



# LEARN C

## THE HARD WAY

BETA  
EDITION

## Preface

This is a rough in-progress dump of the book. The grammar will probably be bad, there will be sections missing, but you get to watch me write the book and see how I do things.

You can also ask for help from me at

[help@learncodethehardway.org](mailto:help@learncodethehardway.org) any time and I'll usually answer within 1 or 2 days.

*This list is a discussion list, not an announce-only list. It's for discussing the book and asking questions.*

Finally, don't forget that I have [Learn Python The Hard Way](#) which you should read if you can't code yet. LCTHW will *not* be for beginners, but for people who have at least read LPTHW or know one other programming language.

1 —  
2 —  
3 — MAIN  
4 —

◀  
PREVIOUS

▶ NEXT

🚑 HELP

Follow  
[@l2sthw](#)

### Interested In Python?

Python is  
also a great  
language.

**Learn  
Python  
The Hard  
Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard  
Way**



# LEARN C

## THE HARD WAY

BETA  
EDITION

## Introduction: The Cartesian Dream Of C

Whatever I have up till now accepted as most true and assured I have gotten either from the senses or through the senses. But from time to time I have found that the senses deceive, and it is prudent never to trust completely those who have deceived us even once.

—Rene Descartes, Meditations On First Philosophy

If there ever were a quote that described programming with C, it would be this. To many programmers, this makes C scary and evil. It is the Devil, Satan, the trickster Loki come to destroy your productivity with his seductive talk of pointers and direct access to the machine. Then, once this computational Lucifer has you hooked, he destroys your world with the evil "segfault" and laughs as he reveals the trickery in your bargain with him.

But, C is not to blame for this state of affairs. No my friends, your computer and the Operating System controlling it are the real tricksters. They conspire to hide their true inner workings from you so that you can never really know what is going on. The C programming language's only failing is giving you access to what is really there, and telling you the cold hard raw truth. C gives you the red pill. C pulls the curtain back to show you the wizard. *C is truth.*

Why use C then if it's so dangerous? Because C gives you power over the false reality of abstraction and liberates you from stupidity.

## What You Will Learn

The purpose of this book is to get you strong enough in C that you'll be able to write your own software in it, or modify someone else's code. At the end of the book we actually take

1 —  
2 —  
3 — MAIN  
4 —

◀  
PREVIOUS

▶ NEXT

🚑  
HELP

Follow  
@l2sthw

code from a more famous book called [K&R C](#) and code review it using what you've learned. To get to this stage you'll have to learn a few things:

- The basics of C syntax and idioms.
- Compilation, make files, linkers.
- Finding bugs and preventing them.
- Defensive coding practices.
- Breaking C code.
- Writing basic Unix systems software.

By the final chapter you will have more than enough ammunition to tackle basic systems software, libraries, and other smaller projects.

## How To Read This Book

This book is intended for programmers who have learned at least one other programming language. I refer you to [Learn Python The Hard Way](#) if you haven't learned a programming language yet. This book is meant for total beginners and works very well as a first book on programming. Once you've done those then you can come back and start this book.

For those who've already learned to code, this book may seem strange at first. It's not like other books where you read paragraph after paragraph of prose and then type in a bit of code here and there. Instead I have you coding right away and then I explain what you just did. This works better because it's easier to explain something you've already experienced.

Because of this structure, there are a few rules you *must* follow in this book:

- Type in all of the code. Do not copy-paste!
- Type the code in exactly, even the comments.
- Get it to run and make sure it prints the same output.
- If there are bugs fix them.
- Do the extra credit but it's alright to skip ones you can't figure out.
- Always try to figure it out first before trying to get help.

If you follow these rules, do everything in the book, and still can't code C then you at least tried. It's not for everyone, but the act of trying will make you a better programmer.

# The Core Competencies

I'm going to guess that you come from a language for weaklings. One of those "usable" languages that lets you get away with sloppy thinking and half-assed hackery like Python or Ruby. Or, maybe you use a language like Lisp that pretends the computer is some purely functional fantasy land with padded walls for little babies. Maybe you've learned Prolog and you think the entire world should just be a database that you walk around in looking for clues. Even worse, I'm betting you've been using an IDE, so your brain is riddled with memory holes and you can't even type out an entire function's name without hitting CTRL-SPACE every 3 characters you type.

No matter what your background, you are probably bad at four skills:

## Reading And Writing

This is especially true if you use an IDE, but generally I find programmers do too much "skimming" and have problems reading for comprehension. They'll skim code they need to understand in detail and think they understand it when they really don't. Other languages provide tools that also let them avoid actually writing any code, so when faced with a language like C they break down. Simplest thing to do is just understand *everyone* has this problem, and you can fix it by forcing yourself to slow down and be meticulous about your reading and writing. At first it'll feel painful and annoying, but take frequent breaks, and then eventually it'll be easy to do.

## Attention To Detail

Everyone is bad at this, and it's the biggest cause of bad software. Other languages let you get away with not paying attention, but C demands your full attention because it is right in the machine and the machine is very picky. With C there is no "kind of similar" or "close enough", so you need to pay attention. Double check your work. Assume everything you write is wrong until you prove it's right.

## Spotting Differences

A key problem people from other languages have is their brain has been trained to spot differences in

*that* language, not in C. When you compare code you've written to my exercise code your eyes will jump right over characters you think don't matter or that aren't familiar. I'll be giving you strategies that force you to see your mistakes, but keep in mind that if your code is not *exactly* like the code in this book it is wrong.

### Planning And Debugging

I love other easier languages because I can just hang out. I type the ideas I have into their interpreter and see results immediately. They're great for just hacking out ideas, but have you noticed that if you keep doing "hack until it works" eventually nothing works? C is harder on you because it requires you to plan out what you'll create first. Sure, you can hack for a bit, but you have to get serious much earlier in C than other languages. I'll be teaching you ways to plan out key parts of your program before you start coding, and this will hopefully make you a better programmer at the same time. Even just a little planning can smooth things out down the road.

Learning C makes you a better programmer because you are forced to deal with these issues earlier and more frequently. You can't be sloppy and half-assed about what you write or nothing will work. The advantage of C is it's a simple language you can figure out on your own, which makes it a great language for learning about the machine and getting stronger in these core programmer skills.

C is harder than some other languages, but that's only because C's not hiding things from you that those other languages try and fail to obfuscate.

## License

This book is free for you to read, but until I'm done you can't distribute it or modify it. I need to make sure that unfinished copies of it do not get out and mess up a student on accident.

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**





# LEARN C

## THE HARD WAY

BETA  
EDITION

## Exercise 0: The Setup

In this chapter you get your system setup to do C programming. The good news for anyone using Linux or Mac OSX is that you are on a system designed *for* programming in C. The authors of the C language were also instrumental in the creation of the Unix operating system, and both Linux and OSX are based on Unix. In fact, the install will be incredibly easy.

I have some bad news for users of Windows: learning C on Windows is painful. You can write C code for Windows, that's not a problem. The problem is all of the libraries, functions, and tools are just a little "off" from everyone else in the C world. C came from Unix and is much easier on a Unix platform. It's just a fact of life that you'll have to accept I'm afraid.

I wanted to get this bad news out right away so that you don't panic. I'm not saying to avoid Windows entirely. I am however saying that, if you want to have the easiest time learning C, then it's time to bust out some Unix and get dirty. This could also be really good for you, since knowing a little bit of Unix will also teach you some of the idioms of C programming and expand your skills.

This also means that for everyone you'll be using the *command line*. Yep, I said it. You've gotta get in there and type commands at the computer. Don't be afraid though because I'll be telling you what to type and what it should look like, so you'll actually be learning quite a few mind expanding skills at the same time.

## Linux

On most Linux systems you just have to install a few packages. For Debian based systems, like Ubuntu you should just have to install a few things using these commands:

1 —  
2 —  
3 —  
4 — MAIN

◀  
PREVIOUS

▶ NEXT

🚑  
HELP

Follow  
@l2sthw

```
$ sudo apt-get install build-essential
```

The above is an example of a command line prompt, so to get to where you can run that, find your "Terminal" program and run it first. Then you'll get a shell prompt similar to the `$` above and can type that command into it. *Do not type the ``\$``, just the stuff after it.*

Here's how you would install the same setup on an RPM based Linux like Fedora:

```
$ su -c "yum groupinstall development-tools"
```

Once you've run that, you should be able to do the first Exercise in this book and it'll work. If not then let me know.

## Mac OSX

On Mac OSX the install is even easier. First, you'll need to either download the latest `XCode` from Apple, or find your install DVD and install it from there. The download will be massive and could take forever, so I recommend installing from the DVD. Also, search online for "installing xcode" for instructions on how to do it.

Once you're done installing XCode, and probably restarting your computer if it didn't make you do that, you can go find your Terminal program and get it put into your Dock. You'll be using Terminal a lot in the book, so it's good to put it in a handy location.

## Windows

For Windows users I'll show you how to get a basic Ubuntu Linux system up and running in a virtual machine so that you can still do all of my exercises, but avoid all the painful Windows installation problems.

... have to figure this one out.

## Text Editor

The choice of text editor for a programmer is a tough one. For beginners I tell them to just use `Gedit` since it's simple and



works for code. However, it doesn't work in certain internationalized situations, and chances are you already have a favorite text editor if you've been programming for a while.

With this in mind, I want you to try out a few of the standard programmer text editors for your platform and then stick with the one that you like best. If you've been using GEdit and like it then stick with it. If you want to try something different, then try it out real quick and pick one.

The most important thing is *do not get stuck picking the perfect editor*. Text editors all just kind of suck in odd ways. Just pick one, stick with it, and if you find something else you like try it out. Don't spend days on end configuring it and making it perfect.

Some text editors to try out are:

- **Gedit** on Linux and OSX.
- **TextWrangler** on OSX.
- **Nano** which runs in Terminal and works nearly everywhere.
- **Emacs** and **Emacs for OSX**. Be prepared to do some learning though.
- **Vim** and **MacVim**

There is probably a different editor for every person out there, but these are just a few of the free ones that I know work. Try a few out, and maybe some commercial ones until you find one that you like.

## WARNING: Do Not Use An IDE

An IDE, or "Integrated Development Environment" will turn you stupid. They are the worst tools if you want to be a good programmer because they hide what's going on from you, and your job is to know what's going on. They are useful if you're trying to get something done and the platform is designed around a particular IDE, but for learning to code C (and many other languages) they are pointless.

### NOTE

If you've played guitar then you know what tablature is, but for everyone else

let me explain. In music there's an established notation called the "staff notation". It's a generic, very old, and universal way to write down what someone should play on an instrument. If you play piano this notation is fairly easy to use, since it was created mostly for piano and composers.

Guitar however is a weird instrument that doesn't really work with notation, so guitarists have an alternative notation called "tablature". What tablature does is, rather than tell you the note to play, it tells you the fret and string you should play at that time. You could learn whole songs without ever knowing about a single thing you're playing. Many people do it this way, but if you want to know *what* you're playing, then tablature is pointless.

It may be harder than tablature, but traditional notation tells you how to play the *music* rather than just how to play the guitar. With traditional notation I can walk over to a piano and play the same song. I can play it on a bass. I can put it into a computer and design whole scores around it. With tablature I can just play it on a guitar.

IDEs are like tablature. Sure, you can code pretty quickly, but you can only code in that one language on that one platform. This is why companies love selling them to you. They know you're lazy, and since it only works on their platform they've got you locked in because you are lazy.

The way you break the cycle is you suck it up and finally learn to code without an IDE. A plain editor, or a programmer's editor like Vim or Emacs, makes you work with the code. It's a little harder, but the end result is you can work with *any* code,

on any computer, in any language, and  
you know what's going on.

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Learn  
Python  
The Hard  
Way**

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**



## Exercise 1: Dust Off That Compiler

Here is a simple first program you can make in C:

```
int main(int argc, char *argv[])
{
    puts("Hello world.");

    return 0;
}
```

You can put this into a `ex1.c` then type:

```
$ make ex1
cc      ex1.c  -o ex1
```

Your computer may use a slightly different command, but the end result should be a file named `ex1` that you can run.

## What You Should See

You can now run the program and see the output.

```
$ ./ex1
Hello world.
```

If you don't then go back and fix it.

## How To Break It

In this book I'm going to have a small section for each program on how to break the program. I'll have you do odd things to the programs, run them in weird ways, or change code so that you can see crashes and compiler errors.

For this program, rebuild it with all compiler warnings on:

```
$ rm ex1
$ CFLAGS="-Wall" make ex1
cc -Wall    ex1.c  -o ex1
ex1.c: In function 'main':
ex1.c:3: warning: implicit declaration of function
'puts'
$ ./ex1
Hello world.
$
```

Now you are getting a warning that says the function "puts" is implicitly declared. The C compiler is smart enough to figure out what you want, but you should be getting rid of all compiler warnings when you can. How you do this is add the following line to the top of `ex1.c` and recompile:

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssth

```
#include <stdio.h>
```

Now do the make again like you just did and you'll see the warning go away.

## Extra Credit

- Open the `ex1` file in your text editor and change or delete random parts. Try running it and see what happens.
- Print out 5 more lines of text or something more complex than hello world.
- Run `man 3 puts` and read about this function and many others.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 2: Make Is Your Python Now

In `Python` you ran programs by just typing `python` and the code you wanted to run. The Python interpreter would just run them, and import any other libraries and things you needed on the fly as it ran. C is a different beast completely where you have to *compile* your source files and manually stitch them together into a binary that can run on its own. Doing this manually is a pain, and in the last exercise you just ran `make` to do it.

In this exercise, you're going to get a crash course in GNU make, and you'll be learning to use it as you learn C. Make will for the rest of this book, be your Python. It will build your code, and run your tests, and set things up and do all the stuff for you that Python normally does.

The difference is, I'm going to show you smarter Makefile wizardry, where you don't have to specify every stupid little thing about your C program to get it to build. I won't do that in this exercise, but after you've been using "baby make" for a while, I'll show you "master make".

## Using Make

The first stage of using make is to just use it to build programs it already knows how to build. Make has decades of knowledge on building a wide variety of files from other files. In the last exercise you did this already using commands like this:

```
$ make ex1
# or this one too
$ CFLAGS="-Wall" make ex1
```

In the first command you're telling make, "I want a file named `ex1` to be created." Make then does the following:

- Does the file `ex1` exist already?
- No. Ok, is there another file that starts with `ex1`?
- Yes, it's called `ex1.c`. Do I know how to build `.c` files?
- Yes, I run this command `cc ex1.c -o ex1` to build them.
- I shall make you one `ex1` by using `cc` to build it from `ex1.c`.

The second command in the listing above is a way to pass "modifiers" to the make command. If you're not familiar with how the Unix shell works, you can create these "environment variables" which will get picked up by programs you run. Sometimes you do this with a command like `export CFLAGS="-Wall"` depending on the shell you use. You can however also just put them before the command you want to run, and that environment variable will be set only while that command runs.

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

⚙  
HELP

Follow  
@l2sthw

In this example I did `CFLAGS="-Wall" make ex1` so that it would add the command line option `-Wall` to the `cc` command that `make` normally runs. That command line option tells the compiler `cc` to report all warnings (which in a sick twist of fate isn't actually all the warnings possible).

You can actually get pretty far with just that way of using `make`, but let's get into making a `Makefile` so you can understand make a little better. To start off, create a file with just this in it:

```
CFLAGS=-Wall -g

clean:
    rm -f ex1
```

Save this file as `Makefile` in your current directory. Make automatically assumes there's a file called `Makefile` and will just run it. Also, *WARNING: Make sure you are only entering TAB characters, not mixtures of TAB and spaces.*

This `Makefile` is showing you some new stuff with make. First we set `CFLAGS` in the file so we never have to set it again, as well as adding the `-g` flag to get debugging. Then we have a section named `clean` which tells make how to clean up our little project.

Make sure it's in the same directory as your `ex1.c` file, and then run these commands:

```
$ make clean
$ make ex1
```

## What You Should See

If that worked then you should see this:

```
$ make clean
rm -f ex1
$ make ex1
cc -Wall -g    ex1.c    -o ex1
ex1.c: In function 'main':
ex1.c:3: warning: implicit declaration of function 'puts'
$
```

Here you can see that I'm running `make clean` which tells make to run our `clean` target. Go look at the Makefile again and you'll see that under this I indent and then I put the shell commands I want `make` to run for me. You could put as many commands as you wanted in there, so it's a great automation tool.

### NOTE

If you fixed `ex1.c` to have `#include <stdio.h>` then your output will not have the warning (which should really be an error) about puts. I have the error here because I didn't fix it.

Notice also that, even though we don't mention `ex1` in the `Makefile`, `make` still knows how to build it *plus* use our special settings.

## How To Break It

That should be enough to get you started, but first let's break this make file in a particular way so you can see what happens. Take the line `rm -f ex1` and dedent it (move it all the way left) so you can see what happens. Rerun `make clean` and you should get something like this:

```
$ make clean
Makefile:4: *** missing separator.  Stop.
```

Always remember to indent, and if you get weird errors like this then double check you're consistently using tab characters since some make variants are very picky.

## Extra Credit

- Create an `all: ex1` target that will build `ex1` with just the command `make`.
- Read `man make` to find out more information on how to run it.
- Read `man cc` to find out more information on what the flags `-Wall` and `-g` do.
- Research Makefiles online and see if you can improve this one even more.
- Find a `Makefile` in another C project and try to understand what it's doing.

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 3: Formatted Printing

Keep that `Makefile` around since it'll help you spot errors and we'll be adding to it when we need to automate more things.

Many programming languages use the C way of formatting output, so let's try it:

```
#include <stdio.h>

int main()
{
    int age = 10;
    int height = 72;

    printf("I am %d years old.\n", age);
    printf("I am %d inches tall.\n", height);

    return 0;
}
```

Once you have that, do the usual `make ex3` to build it and run it. Make sure you *fix all warnings*.

This exercise has a whole lot going on in a small amount of code so let's break it down:

- First you're including another "header file" called `stdio.h`. This tells the compiler that you're going to use the "standard Input/Output functions". One of those is `printf`.
- Then you're using a variable named `age` and setting it to 10.
- Next you're using a variable `height` and setting it to 72.
- Then you use the `printf` function to print the age and height of the tallest 10 year old on the planet.
- In the `printf` you'll notice you're passing in a string, and it's a format string like in many other languages.
- After this format string, you put the variables that should be "replaced" into the format string by `printf`.

The result of doing this is you are handing `printf` some variables and it is constructing a new string then printing that new string to the terminal.

## What You Should See

When you do the whole build you should see something like this:

```
$ make ex3
cc -Wall -g    ex3.c    -o ex3
$ ./ex3
I am 10 years old.
I am 72 inches tall.
$
```

Pretty soon I'm going to stop telling you to run `make` and what the build

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssth

looks like, so please make sure you're getting this right and that it's working.

## External Research

In the *Extra Credit* section of each exercise I may have you go find information on your own and figure things out. This is an important part of being a self-sufficient programmer. If you constantly run to ask someone a question before trying to figure it out first then you never learn to solve problems independently. This leads to you never building confidence in your skills and always needing someone else around to do your work.

The way you break this habit is to *force* yourself to try to answer your own questions first, and to confirm that your answer is right. You do this by trying to break things, experimenting with your possible answer, and doing your own research.

For this exercise I want you to go online and find out *all* of the `printf` escape codes and format sequences. Escape codes are `\n` or `\t` that let you print a newline or tab (respectively). Format sequences are the `%s` or `%d` that let you print a string or a integer. Find all of the ones available, how you can modify them, and what kind of "precisions" and widths you can do.

From now on, these kinds of tasks will be in the Extra Credit and you should do them.

## How To Break It

Try a few of these ways to break this program, which may or may not cause it to crash on your computer:

- Take the `age` variable out of the first `printf` call then recompile. You should get a couple of warnings.
- Run this new program and it will either crash, or print out a really crazy age.
- Put the `printf` back the way it was, and then don't set `age` to an initial value by changing that line to `int age;` then rebuild and run again.

```
# edit ex3.c to break printf
$ make ex3
cc -Wall -g    ex3.c    -o ex3
ex3.c: In function 'main':
ex3.c:8: warning: too few arguments for format
ex3.c:5: warning: unused variable 'age'
$ ./ex3
I am -919092456 years old.
I am 72 inches tall.
# edit ex3.c again to fix printf, but don't init age
$ make ex3
cc -Wall -g    ex3.c    -o ex3
ex3.c: In function 'main':
ex3.c:8: warning: 'age' is used uninitialized in this function
$ ./ex3
I am 0 years old.
I am 72 inches tall.
$
```

## Extra Credit

- Find as many other ways to break `ex3.c` as you can.
- Run `man 3 printf` and read about the other '%' format characters you can use. These should look familiar if you used them in other languages (`printf` is where they come from).
- Add `ex3` to your `Makefile`'s `all` list. Use this to `make clean all` and build all your exercises so far.
- Add `ex3` to your `Makefile`'s `clean` list as well. Now use `make clean` will remove it when you need to.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 4: Introducing Valgrind

It's time to learn about another tool you will live and die by as you learn C called `Valgrind`. I'm introducing `Valgrind` to you now because you're going to use it from now on in the "How To Break It" sections of each exercise. `Valgrind` is a program that runs your programs, and then reports on all of the horrible mistakes you made. It's a wonderful free piece of software that I use constantly while I write C code.

Remember in the last exercise that I told you to break your code by removing one of the arguments to `printf`? It printed out some funky results, but I didn't tell you why it printed those results out. In this exercise we're going to use `Valgrind` to find out why.

### NOTE

These first few exercises are mixing some essential tools the rest of the book needs with learning a little bit of code. The reason is that most of the folks who read this book are not familiar with compiled languages, and definitely not with automation and helpful tools. By getting you to use `make` and `Valgrind` right now I can then use them to teach you C faster and help you find all your bugs early.

After this exercise we won't do many more tools, it'll be mostly code and syntax for a while. But, we'll also have a few tools we can use to really see what's going on and get a good understanding of common mistakes and problems.

## Installing Valgrind

You could install `Valgrind` with the package manager for your OS, but I want you to learn to install things from source. This involves the following process:

- Download a source archive file to get the source.
- Unpack the archive to extract the files onto your computer.
- Run `./configure` to setup build configurations.
- Run `make` to make it build, just like you've been doing.
- Run `sudo make install` to install it onto your computer.

Here's a script of me doing this very process, which I want you to try to replicate:

```
# 1) Download it (use wget if you don't have curl)
curl -O http://valgrind.org/downloads/valgrind-3.6.1.tar.bz2

# use md5sum to make sure it matches the one on the site
md5sum valgrind-3.6.1.tar.bz2

# 2) Unpack it.
tar -xjvf valgrind-3.6.1.tar.bz2

# cd into the newly created directory
cd valgrind-3.6.1

# 3) configure it
./configure

# 4) make it
make

# 5) install it (need root)
sudo make install
```

Follow this, but obviously update it for new Valgrind versions. If it doesn't build then try digging into why as well.

## Using Valgrind

Using `Valgrind` is easy, you just run `valgrind theprogram` and it runs your program, then prints out all the errors your program made while it was running. In this exercise we'll break down one of the error outputs and you can get an instant crash course in "Valgrind hell". Then we'll fix the program.

First, here's a purposefully broken version of the `ex3.c` code for you to build, now called `ex4.c`. For practice, type it in again:

```
#include <stdio.h>

/* Warning: This program is wrong on purpose. */

int main()
{
    int age = 10;
    int height;

    printf("I am %d years old.\n");
    printf("I am %d inches tall.\n", height);

    return 0;
}
```

You'll see it's the same except I've made two classic mistakes:

- I've failed to initialize the `height` variable.
- I've forgot to give the first `printf` the `age` variable.

## What You Should See

Now we will build this just like normal, but instead of running it directly, we'll run it with `Valgrind` (see Source: "Building and running ex4.c with

Valgrind"):

```
$ make ex4
cc -Wall -g    ex4.c    -o ex4
ex4.c: In function 'main':
ex4.c:10: warning: too few arguments for format
ex4.c:7: warning: unused variable 'age'
ex4.c:11: warning: 'height' is used uninitialized in this function
$ valgrind ./ex4
==3082== Memcheck, a memory error detector
==3082== Copyright (C) 2002-2010, and GNU GPL'd, by Julian Seward et al.
==3082== Using Valgrind-3.6.0.SVN-Debian and LibVEX; rerun with -h for copyrig
ht info
==3082== Command: ./ex4
==3082==
I am -16775432 years old.
==3082== Use of uninitialised value of size 8
==3082==    at 0x4E730EB: _itoa_word (_itoa.c:195)
==3082==    by 0x4E743D8: vfprintf (vfprintf.c:1613)
==3082==    by 0x4E7E6F9: printf (printf.c:35)
==3082==    by 0x40052B: main (ex4.c:11)
==3082==
==3082== Conditional jump or move depends on uninitialised value(s)
==3082==    at 0x4E730F5: _itoa_word (_itoa.c:195)
==3082==    by 0x4E743D8: vfprintf (vfprintf.c:1613)
==3082==    by 0x4E7E6F9: printf (printf.c:35)
==3082==    by 0x40052B: main (ex4.c:11)
==3082==
==3082== Conditional jump or move depends on uninitialised value(s)
==3082==    at 0x4E7633B: vfprintf (vfprintf.c:1613)
==3082==    by 0x4E7E6F9: printf (printf.c:35)
==3082==    by 0x40052B: main (ex4.c:11)
==3082==
==3082== Conditional jump or move depends on uninitialised value(s)
==3082==    at 0x4E744C6: vfprintf (vfprintf.c:1613)
==3082==    by 0x4E7E6F9: printf (printf.c:35)
==3082==    by 0x40052B: main (ex4.c:11)
==3082==
I am 0 inches tall.
==3082==
==3082== HEAP SUMMARY:
==3082==    in use at exit: 0 bytes in 0 blocks
==3082==    total heap usage: 0 allocs, 0 frees, 0 bytes allocated
==3082==
==3082== All heap blocks were freed -- no leaks are possible
==3082==
==3082== For counts of detected and suppressed errors, rerun with: -v
==3082== Use --track-origins=yes to see where uninitialised values come from
==3082== ERROR SUMMARY: 4 errors from 4 contexts (suppressed: 4 from 4)
$
```

## NOTE

If you run valgrind and it says something like `by 0x4052112: (below main) (libc-start.c:226)` instead of a line number in `main.c` then add run your valgrind command like this `valgrind --track-origins=yes ./ex4` to make it work. For some reason the Debian or Ubuntu version of `valgrind` does this but not other versions.

This one is huge because `Valgrind` is telling you exactly where every problem in your program is. Starting at the top here's what you're reading, line by line (line numbers are on the left so you can follow):

1

You do the usual `make ex4` and that builds it. Make sure the `cc` command you see is the same and has the `-g` option or your `Valgrind` output won't have line numbers.

2-6

Notice that the compiler is also yelling at you about this source file and it warns you that you have "too few arguments for format". That's where you forgot to include the `age` variable.

7

Then you run your program using `valgrind ./ex4`.

8

Then `Valgrind` goes crazy and yells at you for:

14-18

On line `main (ex4.c:11)` (read as "in the main function in file ex4.c at line 11) you have "Use of uninitialised value of size 8". You find this by looking at the error, then you see what's called a "stack trace" right under that. The line to look at first (ex4.c:11) is the bottom one, and if you don't see what's going wrong then you go up, so you'd try `printf.c:35`. Typically it's the bottom most line that matters (in this case, on line 18).

20-24

Next error is yet another one on line ex4.c:11 in the main function. `Valgrind` hates this line. This error says that some kind of if-statement or while-loop happened that was based on an uninitialized variable, in this case `height`.

25-35

The remaining errors are more of the same because the variable keeps getting used.

37-46

Finally the program exits and `Valgrind` tells you a summary of how bad your program is.

That is quite a lot of information to take in, but here's how you deal with it:

- Whenever you run your C code and get it working, rerun it under `Valgrind` to check it.
- For each error that you get, go to the source:line indicated and fix it. You may have to search online for the error message to figure out what it means.
- Once your program is "Valgrind pure" then it should be good, and you have probably learned something about how you write code.

In this exercise I'm not expecting you to fully grasp `Valgrind` right away, but instead get it installed and learn how to use it real quick so we can apply it to all the later exercises.

## Extra Credit

- Fix this program using `Valgrind` and the compiler as your guide.
- Read up on `Valgrind` on the internet.
- Download other software and build it by hand. Try something you already use but never built for yourself.
- Look at how the `Valgrind` source files are laid out in the source directory and read its Makefile. Don't worry, none of that makes sense to me either.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 5: The Structure Of A C Program

You know how to use `printf` and have a couple basic tools at your disposal, so let's break down a simple C program line-by-line so you know how one is structured. In this program you're going to type in a few more things that you're unfamiliar with, and I'm going to lightly break them down. Then in the next few exercises we're going to work with these concepts.

```
#include <stdio.h>

/* This is a comment. */
int main(int argc, char *argv[])
{
    int distance = 100;

    // this is also a comment
    printf("You are %d miles away.\n", distance);

    return 0;
}
```

Type this code in, make it run, and make sure you get *no Valgrind errors*. You probably won't but get in the habit of checking it.

## What You Should See

This has pretty boring output, but the point of this exercise is to analyze the code:

```
$ make ex5
cc -Wall -g    ex5.c    -o ex5
$ ./ex5
You are 100 miles away.
$
```

## Breaking It Down

There's a few features of the C language in this code that you might have only slightly figured out while you were typing code. Let's break this down line-by-line quickly, and then we can do exercises to understand each part better:

ex5.c:1

An `include` and it is the way to import the contents of one file into this source file. C has a convention of using `.h` extensions for "header" files, which then contain lists of functions you want to use in your program.

ex5.c:3

This is a multi-line `comment` and you could put as many lines of

1  
2  
3  
4 MAIN

◀  
PREVIOUS

▶ NEXT

🛠️ HELP

Follow  
@lzssth

text between the `/*` and closing `*/` characters as you want.

ex5.c:4

A more complex version of the `main function` you've been using blindly so far. How C programs work is the operating system loads your program, and then runs the function named `main`. For the function to be totally complete it needs to return an `int` and take two parameters, an `int` for the argument count, and an array of `char *` strings for the arguments. Did that just fly over your head? Do not worry, we'll cover this soon.

ex5.c:5

To start the body of any function you write a `{` character that indicates the beginning of a "block". In Python you just did a `:` and indented. In other languages you might have a `begin` or `do` word to start.

ex5.c:6

A variable declaration and assignment at the same time. This is how you create a variable, with the syntax `type name = value;`. In C statements (except for logic) end in a `;` (semicolon) character.

ex5.c:8

Another kind of comment, and it works like Python or Ruby comments where it starts at the `//` and goes until the end of the line.

ex5.c:9

A call to your old friend `printf`. Like in many languages function calls work with the syntax `name(arg1, arg2);` and can have no arguments, or any number. The `printf` function is actually kind of weird and can take multiple arguments. We'll see that later.

ex5.c:11

A return from the main function, which gives the OS your exit value. You may not be familiar with how Unix software uses return codes, so we'll cover that as well.

ex5.c:12

Finally, we end the main function with a closing brace `}` character and that's the end of the program.

There's a lot of information in this break-down, so study it line-by-line and make sure you at least have a little grasp of what's going on. You won't know everything, but you can probably guess before we continue.

## Extra Credit

- For each line, write out the symbols you don't understand and see if you can guess what they mean. Write a little chart on paper with your guess that you can use to check later and see if you get it right.
- Go back to the source code from the previous exercises and do a similar break-down to see if you're getting it. Write down what you don't know and can't explain to yourself.

**Interested  
In  
Python?**

Python is also a

**Interested In Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard Way**

great language.

**Learn  
Python The  
Hard Way**



## Exercise 6: Types Of Variables

You should be getting a grasp of how a simple C program is structured, so let's do the next simplest thing which is making some variables of different types:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int distance = 100;
    float power = 2.345f;
    double super_power = 56789.4532;
    char initial = 'A';
    char first_name[] = "Zed";
    char last_name[] = "Shaw";

    printf("You are %d miles away.\n", distance);
    printf("You have %f levels of power.\n", power);
    printf("You have %f awesome super powers.\n", super_power);
    printf("I have an initial %c.\n", initial);
    printf("I have a first name %s.\n", first_name);
    printf("I have a last name %s.\n", last_name);
    printf("My whole name is %s %c. %s.\n",
           first_name, initial, last_name);

    return 0;
}
```

In this program we're declaring variables of different types and then printing them with different `printf` format strings.

## What You Should See

Your output should look like mine, and you can start to see how the format strings for C are similar to Python and other languages. They've been around for a long time.

```
$ make ex6
cc -Wall -g    ex6.c    -o ex6
$ ./ex6
You are 100 miles away.
You have 2.345000 levels of power.
You have 56789.453200 awesome super powers.
I have an initial A.
I have a first name Zed.
I have a last name Shaw.
My whole name is Zed A. Shaw.
$
```

What you can see is we have a set of "types", which are ways of telling the C compiler what each variable should represent, and then format strings to match different types. Here's the breakdown of how they match up:

Integers

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssth

You declare Integers with the `int` keyword, and print them with `%d`.

#### Floating Point

Declared with `float` or `double` depending on how big they need to be (double is bigger), and printed with `%f`.

#### Character

Declared with `char`, written with a `'` (single-quote) character around the char, and then printed with `%c`.

#### String (Array of Characters)

Declared with `char name[]`, written with `"` characters, and printed with `%s`.

You'll notice that C makes a distinction between single-quote for `char` and double-quote for `char[]` or strings.

### NOTE

When talking about C types, I will typically write in English `char[]` instead of the whole `char SOMENAME[]`. This is not valid C code, just a simpler way to talk about types when writing English.

## How To Break It

You can easily break this program by passing the wrong thing to the `printf` statements. For example, if you take the line that prints my name, but put the `initial` variable before the `first_name` in the arguments, you'll get a bug. Make that change and the compiler will yell at you, then when you run it you might get a "Segmentation fault" like I did:

```
$ make ex6
cc -Wall -g    ex6.c    -o ex6
ex6.c: In function 'main':
ex6.c:19: warning: format '%s' expects type 'char *', but argument 2 has type 'int'
ex6.c:19: warning: format '%c' expects type 'int', but argument 3 has type 'char *'
$ ./ex6
You are 100 miles away.
You have 2.345000 levels of power.
You have 56789.453125 awesome super powers.
I have an initial A.
I have a first name Zed.
I have a last name Shaw.
Segmentation fault
$
```

Run this change under Valgrind too to see what it tells you about the error "Invalid read of size 1".

## Extra Credit

- Come up with other ways to break this C code by changing the `printf`, then fix them.

- Go search for "printf formats" and try using a few of the more exotic ones.
- Research how many different ways you can write a number. Try octal, hexadecimal, and others you can find.
- Try printing an empty string that's just `""`.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 7: More Variables, Some Math

Let's get familiar with more things you can do with variables by declaring various `ints`, `floats`, `chars`, and `doubles`. We'll then use these in various math expressions so you get introduced to C's basic math.

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int bugs = 100;
    double bug_rate = 1.2;

    printf("You have %d bugs at the imaginary rate of %f.\n",
           bugs, bug_rate);

    long universe_of_defects = 1L * 1024L * 1024L * 1024L;
    printf("The entire universe has %ld bugs.\n",
           universe_of_defects);

    double expected_bugs = bugs * bug_rate;
    printf("You are expected to have %f bugs.\n",
           expected_bugs);

    double part_of_universe = expected_bugs / universe_of_defects;
    printf("That is only a %e portion of the universe.\n",
           part_of_universe);

    // this makes no sense, just a demo of something weird
    char nul_byte = '\0';
    int care_percentage = bugs * nul_byte;
    printf("Which means you should care %d%%.\n",
           care_percentage);

    return 0;
}
```

Here's what's going on in this little bit of nonsense:

ex7.c:1-4

The usual start of a C program.

ex7.c:5-6

Declare an `int` and `double` for some fake bug data.

ex7.c:8-9

Print out those two, so nothing new here.

ex7.c:11

Declare a huge number using a new type `long` for storing big numbers.

ex7.c:12-13

Print out that number using `%ld` which adds a modifier to the usual `%d`. Adding 'l' (the letter ell) means "print this as a long decimal".

ex7.c:15-17

Just more math and printing.

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssthw

ex7.c:19-21

Craft up a depiction of your bug rate compared to the bugs in the universe, which is a completely inaccurate calculation. It's so small though that we have to use `%e` to print it in scientific notation.

ex7.c:24

Make a character, with a special syntax `'\0'` which creates a 'nul byte' character. This is effectively the number 0.

ex7.c:25

Multiply bugs by this character, which produces 0 for how much you should care. This demonstrates an ugly hack you find sometimes.

ex7.c:26-27

Print that out, and notice I've got a `'%%'` (two percent chars) so I can print a `'%'` (percent) character.

ex7.c:28-30

The end of the `main` function.

This bit of source is entirely just an exercise, and demonstrates how some math works. At the end, it also demonstrates something you see in C, but not in many other languages. To C, a "character" is just an integer. It's a really small integer, but that's all it is. This means you can do math on them, and a lot of software does just that, for good or bad.

This last bit is your first glance at how C gives you direct access to the machine. We'll be exploring that more in later exercises.

## What You Should See

As usual, here's what you should see for the output:

```
$ make ex7
cc -Wall -g    ex7.c    -o ex7
$ ./ex7
You have 100 bugs at the imaginary rate of 1.200000.
The entire universe has 1073741824 bugs.
You are expected to have 120.000000 bugs.
That is only a 1.117587e-07 portion of the universe.
Which means you should care 0%.
$
```

## How To Break It

Again, go through this and try breaking the `printf` by passing in the wrong arguments. See what happens when you try to print out that `nul_byte` variable too with `%s` vs. `%c`. When you break it, run it under `Valgrind` to see what it says about your breaking attempts.

## Extra Credit

- Make the number you assign to `universe_of_defects` various sizes until you get a warning from the compiler.
- What do these really huge numbers actually print out?
- Change `long` to `unsigned long` and try to find the number that makes that one too big.
- Go search online to find out what `unsigned` does.



- Try to explain to yourself (before I do in the next exercise) why you can multiply a `char` and an `int`.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 8: Sizes And Arrays

In the last exercise you did math, but with a `'\0'` (nul) character. This may be odd coming from other languages, since they try to treat "strings" and "byte arrays" as different beasts. C however treats strings as just arrays of bytes, and it's only the different printing functions that know there's a difference.

Before I can really explain the significance of this, I have to introduce a few more concepts: `sizeof` and arrays. Here's the code we'll be talking about:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int areas[] = {10, 12, 13, 14, 20};
    char name[] = "Zed";
    char full_name[] = {
        'Z', 'e', 'd',
        ' ', 'A', ' ',
        'S', 'h', 'a', 'w', '\0'
    };

    // WARNING: On some systems you may have to change the
    // %ld in this code to a %u since it will use unsigned ints
    printf("The size of an int: %ld\n", sizeof(int));
    printf("The size of areas (int[]): %ld\n",
        sizeof(areas));
    printf("The number of ints in areas: %ld\n",
        sizeof(areas) / sizeof(int));
    printf("The first area is %d, the 2nd %d.\n",
        areas[0], areas[1]);

    printf("The size of a char: %ld\n", sizeof(char));
    printf("The size of name (char[]): %ld\n",
        sizeof(name));
    printf("The number of chars: %ld\n",
        sizeof(name) / sizeof(char));

    printf("The size of full_name (char[]): %ld\n",
        sizeof(full_name));
    printf("The number of chars: %ld\n",
        sizeof(full_name) / sizeof(char));

    printf("name=\"%s\" and full_name=\"%s\"\n",
        name, full_name);

    return 0;
}
```

In this code we create a few arrays with different data types in them. Because arrays of data are so central to how C works, there's a huge number of ways to create them. For now, just use the syntax `type name[] = {initializer};` and we'll explore more. What this syntax means is, "I want an array of type that is initialized to {..}." When C sees this it does the following:

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzzsthw

- Look at the type, in this first case it's `int`.
- Look at the `[]` and see that there's no length given.
- Look at the initializer, `{10, 12, 13, 14, 20}` and figure out that you want those 5 ints in your array.
- Create a piece of memory in the computer, that can hold 5 integers one after another.
- Take the name you want, `areas` and assign it this location.

In the case of `areas` it's creating an array of 5 ints that contain those numbers. When it gets to `char name[] = "Zed";` it's doing the same thing, except it's creating an array of 3 chars and assigning that to `name`. The final array we make is `full_name`, but we use the annoying syntax of spelling it out, one character at a time. To C, `name` and `full_name` are identical methods of creating a char array.

The rest of the file, we're using a keyword called `sizeof` to ask C how big things are in *bytes*. C is all about the size and location of pieces of memory and what you do with them. To help you keep that straight, it gives you `sizeof` so you can ask how big something is before you work with it.

This is where stuff gets tricky, so first let's run this and then explain further.

## What You Should See

```
$ make ex8
cc -Wall -g    ex8.c    -o ex8
$ ./ex8
The size of an int: 4
The size of areas (int[]): 20
The number of ints in areas: 5
The first area is 10, the 2nd 12.
The size of a char: 1
The size of name (char[]): 4
The number of chars: 4
The size of full_name (char[]): 12
The number of chars: 12
name="Zed" and full_name="Zed A. Shaw"
$
```

Now you see the output of these different `printf` calls and start to get a glimpse of what C is doing. Your output could actually be totally different from mine, since your computer might have different size integers. I'll go through my output:

5

My computer thinks an `int` is 4 bytes in size. Your computer might use a different size if it's a 32-bit vs. 64-bit.

6

The `areas` array has 5 integers in it, so it makes sense that my computer requires 20 bytes to store it.

7

If we divide the size of `areas` by size of an `int` then we get 5 elements. Looking at the code, this matches what we put in the initializer.

8

We then did an array access to get `areas[0]` and `areas[1]` which means C is "zero indexed" like Python and Ruby.

9-11

We repeat this for the `name` array, but notice something odd about the size of the array? It says it's 4 bytes long, but we only typed "Zed" for 3 characters. Where's the 4th one coming from?

12-13

We do the same thing with `full_name` and notice it gets this correct.

13

Finally we just print out the `name` and `full_name` to prove that they actually are "strings" according to printf.

Make sure you can go through and see how these output lines match what was created. We'll be building on this and exploring more about arrays and storage next.

## How To Break It

Breaking this program is fairly easy. Try some of these:

- Get rid of the `'\0'` at the end of `full_name` and re-run it. Run it under Valgrind too. Now, move the definition of `full_name` to the top of `main` before `areas`. Try running it under Valgrind a few times and see if you get some new errors. In some cases, you might still get lucky and not catch any errors.
- Change it so that instead of `areas[0]` you try to print `areas[10]` and see what Valgrind thinks of that.
- Try other versions of these, doing it to `name` and `full_name` too.

## Extra Credit

- Try assigning to elements in the `areas` array with `areas[0] = 100;` and similar.
- Try assigning to elements of `name` and `full_name`.
- Try setting one element of `areas` to a character from `name`.
- Go search online for the different sizes used for integers on different CPUs.

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 9: Arrays And Strings

In the last exercise you went through an introduction to creating basic arrays and how they map to strings. In this exercise we'll more completely show the similarity between arrays and strings, and get into more about memory layouts.

This exercise shows you that C stores its strings simply as an array of bytes, terminated with the `'\0'` (nul) byte. You probably clued into this in the last exercise since we did it manually. Here's how we do it in another way to make it even more clear by comparing it to an array of numbers:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int numbers[4] = {0};
    char name[4] = {'a'};

    // first, print them out raw
    printf("numbers: %d %d %d %d\n",
           numbers[0], numbers[1],
           numbers[2], numbers[3]);

    printf("name each: %c %c %c %c\n",
           name[0], name[1],
           name[2], name[3]);

    printf("name: %s\n", name);

    // setup the numbers
    numbers[0] = 1;
    numbers[1] = 2;
    numbers[2] = 3;
    numbers[3] = 4;

    // setup the name
    name[0] = 'Z';
    name[1] = 'e';
    name[2] = 'd';
    name[3] = '\0';

    // then print them out initialized
    printf("numbers: %d %d %d %d\n",
           numbers[0], numbers[1],
           numbers[2], numbers[3]);

    printf("name each: %c %c %c %c\n",
           name[0], name[1],
           name[2], name[3]);

    // print the name like a string
    printf("name: %s\n", name);

    // another way to use name
    char *another = "Zed";

    printf("another: %s\n", another);
```

```

printf("another each: %c %c %c %c\n",
      another[0], another[1],
      another[2], another[3]);

return 0;
}

```

In this code, we setup some arrays the tedious way, by assigning a value to each element. In `numbers` we are setting up numbers, but in `name` we're actually building a string manually.

## What You Should See

When you run this code you should see first the arrays printed with their contents initialized to zero, then in its initialized form:

```

$ make ex9
cc -Wall -g    ex9.c    -o ex9
$ ./ex9
numbers: 0 0 0 0
name each: a
name: a
numbers: 1 2 3 4
name each: Z e d
name: Zed
another: Zed
another each: Z e d
$

```

You'll notice some interesting things about this program:

- I didn't have to give all 4 elements of the arrays to initialize them. This is a short-cut that C has where, if you set just one element, it'll fill the rest in with 0.
- When each element of `numbers` is printed they all come out as 0.
- When each element of `name` is printed, only the first element 'a' shows up because the `'\0'` character is special and won't display.
- Then the first time we print `name` it only prints "a" because, since the array will be filled with 0 after the first 'a' in the initializer, then the string is correctly terminated by a `'\0'` character.
- We then setup the arrays with a tedious manual assignment to each thing and print them out again. Look at how they changed. Now the numbers are set, but see how the `name` string prints my name correctly?
- There's also two syntaxes for doing a string: `char name[4] = {'a'}` on line 6 vs. `char *another = "name"` on line 44. The first one is less common and the second is what you should use for string literals like this.

Notice that I'm using the same syntax and style of code to interact with both an array of integers and an array of characters, but that `printf` thinks that the `name` is just a string. Again, this is because to the C language there's no difference between a string and an array of characters.

Finally, when you make string literals you should usually use the `char`

`*another = "Literal"` syntax. This works out to be the same thing, but it's more idiomatic and easier to write.

## How To Break It

The source of almost all bugs in C come from forgetting to have enough space, or forgetting to put a `'\0'` at the end of a string. In fact it's so common and hard to get right that the majority of good C code just doesn't use C style strings. In later exercises we'll actually learn how to avoid C strings completely.

In this program the key to breaking it is to forget to put the `'\0'` character at the end of the strings. There's a few ways to do this:

- Get rid of the initializers that setup `name`.
- Accidentally set `name[3] = 'A';` so that there's no terminator.
- Set the initializer to `{'a', 'a', 'a', 'a'}` so there's too many 'a' characters and no space for the `'\0'` terminator.

Try to come up with some other ways to break this, and as usual run all of these under Valgrind so you can see exactly what is going on and what the errors are called. Sometimes you'll make these mistakes and even Valgrind can't find them, but try moving where you declare the variables to see if you get the error. This is part of the voodoo of C, that sometimes just where the variable is located changes the bug.

## Extra Credit

- Assign the characters into `numbers` and then use `printf` to print them a character at a time. What kind of compiler warnings did you get?
- Do the inverse for `name`, trying to treat it like an array of `int` and print it out one `int` at a time. What does Valgrind think of that?
- How many other ways can you print this out?
- If an array of characters is 4 bytes long, and an integer is 4 bytes long, then can you treat the whole `name` array like it's just an integer? How might you accomplish this crazy hack?
- Take out a piece of paper and draw out each of these arrays as a row of boxes. Then do the operations you just did on paper to see if you get them right.
- Convert `name` to be in the style of `another` and see if the code keeps working.

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 10: Arrays Of Strings, Looping

You can make an array of various types, and have the idea down that a "string" and an "array of bytes" are the same thing. The next thing is to take this one step further and do an array that has strings in it. We'll also introduce your first looping construct, the `for-loop` to help print out this new data structure.

The fun part of this is that there's been an array of strings hiding in your programs for a while now, the `char *argv[]` in the `main` function arguments. Here's code that will print out any command line arguments you pass it:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int i = 0;

    // go through each string in argv
    // why am I skipping argv[0]?
    for(i = 1; i < argc; i++) {
        printf("arg %d: %s\n", i, argv[i]);
    }

    // let's make our own array of strings
    char *states[] = {
        "California", "Oregon",
        "Washington", "Texas"
    };
    int num_states = 4;

    for(i = 0; i < num_states; i++) {
        printf("state %d: %s\n", i, states[i]);
    }

    return 0;
}
```

The format of a `for-loop` is this:

```
for(INITIALIZER; TEST; INCREMENTER) {
    CODE;
}
```

Here's how the `for-loop` works:

- The `INITIALIZER` is code that is run to setup the loop, in this case `i = 0`.
- Next the `TEST` boolean expression is checked, and if it's false (0) then `CODE` is skipped, doing nothing.
- The `CODE` runs, does whatever it does.
- After the `CODE` runs, the `INCREMENTER` part is run, usually incrementing something, like in `i++`.

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

🏠  
HELP

Follow  
@lzssth



- And it continues again with Step 2 until the `TEST` is false (0).

This `for-loop` is going through the command line arguments using `argc` and `argv` like this:

- The OS passes each command line argument as a string in the `argv` array. The program's name (`./ex10`) is at 0, with the rest coming after it.
- The OS also sets `argc` to the number of arguments in the `argv` array so you can process them without going past the end. Remember that if you give one argument, the program's name is the first, so `argc` is 2.
- The `for-loop` sets up with `i = 1` in the initializer.
- It then tests that `i` is less than `argc` with the test `i < argc`. Since initially `$1 < 2$` it will pass.
- It then runs the code which just prints out the `i` and uses `i` to index into `argv`.
- The incrementer is then run using the `i++` syntax, which is a handy way of writing `i = i + 1`.
- This then repeats until `i < argc` is finally false (0) when the loop exits and the program continues on.

## What You Should See

To play with this program you have to run it two ways. The first way is to pass in some command line arguments so that `argc` and `argv` get set. The second is to run it with no arguments so you can see that the first `for-loop` doesn't run since `i < argc` will be false.

```
$ make ex10
cc -Wall -g    ex10.c    -o ex10
$ ./ex10 i am a bunch of arguments
arg 1: i
arg 2: am
arg 3: a
arg 4: bunch
arg 5: of
arg 6: arguments
state 0: California
state 1: Oregon
state 2: Washington
state 3: Texas
$
$ ./ex10
state 0: California
state 1: Oregon
state 2: Washington
state 3: Texas
$
```

## Understanding Arrays Of Strings

From this you should be able to figure out that in C you make an "array of strings" by combining the `char *str = "blah"` syntax with the `char str[] = {'b', 'l', 'a', 'h'}` syntax to construct a 2-dimensional array. The syntax `char *states[] = {...}` on line 14 is this 2-dimension combination, with each string being one element, and each character in the string being another.

Confusing? The concept of multiple dimensions is something most people never think about so what you should do is build this array of strings on paper:

- Make a grid with the index of each *string* on the left.
- Then put the index of each *character* on the top.
- Then, fill in the squares in the middle with what single character goes in that cell.
- Once you have the grid, trace through the code manually using this grid of paper.

Another way to figure this is out is to build the same structure in a programming language you are more familiar with like Python or Ruby.

## How To Break It

- Take your favorite other language, and use it to run this program, but with as many command line arguments as possible. See if you can bust it by giving it way too many arguments.
- Initialize `i` to 0 and see what that does. Do you have to adjust `argc` as well or does it just work? Why does 0-based indexing work here?
- Set `num_states` wrong so that it's a higher value and see what it does.

## Extra Credit

- Figure out what kind of code you can put into the parts of a `for-loop`.
- Look up how to use the `,` (comma) character to separate multiple statements in the parts of the `for-loop`, but between the `;` (semicolon) characters.
- Read what a `NULL` is and try to use it in one of the elements of the `states` array to see what it'll print.
- See if you can assign an element from the `states` array to the `argv` array before printing both. Try the inverse.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 11: While-Loop And Boolean Expressions

You've had your first taste of how C does loops, but the boolean expression `i < argc` might have not been clear to you. Let me explain something about it before we see how a `while-loop` works.

In C, there's not really a "boolean" type, and instead any integer that's 0 is "false" and otherwise it's "true". In the last exercise the expression `i < argc` actually resulted in 1 or 0, not an explicit `True` or `False` like in Python. This is another example of C being closer to how a computer works, because to a computer truth values are just integers.

Now you'll take and implement the same program from the last exercise but use a `while-loop` instead. This will let you compare the two so you can see how one is related to another.

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    // go through each string in argv

    int i = 0;
    while(i < argc) {
        printf("arg %d: %s\n", i, argv[i]);
        i++;
    }

    // let's make our own array of strings
    char *states[] = {
        "California", "Oregon",
        "Washington", "Texas"
    };

    int num_states = 4;
    i = 0; // watch for this
    while(i < num_states) {
        printf("state %d: %s\n", i, states[i]);
        i++;
    }

    return 0;
}
```

You can see from this that a `while-loop` is simpler:

```
while(TEST) {
    CODE;
}
```

It simply runs the `CODE` as long as `TEST` is true (1). This means that to replicate how the `for-loop` works we need to do our own initializing and incrementing of `i`.

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

⚡  
HELP

Follow  
@l2sthw

# What You Should See

The output is basically the same, so I just did it a little different so you can see another way it runs.

```
$ make ex11
cc -Wall -g    ex11.c    -o ex11
$ ./ex11
arg 0: ./ex11
state 0: California
state 1: Oregon
state 2: Washington
state 3: Texas
$
$ ./ex11 test it
arg 0: ./ex11
arg 1: test
arg 2: it
state 0: California
state 1: Oregon
state 2: Washington
state 3: Texas
$
```

## How To Break It

In your own code you should favor `for-loop` constructs over `while-loop` because a `for-loop` is harder to break. Here's a few common ways:

- Forget to initialize the first `int i;` so have it loop wrong.
- Forget to initialize the second loop's `i` so that it retains the value from the end of the first loop. Now your second loop might or might not run.
- Forget to do a `i++` increment at the end of the loop and you get a "forever loop", one of the dreaded problems of the first decade or two of programming.

## Extra Credit

- Make these loops count backward by using `i--` to start at `argc` and count down to 0. You may have to do some math to make the array indexes work right.
- Use a while loop to *copy* the values from `argv` into `states`.
- Make this copy loop never fail such that if there's too many `argv` elements it won't put them all into `states`.
- Research if you've really copied these strings. The answer may surprise and confuse you though.

**Interested  
In  
Python?**

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

**Interested In Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 12: If, Else-If, Else

Something common in every language is the `if-statement`, and C has one. Here's code that uses an `if-statement` to make sure you enter only 1 or 2 arguments:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int i = 0;

    if(argc == 1) {
        printf("You only have one argument. You suck.\n");
    } else if(argc > 1 && argc < 4) {
        printf("Here's your arguments:\n");

        for(i = 0; i < argc; i++) {
            printf("%s ", argv[i]);
        }
        printf("\n");
    } else {
        printf("You have too many arguments. You suck.\n");
    }

    return 0;
}
```

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@l2sthw

The format for the `if-statement` is this:

```
if(TEST) {
    CODE;
} else if(TEST) {
    CODE;
} else {
    CODE;
}
```

This is like most other languages except for some specific C differences:

- As mentioned before, the `TEST` parts are false if they evaluate to 0, and true otherwise.
- You have to put parenthesis around the `TEST` elements, while some other languages let you skip that.
- You don't need the `{ }` braces to enclose the code, but it is very bad form to not use them. The braces make it clear where one branch of code begins and ends. If you don't include it then obnoxious errors come up.

Other than that, they work like others do. You don't need to have either `else if` or `else` parts.

## What You Should See

This one is pretty simple to run and try out:

```
$ make ex12
cc -Wall -g    ex12.c    -o ex12
$ ./ex12
You only have one argument. You suck.
$ ./ex12 one
Here's your arguments:
./ex12 one
$ ./ex12 one two
Here's your arguments:
./ex12 one two
$ ./ex12 one two three
You have too many arguments. You suck.
$
```

## How To Break It

This one isn't easy to break because it's so simple, but try messing up the tests in the `if-statement`.

- Remove the `else` at the end and it won't catch the edge case.
- Change the `&&` to a `||` so you get an "or" instead of "and" test and see how that works.

## Extra Credit

- You were briefly introduced to `&&`, which does an "and" comparison, so go research online the different "boolean operators".
- Write a few more test cases for this program to see what you can come up with.
- Go back to Exercises 10 and 11, and use `if-statements` to make the loops exit early. You'll need the `break` statement to do that. Go read about it.
- Is the first test really saying the right thing? To you the "first argument" isn't the same first argument a user entered. Fix it.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 13: Switch Statement

In other languages like Ruby you have a `switch-statement` that can take any expression. Some languages like Python just don't have a `switch-statement` since an `if-statement` with boolean expressions is about the same thing. For these languages, `switch-statements` are more alternatives to `if-statements` and work the same internally.

The `switch-statement` is actually entirely different and is really a "jump table". Instead of random boolean expressions, you can only put expressions that result in integers, and these integers are used to calculate jumps from the top of the `switch` to the part that matches that value. Here's some code that we'll break down to understand this concept of "jump tables":

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    if(argc != 2) {
        printf("ERROR: You need one argument.\n");
        // this is how you abort a program
        return 1;
    }

    int i = 0;
    for(i = 0; argv[1][i] != '\0'; i++) {
        char letter = argv[1][i];

        switch(letter) {
            case 'a':
            case 'A':
                printf("%d: 'A'\n", i);
                break;

            case 'e':
            case 'E':
                printf("%d: 'E'\n", i);
                break;

            case 'i':
            case 'I':
                printf("%d: 'I'\n", i);
                break;

            case 'o':
            case 'O':
                printf("%d: 'O'\n", i);
                break;

            case 'u':
            case 'U':
                printf("%d: 'U'\n", i);
                break;

            case 'y':
            case 'Y':
                if(i > 2) {
                    // it's only sometimes Y
                }
            }
        }
    }
}
```

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@l2sthw

```

        printf("%d: 'Y'\n", i);
    }
    break;

    default:
        printf("%d: %c is not a vowel\n", i, letter);
    }
}

return 0;
}

```

In this program we take a single command line argument and print out all of the vowels in an incredibly tedious way to demonstrate a `switch-statement`. Here's how the `switch-statement` works:

- The compiler marks the place in the program where the `switch-statement` starts, let's call this location Y.
- It then evaluates the expression in `switch(letter)` to come up with a number. In this case the number will be the raw ASCII code of the letter in `argv[1]`.
- The compiler has also translated each of the `case` blocks like `case 'A':` into a location in the program that is that far away. So the code under `case 'A'` is at Y+'A' in the program.
- It then does the math to figure out where Y+letter is located in the `switch-statement`, and if it's too far then it adjusts it to Y+default.
- Once it knows the location, the program "jumps" to that spot in the code, and then continues running. This is why you have `break` on some of the `case` blocks, but not others.
- If `'a'` is entered, then it jumps to `case 'a'`, there's no break so it "falls through" to the one right under it `case 'A'` which has code and a `break`.
- Finally it runs this code, hits the break then exits out of the `switch-statement` entirely.

This is a deep dive into how the `switch-statement` works, but in practice you just have to remember a few simple rules:

- Always include a `default:` branch so that you catch any missing inputs.
- Don't allow "fall through" unless you really want it, and it's a good idea to add a comment `//fallthrough` so people know it's on purpose.
- Always write the `case` and the `break` before you write the code that goes in it.
- Try to just use `if-statements` instead if you can.

## What You Should See

Here's an example of me playing with this, and also demonstrating various ways to pass the argument in:

```

$ make ex13
cc -Wall -g    ex13.c    -o ex13
$ ./ex13
ERROR: You need one argument.
$
$ ./ex13 Zed

```



```

0: Z is not a vowel
1: 'E'
2: d is not a vowel
$
$ ./ex13 Zed Shaw
ERROR: You need one argument.
$
$ ./ex13 "Zed Shaw"
0: Z is not a vowel
1: 'E'
2: d is not a vowel
3:  is not a vowel
4: S is not a vowel
5: h is not a vowel
6: 'A'
7: w is not a vowel
$

```

Remember that there's that `if-statement` at the top that exits with a `return 1;` when you don't give enough arguments. Doing a return that's not 0 is how you indicate to the OS that the program had an error. Any value that's greater than 0 can be tested for in scripts and other programs to figure out what happened.

## How To Break It

It is *incredibly* easy to break a `switch-statement`. Here's just a few of the ways you can mess one of these up:

- Forget a `break` and it'll run two or more blocks of code you don't want it to run.
- Forget a `default` and have it silently ignore values you forgot.
- Accidentally put in variable into the `switch` that evaluates to something unexpected, like an `int` that becomes weird values.
- Use uninitialized values in the `switch`.

You can also break this program in a few other ways. See if you can bust it yourself.

## Extra Credit

- Write another program that uses math on the letter to convert it to lowercase, and then remove all the extraneous uppercase letters in the switch.
- Use the `','` (comma) to initialize `letter` in the `for-loop`.
- Make it handle all of the arguments you pass it with yet another `for-loop`.
- Convert this `switch-statement` to an `if-statement`. Which do you like better?
- In the case for 'Y' I have the break outside the `if-statement`. What's the impact of this and what happens if you move it inside the `if-statement`. Prove to yourself that you're right.

**Interested  
In  
Python?**

Python is also a  
great language.

**Interested In Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 14: Writing And Using Functions

Until now you've just used functions that are part of the `stdio.h` header file. In this exercise you will write some functions and use some other functions.

```
#include <stdio.h>
#include <ctype.h>

// forward declarations
int can_print_it(char ch);
void print_letters(char arg[]);

void print_arguments(int argc, char *argv[])
{
    int i = 0;

    for(i = 0; i < argc; i++) {
        print_letters(argv[i]);
    }
}

void print_letters(char arg[])
{
    int i = 0;

    for(i = 0; arg[i] != '\0'; i++) {
        char ch = arg[i];

        if(can_print_it(ch)) {
            printf("%c" == %d ", ch, ch);
        }
    }

    printf("\n");
}

int can_print_it(char ch)
{
    return isalpha(ch) || isblank(ch);
}

int main(int argc, char *argv[])
{
    print_arguments(argc, argv);
    return 0;
}
```

In this example you're creating functions to print out the characters and ASCII codes for any that are "alpha" or "blanks". Here's the breakdown:

ex14.c:2

Include a new header file so we can gain access to `isalpha` and `isblank`.

ex14.c:5-6

Tell C that you will be using some functions later in your

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssth

program, without having to actually define them. This is a "forward declaration" and it solves the chicken-and-egg problem of needing to use a function before you've defined it.

ex14.c:8-15

Define the `print_arguments` which knows how to print the same array of strings that `main` typically gets.

ex14.c:17-30

Define the next function `print_letters` that is called by `print_arguments` and knows how to print each of the characters and their codes.

ex14.c:32-35

Define `can_print_it` which simply returns the truth value (0 or 1) of `isalpha(ch) || isblank(ch)` back to its caller `print_letters`.

ex14.c:38-42

Finally `main` simply calls `print_arguments` to make the whole chain of function calls go.

I shouldn't have to describe what's in each function because it's all things you've ran into before. What you should be able to see though is that I've simply defined functions the same way you've been defining `main`. The only difference is you have to help C out by telling it ahead of time if you're going to use functions it hasn't encountered yet in the file. That's what the "forward declarations" at the top do.

## What You Should See

To play with this program you just feed it different command line arguments, which get passed through your functions. Here's me playing with it to demonstrate:

```
$ make ex14
cc -Wall -g    ex14.c    -o ex14

$ ./ex14
'e' == 101 'x' == 120

$ ./ex14 hi this is cool
'e' == 101 'x' == 120
'h' == 104 'i' == 105
't' == 116 'h' == 104 'i' == 105 's' == 115
'i' == 105 's' == 115
'c' == 99 'o' == 111 'o' == 111 'l' == 108

$ ./ex14 "I go 3 spaces"
'e' == 101 'x' == 120
'I' == 73 ' ' == 32 'g' == 103 'o' == 111 ' ' == 32 ' ' == 32 's' == 115 'p' =
= 112 'a' == 97 'c' == 99 'e' == 101 's' == 115
$
```

The `isalpha` and `isblank` do all the work of figuring out if the given character is a letter or a blank. When I do the last run it prints everything but the '3' character, since that is a digit.

## How To Break It

There's two different kinds of "breaking" in this program:

- Confuse the compiler by removing the forward declarations so it complains about `can_print_it` and `print_letters`.
- When you call `print_arguments` inside `main`, try adding 1 to `argc` so that it goes past the end of the `argv` array.

## Extra Credit

- Rework these functions so that you have fewer functions. For example, do you really need `can_print_it`?
- Have `print_arguments` figure how long each argument string is using the `strlen` function, and then pass that length to `print_letters`. Then, rewrite `print_letters` so it only processes this fixed length and doesn't rely on the `'\0'` terminator. You will need the `#include <string.h>` for this.
- Use `man` to lookup information on `isalpha` and `isblank`. Use the other similar functions to print out only digits or other characters.
- Go read about how different people like to format their functions. Never use the "K&R syntax" as it's antiquated and confusing, but understand what it's doing in case you run into someone who likes it.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 15: Pointers Dreaded Pointers

Pointers are famous mystical creatures in C that I will attempt to demystify by teaching you the vocabulary used to deal with them. They actually aren't that complex, it's just they are frequently abused in weird ways that make them hard to use. If you avoid the stupid ways to use pointers then they're fairly easy.

To demonstrate pointers in a way we can talk about them, I've written a frivolous program that prints a group of people's ages in three different ways:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    // create two arrays we care about
    int ages[] = {23, 43, 12, 89, 2};
    char *names[] = {
        "Alan", "Frank",
        "Mary", "John", "Lisa"
    };

    // safely get the size of ages
    int count = sizeof(ages) / sizeof(int);
    int i = 0;

    // first way using indexing
    for(i = 0; i < count; i++) {
        printf("%s has %d years alive.\n",
            names[i], ages[i]);
    }

    printf("---\n");

    // setup the pointers to the start of the arrays
    int *cur_age = ages;
    char **cur_name = names;

    // second way using pointers
    for(i = 0; i < count; i++) {
        printf("%s is %d years old.\n",
            *(cur_name+i), *(cur_age+i));
    }

    printf("---\n");

    // third way, pointers are just arrays
    for(i = 0; i < count; i++) {
        printf("%s is %d years old again.\n",
            cur_name[i], cur_age[i]);
    }

    printf("---\n");

    // fourth way with pointers in a stupid complex way
    for(cur_name = names, cur_age = ages;
```

```

        (cur_age - ages) < count;
        cur_name++, cur_age++)
    {
        printf("%s lived %d years so far.\n",
               *cur_name, *cur_age);
    }

    return 0;
}

```

Before explaining how pointers work, let's break this program down line-by-line so you get an idea of what's going on. As you go through this detailed description, try to answer the questions for yourself on a piece of paper, then see if what you guessed was going on matches my description of pointers later.

ex15.c:6-10

Create two arrays, `ages` storing some `int` data, and `names` storing an array of strings.

ex15.c:12-13

Some variables for our `for-loops` later.

ex15.c:16-19

You know this is just looping through the two arrays and printing how old each person is. This is using `i` to index into the array.

ex15.c:24

Create a pointer that points at `ages`. Notice the use of `int *` to create a "pointer to integer" type of pointer. That's similar to `char *`, which is a "pointer to char", and a string is an array of chars. Seeing the similarity yet?

ex15.c:25

Create a pointer that points at `names`. A `char *` is already a "pointer to char", so that's just a string. You however need 2 levels, since `names` is 2-dimensional, that means you need `char **` for a "pointer to (a pointer to char)" type. Study that too, explain it to yourself.

ex15.c:28-31

Loop through `ages` and `names` but instead use the pointers *plus an offset of i*. Writing `*(cur_name+i)` is the same as writing `name[i]`, and you read it as "the value of (pointer `cur_name` plus i)".

ex15.c:35-39

This shows how the syntax to access an element of an array is the same for a pointer and an array.

ex15.c:44-50

Another admittedly insane loop that does the same thing as the other two, but instead it uses various pointer arithmetic methods:

ex15.c:44

Initialize our `for-loop` by setting `cur_name` and `cur_age` to the beginning of the `names` and `ages` arrays.

ex15.c:45

The test portion of the `for-loop` then compares the *distance* of the pointer `cur_age` from the start of `ages`. Why does that work?

ex15.c:46

The increment part of the `for-loop` then increments both `cur_name` and `cur_age` so that they point at the *next* element of the `name` and `age` arrays.

ex15.c:48-49

The pointers `cur_name` and `cur_age` are now pointing at one element of the arrays they work on, and we can print them out using just `*cur_name` and `*cur_age`, which means "the value of wherever `cur_name` is pointing".

This seemingly simple program has a large amount of information, and the goal is to get you to attempt figuring pointers out for yourself before I explain them. *Don't continue until you've written down what you think a pointer does.*

## What You Should See

After you run this program try to trace back each line printed out to the line in the code that produced it. If you have to, alter the `printf` calls to make sure you got the right line number.

```
$ make ex15
cc -Wall -g    ex15.c    -o ex15
$ ./ex15
Alan has 23 years alive.
Frank has 43 years alive.
Mary has 12 years alive.
John has 89 years alive.
Lisa has 2 years alive.
---
Alan is 23 years old.
Frank is 43 years old.
Mary is 12 years old.
John is 89 years old.
Lisa is 2 years old.
---
Alan is 23 years old again.
Frank is 43 years old again.
Mary is 12 years old again.
John is 89 years old again.
Lisa is 2 years old again.
---
Alan lived 23 years so far.
Frank lived 43 years so far.
Mary lived 12 years so far.
John lived 89 years so far.
Lisa lived 2 years so far.
$
```

## Explaining Pointers

When you type something like `ages[i]` you are "indexing" into the array `ages`, and you're using the number that's held in `i` to do it. If `i` is set to 0 then it's the same as typing `ages[0]`. We've been calling this number `i` an "index" since it's a location inside `ages` that we want. It could also be called an "address", that's a way of saying "I want the integer in `ages` that is at address `i`".

If `i` is an index, then what's `ages`? To C `ages` is a location in the



computer's memory where all of these integers start. It is *also* an address, and the C compiler will replace anywhere you type `ages` with the address of the very first integer in `ages`. Another way to think of `ages` is it's the "address of the first integer in `ages`". But, the trick is `ages` is an address inside the *entire computer*. It's not like `i` which was just an address inside `ages`. The `ages` array name is actually an address in the computer.

That leads to a certain realization: C thinks your whole computer is one massive array of bytes. Obviously this isn't very useful, but then C layers on top of this massive array of bytes the concept of *types* and *sizes* of those types. You already saw how this worked in previous exercises, but now you can start to get an idea that C is somehow doing the following with your arrays:

- Creating a block of memory inside your computer.
- "Pointing" the name `ages` at the beginning of that block.
- "Indexing" into the block by taking the base address of `ages` and getting the element that's `i` away from there.
- Converting that address at `ages+i` into a valid `int` of the right size, such that the index works to return what you want: the int at index `i`.

If you can take a base address, like `ages`, and then "add" to it with another address like `i` to produce a new address, then can you just make something that points right at this location all the time? Yes, and that thing is called a "pointer". This is what the pointers `cur_age` and `cur_name` are doing. They are variables pointing at the location where `ages` and `names` live in your computer's memory. The example program is then moving them around or doing math on them to get values out of the memory. In one instance, they just add `i` to `cur_age`, which is the same as what it does with `array[i]`. In the last `for-loop` though these two pointers are being moved on their own, without `i` to help out. In that loop, the pointers are treated like a combination of array and integer offset rolled into one.

A pointer is simply an address pointing somewhere inside the computer's memory, with a type specifier so you get the right size of data with it. It is kind of like a combined `ages` and `i` rolled into one data type. C knows where pointers are pointing, knows the data type they point at, the size of those types, and how to get the data for you. Just like `i` you can increment them, decrement them, subtract or add to them. But, just like `ages` you can also get values out with them, put new values in, and all the array operations.

The purpose of a pointer is to let you manually index into blocks or memory when an array won't do it right. In almost all other cases you actually want to use an array. But, there are times when you *have* to work with a raw block of memory and that's where a pointer comes in. A pointer gives you raw, direct access to a block of memory so you can work with it.

The final thing to grasp at this stage is that you can use either syntax for most array or pointer operations. You can take a pointer to something, but use the array syntax for accessing it. You can take an array and do pointer arithmetic with it.

# Practical Pointer Usage

There are four primary useful things you do with pointers in C code:

- Ask the OS for a chunk of memory and use a pointer to work with it. This includes strings and something you haven't seen yet, `structs`.
- Passing large blocks of memory (like large structs) to functions with a pointer so you don't have to pass the whole thing to them.
- Taking the address of a function so you can use it as a dynamic callback.
- Complex scanning of chunks of memory such as converting bytes off a network socket into data structures or parsing files.

For nearly everything else you see people use pointers, they should be using arrays. In the early days of C programming people used pointers to speed up their programs because the compilers were really bad at optimizing array usage. These days the syntax to access an array vs. a pointer are translated into the same machine code and optimized the same, so it's not as necessary. Instead, you go with arrays every time you can, and then only use pointers as a performance optimization if you absolutely have to.

## The Pointer Lexicon

I'm now going to give you a little lexicon to use for reading and writing pointers. Whenever you run into a complex pointer statement, just refer to this and break it down bit by bit (or just don't use that code since it's probably not good code):

`type *ptr`

"a pointer of type named ptr"

`*ptr`

"the value of whatever ptr is pointed at"

`*(ptr + i)`

"the value of (whatever ptr is pointed at plus i)"

`&thing`

"the address of thing"

`type *ptr = &thing`

"a pointer of type named ptr set to the address of thing"

`ptr++`

"increment where ptr points"

We'll be using this simple lexicon to break down all of the pointers we use from now on in the book.

## Pointers Are Not Arrays

No matter what, you should never think that pointers and arrays are the same thing. They are not the same thing, even though C lets you work with them in many of the same ways. For example, if you do `sizeof(cur_age)` in the code above, you would get the size of the *pointer*, not the size of what it points at. If you want the size of the full array, you

have to use the array's name, `age` as I did on line 12.

TODO: expand on this some more with what doesn't work on both the same.

## How To Break It

You can break this program by simply pointing the pointers at the wrong things:

- Try to make `cur_age` point at `names`. You'll need to use a C cast to force it, so go look that up and try to figure it out.
- In the final `for-loop` try getting the math wrong in weird ways.
- Try rewriting the loops so they start at the end of the arrays and go to the beginning. This is harder than it looks.

## Extra Credit

- Rewrite all the array usage in this program so that it's pointers.
- Rewrite all the pointer usage so they're arrays.
- Go back to some of the other programs that use arrays and try to use pointers instead.
- Process command line arguments using just pointers similar to how you did `names` in this one.
- Play with combinations of getting the value of and the address of things.
- Add another `for-loop` at the end that prints out the addresses these pointers are using. You'll need the `%p` format for `printf`.
- Rewrite this program to use a function for each of the ways you're printing out things. Try to pass pointers to these functions so they work on the data. Remember you can declare a function to accept a pointer, but just use it like an array.
- Change the `for-loops` to `while-loops` and see what works better for which kind of pointer usage.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 16: Structs And Pointers To Them

In this exercise you'll learn how to make a `struct`, point a pointer at them, and use them to make sense of internal memory structures. I'll also apply the knowledge of pointers from the last exercise and get you constructing these structures from raw memory using `malloc`.

As usual, here's the program we'll talk about, so type it in and make it work:

```
#include <stdio.h>
#include <assert.h>
#include <stdlib.h>
#include <string.h>

struct Person {
    char *name;
    int age;
    int height;
    int weight;
};

struct Person *Person_create(char *name, int age, int height, int weight)
{
    struct Person *who = malloc(sizeof(struct Person));
    assert(who != NULL);

    who->name = strdup(name);
    who->age = age;
    who->height = height;
    who->weight = weight;

    return who;
}

void Person_destroy(struct Person *who)
{
    assert(who != NULL);

    free(who->name);
    free(who);
}

void Person_print(struct Person *who)
{
    printf("Name: %s\n", who->name);
    printf("\tAge: %d\n", who->age);
    printf("\tHeight: %d\n", who->height);
    printf("\tWeight: %d\n", who->weight);
}

int main(int argc, char *argv[])
{
    // make two people structures
    struct Person *joe = Person_create(
        "Joe Alex", 32, 64, 140);

    struct Person *frank = Person_create(
```

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

🔧  
HELP

Follow  
@l2sthw

```

        "Frank Blank", 20, 72, 180);

    // print them out and where they are in memory
    printf("Joe is at memory location %p:\n", joe);
    Person_print(joe);

    printf("Frank is at memory location %p:\n", frank);
    Person_print(frank);

    // make everyone age 20 years and print them again
    joe->age += 20;
    joe->height -= 2;
    joe->weight += 40;
    Person_print(joe);

    frank->age += 20;
    frank->weight += 20;
    Person_print(frank);

    // destroy them both so we clean up
    Person_destroy(joe);
    Person_destroy(frank);

    return 0;
}

```

To describe this program, I'm going to use a different approach than before. I'm not going to give you a line-by-line breakdown of the program, but I'm going to make *you* write it. I'm going to give you a guide through the program based on the parts it contains, and your job is to write out what each line does.

includes

I include some new header files here to gain access to some new functions. What does each give you?

`struct Person`

This is where I'm creating a structure that has 4 elements to describe a person. The final result is a new compound type that lets me reference these elements all as one, or each piece by name. It's similar to a row of a database table or a class in an OOP language.

function `Person_create`

I need a way to create these structures so I've made a function to do that. Here's the important things this function is doing:

- I use `malloc` for "memory allocate" to ask the OS to give me a piece of raw memory.
- I pass to `malloc` the `sizeof(struct Person)` which calculates the total size of the struct, given all the fields inside it.
- I use `assert` to make sure that I have a valid piece of memory back from malloc. There's a special constant called `NULL` that you use to mean "unset or invalid pointer". This `assert` is basically checking that malloc didn't return a NULL invalid pointer.
- I initialize each field of `struct Person` using the `x->y` syntax, to say what part of the struct I want to set.
- I use the `strdup` function to duplicate the string for the name, just to make sure that this structure actually owns it. The `strdup` actually is like `malloc` and it also copies the original string into the memory it creates.

function `Person_destroy`

If I have a create, then I always need a destroy function, and this is what destroys `Person` structs. I again use `assert` to make sure I'm not getting bad input. Then I use the function `free` to return the memory I got with `malloc` and `strdup`. If you don't do this you get a "memory leak".

function `Person_print`

I then need a way to print out people, which is all this function does. It uses the same `x->y` syntax to get the field from the struct to print it.

function `main`

In the main function I use all the previous functions and the `struct Person` to do the following:

- Create two people, `joe` and `frank`.
- Print them out, but notice I'm using the `%p` format so you can see *where* the program has actually put your struct in memory.
- Age both of them by 20 years, with changes to their body too.
- Print each one after aging them.
- Finally destroy the structures so we can clean up correctly.

Go through this description carefully, and do the following:

- Look up every function and header file you don't know about. Remember that you can usually do `man 2 function` or `man 3 function` and it'll tell you about it. You can also search online for the information.
- Write a *comment* above each and every single line saying what the line does in English.
- Trace through each function call and variable so you know where it comes from in the program.
- Look up any symbols you don't know as well.

## What You Should See

After you augment the program with your description comments, make sure it really runs and produces this output:

```
$ make ex16
cc -Wall -g    ex16.c    -o ex16

$ ./ex16
Joe is at memory location 0xeba010:
Name: Joe Alex
  Age: 32
  Height: 64
  Weight: 140
Frank is at memory location 0xeba050:
Name: Frank Blank
  Age: 20
  Height: 72
  Weight: 180
Name: Joe Alex
  Age: 52
  Height: 62
```

```
Weight: 180
Name: Frank Blank
Age: 40
Height: 72
Weight: 200
```

## Explaining Structures

If you've done the work I asked you then structures should be making sense, but let me explain them explicitly just to make sure you've understood it.

A structure in C is a collection of other data types (variables) that are stored in one block of memory but let you access each variable independently by name. They are similar to a record in a database table, or a very simplistic class in an object oriented language. We can break one down this way:

- In the above code, you make a `struct` that has the fields you'd expect for a person: name, age, weight, height.
- Each of those fields has a type, like `int`.
- C then packs those together so they can all be contained in one single `struct`.
- The `struct Person` is now a *compound data type*, which means you can now refer to `struct Person` in the same kinds of expressions you would other data types.
- This lets you pass the whole cohesive grouping to other functions, as you did with `Person_print`.
- You can then access the individual parts of a `struct` by their names using `x->y` if you're dealing with a pointer.
- There's also a way to make a struct that doesn't need a pointer, and you use the `x.y` (period) syntax to work with it. You'll do this in the Extra Credit.

If you didn't have `struct` you'd need to figure out the size, packing, and location of pieces of memory with contents like this. In fact, in most early assembler code (and even some now) this is what you do. With C you can let C handle the memory structuring of these compound data types and then focus on what you do with them.

## How To Break It

With this program the ways to break it involve how you use the pointers and the `malloc` system:

- Try passing `NULL` to `Person_destroy` to see what it does. If it doesn't abort then you must not have the `-g` option in your Makefile's `CFLAGS`.
- Forget to call `Person_destroy` at the end, then run it under `Valgrind` to see it report that you forgot to free the memory. Figure out the options you need to pass to `Valgrind` to get it to print how you leaked this memory.
- Forget to free `who->name` in `Person_destroy` and compare the output. Again, use the right options to see how `Valgrind` tells you exactly where you messed up.
- This time, pass `NULL` to `Person_print` and see what `Valgrind`

thinks of that.

- You should be figuring out that `NULL` is a quick way to crash your program.

## Extra Credit

In this exercise I want you to attempt something difficult for the extra credit: Convert this program to *not* use pointers and `malloc`. This will be hard, so you'll want to research the following:

- How to create a `struct` on the *stack*, which means just like you've been making any other variable.
- How to initialize it using the `x.y` (period) character instead of the `x->y` syntax.
- How to pass a structure to other functions without using a pointer.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 17: Heap And Stack Memory Allocation

In this exercise you're going to make a big leap in difficulty and create an entire small program to manage a database. This database isn't very efficient and doesn't store very much, but it does demonstrate most of what you've learned so far. It also introduces memory allocation more formally and gets you started working with files. We use some file I/O functions, but I won't be explaining them too well so you can try to figure them out first.

As usual, type this whole program in and get it working, then we'll discuss:

```
#include <stdio.h>
#include <assert.h>
#include <stdlib.h>
#include <errno.h>
#include <string.h>

#define MAX_DATA 512
#define MAX_ROWS 100

struct Address {
    int id;
    int set;
    char name[MAX_DATA];
    char email[MAX_DATA];
};

struct Database {
    struct Address rows[MAX_ROWS];
};

struct Connection {
    FILE *file;
    struct Database *db;
};

void die(const char *message)
{
    if(errno) {
        perror(message);
    } else {
        printf("ERROR: %s\n", message);
    }

    exit(1);
}

void Address_print(struct Address *addr)
{
    printf("%d %s %s\n",
        addr->id, addr->name, addr->email);
}

void Database_load(struct Connection *conn)
{
```

1 \_\_\_\_\_  
2 \_\_\_\_\_  
3 \_\_\_\_\_  
4 \_\_\_\_\_ MAIN

◀  
PREVIOUS

▶ NEXT

⚙️ HELP

Follow  
@lzssth

```

    int rc = fread(conn->db, sizeof(struct Database), 1, conn->file);
    if(rc != 1) die("Failed to load database.");
}

struct Connection *Database_open(const char *filename, char mode)
{
    struct Connection *conn = malloc(sizeof(struct Connection));
    if(!conn) die("Memory error");

    conn->db = malloc(sizeof(struct Database));
    if(!conn->db) die("Memory error");

    if(mode == 'c') {
        conn->file = fopen(filename, "w");
    } else {
        conn->file = fopen(filename, "r+");

        if(conn->file) {
            Database_load(conn);
        }
    }

    if(!conn->file) die("Failed to open the file");

    return conn;
}

void Database_close(struct Connection *conn)
{
    if(conn) {
        if(conn->file) fclose(conn->file);
        if(conn->db) free(conn->db);
        free(conn);
    }
}

void Database_write(struct Connection *conn)
{
    rewind(conn->file);

    int rc = fwrite(conn->db, sizeof(struct Database), 1, conn->file);
    if(rc != 1) die("Failed to write database.");

    rc = fflush(conn->file);
    if(rc == -1) die("Cannot flush database.");
}

void Database_create(struct Connection *conn)
{
    int i = 0;

    for(i = 0; i < MAX_ROWS; i++) {
        // make a prototype to initialize it
        struct Address addr = {.id = i, .set = 0};
        // then just assign it
        conn->db->rows[i] = addr;
    }
}

void Database_set(struct Connection *conn, int id, const char *name, const char *email)
{
    struct Address *addr = &conn->db->rows[id];
    if(addr->set) die("Already set, delete it first");

    addr->set = 1;
    // WARNING: bug, read the "How To Break It" and fix this
    char *res = strncpy(addr->name, name, MAX_DATA);
    // demonstrate the strncpy bug
    if(!res) die("Name copy failed");
}

```

```

        res = strncpy(addr->email, email, MAX_DATA);
        if(!res) die("Email copy failed");
    }

void Database_get(struct Connection *conn, int id)
{
    struct Address *addr = &conn->db->rows[id];

    if(addr->set) {
        Address_print(addr);
    } else {
        die("ID is not set");
    }
}

void Database_delete(struct Connection *conn, int id)
{
    struct Address addr = {.id = id, .set = 0};
    conn->db->rows[id] = addr;
}

void Database_list(struct Connection *conn)
{
    int i = 0;
    struct Database *db = conn->db;

    for(i = 0; i < MAX_ROWS; i++) {
        struct Address *cur = &db->rows[i];

        if(cur->set) {
            Address_print(cur);
        }
    }
}

int main(int argc, char *argv[])
{
    if(argc < 3) die("USAGE: ex17 <dbfile> <action> [action params]");

    char *filename = argv[1];
    char action = argv[2][0];
    struct Connection *conn = Database_open(filename, action);
    int id = 0;

    if(argc > 3) id = atoi(argv[3]);
    if(id >= MAX_ROWS) die("There's not that many records.");

    switch(action) {
        case 'c':
            Database_create(conn);
            Database_write(conn);
            break;

        case 'g':
            if(argc != 4) die("Need an id to get");

            Database_get(conn, id);
            break;

        case 's':
            if(argc != 6) die("Need id, name, email to set");

            Database_set(conn, id, argv[4], argv[5]);
            Database_write(conn);
            break;

        case 'd':
            if(argc != 4) die("Need id to delete");
    }
}

```

```

        Database_delete(conn, id);
        Database_write(conn);
        break;

    case 'l':
        Database_list(conn);
        break;
    default:
        die("Invalid action, only: c=create, g=get, s=set, d=del, l=list")
;
}

Database_close(conn);

return 0;
}

```

In this program I am using a set of structures to create a simple database for an address book. In it I'm using some things you've never seen, so you should go through it line-by-line, explain what each line does, and look up any functions you do not recognize. There are few key things I'm doing that you should pay attention to as well:

`#define` for constants

I use another part of the "C Pre-Processor" to create constant settings of `MAX_DATA` and `MAX_ROWS`. I'll cover more of what the CPP does, but this is a way to create a constant that will work reliably. There's other ways but they don't apply in certain situations.

Fixed Sized Structs

The `Address` struct then uses these constants to create a piece of data that is fixed in size making it less efficient, but easier to store and read. The `Database` struct is then also fixed size because it is a fixed length array of `Address` structs. That lets you write the whole thing to disk in one move later on.

`die` function to abort with an error

In a small program like this you can make a single function that kills the program with an error if there's anything wrong. I call this `die`, and it's used after any failed function calls or bad inputs to exit with an error using `exit`.

`errno` and `perror()` for error reporting

When you have an error return from a function, it will usually set an "external" variable called `errno` to say exactly what error happened. These are just numbers, so you can use `perror` to "print the error message".

FILE functions

I'm using all new functions like `fopen`, `fread`, `fclose`, and `rewind` to work with files. Each of these functions works on a `FILE` struct that's just like your structs, but it's defined by the C standard library.

nested struct pointers

There's use of nested structures and getting the address of array elements that you should study. Specifically code like `&conn->db->rows[i]` which reads "get the `i` element of `rows`, which is in `db`, which is in `conn`, then get the address of (`&`) it".

copying struct prototypes

best shown in `Database_delete`, you can see I'm using a temporary local `Address`, initializing its `id` and `set` fields, and

then simply copying it into the `rows` array by assigning it to the element I want. This trick makes sure that all fields but `set` and `id` are initialized to 0s and is actually easier to write. Incidentally, you shouldn't be using `memcpy` to do these kinds of struct copying operations. Modern C allows you to simply assign one struct to another and it'll handle the copying for you.

processing complex arguments

I'm doing some more complex argument parsing, but this isn't really the best way to do it. We'll get into better option parsing later in the book.

converting strings to ints

I use the `atoi` function to take the string for the id on the command line and convert it to the `int id` variable. Read up on this function and similar ones.

allocating large data on the "heap"

The whole point of this program is that I'm using `malloc` to ask the OS for a large amount of memory to work with when I create the `Database`. I cover this in more detail below.

NULL is 0 so boolean works

In many of the checks I'm testing that a pointer is not NULL by simply doing `if(!ptr) die("fail!")`; this is valid because NULL will evaluate to false. You could be explicit and say `if(ptr == NULL) die("fail!")` as well. On some rare systems NULL will be stored in the computer (represented) as something not 0, but the C standard says you should still be able to write code as if it has a 0 value. From now on when I say "NULL is 0" I mean its value for anyone who is overly pedantic.

## What You Should See

You should spend as much time as you can testing that it works, and running it with `valgrind` to confirm you've got all the memory usage right. Here's a session of me testing it normally and then using `Valgrind` to check the operations:

```
$ make ex17
cc -Wall -g    ex17.c    -o ex17
$ ./ex17 db.dat c
$ ./ex17 db.dat s 1 zed zed@zedshaw.com
$ ./ex17 db.dat s 2 frank frank@zedshaw.com
$ ./ex17 db.dat s 3 joe joe@zedshaw.com
$
$ ./ex17 db.dat l
1 zed zed@zedshaw.com
2 frank frank@zedshaw.com
3 joe joe@zedshaw.com
$ ./ex17 db.dat d 3
$ ./ex17 db.dat l
1 zed zed@zedshaw.com
2 frank frank@zedshaw.com
$ ./ex17 db.dat g 2
2 frank frank@zedshaw.com
$
$ valgrind --leak-check=yes ./ex17 db.dat g 2
# cut valgrind output...
$
```

The actual output of `valgrind` is taken out since you should be able to detect it.

## NOTE

`Valgrind` will report that you're leaking small blocks of memory, but sometimes it's just over-reporting from OSX's internal APIs. If you see it showing leaks that aren't inside your code then just ignore them.

## Heap vs. Stack Allocation

You kids these days have it great. You play with your Ruby or Python and just make objects and variables without any care for where they live. You don't care if it's on the "stack", and the heap? Fuggedaboutit. You don't even know, and you know what, chances are your language of choice doesn't even put the variables on stack at all. It's all heap, and you don't even *know* if it is.

C is different because it's using the real CPU's actual machinery to do its work, and that involves a chunk of ram called the stack and another called the heap. What's the difference? It all depends on where you get the storage.

The heap is easier to explain as it's just all the remaining memory in your computer, and you access it with the function `malloc` to get more. Each time you call `malloc`, the OS uses internal functions to register that piece of memory to you, and then returns a pointer to it. When you're done with it, you use `free` to return it to the OS so that it can be used by other programs. Failing to do this will cause your program to "leak" memory, but `Valgrind` will help you track these leaks down.

The stack is a special region of memory that stores temporary variables each function creates as locals to that function. How it works is each argument to a function is "pushed" onto the stack, and then used inside the function. It is really a stack data structure, so the last thing in is the first thing out. This also happens with all local variables like `char action` and `int id` in `main`. The advantage of using a stack for this is simply that, when the function exits, the C compiler "pops" these variables off the stack to clean up. This is simple and prevents memory leaks if the variable is on the stack.

The easiest way to keep this straight is with this mantra: If you didn't get it from `malloc` or a function that got it from `malloc`, then it's on the stack.

There's three primary problems with stacks and heaps to watch for:

- If you get a block of memory from `malloc`, and have that pointer on the stack, then when the function exits, the pointer will get popped off and lost.
- If you put too much data on the stack (like large structs and arrays) then you can cause a "stack overflow" and the program will abort. In this case, use the heap with `malloc`.

- If you take a pointer to something on the stack, and then pass that or return it from your function, then the function receiving it will "segmentation fault" (segfault) because the actual data will get popped off and disappear. You'll be pointing at dead space.

This is why in the program I've created a `Database_open` that allocates memory or dies, and then a `Database_close` that frees everything. If you create a "create" function, that makes the whole thing or nothing, and then a "destroy" function that cleans up everything safely, then it's easier to keep it all straight.

Finally, when a program exits the OS will clean up all the resources for you, but sometimes not immediately. A common idiom (and one I use in this exercise) is to just abort and let the OS clean up on error.

## How To Break It

This program has a lot of places you can break it, so try some of these but also come up with your own:

- The classic way is to remove some of the safety checks such that you can pass in arbitrary data. For example, if you remove the check on line 160 that prevents you from passing in any record number.
- You can also try corrupting the data file. Open it in any editor and change random bytes then close it.
- You could also find ways to pass bad arguments to the program when it's run, such as getting the file and action backwards will make it create a file named after the action, then do an action based on the first character.
- There is a bug in this program because of `strncpy` being poorly designed. Go read about `strncpy` then try to find out what happens when the `name` or `address` you give is *greater* than 512 bytes. Fix this by simply forcing the last character to `'\0'` so that it's always set no matter what (which is what `strncpy` should do).
- In the extra credit I have you augment the program to create arbitrary size databases. Try to see what the biggest database is before you cause the program to die for lack of memory from `malloc`.

## Extra Credit

- The `die` function needs to be augmented to let you pass the `conn` variable so it can close it and clean up.
- Change the code to accept parameters for `MAX_DATA` and `MAX_ROWS`, store them in the `Database` struct, and write that to the file, thus creating a database that can be arbitrarily sized.
- Add more operations you can do on the database, like `find`.
- Read about how C does its struct packing, and then try to see why your file is the size it is. See if you can calculate a new size after adding more fields.
- Add some more fields to the `Address` and make them searchable.
- Write a shell script that will do your testing automatically for

you by running commands in the right order. Hint: Use `set -e` at the top of a `bash` to make it abort the whole script if any command has an error.

- Try reworking the program to use a single global for the database connection. How does this new version of the program compare to the other one?
- Go research "stack data structure" and write one in your favorite language, then try to do it in C.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 18: Pointers To Functions

Functions in C are actually just pointers to a spot in the program where some code exists. Just like you've been creating pointers to structs, strings, and arrays, you can point a pointer at a function too. The main use for this is to pass "callbacks" to other functions, or to simulate classes and objects. In this exercise we'll do some callbacks, and in the next one we'll make a simple object system.

The format of a function pointer goes like this:

```
int (*POINTER_NAME)(int a, int b)
```

A way to remember how to write one is to do this:

- Write a normal function declaration: `int callme(int a, int b)`
- Wrap function name with pointer syntax: `int (*callme)(int a, int b)`
- Change the name to the pointer name: `int (*compare_cb)(int a, int b)`

The key thing to remember is, when you're done with this, the *variable* name for the pointer is called `compare_cb` and then you use it just like it's a function. This is similar to how pointers to arrays can be used just like the arrays they point to. Pointers to functions can be used like the functions they point to but with a different name.

```
int (*tester)(int a, int b) = sorted_order;  
printf("TEST: %d is same as %d\n", tester(2, 3),  
sorted_order(2, 3));
```

This will work even if the function pointer returns a pointer to something:

- Write it: `char *make_coolness(int awesome_levels)`
- Wrap it: `char *(*make_coolness)(int awesome_levels)`
- Rename it: `char *(*coolness_cb)(int awesome_levels)`

The next problem to solve with using function pointers is that it's hard to give them as parameters to a function, like when you want to pass the function callback to another function. The solution to this is to use `typedef` which is a C keyword for making new names for other more complex types. The only thing you need to do is put `typedef` before the same function pointer syntax, and then after that you can use the name like it's a type. I demonstrate this in the following exercise code:

```
#include <stdio.h>  
#include <stdlib.h>  
#include <errno.h>  
#include <string.h>  
  
/** Our old friend die from ex17. */  
void die(const char *message)  
{  
    if(errno) {
```

```

        perror(message);
    } else {
        printf("ERROR: %s\n", message);
    }

    exit(1);
}

// a typedef creates a fake type, in this
// case for a function pointer
typedef int (*compare_cb)(int a, int b);

/**
 * A classic bubble sort function that uses the
 * compare_cb to do the sorting.
 */
int *bubble_sort(int *numbers, int count, compare_cb cmp)
{
    int temp = 0;
    int i = 0;
    int j = 0;
    int *target = malloc(count * sizeof(int));

    if(!target) die("Memory error.");

    memcpy(target, numbers, count * sizeof(int));

    for(i = 0; i < count; i++) {
        for(j = 0; j < count - 1; j++) {
            if(cmp(target[j], target[j+1]) > 0) {
                temp = target[j+1];
                target[j+1] = target[j];
                target[j] = temp;
            }
        }
    }

    return target;
}

int sorted_order(int a, int b)
{
    return a - b;
}

int reverse_order(int a, int b)
{
    return b - a;
}

int strange_order(int a, int b)
{
    if(a == 0 || b == 0) {
        return 0;
    } else {
        return a % b;
    }
}

/**
 * Used to test that we are sorting things correctly
 * by doing the sort and printing it out.
 */
void test_sorting(int *numbers, int count, compare_cb cmp)
{
    int i = 0;
    int *sorted = bubble_sort(numbers, count, cmp);

    if(!sorted) die("Failed to sort as requested.");
}

```

```

    for(i = 0; i < count; i++) {
        printf("%d ", sorted[i]);
    }
    printf("\n");

    free(sorted);
}

int main(int argc, char *argv[])
{
    if(argc < 2) die("USAGE: ex18 4 3 1 5 6");

    int count = argc - 1;
    int i = 0;
    char **inputs = argv + 1;

    int *numbers = malloc(count * sizeof(int));
    if(!numbers) die("Memory error.");

    for(i = 0; i < count; i++) {
        numbers[i] = atoi(inputs[i]);
    }

    test_sorting(numbers, count, sorted_order);
    test_sorting(numbers, count, reverse_order);
    test_sorting(numbers, count, strange_order);

    free(numbers);

    return 0;
}

```

In this program you're creating a dynamic sorting algorithm that can sort an array of integers using a comparison callback. Here's the breakdown of this program so you can clearly understand it:

ex18.c:1-6

The usual includes needed for all the functions we call.

ex18.c:7-17

This is the `die` function from the previous exercise which I'll use to do error checking.

ex18.c:21

This is where the `typedef` is used, and later I use `compare_cb` like it's a type similar to `int` or `char` in `bubble_sort` and `test_sorting`.

ex18.c:27-49

A bubble sort implementation, which is a very inefficient way to sort some integers. This function contains:

ex18.c:27

Here's where I use the `typedef` for `compare_cb` as the last parameter `cmp`. This is now a function that will return a comparison between two integers for sorting.

ex18.c:29-34

The usual creation of variables on the stack, followed by a new array of integers on the heap using `malloc`. Make sure you understand what `count * sizeof(int)` is doing.

ex18.c:38

The outer-loop of the bubble sort.

ex18.c:39

The inner-loop of the bubble sort

ex18.c:40

Now I call the `cmp` callback just like it's a normal function, but instead of being the name of something we defined, it's just a pointer to it. This lets the caller pass in anything they want as long as it matches the "signature" of the `compare_cb` `typedef`.

ex18.c:41-43

The actual swapping operation a bubble sort needs to do what it does.

ex18.c:48

Finally return the newly created and sorted result array `target`.

ex18.c:51-68

Three different versions of the `compare_cb` function type, which needs to have the same definition as the `typedef` we created. The C compiler will complain to you if you get this wrong and say the types don't match.

ex18.c:74-87

This is a tester for the `bubble_sort` function. You can see now how I'm also using `compare_cb` to then pass to `bubble_sort` demonstrating how these can be passed around like any other pointers.

ex18.c:90-103

A simple main function that sets up an array based on integers you pass on the command line, then calls the `test_sorting` function.

ex18.c:105-107

Finally, you get to see how the `compare_cb` function pointer `typedef` is used. I simply call `test_sorting` but give it the name of `sorted_order`, `reverse_order`, and `strange_order` as the function to use. The C compiler then finds the address of those functions, and makes it a pointer for `test_sorting` to use. If you look at `test_sorting` you'll see it then passes each of these to `bubble_sort` but it actually has no idea what they do, only that they match the `compare_cb` prototype and should work.

ex18.c:109

Last thing we do is free up the array of numbers we made.

## What You Should See

Running this program is simple, but try different combinations of numbers, and try even non-numbers to see what it does.

```
$ make ex18
cc -Wall -g    ex18.c    -o ex18
$ ./ex18 4 1 7 3 2 0 8
0 1 2 3 4 7 8
8 7 4 3 2 1 0
3 4 2 7 1 0 8
$
```

## How To Break It

I'm going to have you do something kind of weird to break this. These function pointers are pointers like every other pointer, so they point at blocks of memory. C has this ability to take one pointer and convert it to another so you can process the data in different ways. It's usually not necessary, but to show you how to hack your computer, I want you to add this at the end of `test_sorting`:

```
unsigned char *data = (unsigned char *)cmp;

for(i = 0; i < 25; i++) {
    printf("%02x:", data[i]);
}

printf("\n");
```

This loop is sort of like converting your function to a string and then printing out it's contents. This won't break your program unless the CPU and OS you're on has a problem with you doing this. What you'll see is a string of hexadecimal numbers after it prints the sorted array:

```
55:48:89:e5:89:7d:fc:89:75:f8:8b:55:fc:8b:45:f8:29:d0:
c9:c3:55:48:89:e5:89:
```

That should be the raw assembler byte code of the function itself, and you should see they start the same, but then have different endings. It's also possible that this loop isn't getting all of the function or is getting too much and stomping on another piece of the program. Without more analysis you wouldn't know.

## Extra Credit

- Get a hex editor and open up `ex18`, then find this sequence of hex digits that start a function to see if you can find the function in the raw program.
- Find other random things in your hex editor and change them. Rerun your program and see what happens. Changing strings you find are the easiest things to change.
- Pass in the wrong function for the `compare_cb` and see what the C compiler complains about.
- Pass in NULL and watch your program seriously bite it. Then run `Valgrind` and see what that reports.
- Write another sorting algorithm, then change `test_sorting` so that it takes *both* an arbitrary sort function and the sort function's callback comparison. Use it to test both of your algorithms.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 19: A Simple Object System

I learned C before I learned Object Oriented Programming, so it helped me to build an OOP system in C to understand the basics of what OOP meant. You are probably the kind of person who learned an OOP language before you learned C, so this kind of bridge might help you as well. In this exercise, you will build a simple object system, but also learn more about the C Pre-Processor or CPP.

This exercise will build a simple game where you kill a Minotaur in a small little castle. Nothing fancy, just four rooms and a bad guy. This project will also be a multi-file project, and look more like a real C software project than your previous ones. This is why I'm introducing the CPP here because you need it to start using multiple files in your own software.

## How The CPP Works

The C Pre-Processor is a template processing system. It's a highly targeted one that helps make C easier to work with, but it does this by having a syntax aware templating mechanism. Traditionally people just used the CPP to store constants and make "macros" to simplify repetitive coding. In modern C you'll actually use the CPP as a code generator to create templated pieces of code.

How the CPP works is you give it one file, usually a .c file, and it processes various bits of text starting with the `#` (octothorpe) character. When it encounters one of these it performs a specific replacement on the text of the input file. It's main advantage though is it can *include* other files, and then augment its list of macros based on that file's contents.

A quick way to see what the CPP does is take the last exercise and run this:

```
cpp ex18.c | less
```

It will be a huge amount of output, but scroll through it and you'll see the contents of the other files you included with `#include`. Scroll down to the original code and you can see how the `cpp` is altering the source based on various `#define` macros in the header files.

The C compiler is so tightly integrated with `cpp` that it just runs this for you and understands how it works intimately. In modern C, the `cpp` system is so integral to C's function that you might as well just consider it to be part of the language.

In the remaining sections, we'll be using more of the CPP syntax and explaining it as we go.

# The Prototype Object System

The OOP system we'll create is a simple "prototype" style object system more like JavaScript. Instead of classes, you start with prototypes that have fields set, and then use those as the basis of creating other object instances. This "classless" design is much easier to implement and work with than a traditional class based one.

## The Object Header File

I want to put the data types and function declarations into a separate header file named `object.h`. This is standard C practice and it lets you ship binary libraries but still let the programmer compile against it. In this file I have several advanced CPP techniques I'm going to quickly describe and then have you see in action later:

```
#ifndef _object_h
#define _object_h

typedef enum {
    NORTH, SOUTH, EAST, WEST
} Direction;

typedef struct {
    char *description;
    int (*init)(void *self);
    void (*describe)(void *self);
    void (*destroy)(void *self);
    void (*move)(void *self, Direction direction);
    int (*attack)(void *self, int damage);
} Object;

int Object_init(void *self);
void Object_destroy(void *self);
void Object_describe(void *self);
void *Object_move(void *self, Direction direction);
int Object_attack(void *self, int damage);
void *Object_new(size_t size, Object proto, char *description);

#define NEW(T, N) Object_new(sizeof(T), T##Proto, N)
#define _(N) proto.N

#endif
```

Taking a look at this file, you can see we have a few new pieces of syntax you haven't encountered before:

`#ifndef`

You've seen a `#define` for making simple constants, but the CPP can also do logic and remove sections of code. This `#ifndef` is "if not defined" and checks if there's already a `#define _object_h` and if there is it skips all of this code. I do this so that we can include this file any time we want and not worry about it defining things multiple times.

`#define`

With the above `#ifndef` shielding this file from we then add the `_object_h` define so that any attempts to include it later cause the above to skip.

`#define NEW(T,N)`

This makes a macro, and it works like a template function that

spits out the code on the right, whenever you write use the macro on the left. This one is simply making a short version of the normal way we'll call `Object_new` and avoids potential errors with calling it wrong. The way the macro works is the `T` and `N` parameters to `NEW` are "injected" into the line of code on the right. The syntax `T##Proto` says to "concat Proto at the end of T", so if you had `NEW(Room, "Hello.")` then it'd make `RoomProto` there.

```
#define _(N)
```

This macro is a bit of "syntactic sugar" for the object system and basically helps you write `obj->proto.blah` as simply `obj->_(blah)`. It's not necessary, but it's a fun little trick that I'll use later.

## The Object Source File

The `object.h` file is declaring functions and data types that are defined (created) in the `object.c`, so that's next:

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include "object.h"
#include <assert.h>

void Object_destroy(void *self)
{
    Object *obj = self;

    if(obj) {
        if(obj->description) free(obj->description);
        free(obj);
    }
}

void Object_describe(void *self)
{
    Object *obj = self;
    printf("%s.\n", obj->description);
}

int Object_init(void *self)
{
    // do nothing really
    return 1;
}

void *Object_move(void *self, Direction direction)
{
    printf("You can't go that direction.\n");
    return NULL;
}

int Object_attack(void *self, int damage)
{
    printf("You can't attack that.\n");
    return 0;
}

void *Object_new(size_t size, Object proto, char *description)
{
    // setup the default functions in case they aren't set
    if(!proto.init) proto.init = Object_init;
    if(!proto.describe) proto.describe = Object_describe;
    if(!proto.destroy) proto.destroy = Object_destroy;
```



```

if(!proto.attack) proto.attack = Object_attack;
if(!proto.move) proto.move = Object_move;

// this seems weird, but we can make a struct of one size,
// then point a different pointer at it to "cast" it
Object *el = calloc(1, size);
*el = proto;

// copy the description over
el->description = strdup(description);

// initialize it with whatever init we were given
if(!el->init(el)) {
    // looks like it didn't initialize properly
    el->destroy(el);
    return NULL;
} else {
    // all done, we made an object of any type
    return el;
}
}

```

There's really nothing new in this file, except one *tiny* little trick. The function `Object_new` uses an aspect of how `structs` work by putting the base prototype at the beginning of the struct. When you look at the `ex19.h` header later, you'll see how I make the first field in the struct an `Object`. Since C puts the fields in a struct in order, and since a pointer just points at a chunk of memory, I can "cast" a pointer to anything I want. In this case, even though I'm taking a potentially larger block of memory from `calloc`, I'm using a `Object` pointer to work with it.

I explain this a bit better when we write the `ex19.h` file since it's easier to understand when you see it being used.

That creates your base object system, but you'll need a way to compile it and link it into your `ex19.c` file to create a complete program. The `object.c` file on its own doesn't have a `main` so it isn't enough to make a full program. Here's a `Makefile` that will do this based on the one you've been using:

```

CFLAGS=-Wall -g

all: ex19

ex19: object.o

clean:
    rm -f ex19

```

This `Makefile` is doing nothing more than saying that `ex19` depends on `object.o`. Remember how `make` knows how to build different kinds of files by their extensions? Doing this tells make the following:

- When I say run `make` the default `all` should just build `ex19`.
- When you build `ex19`, you need to also build `object.o` and include it in the build.
- `make` can't see anything in the file for `object.o`, but it does see an `object.c` file, and it knows how to turn a `.c` into a `.o`, so it does that.
- Once it has `object.o` built it then runs the correct compile command to build `ex19` from `ex19.c` and `object.o`.

# The Game Implementation

Once you have those files you just need to implement the actual game using the object system, and first step is putting all the data types and function declarations in a `ex19.h` file:

```
#ifndef _ex19_h
#define _ex19_h

#include "object.h"

struct Monster {
    Object proto;
    int hit_points;
};

typedef struct Monster Monster;

int Monster_attack(void *self, int damage);
int Monster_init(void *self);

struct Room {
    Object proto;

    Monster *bad_guy;

    struct Room *north;
    struct Room *south;
    struct Room *east;
    struct Room *west;
};

typedef struct Room Room;

void *Room_move(void *self, Direction direction);
int Room_attack(void *self, int damage);
int Room_init(void *self);

struct Map {
    Object proto;
    Room *start;
    Room *location;
};

typedef struct Map Map;

void *Map_move(void *self, Direction direction);
int Map_attack(void *self, int damage);
int Map_init(void *self);

#endif
```

That sets up three new Objects you'll be using: `Monster`, `Room`, and `Map`.

Taking a look at `object.c:52` you can see where I use a pointer `Object *el = calloc(1, size)`. Go back and look at the `NEW` macro in `object.h` and you can see that it is getting the `sizeof` another struct, say `Room`, and I allocate that much. However, because I've pointed a `Object` pointer at this block of memory, and because I put an `Object proto` field at the front of `Room`, I'm able to treat a `Room` like it's an `Object`.

The way to break this down is like so:

- I call `NEW(Room, "Hello.")` which the CPP expands as a macro

- into `Object_new(sizeof(Room), RoomProto, "Hello.")`.
- This runs, and inside `Object_new` I allocate a piece of memory that's `Room` in size, *but* point a `Object *el` pointer at it.
- Since C puts the `Room.proto` field first, that means the `el` pointer is really only pointing at enough of the block of memory to see a full `Object` struct. It has no idea that it's even called `proto`.
- It then uses this `Object *el` pointer to set the contents of the piece of memory correctly with `*el = proto;`. Remember that you can copy structs, and that `*el` means "the value of whatever el points at", so this means "assign the proto struct to whatever el points at".
- Now that this mystery struct is filled in with the right data from `proto`, the function can then call `init` or `destroy` on the `Object`, but the cool part is whoever called this function can *change* these out for whatever ones they want.

And with that, we have a way to get this one function to construct new types, and give them new functions to change their behavior. This may seem like "hackery" but it's stock C and totally valid. In fact there's quite a few standard system functions that work this same way, and we'll be using some of them for converting addresses in network code.

With the function definitions and data structures written out I can now actually implement the game with four rooms and a minotaur to beat up:

```
#include <stdio.h>
#include <errno.h>
#include <stdlib.h>
#include <string.h>
#include <time.h>
#include "ex19.h"

int Monster_attack(void *self, int damage)
{
    Monster *monster = self;

    printf("You attack %s!\n", monster->(description));

    monster->hit_points -= damage;

    if(monster->hit_points > 0) {
        printf("It is still alive.\n");
        return 0;
    } else {
        printf("It is dead!\n");
        return 1;
    }
}

int Monster_init(void *self)
{
    Monster *monster = self;
    monster->hit_points = 10;
    return 1;
}

Object MonsterProto = {
    .init = Monster_init,
    .attack = Monster_attack
};
```

```

void *Room_move(void *self, Direction direction)
{
    Room *room = self;
    Room *next = NULL;

    if(direction == NORTH && room->north) {
        printf("You go north, into:\n");
        next = room->north;
    } else if(direction == SOUTH && room->south) {
        printf("You go south, into:\n");
        next = room->south;
    } else if(direction == EAST && room->east) {
        printf("You go east, into:\n");
        next = room->east;
    } else if(direction == WEST && room->west) {
        printf("You go west, into:\n");
        next = room->west;
    } else {
        printf("You can't go that direction.");
        next = NULL;
    }

    if(next) {
        next->_(describe)(next);
    }

    return next;
}

int Room_attack(void *self, int damage)
{
    Room *room = self;
    Monster *monster = room->bad_guy;

    if(monster) {
        monster->_(attack)(monster, damage);
        return 1;
    } else {
        printf("You flail in the air at nothing. Idiot.\n");
        return 0;
    }
}

Object RoomProto = {
    .move = Room_move,
    .attack = Room_attack
};

void *Map_move(void *self, Direction direction)
{
    Map *map = self;
    Room *location = map->location;
    Room *next = NULL;

    next = location->_(move)(location, direction);

    if(next) {
        map->location = next;
    }

    return next;
}

int Map_attack(void *self, int damage)
{
    Map* map = self;
    Room *location = map->location;

```

```

        return location->_(attack)(location, damage);
    }

int Map_init(void *self)
{
    Map *map = self;

    // make some rooms for a small map
    Room *hall = NEW(Room, "The great Hall");
    Room *throne = NEW(Room, "The throne room");
    Room *arena = NEW(Room, "The arena, with the minotaur");
    Room *kitchen = NEW(Room, "Kitchen, you have the knife now");

    // put the bad guy in the arena
    arena->bad_guy = NEW(Monster, "The evil minotaur");

    // setup the map rooms
    hall->north = throne;

    throne->west = arena;
    throne->east = kitchen;
    throne->south = hall;

    arena->east = throne;
    kitchen->west = throne;

    // start the map and the character off in the hall
    map->start = hall;
    map->location = hall;

    return 1;
}

Object MapProto = {
    .init = Map_init,
    .move = Map_move,
    .attack = Map_attack
};

int process_input(Map *game)
{
    printf("\n> ");

    char ch = getchar();
    getchar(); // eat ENTER

    int damage = rand() % 4;

    switch(ch) {
        case -1:
            printf("Giving up? You suck.\n");
            return 0;
            break;

        case 'n':
            game->_(move)(game, NORTH);
            break;

        case 's':
            game->_(move)(game, SOUTH);
            break;

        case 'e':
            game->_(move)(game, EAST);
            break;

        case 'w':
            game->_(move)(game, WEST);

```

```

        break;

    case 'a':

        game->_(attack)(game, damage);
        break;
    case 'l':
        printf("You can go:\n");
        if(game->location->north) printf("NORTH\n");
        if(game->location->south) printf("SOUTH\n");
        if(game->location->east) printf("EAST\n");
        if(game->location->west) printf("WEST\n");
        break;

    default:
        printf("What?: %d\n", ch);
}

return 1;
}

int main(int argc, char *argv[])
{
    // simple way to setup the randomness
    srand(time(NULL));

    // make our map to work with
    Map *game = NEW(Map, "The Hall of the Minotaur.");

    printf("You enter the ");
    game->location->_(describe)(game->location);

    while(process_input(game)) {
    }

    return 0;
}

```

Honestly there isn't much in this that you haven't seen, and only you might need to understand how I'm using the macros I made from the headers files. Here's the important key things to study and understand:

- Implementing a prototype involves creating its version of the functions, and then creating a single struct ending in "Proto". Look at `MonsterProto`, `RoomProto` and `MapProto`.
- Because of how `Object_new` is implemented, if you don't set a function in your prototype, then it will get the default implementation created in `object.c`.
- In `Map_init` I create the little world, but more importantly I use the `NEW` macro from `object.h` to build all of the objects. To get this concept in your head, try replacing the `NEW` usage with direct `Object_new` calls to see how it's being translated.
- Working with these objects involves calling functions on them, and the `_(N)` macro does this for me. If you look at the code `monster->_(attack)(monster, damage)` you see that I'm using the macro, which gets replaced with `monster->proto.attack(monster, damage)`. Study this transformation again by rewriting these calls back to their original. Also, if you get stuck then run `cpp` manually to see what it's going to do.
- I'm using two new functions `srand` and `rand`, which setup a simple random number generator good enough for the game. I also use `time` to initialize the random number generator. Research those.
- I use a new function `getchar` that gets a single character from

## What You Should See

Here's me playing my own game:

```
$ make ex19
cc -Wall -g -c -o object.o object.c
cc -Wall -g ex19.c object.o -o ex19
$ ./ex19
You enter the The great Hall.

> l
You can go:
NORTH

> n
You go north, into:
The throne room.

> l
You can go:
SOUTH
EAST
WEST

> e
You go east, into:
Kitchen, you have the knife now.

> w
You go west, into:
The throne room.

> s
You go south, into:
The great Hall.

> n
You go north, into:
The throne room.

> w
You go west, into:
The arena, with the minotaur.

> a
You attack The evil minotaur!
It is still alive.

> a
You attack The evil minotaur!
It is dead!

> ^D
Giving up? You suck.
$
```

## Auditing The Game

As an exercise for you I have left out all of the `assert` checks I normally put into a piece of software. You've seen me use `assert` to make sure a program is running correctly, but now I want you to go back and do the following:

- Look at each function you've defined, one file at a time.
- At the top of each function, add `asserts` that make sure the input parameters are correct. For example, in `Object_new` you want a `assert(description != NULL)`.
- Go through each line of the function, and find any functions being called. Read the documentation (man page) for that function, and confirm what it returns for an error. Add another assert to check that the error didn't happen. For example, in `Object_new` you need one after the call to `calloc` that does `assert(e1 != NULL)`.
- If a function is expected to return a value, either make sure it returns an error value (like NULL), or have an assert to make sure that the returned variable isn't invalid. For example, in `Object_new`, you need to have `assert(e1 != NULL)` again before the last return since that part can never be NULL.
- For every `if-statement` you write, make sure there's an else clause unless that if is an error check that causes an exit.
- For every `switch-statement` you write, make sure that there's a `default` case that handles anything you didn't anticipate.

Take your time going through every line of the function and find any errors you make. Remember that the point of this exercise is to stop being a "coder" and switch your brain into being a "hacker". Try to see how you could break it, then write code to prevent it or abort early if you can.

## Extra Credit

- Update the `Makefile` so that when you do `make clean` it will also remove the `object.o` file.
- Write a test script that works the game in different ways and augment the `Makefile` so you can run `make test` and it'll thrash the game with your script.
- Add more rooms and monsters to the game.
- Put the game mechanics into a third file, compile it to `.o`, and then use that to write another little game. If you're doing it right you should only have a new `Map` and a `main` function in the new game.

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 20: Zed's Awesome Debug Macros

There is a constant problem in C that you have been dancing around but which I am going to solve in this exercise using a set of macros I developed. You can thank me later when you realize how insanely awesome these macros are. Right now you won't realize how awesome they are, so you'll just have to use them and then you can walk up to me one day and say, "Zed, those Debug Macros were the bomb. I owe you my first born child because you saved me a decade of heartache and prevented me from killing myself more than once. Thank you good sir, here's a million dollars and the original Snakehead Telecaster prototype signed by Leo Fender."

Yes, they are that awesome.

## The C Error Handling Problem

In almost every programming language handling errors is a difficult activity. There's entire programming languages that try as hard as they can to avoid even the concept of an error. Other languages invent complex control structures like exceptions to pass error conditions around. The problem exists mostly because programmers assume errors don't happen and this optimism infects the type of languages they use and create.

C tackles the problem by returning error codes and setting a global `errno` value that you check. This makes for complex code that simply exists to check if something you did had an error. As you write more and more C code you'll write code with the pattern:

- Call a function.
- If the return value is an error (must look that up each time too).
- Then cleanup all the resource created so far.
- and print out an error message that hopefully helps.

This means for every function call (and yes, *every* function) you are potentially writing 3-4 more lines just to make sure it worked. That doesn't include the problem of cleaning up all of the junk you've built to that point. If you have 10 different structures, 3 files, and a database connection, when you get an error then you would have 14 more lines.

In the past this wasn't a problem because C programs did what you've been doing when there's an error: die. No point in bothering with cleanup when the OS will do it for you. Today though many C programs need to run for weeks, months, or years and handle errors from many different sources gracefully. You can't just have your webserver die at the slightest touch, and you definitely can't have a library you've written nuke

a the program its used in. That's just rude.

Other languages solve this problem with exceptions, but those have problems in C (and in other languages too). In C you only have one return value, but exceptions are an entire stack based return system with arbitrary values. Trying to marshal exceptions up the stack in C is difficult, and no other libraries will understand it.

## The Debug Macros

The solution I've been using for years is a small set of "debug macros" that implement a basic debugging and error handling system for C. This system is easy to understand, works with every library, and makes C code more solid and clearer.

It does this by adopting the convention that whenever there's an error, your function will jump to an "error:" part of the function that knows how to cleanup everything and return an error code. You use a macro called `check` to check return codes, print an error message, and then jump to the cleanup section. You combine that with a set of logging functions for printing out useful debug messages.

I'll now show you the entire contents of the most awesome set of brilliance you've ever seen:

```
#ifndef __dbg_h__
#define __dbg_h__

#include <stdio.h>
#include <errno.h>
#include <string.h>

#ifdef NDEBUG
#define debug(M, ...)
#else
#define debug(M, ...) fprintf(stderr, "DEBUG %s:%d: " M "\n", __FILE__, __LINE__, ##__VA_ARGS__)
#endif

#define clean_errno() (errno == 0 ? "None" : strerror(errno))

#define log_err(M, ...) fprintf(stderr, "[ERROR] (%s:%d: errno: %s) " M "\n", __FILE__, __LINE__, clean_errno(), ##__VA_ARGS__)

#define log_warn(M, ...) fprintf(stderr, "[WARN] (%s:%d: errno: %s) " M "\n", __FILE__, __LINE__, clean_errno(), ##__VA_ARGS__)

#define log_info(M, ...) fprintf(stderr, "[INFO] (%s:%d) " M "\n", __FILE__, __LINE__, ##__VA_ARGS__)

#define check(A, M, ...) if(!(A)) { log_err(M, ##__VA_ARGS__); errno=0; goto error; }

#define sentinel(M, ...) { log_err(M, ##__VA_ARGS__); errno=0; goto error; }

#define check_mem(A) check((A), "Out of memory.")

#define check_debug(A, M, ...) if(!(A)) { debug(M, ##__VA_ARGS__); errno=0; goto error; }

#endif
```

Yes, that's it, and here's what every line does:

dbg.h:1-2

The usual defense against accidentally including the file twice, which you saw in the last exercise.

dbg.h:4-6

Includes for the functions that these macros need.

dbg.h:8

The start of a `#ifdef` which lets you recompile your program so that all the debug log messages are removed.

dbg.h:9

If you compile with `NDEBUG` defined, then "no debug" messages will remain. You can see in this case the `#define debug()` is just replaced with nothing (the right side is empty).

dbg.h:10

The matching `#else` for the above `#ifdef`.

dbg.h:11

The alternative `#define debug` that translates any use of `debug("format", arg1, arg2)` into an `fprintf` call to `stderr`. Many C programmers don't know, but you can create macros that actually work like `printf` and take variable arguments. Some C compilers (actually cpp) don't support this, but the ones that matter do. The magic here is the use of `##_VA_ARGS__` which says "put whatever they had for extra arguments (...) here". Also notice the use of `__FILE__` and `__LINE__` to get the current `file:line` for the debug message. Very helpful.

dbg.h:12

The end of the `#ifdef`.

dbg.h:14

The `clean_errno` macro that's used in the others to get a safe readable version of `errno`. That strange syntax in the middle is a "ternary operator" and you'll learn what it does later.

dbg.h:16-20

The `log_err`, `log_warn`, and `log_info` macros for logging messages meant for the end user. Works like `debug` but can't be compiled out.

dbg.h:22

The best macro ever, `check` will make sure the condition `A` is true, and if not logs the error `M` (with variable arguments for `log_err`), then jumps to the function's `error:` for cleanup.

dbg.h:24

The 2nd best macro ever, `sentinel` is placed in any part of a function that shouldn't run, and if it does prints an error message then jumps to the `error:` label. You put this in `if-statements` and `switch-statements` to catch conditions that shouldn't happen, like the `default:`.

dbg.h:26

A short-hand macro `check_mem` that makes sure a pointer is valid, and if it isn't reports it as an error with "Out of memory."

dbg.h:28

An alternative macro `check_debug` that still checks and handles an error, but if the error is common then you don't want to bother reporting it. In this one it will use `debug` instead of `log_err` to report the message, so when you define `NDEBUG` the check still happens, the error jump goes off, but the message isn't printed.

# Using dbg.h

Here's an example of using all of `dbg.h` in a small program. This doesn't actually do anything but demonstrate how to use each macro, but we'll be using these macros in all of the programs we write from now on, so be sure to understand how to use them.

```
#include "dbg.h"
#include <stdlib.h>
#include <stdio.h>

void test_debug()
{
    // notice you don't need the \n
    debug("I have Brown Hair.");

    // passing in arguments like printf
    debug("I am %d years old.", 37);
}

void test_log_err()
{
    log_err("I believe everything is broken.");
    log_err("There are %d problems in %s.", 0, "space");
}

void test_log_warn()
{
    log_warn("You can safely ignore this.");
    log_warn("Maybe consider looking at: %s.", "/etc/passwd");
}

void test_log_info()
{
    log_info("Well I did something mundane.");
    log_info("It happened %f times today.", 1.3f);
}

int test_check(char *file_name)
{
    FILE *input = NULL;
    char *block = NULL;

    block = malloc(100);
    check_mem(block); // should work

    input = fopen(file_name, "r");
    check(input, "Failed to open %s.", file_name);

    free(block);
    fclose(input);
    return 0;

error:
    if(block) free(block);
    if(input) fclose(input);
    return -1;
}

int test_sentinel(int code)
{
    char *temp = malloc(100);
    check_mem(temp);

    switch(code) {
        case 1:
```

```

        log_info("It worked.");
        break;
    default:
        sentinel("I shouldn't run.");
    }

    free(temp);
    return 0;

error:
    if(temp) free(temp);
    return -1;
}

int test_check_mem()
{
    char *test = NULL;
    check_mem(test);

    free(test);
    return 1;

error:
    return -1;
}

int test_check_debug()
{
    int i = 0;
    check_debug(i != 0, "Oops, I was 0.");

    return 0;
error:
    return -1;
}

int main(int argc, char *argv[])
{
    check(argc == 2, "Need an argument.");

    test_debug();
    test_log_err();
    test_log_warn();
    test_log_info();

    check(test_check("ex20.c") == 0, "failed with ex20.c");
    check(test_check(argv[1]) == -1, "failed with argv");
    check(test_sentinel(1) == 0, "test_sentinel failed.");
    check(test_sentinel(100) == -1, "test_sentinel failed.");
    check(test_check_mem() == -1, "test_check_mem failed.");
    check(test_check_debug() == -1, "test_check_debug failed.");

    return 0;

error:
    return 1;
}

```

Pay attention to how `check` is used, and how when it is `false` it will jump to the `error:` label to do a cleanup. The way to read those lines is, "check that A is true and if not say M and jump out."

## What You Should See

When you run this, give it some bogus first parameter and you should see this:

```
$ make ex20
cc -Wall -g -DNDEBUG      ex20.c      -o ex20
$ ./ex20 test
[ERROR] (ex20.c:16: errno: None) I believe everything is broken.
[ERROR] (ex20.c:17: errno: None) There are 0 problems in space.
[WARN] (ex20.c:22: errno: None) You can safely ignore this.
[WARN] (ex20.c:23: errno: None) Maybe consider looking at: /etc/passwd.
[INFO] (ex20.c:28) Well I did something mundane.
[INFO] (ex20.c:29) It happened 1.300000 times today.
[ERROR] (ex20.c:38: errno: No such file or directory) Failed to open test.
[INFO] (ex20.c:57) It worked.
[ERROR] (ex20.c:60: errno: None) I shouldn't run.
[ERROR] (ex20.c:74: errno: None) Out of memory.
```

See how it reports the exact line number where the `check` failed? That's going to save you hours of debugging later. See also how it prints the error message for you when `errno` is set? Again, that will save you hours of debugging.

## How The CPP Expands Macros

It's now time for you to get a small introduction to the CPP so that you know how these macros actually work. To do this, I'm going to break down the most complex macro from `dbg.h` and have you run `cpp` so you can see what it's actually doing.

Imagine I have a function called `dosomething()` that return the typical 0 for success and -1 for an error. Every time I call `dosomething` I have to check for this error code, so I'd write code like this:

```
int rc = dosomething();

if(rc != 0) {
    fprintf(stderr, "There was an error: %s\n",
    strerror());
    goto error;
}
```

What I want to use the CPP for is to encapsulate this `if-statement` I have to use all the time into a more readable and memorable line of code. I want what you've been doing in `dbg.h` with the `check` macro:

```
int rc = dosomething();
check(rc == 0, "There was an error.");
```

This is *much* clearer and explains exactly what's going on: check that the function worked, and if not report an error. To do this, we need some special CPP "tricks" that make the CPP useful as a code generation tool. Take a look at the `check` and `log_err` macros again:

```
#define log_err(M, ...) fprintf(stderr, "[ERROR]
(%s:%d: errno: %s) " M "\n", __FILE__, __LINE__,
clean_errno(), ##__VA_ARGS__)
#define check(A, M, ...) if(!(A)) { log_err(M,
##__VA_ARGS__); errno=0; goto error; }
```

The first macro, `log_err` is simpler and simply replace itself with a call to `fprintf` to `stderr`. The only tricky part of this macro is the use of `...` in the definition `log_err(M, ...)`. What this does is let you pass variable arguments to the macro, so you can pass in the arguments that

should go to `fprintf`. How do they get injected into the `fprintf` call? Look at the end to the `##_VA_ARGS__` and that's telling the CPP to take the args entered where the `...` is, and inject them at that part of the `fprintf` call. You can then do things like this:

```
log_err("Age: %d, name: %s", age, name);
```

The arguments `age, name` are the `...` part of the definition, and those get injected into the `fprintf` output to become:

```
fprintf(stderr, "[ERROR] (%s:%d: errno: %s) Age %d:
name %d\n",
    __FILE__, __LINE__, clean_errno(), age, name);
```

See the `age, name` at the end? That's how `...` and `##_VA_ARGS__` work together, and it will work in macros that call other variable argument macros. Look at the `check` macro now and see it calls `log_err`, but `check` is *also* using the `...` and `##_VA_ARGS__` to do the call. That's how you can pass full `printf` style format strings to `check`, which go to `log_err`, and then make both work like `printf`.

Next thing to study is how `check` crafts the `if-statement` for the error checking. If we strip out the `log_err` usage we see this:

```
if(!(A)) { errno=0; goto error; }
```

Which means, if `A` is false, then clear `errno` and goto the error label. That has `check` macro being replaced with the `if-statement` so if we manually expanded out the macro `check(rc == 0, "There was an error.")` we'd get:

```
if(!(rc == 0)) {
    log_err("There was an error.");
    errno=0;
    goto error;
}
```

What you should be getting from this trip through these two macros is that the CPP replaces macros with the expanded version of their definition, but that it will do this *recursively*, expanding all the macros in macros. The CPP then is just a recursive templating system, as I mentioned before. Its power comes from its ability to generate whole blocks of parameterized code thus becoming a handy code generation tool.

That leaves one question: Why not just use a function like `die`? The reason is you want `file:line` numbers and the `goto` operation for an error handling exit. If you did this inside a function, you wouldn't get a line number for where the error actually happened, and the `goto` would be much more complicated.

Another reason is you still have to write the raw `if-statement`, which looks like all the other `if-statements` in your code, so it's not as clear that this one is an error check. By wrapping the `if-statement` in a macro called `check` you make it clear that this is just error checking, and not part of the main flow.

Finally, CPP has the ability to *conditionally compile* portions of code, so you can have code that's only present when you build a developer or

debug version of the program. You can see this already in the `dbg.h` file where the `debug` macro has a body only if it's asked for by the compiler. Without this ability, you'd need a wasted `if-statement` that checks for "debug mode", and then still wastes CPU doing that check for no value.

## Extra Credit

- Put `#define NDEBUG` at the top of the file and check that all the debug messages go away.
- Undo that line, and add `-DNDEBUG` to `CFLAGS` at the top of the `Makefile` then recompile to see the same thing.
- Modify the logging so that it include the function name as well as the `file:line`.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 21: Advanced Data Types And Flow Control

This exercise will be a complete compendium of the available C data types and flow control structures you can use. It will work as a reference to complete your knowledge, and won't have any code for you to enter. I'll have you memorize some of the information by creating flash cards so you can get the important concepts solid in your mind.

For this exercise to be useful, you should spend at least a week hammering the content and filling out all the element I have missing here. You'll be writing out what each one means, and then writing a program to confirm what you've researched.

### Available Data Types

`int`

Stores a regular integer, defaulting to 32 bits in size.

`double`

Holds a large floating point number.

`float`

Holds a smaller floating point number.

`char`

Holds a single 1 byte character.

`void`

Indicates "no type" and used to say a function returns nothing, or a pointer has no type as in `void *thing`.

`enum`

Enumerated types, work as integers, convert to integers, but give you symbolic names for sets. Some compilers will warn you when you don't cover all elements of an enum in `switch-statements`.

### Type Modifiers

### unsigned

Changes the type so that it does not have negative numbers, giving you a larger upper bound but nothing lower than 0.

### signed

Gives you negative and positive numbers, but halves your upper bound in exchange for the same lower bound negative.

### long

Uses a larger storage for the type so that it can hold bigger numbers, usually doubling the current size.

### short

Uses smaller storage for the type so it stores less, but takes half the space.

## Type Qualifiers

### const

Indicates the variable won't change after being initialized.

### volatile

Indicates that all bets are off, and the compiler should leave this alone and try not to do any fancy optimizations to it. You usually only need this if you're doing really weird stuff to your variables.

### register

Forces the compiler to keep this variable in a register, and the compiler can just ignore you. These days compilers are better at figuring out where to put variables, so only use this if you actually can measure it improving the speed.

## Type Conversion

C uses a sort of "stepped type promotion" mechanism, where it looks at two operands on either side of an expression, and promotes the smaller side to match the larger side before doing the operation. If one side of an expression is on this list, then the other side is converted to that type before the operation is done, and this goes in this order:

- long double
- double
- float
- int (but only `char` and `short int`);

- long

If you find yourself trying to figure out how your conversions are working in an expression, then don't leave it to the compiler. Use explicit casting operations to make it exactly what you want. For example, if you have:

```
long + char - int * double
```

Rather than trying to figure out if it will be converted to double correctly, just use casts:

```
(double)long - (double)char - (double)int * double
```

Putting the type you want in parenthesis before the variable name is how you force it into the type you really need. The important thing though is *always promote up, not down*. Don't cast `long` into `char` unless you know what you're doing.

## Type Sizes

The `stdint.h` defines both a set of `typedefs` for exact sized integer types, as well as a set of macros for the sizes of all the types. This is easier to work with than the older `limits.h` since it is consistent. The types defined are:

```
int8_t
```

8 bit signed integer.

```
uint8_t
```

8 bit unsigned integer.

```
int16_t
```

16 bit signed integer.

```
uint16_t
```

16 bit unsigned integer.

```
int32_t
```

32 bit signed integer.

```
uint32_t
```

32 bit unsigned integer.

```
int64_t
```

64 bit signed integer.

```
uint64_t
```

64 bit unsigned integer.

The pattern here is of the form (u)int(BITS)\_t where a *u* is put in front to indicate "unsigned", then *BITS* is a number for the number of bits. This pattern is then repeated for macros that return the maximum values of these types:

`INT(N)_MAX`

Maximum positive number of the signed integer of bits ( $N$ ), such as `INT16_MAX`.

`INT(N)_MIN`

Minimum negative number of signed integer of bits ( $N$ ).

`UINT(N)_MAX`

Maximum positive number of unsigned integer of bits ( $N$ ). Since it's unsigned the minimum is 0 and can't have a negative value.

## WARNING

Pay attention! Do not go looking for a literal `INT(N)_MAX` definition in any header file. I'm using the `(N)` as a placeholder for any number of bits your platform currently supports. This `(N)` could be any number, 8, 16, 32, 64, even maybe 128. I use this notation in this exercise so that I don't have to literally write out every possible combination.

There are also macros in `stdint.h` for sizes of the `size_t` type, integers large enough to hold pointers, and other handy size defining macros. Compilers have to at least have these, and then they can allow other larger types.

Here is a full list should be in `stdint.h`:

`int_least(N)_t`

holds at least ( $N$ ) bits.

`uint_least(N)_t`

holds at least ( $N$ ) bits unsigned.

`INT_LEAST(N)_MAX`

max value of the matching least ( $N$ ) type.

`INT_LEAST(N)_MIN`

min value of the matching least ( $N$ ) type.

`UINT_LEAST(N)_MAX`

unsigned maximum of the matching ( $N$ ) type.

`int_fast(N)_t`

similar to `int_least*N*_t` but asking for the "fastest" with at least that precision.

`uint_fast(N)_t`

unsigned fastest least integer.

`INT_FAST(N)_MAX`

max value of the matching fastest (*N*) type.

`INT_FAST(N)_MIN`

min value of the matching fastest (*N*) type.

`UINT_FAST(N)_MAX`

unsigned max value of the matching fastest (*N*) type.

`intptr_t`

a *signed* integer large enough to hold a pointer.

`uintptr_t`

an *unsigned* integer large enough to hold a pointer.

`INTPTR_MAX`

max value of a `intptr_t`.

`INTPTR_MIN`

min value of a `intptr_t`.

`UINTPTR_MAX`

unsigned max value of a `uintptr_t`.

`intmax_t`

biggest number possible on that system.

`uintmax_t`

biggest unsigned number possible.

`INTMAX_MAX`

largest value for the biggest signed number.

`INTMAX_MIN`

smallest value for the biggest signed number.

`UINTMAX_MAX`

largest value for the biggest unsigned number.

`PTRDIFF_MIN`

minimum value of `ptrdiff_t`.

`PTRDIFF_MAX`

maximum value of `ptrdiff_t`.

`SIZE_MAX`

maximum of a `size_t`.

## Available Operators

This is a comprehensive list of all the operators you have in the C language. In this list, I'm indicating the following:

(binary)

The operator has a left and right: `X + Y`.

(unary)

The operator is on its own: `-X`.

(prefix)

The operator comes before the variable: `++X`.

(postfix)

Usually the same as the `(prefix)` version, but placing it

after gives it a different meaning: `X++`.

(ternary)

There's only one of these, so it's actually called the

ternary but it means "three operands": `X ? Y : Z`.

## Math Operators

These are your basic math operations, plus I put `()` in with these since it calls a function and is close to a "math" operation.

`()`

Function call.

`*`

(binary)

multiply.

`/`

divide.

`+`

(binary)

addition.

`+`

(unary)

positive number.

`++`

(postfix)

read, then increment.

`++`

(prefix)

increment, then read.

`--`

(postfix)

read, then decrement.

`--`

(prefix)

decrement, then read.

`-`

(binary)

subtract.

`-`

(unary)

negative number.

## Data Operators

These are used to access data in different ways and forms.

`->`

struct pointer access.

`.`

struct value access.

`[]`

Array index.

`sizeof`

size of a type or variable.

`&` (unary)

Address of.

`*` (unary)

Value of.

## Logic Operators

These handle testing equality and inequality of variables.

`!=`

does not equal.

`<`

less than.

`<=`

less than or equal.

`==`

equal (not assignment).

`>`

greater than.

`>=`

greater than or equal.

## Bit Operators

These are more advanced and for shifting and modifying the raw bits in integers.

`&` (binary)

Bitwise and.

`<<`

Shift left.

`>>`

Shift right.

`^`

bitwise xor (exclusive or).

|

bitwise or.

~

compliment (flips all the bits).

## Boolean Operators

Used in truth testing. Study the ternary operator carefully, it is very handy.

!

not.

&&

and.

||

or.

?:

Ternary truth test, read `X ? Y : Z` as "if X then Y else Z".

## Assignment Operators

Compound assignment operators that assign a value, and/or perform an operation at the same time. Most of the above operations can also be combined into a compound assignment operator.

=

assign.

%=

modulus assign.

&=

bitwise and assign.

\*=

multiply assign.

+=

plus assign.

-=

minus assign.

/=

divide assign.

<<=

shift left, assign.

>>=

shift right, assign.



`^=`

bitwise xor, assign.

`|=`

bitwise or, assign.

## Available Control Structures

There's a few control structures you haven't encountered yet:

`do-while`

`do { ... } while(X);` First does the code in the block, then

tests the `X` expression before exiting.

`break`

Put this in a loop, and it breaks out ending it early.

`continue`

Stops the body of a loop and jumps to the test so it can continue.

`goto`

Jumps to a spot in the code where you've placed a `label:`, and

you've been using this in the `dbg.h` macros to go to the `error:` label.

## Extra Credit

- Read `stdint.h` or a description of it and write out all the possible available size identifiers.
- Go through each item here and write out what it does in code. Research it so you know you got it right by looking it up online.
- Get this information solid as well by making flash cards and spending 15 minutes a day memorizing it.
- Create a program that prints out examples of each type and confirm that your research is right.

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**





## Exercise 22: The Stack, Scope, And Globals

The concept of "scope" seems to confuse quite a few people when they first start programming. Originally it came from the use of the system stack (which we lightly covered earlier) and how it was used to store temporary variables. In this exercise, we'll learn about scope by learning about how a stack data structure works, and then feeding that concept back in to how modern C does scoping.

The real purpose of this exercise though is to learn where the hell things live in C. When someone doesn't grasp the concept of scope, it's almost always a failure in understanding where variables are created, exist, and die. Once you know where things are, the concept of scope becomes easier.

This exercise will require three files:

ex22.h

A header file that sets up some external variables and some functions.

ex22.c

Not your main like normal, but instead a source file that will become a object file `ex22.o` which will have some functions and variables in it defined from `ex22.h`.

ex22\_main.c

The actual `main` that will include the other two and demonstrate what they contain as well as other scope concepts.

### ex22.h and ex22.c

Your first step is to create your own header file named `ex22.h` which defines the functions and "extern" variables you need:

```
#ifndef _ex22_h
#define _ex22_h

// makes THE_SIZE in ex22.c available to other .c files
extern int THE_SIZE;

// gets and sets an internal static variable in ex22.c
int get_age();
void set_age(int age);

// updates a static variable that's inside update_ratio
double update_ratio(double ratio);

void print_size();

#endif
```

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@l2sthw

The important thing to see is the use of `extern int THE_SIZE`, which I'll explain after you also create the matching `ex22.c`:

```
#include <stdio.h>
#include "ex22.h"
#include "dbg.h"

int THE_SIZE = 1000;

static int THE_AGE = 37;

int get_age()
{
    return THE_AGE;
}

void set_age(int age)
{
    THE_AGE = age;
}

double update_ratio(double new_ratio)
{
    static double ratio = 1.0;

    double old_ratio = ratio;
    ratio = new_ratio;

    return old_ratio;
}

void print_size()
{
    log_info("I think size is: %d", THE_SIZE);
}
```

These two files introduce some new kinds of storage for variables:

#### `extern`

This keyword is a way to tell the compiler "the variable exists, but it's in another 'external' location". Typically this means that one .c file is going to use a variable that's been defined in another .c file. In this case, we're saying `ex22.c` has a variable `THE_SIZE` that will be accessed from `ex22_main.c`.

#### `static` (file)

This keyword is kind of the inverse of `extern` and says that the variable is only used in this .c file, and should not be available to other parts of the program. Keep in mind that `static` at the file level (as with `THE_AGE` here) is different than in other places.

#### `static` (function)

If you declare a variable in a function `static`, then that variable acts like a `static` defined in the file, but it's only accessible from that function. It's a way of creating constant state for a function, but in reality it's *rarely* used in modern C programming because they are hard to use with threads.

In these two files then, you have the following variables and functions that you should understand:

#### `THE_SIZE`

This is the variable you declared `extern` that you'll play with from `ex22_main.c`.

`get_age` and `set_age`

These are taking the static variable `THE_AGE`, but exposing it to other parts of the program through functions. You couldn't access `THE_AGE` directly, but these functions can.

`update_ratio`

This takes a new `ratio` value, and returns the old one. It uses a function level static variable `ratio` to keep track of what the ratio currently is.

`print_size`

Prints out what `ex22.c` thinks `THE_SIZE` is currently.

## ex22\_main.c

Once you have that file written, you can then make the main function which uses all of these and demonstrates some more scope conventions:

```
#include "ex22.h"
#include "dbg.h"

const char *MY_NAME = "Zed A. Shaw";

void scope_demo(int count)
{
    log_info("count is: %d", count);

    if(count > 10) {
        int count = 100; // BAD! BUGS!

        log_info("count in this scope is %d", count);
    }

    log_info("count is at exit: %d", count);

    count = 3000;

    log_info("count after assign: %d", count);
}

int main(int argc, char *argv[])
{
    // test out THE_AGE accessors
    log_info("My name: %s, age: %d", MY_NAME, get_age());

    set_age(100);

    log_info("My age is now: %d", get_age());

    // test out THE_SIZE extern
    log_info("THE_SIZE is: %d", THE_SIZE);
    print_size();

    THE_SIZE = 9;

    log_info("THE SIZE is now: %d", THE_SIZE);
    print_size();

    // test the ratio function static
    log_info("Ratio at first: %f", update_ratio(2.0));
    log_info("Ratio again: %f", update_ratio(10.0));
    log_info("Ratio once more: %f", update_ratio(300.0));

    // test the scope demo
    int count = 4;
    scope_demo(count);
}
```

```

scope_demo(count * 20);

log_info("count after calling scope_demo: %d", count);

return 0;
}

```

I'll break this file down line-by-line, and as I do you should find each variable I mention and where it lives.

ex22\_main.c:4

Making a `const` which stands for constant and is an alternative to using a `define` to create a constant variable.

ex22\_main.c:6

A simple function that demonstrates more scope issues in a function.

ex22\_main.c:8

Prints out the value of `count` as it is at the top of the function.

ex22\_main.c:10

An `if-statement` that starts a new *scope block*, and then has another `count` variable in it. This version of `count` is actually a whole new variable. It's kind of like the `if-statement` started a new "mini function".

ex22\_main.c:11

The `count` that is local to this block is actually different from the one in the function's parameter list. What happens as we continue.

ex22\_main.c:13

Prints it out so you can see it's actually 100 here, not what was passed to `scope_demo`.

ex22\_main.c:16

Now for the freaky part. You have `count` in two places: the parameters to this function, and in the `if-statement`. The `if-statement` created a new block, so the `count` on line 11 *does not impact the parameter with the same name*. This line prints it out and you'll see that it prints the value of the parameter, not 100.

ex22\_main.c:18-20

Then I set the parameter `count` to 3000 and print that out, which will demonstrate that you can change function parameters and they don't impact the caller's version of the variable.

Make sure you trace through this function, but don't think that you understand scope quite yet. Just start to realize that if you make a variable inside a block (as in `if-statements` or `while-loops`), then those variables are *new* variables that exist only in that block. This is crucial to understand, and is also a *source of many bugs*. We'll address why you shouldn't do this shortly.

The rest of the `ex22_main.c` then demonstrates all of these by manipulating and printing them out:

ex22\_main.c:26

Prints out the current values of `MY_NAME` and gets `THE_AGE` from `ex22.c` using the accessor function `get_age`.

ex22\_main.c:27-30

Uses `set_age` in `ex22.c` to change `THE_AGE` and then print it

out.

ex22\_main.c:33-39

Then I do the same thing to `THE_SIZE` from `ex22.c`, but this time I'm accessing it directly, and also demonstrating that it's actually changing in that file by printing it here and with `print_size`.

ex22\_main.c:42-44

Show how the static variable `ratio` inside `update_ratio` is maintained between function calls.

ex22\_main.c:46-51

Finally running `scope_demo` a few times so you can see the scope in action. Big thing to notice is that the local `count` variable remains unchanged. You *must* get that passing in a variable like this will not let you change it in the function. To do that you need our old friend the pointer. If you were to pass a pointer to this `count`, then the called function has the address of it and can change it.

That explains what's going on in all of these files, but you should trace through them and make sure you know where everything is as you study it.

## What You Should See

This time, instead of using your `Makefile`, I want you to build these two files manually so you can see how they are actually put together by the compiler. Here's what you should do and what you should see for output.

```
$ cc -Wall -g -DNDEBUG -c -o ex22.o ex22.c
$ cc -Wall -g -DNDEBUG ex22_main.c ex22.o -o ex22_main
$ ./ex22_main
[INFO] (ex22_main.c:26) My name: Zed A. Shaw, age: 37
[INFO] (ex22_main.c:30) My age is now: 100
[INFO] (ex22_main.c:33) THE_SIZE is: 1000
[INFO] (ex22.c:32) I think size is: 1000
[INFO] (ex22_main.c:38) THE SIZE is now: 9
[INFO] (ex22.c:32) I think size is: 9
[INFO] (ex22_main.c:42) Ratio at first: 1.000000
[INFO] (ex22_main.c:43) Ratio again: 2.000000
[INFO] (ex22_main.c:44) Ratio once more: 10.000000
[INFO] (ex22_main.c:8) count is: 4
[INFO] (ex22_main.c:16) count is at exit: 4
[INFO] (ex22_main.c:20) count after assign: 3000
[INFO] (ex22_main.c:8) count is: 80
[INFO] (ex22_main.c:13) count in this scope is 100
[INFO] (ex22_main.c:16) count is at exit: 80
[INFO] (ex22_main.c:20) count after assign: 3000
[INFO] (ex22_main.c:51) count after calling scope_demo: 4
```

Make sure you trace how each variable is changing and match it to the line that gets output. I'm using `log_info` from the `dbg.h` macros so you can get the exact line number where each variable is printed and find it in the files for tracing.

## Scope, Stack, And Bugs

If you've done this right you should now see many of the different ways you can place variables in your C code. You can use `extern` or access

functions like `get_age` to create globals. You can make new variables inside any blocks, and they'll retain their own values until that block exits, leaving the outer variables alone. You also can pass a value to a function, and change the parameter but not change the caller's version of it.

The most important thing to realize though is that all of this causes bugs. C's ability to place things in many places in your machine and then let you access it in those places means you get confused easily about where something lives. If you don't where it lives then there's a chance you'll not manage it properly.

With that in mind, here's some rules to follow when writing C code so you avoid bugs related to the stack:

- Do not "shadow" a variable like I've done here with `count` in `scope_demo`. It leaves you open to subtle and hidden bugs where you *think* you're changing a value and you actually aren't.
- Avoid too many globals, especially if across multiple files. If you have to then use accessor functions like I've done with `get_age`. This doesn't apply to constants, since those are read-only. I'm talking about variables like `THE_SIZE`. If you want people to modify or set this, then make accessor functions.
- When in doubt, put it on the heap. Don't rely on the semantics of the stack or specialized locations and instead just create things with `malloc`.
- Don't use function static variables like I did in `update_ratio`. They're rarely useful and end up being a huge pain when you need to make your code concurrent in threads. They are also hard as hell to find compared to a well done global variable.
- Avoid reusing function parameters as it's confusing whether you're just reusing it or if you think you're changing the *caller's* version of it.

As with all things, these rules can be broken when it's practical. In fact, I guarantee you'll run into code that breaks all of these rules and is perfectly fine. The constraints of different platforms makes it necessary sometimes.

## How To Break It

For this exercise, breaking the program involves trying to access or change things you can't:

- Try to directly access variables in `ex22.c` from `ex22_main.c` that you think you can't. For example, you can't get at `ratio` inside `update_ratio`? What if you had a pointer to it?
- Ditch the `extern` declaration in `ex22.h` to see what you get for errors or warnings.
- Add `static` or `const` specifiers to different variables and then try to change them.

## Extra Credit

- Research the concept of "pass by value" vs. "pass by reference". Write an example of both.



- Use pointers to gain access to things you shouldn't have access to.
- Use valgrind to see what this kind of access looks like when you do it wrong.
- Write a recursive function that causes a stack overflow. Don't know what a recursive function is? Try calling `scope_demo` at the bottom of `scope_demo` itself so that it loops.
- Rewrite the `Makefile` so that it can build this.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 23: Meet Duff's Device

This exercise is a brain teaser where I introduce you to one of the most famous hacks in C called "Duff's Device", named after Tom Duff the "inventor". This little slice of awesome (evil?) has nearly everything you've been learning wrapped in one tiny little package. Figuring out how it works is also a good fun puzzle.

### NOTE

Part of the fun of C is that you can come up with crazy hacks like this, but this is also what makes C annoying to use. It's good to learn about these tricks because it gives you a deeper understanding of the language and your computer. But, you should never use this. Always strive for easy to read code.

Duff's device was "discovered" (created?) by Tom Duff and is a trick with the C compiler that actually shouldn't work. I won't tell you what it does yet since this is meant to be a puzzle for you to ponder and try to solve. You are to get this code running and then try to figure out what it does, and *why* it does it this way.

```
#include <stdio.h>
#include <string.h>
#include "dbg.h"

int normal_copy(char *from, char *to, int count)
{
    int i = 0;

    for(i = 0; i < count; i++) {
        to[i] = from[i];
    }

    return i;
}

int duffs_device(char *from, char *to, int count)
{
    {
        int n = (count + 7) / 8;

        switch(count % 8) {
            case 0: do { *to++ = *from++;
                        case 7: *to++ = *from++;
                        case 6: *to++ = *from++;
                        case 5: *to++ = *from++;
                        case 4: *to++ = *from++;
                        case 3: *to++ = *from++;
```

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

🛠  
HELP

Follow  
@l2sthw

```

        case 2: *to++ = *from++;
        case 1: *to++ = *from++;
    } while(--n > 0);
    }
}

return count;
}

int zeds_device(char *from, char *to, int count)
{
    {
        int n = (count + 7) / 8;

        switch(count % 8) {
            case 0:
                again: *to++ = *from++;

                case 7: *to++ = *from++;
                case 6: *to++ = *from++;
                case 5: *to++ = *from++;
                case 4: *to++ = *from++;
                case 3: *to++ = *from++;
                case 2: *to++ = *from++;
                case 1: *to++ = *from++;
                    if(--n > 0) goto again;
            }
        }

        return count;
    }
}

int valid_copy(char *data, int count, char expects)
{
    int i = 0;
    for(i = 0; i < count; i++) {
        if(data[i] != expects) {
            log_err("[%d] %c != %c", i, data[i], expects);
            return 0;
        }
    }

    return 1;
}

int main(int argc, char *argv[])
{
    char from[1000] = {'a'};
    char to[1000] = {'c'};
    int rc = 0;

    // setup the from to have some stuff
    memset(from, 'x', 1000);
    // set it to a failure mode
    memset(to, 'y', 1000);
    check(valid_copy(to, 1000, 'y'), "Not initialized right.");

    // use normal copy to
    rc = normal_copy(from, to, 1000);
    check(rc == 1000, "Normal copy failed: %d", rc);
    check(valid_copy(to, 1000, 'x'), "Normal copy failed.");

    // reset
    memset(to, 'y', 1000);

    // duffs version
    rc = duffs_device(from, to, 1000);
    check(rc == 1000, "Duff's device failed: %d", rc);
    check(valid_copy(to, 1000, 'x'), "Duff's device failed copy.");
}

```

```

// reset
memset(to, 'y', 1000);

// my version
rc = zeds_device(from, to, 1000);
check(rc == 1000, "Zed's device failed: %d", rc);
check(valid_copy(to, 1000, 'x'), "Zed's device failed copy.");

return 0;
error:
return 1;
}

```

In this code I have three versions of a copy function:

`normal_copy`

Which is just a plain `for-loop` that copies characters from one array to another.

`duffs_device`

This is the brain teaser called "Duff's Device", named after Tom Duff, the person to blame for this delicious evil.

`zeds_device`

A version of "Duff's Device" that just uses a `goto` so you can get a clue about what's happening with the weird `do-while` placement in `duffs_device`.

Study these three functions before continuing. Try to explain what's going on to yourself before continuing.

## What You Should See

There's no output from this program, it just runs and exits. You should run it under `valgrind` and make sure there are no errors.

## Solving The Puzzle

The first thing to understand is that C is rather loose regarding some of its syntax. This is why you can put half of a `do-while` in one part of a `switch-statement`, then the other half somewhere else and it will still work. If you look at my version with the `goto again` it's actually more clear what's going on, but make sure you understand how that part works.

The second thing is how the default fallthrough semantics of `switch-statements` means you can jump to a particular case, and then it will just keep running until the end of the switch.

The final clue is the `count % 8` and the calculation of `n` at the top.

Now, to solve how these functions work, do the following:

- Print this code out so you can write on some paper.
- On a piece of paper, write each of the variables in a table as they are when they get initialized right before the `switch-statement`.
- Follow the logic to the switch, then do the jump to the right case.

- Update the variables, including the `to`, `from`, and the arrays they point at.
- When you get to the `while` part or my `goto` alternative, check your variables and then follow the logic either back to the top of the `do-while` or to where the `again` label is located.
- Follow through this manual tracing, updating the variables, until you are sure you see how this flows.

## Why Bother?

When you've figured out how it actually works, the final question is: Why would you ever want to do this? The purpose of this trick is to manually do "loop unrolling". Large long loops can be slow, so one way to speed them up is to find some fixed chunk of the loop, and then just duplicate the code in the loop out that many times sequentially. For example, if you know a loop runs a minimum of 20 times, then you can put the contents of the loop 20 times in the source code.

Duff's device is basically doing this automatically by chunking up the loop into 8 iteration chunks. It's clever and actually works, but these days a good compiler will do this for you. You shouldn't need this except in the rare case where you have *proven* it would improve your speed.

## Extra Credit

- Never use this again.
- Go look at the Wikipedia entry for "Duff's Device" and see if you can spot the error. Compare it to the version I have here and read the article carefully to try to understand why the Wikipedia code won't work for you but worked for Tom Duff.
- Create a set of macros that lets you create any length device like this. For example, what if you wanted to have 32 case statements and didn't want to write out all of them? Can you do a macro that lays down 8 at a time?
- Change the `main` to conduct some speed tests to see which one is really the fastest.
- Read about `memcpy`, `memmove`, `memset`, and also compare their speed.
- Never use this again!

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 24: Input, Output, Files

You've been using `printf` to print things, and that's great and all, but you need more. In this exercise program you're using the functions `fscanf` and `fgets` to build information about a person in a structure. After this simple introduction to reading input, you'll get a full list of the functions that C has for I/O. Some of these you've already seen and used, so this will be another memorization exercise.

```
#include <stdio.h>
#include "dbg.h"

#define MAX_DATA 100

typedef enum EyeColor {
    BLUE_EYES, GREEN_EYES, BROWN_EYES,
    BLACK_EYES, OTHER_EYES
} EyeColor;

const char *EYE_COLOR_NAMES[] = {
    "Blue", "Green", "Brown", "Black", "Other"
};

typedef struct Person {
    int age;
    char first_name[MAX_DATA];
    char last_name[MAX_DATA];
    EyeColor eyes;
    float income;
} Person;

int main(int argc, char *argv[])
{
    Person you = {.age = 0};
    int i = 0;
    char *in = NULL;

    printf("What's your First Name? ");
    in = fgets(you.first_name, MAX_DATA-1, stdin);
    check(in != NULL, "Failed to read first name.");

    printf("What's your Last Name? ");
    in = fgets(you.last_name, MAX_DATA-1, stdin);
    check(in != NULL, "Failed to read last name.");

    printf("How old are you? ");
    int rc = fscanf(stdin, "%d", &you.age);
    check(rc > 0, "You have to enter a number.");

    printf("What color are your eyes:\n");
    for(i = 0; i <= OTHER_EYES; i++) {
        printf("%d) %s\n", i+1, EYE_COLOR_NAMES[i]);
    }
    printf("> ");

    int eyes = -1;
    rc = fscanf(stdin, "%d", &eyes);
    check(rc > 0, "You have to enter a number.");
```

1  
2  
3  
4 MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssthw

```

    you.eyes = eyes - 1;
    check(you.eyes <= OTHER_EYES && you.eyes >= 0, "Do it right, that's not an
option.");

    printf("How much do you make an hour? ");
    rc = fscanf(stdin, "%f", &you.income);
    check(rc > 0, "Enter a floating point number.");

    printf("----- RESULTS -----\\n");

    printf("First Name: %s", you.first_name);
    printf("Last Name: %s", you.last_name);
    printf("Age: %d\\n", you.age);
    printf("Eyes: %s\\n", EYE_COLOR_NAMES[you.eyes]);
    printf("Income: %f\\n", you.income);

    return 0;
error:

    return -1;
}

```

This program is deceptively simple, and introduces a function called `fscanf` which is the "file scanf". The `scanf` family of functions are the inverse of the `printf` versions. Where `printf` printed out data based on a format, `scanf` reads (or scans) input based on a format.

There's nothing original in the beginning of the file, so here's what the `main` is doing:

ex24.c:24-28

Set up some variables we'll need.

ex24.c:30-32

Get your first name using the `fgets` function, which reads a string from the input (in this case `stdin`) but makes sure it doesn't overflow the given buffer.

ex24.c:34-36

Same thing for `you.last_name`, again using `fgets`.

ex24.c:38-39

Uses `fscanf` to read an integer from `stdin` and put it into `you.age`. You can see that the same format string is used as `printf` to print an integer. You should also see that you have to give the address of `you.age` so that `fscanf` has a pointer to it and can modify it. This is a good example of using a pointer to a piece of data as an "out parameter".

ex24.c:41-45

Print out all the options available for eye color, with a matching number that works with the `EyeColor` enum above.

ex24.c:47-50

Using `fscanf` again, get a number for the `you.eyes`, but make sure the input is valid. This is important because someone can enter a value outside the `EYE_COLOR_NAMES` array and cause a segfault.

ex24.c:52-53

Get how much you make as a `float` for the `you.income`.

ex24.c:55-61

Print everything out so you can see if you have it right. Notice that `EYE_COLOR_NAMES` is used to print out what the `EyeColor` enum is actually called.

# What You Should See

When you run this program you should see your inputs being properly converted. Make sure you try to give it bogus input too so you can see how it protects against the input.

```
$ make ex24
cc -Wall -g -DNDEBUG    ex24.c    -o ex24
$ ./ex24
What's your First Name? Zed
What's your Last Name? Shaw
How old are you? 37
What color are your eyes:
1) Blue
2) Green
3) Brown
4) Black
5) Other
> 1
How much do you make an hour? 1.2345
----- RESULTS -----
First Name: Zed
Last Name: Shaw
Age: 37
Eyes: Blue
Income: 1.234500
```

## How To Break It

This is all fine and good, but the real important part of this exercise is how `scanf` actually sucks. It's fine for simple conversion of numbers, but fails for strings because it's difficult to tell `scanf` how big a buffer is before you read. There's also a problem with a function like `gets` (not `fgets`, the non-f version) which we avoided. That function has no idea how big the input buffer is at all and will just trash your program.

To demonstrate the problems with `fscanf` and strings, change the lines that use `fgets` so they are `fscanf(stdin, "%50s", you.first_name)` and then try to use it again. Notice it seems to read too much and then eat your enter key? This doesn't do what you think it does, and really rather than deal with weird `scanf` issues, just use `fgets`.

Next, change the `fgets` to use `gets`, then bust out your `valgrind` and do this: `valgrind ./ex24 < /dev/urandom` to feed random garbage into your program. This is called "fuzzing" your program, and it is a good way to find input bugs. In this case, you're feeding garbage from the `/dev/urandom` file, and then watching it crash. On some platforms you may have to do this a few times, or even adjust the `MAX_DATA` define so it's small enough.

The `gets` function is so bad that some platforms actually warn you when the *program* runs that you're using `gets`. You should never use this function, ever.

Finally, take the input for `you.eyes` and remove the check that the number given is within the right range. Then feed it bad numbers like -1 or 1000. Do this under Valgrind too so you can see what happens.



# The I/O Functions

This is a short list of various I/O functions that you should look up and create index cards that have the function name, what it does, and all the variants similar to it.

- `fscanf`
- `fgets`
- `fopen`
- `freopen`
- `fdopen`
- `fclose`
- `fcloseall`
- `fgetpos`
- `fseek`
- `ftell`
- `rewind`
- `fprintf`
- `fwrite`
- `fread`

Go through these and memorize the different variants and what they do. For example, for the card on `fscanf` you'll have `scanf`, `sscanf`, `vscanf`, etc. and then what each of those do on the back.

Finally, to get the information you need for these cards, use `man` to read the help for it. For example, the page for `fscanf` comes from `man fscanf`.

## Extra Credit

- Rewrite this to not use `fscanf` at all. You'll need to use functions like `atoi` to convert the input strings to numbers.
- Change this to use plain `scanf` instead of `fscanf` to see what the difference is.
- Fix it so that the input names get stripped of the trailing newline characters and any whitespace.
- Use `scanf` to write a function that reads a character at a time and files in the names but doesn't go past the end. Make this function generic so it can take a size for the string, and make sure you end the string with `'\0'` no matter what.

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 25: Variable Argument Functions

In C you can create your own versions of functions like `printf` and `scanf` by creating a "variable argument function". These functions use the header `stdarg.h` and with them you can create nicer interfaces to your library. They are handy for certain types of "builder" functions, formatting functions, and anything that takes variable arguments.

Understanding "vararg functions" is *not* essential to creating C programs. I think I've used it maybe a 20 times in my code in the years I've been programming. However, knowing how a vararg function works will help you debug the ones you use and gives you more understanding of the computer.

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

🏠  
HELP

Follow  
@l2sthw

```
/** WARNING: This code is fresh and potentially isn't correct yet. */
```

```
#include <stdlib.h>
#include <stdio.h>
#include <stdarg.h>
#include "dbg.h"
```

```
#define MAX_DATA 100
```

```
int read_string(char **out_string, int max_buffer)
{
    *out_string = calloc(1, max_buffer + 1);
    check_mem(*out_string);

    char *result = fgets(*out_string, max_buffer, stdin);
    check(result != NULL, "Input error.");

    return 0;
}
```

```
error:
    if(*out_string) free(*out_string);
    *out_string = NULL;
    return -1;
}
```

```
int read_int(int *out_int)
{
    char *input = NULL;
    int rc = read_string(&input, MAX_DATA);
    check(rc == 0, "Failed to read number.");

    *out_int = atoi(input);

    free(input);
    return 0;
}
```

```
error:
    if(input) free(input);
    return -1;
}
```

```
int read_scan(const char *fmt, ...)
```

```

int i = 0;
int rc = 0;
int *out_int = NULL;
char *out_char = NULL;
char **out_string = NULL;
int max_buffer = 0;

va_list argp;
va_start(argp, fmt);

for(i = 0; fmt[i] != '\0'; i++) {
    if(fmt[i] == '%') {
        i++;
        switch(fmt[i]) {
            case '\0':
                sentinel("Invalid format, you ended with %%.");
                break;

            case 'd':
                out_int = va_arg(argp, int *);
                rc = read_int(out_int);
                check(rc == 0, "Failed to read int.");
                break;

            case 'c':
                out_char = va_arg(argp, char *);
                *out_char = fgetc(stdin);
                break;

            case 's':
                max_buffer = va_arg(argp, int);
                out_string = va_arg(argp, char **);
                rc = read_string(out_string, max_buffer);
                check(rc == 0, "Failed to read string.");
                break;

            default:
                sentinel("Invalid format.");
        }
    } else {
        fgetc(stdin);
    }

    check(!feof(stdin) && !ferror(stdin), "Input error.");
}

va_end(argp);
return 0;

error:
    va_end(argp);
    return -1;
}

int main(int argc, char *argv[])
{
    char *first_name = NULL;
    char initial = ' ';
    char *last_name = NULL;
    int age = 0;

    printf("What's your first name? ");
    int rc = read_scan("%s", MAX_DATA, &first_name);
    check(rc == 0, "Failed first name.");

    printf("What's your initial? ");
    rc = read_scan("%c\n", &initial);
    check(rc == 0, "Failed initial.");
}

```

```

printf("What's your last name? ");
rc = read_scan("%s", MAX_DATA, &last_name);
check(rc == 0, "Failed last name.");

printf("How old are you? ");
rc = read_scan("%d", &age);

printf("---- RESULTS ----\n");
printf("First Name: %s", first_name);
printf("Initial: '%c'\n", initial);
printf("Last Name: %s", last_name);
printf("Age: %d\n", age);

free(first_name);
free(last_name);
return 0;
error:
return -1;
}

```

This program is similar to the previous exercise, except I have written my own `scanf` style function that handles strings the way I want. The main function should be clear to you, as well as the two functions `read_string` and `read_int` since they do nothing new.

The varargs function is called `read_scan` and it does the same thing that `scanf` is doing using the `va_list` data structure and it's supporting macros and functions. Here's how it works:

- I set as the last parameter of the function the keyword `...` which indicates to C that this function will take any number of arguments after the `fmt` argument. I could put many other arguments before this, but I can't put anymore after this.
- After setting up some variables, I create a `va_list` variable and initialize it with `va_start`. This configures the gear in `stdarg.h` that handles variable arguments.
- I then use a `for-loop` to loop through the format string `fmt` and process the same kind of formats that `scanf` has, but much simpler. I just have integers, characters, and strings.
- When I hit a format, I use the `switch-statement` to figure out what to do.
- Now, to get a variable from the `va_list argp` I use the macro `va_arg(argp, TYPE)` where TYPE is the exact type of what I will assign this function parameter to. The downside to this design is you're flying blind, so if you don't have enough parameters then oh well, you'll most likely crash.
- The interesting difference from `scanf` is I'm assuming that people want `read_scan` to create the strings it reads when it hits a `'s'` format sequence. When you give this sequence, the function takes two parameters off the `va_list argp` stack: the max function size to read, and the output character string pointer. Using that information it just runs `read_string` to do the real work.
- This makes `read_scan` more consistent than `scanf` since you *always* give an address-of `&` on variables to have them set appropriately.
- Finally, if it encounters a character that's not in the format, it just reads one char to skip it. It doesn't care what that char is, just that it should skip it.

# What You Should See

When you run this one it's similar to the last one:

```
$ make ex25
cc -Wall -g -DNDEBUG    ex25.c    -o ex25
$ ./ex25
What's your first name? Zed
What's your initial? A
What's your last name? Shaw
How old are you? 37
---- RESULTS ----
First Name: Zed
Initial: 'A'
Last Name: Shaw
Age: 37
```

## How To Break It

This program should be more robust against buffer overflows, but it doesn't handle the formatted input as well as `scanf`. To try breaking this, change the code that you forget to pass in the initial size for '%s' formats. Try also giving it more data than `MAX_DATA`, and then see how not using `calloc` in `read_string` changes how it works. Finally, there's a problem that `fgets` eats the newlines, so try to fix that using `fgetc` but leave out the `\0` that ends the string.

## Extra Credit

- Make double and triple sure that you know what each of the `out_` variables are doing. Most important is `out_string` and how it's a pointer to a pointer, so getting when you're setting the pointer vs. the contents is important. Break down each of the
- Write a similar function to `printf` that uses the varargs system and rewrite `main` to use it.
- As usual, read the man page on all of this so you know what it does on your platform. Some platforms will use macros and others use functions, and some have these do nothing. It all depends on the compiler and the platform you use.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 26: Write A First Real Program

You are at the half-way mark in the book, so you need to take a mid-term. In this mid-term you're going to recreate a piece of software I wrote specifically for this book called `devpkg`. You'll then extend it in a few key ways and improve the code, most importantly by writing some unit tests for it.

### NOTE

I wrote this exercise before writing some of the exercises you might need to complete this. If you are attempting this one now, please keep in mind that the software may have bugs, that you might have problems because of my mistakes, and that you might not know everything you need to finish it. If so, tell me at [help@learncodethehardway.org](mailto:help@learncodethehardway.org) and then wait until I finish the other exercises.

## What Is `devpkg`?

`Devpkg` is a simple C program that installs other software. I made it specifically for this book as a way to teach you how a real software project is structured, and also how to reuse other people's libraries. It uses a portability library called [The Apache Portable Runtime \(APR\)](#) that has many handy C functions which work on tons of platforms, including Windows. Other than that, it just grabs code from the internet (or local files) and does the usual `./configure ; make ; make install` every programmer does.

Your goal in this exercise is to build `devpkg` from source, finish each *Challenge* I give, and use the source to understand what `devpkg` does and why.

## What We Want To Make

We want a tool that has three commands:

```
devpkg -S
    Sets up a new install on a computer.
devpkg -I
```

1 \_\_\_\_\_  
2 \_\_\_\_\_  
3 \_\_\_\_\_  
4 \_\_\_\_\_  
MAIN

◀  
PREVIOUS

▶ NEXT

⚙️ HELP

Follow  
@l2sthw

- Installs a piece of software from a URL.
- `devpkg -L`
  - Lists all the software that's been installed.
- `devpkg -F`
  - Fetches some source code for manual building.
- `devpkg -B`
  - Builds fetches source code and installs it, even if already installed.

We want `devpkg` to be able to take almost any URL, figure out what kind of project it is, download it, install it, and register that it downloaded that software. We'd also like it to process a simple dependency list so it can install all the software that a project might need as well.

## The Design

To accomplish this goal `devpkg` will have a very simple design:

Use external commands

You'll do most of the work through external commands like `curl`, `git`, and `tar`. This reduces the amount of code `devpkg` needs to get things done.

Simple File Database

You could easily make it more complex, but to start you'll just make a single simple file database at `/usr/local/.devpkg/db` to keep track of what's installed.

`/usr/local` Always

Again you could make this more advanced, but for starters just assume it's always `/usr/local` which is a standard install path for most software on Unix.

configure, make, make install

It's assumed that most software can install with just a `configure; make; make install` and maybe `configure` is optional. If the software can't at a minimum do that, then there's some options to modify the commands, but otherwise `devpkg` won't bother.

The User Can Be root

We'll assume the user can become root using `sudo`, but that they don't want to become root until the end.

This will keep our program small at first and work well enough to get it going, at which point you'll be able to modify it further for this exercise.

## The Apache Portable Runtime

One more thing you'll do is leverage the [The Apache Portable Runtime \(APR\)](#) libraries to get a good set of portable routines for doing this kind of work. The APR isn't necessary, and you could probably write this program without them, but it'd take more code than necessary. I'm also forcing you to use APR now so you get used to linking and using other libraries. Finally, the APR also works on *Windows* so your skills with it are transferable to many other platforms.

You should go get both the `apr-1.4.5` and the `apr-util-1.3` libraries, as well as browse through the documentation available at the [main APR](#)

[site at apr.apache.org](http://site.at.apr.apache.org)

Here's a shell script that will install all the stuff you need. You should write this into a file by hand, and then run it until it can install APR without any errors.

```
set -e

# go somewhere safe
cd /tmp

# get the source to base APR 1.4.6
curl -L -O http://archive.apache.org/dist/apr/apr-1.4.6.tar.gz

# extract it and go into the source
tar -xzvf apr-1.4.6.tar.gz
cd apr-1.4.6

# configure, make, make install
./configure
make
sudo make install

# reset and cleanup
cd /tmp
rm -rf apr-1.4.6 apr-1.4.6.tar.gz

# do the same with apr-util
curl -L -O http://archive.apache.org/dist/apr/apr-util-1.4.1.tar.gz

# extract
tar -xzvf apr-util-1.4.1.tar.gz
cd apr-util-1.4.1

# configure, make, make install
./configure --with-apr=/usr/local/apr
# you need that extra parameter to configure because
# apr-util can't really find it because...who knows.

make
sudo make install

#cleanup
cd /tmp
rm -rf apr-util-1.4.1* apr-1.4.6*
```

I'm having you write this script out because this is basically what we want `devpkg` to do, but with extra options and checks. In fact, you could just do it all in shell with less code, but then that wouldn't be a very good program for a C book would it?

Simply run this script and fix it until it works, then you'll have the libraries you need to complete the rest of this project.

## Project Layout

You need to setup some simple project files to get started. Here's how I usually craft a new project:

```
mkdir devpkg
cd devpkg
touch README Makefile
```

## Other Dependencies



You should have already installed APR and APR-util, so now you need a few more files as basic dependencies:

- `dbg.h` from Exercise 20.
- `bstrlib.h` and `bstrlib.c` from <http://bstring.sourceforge.net/>. Download the .zip file, extract it, and copy just those two files out.
- Type `make bstrlib.o` and if it doesn't work, read the "Fixing bstring" instructions below.

## NOTE

In some platforms the `bstring.c` file will have an error like:

```
bstrlib.c:2762: error:
expected declaration
specifiers or '...' before
numeric constant
```

This is from a bad define the authors added which doesn't work always. You just need to change the line 2759 that reads `#ifdef __GNUC__` and make it:

```
#if defined(__GNUC__) &&
!defined(__APPLE__)
```

Then it should work on Apple Mac OSX.

When that's all done, you should have a `Makefile`, `README`, `dbg.h`, `bstrlib.h`, and `bstrlib.c` ready to go.

## The Makefile

A good place to start is the `Makefile` since this lays out how things are built and what source files you'll be creating.

```
PREFIX?=/usr/local
CFLAGS=-g -Wall -I${PREFIX}/apr/include/apr-1 -I${PREFIX}/apr/include/apr-util-1
LDFLAGS=-L${PREFIX}/apr/lib -lapr-1 -pthread -laprutil-1

all: devpkg

devpkg: bstrlib.o db.o shell.o commands.o

install: all
    install -d $(DESTDIR)/$(PREFIX)/bin/
    install devpkg $(DESTDIR)/$(PREFIX)/bin/

clean:
    rm -f *.o
    rm -f devpkg
    rm -rf *.dSYM
```

There's nothing in this that you haven't seen before, except maybe the strange `?=` syntax, which says "set PREFIX equal to this unless PREFIX is already set".

## NOTE

If you are on more recent versions of Ubuntu and you get errors about `apr_off_t` or

`off64_t` then add `-D_LARGEFILE64_SOURCE=1` to `CFLAGS`.

Another thing is you need to add

`/usr/local/apr/lib` to a file in `/etc/ld.conf.so.d/` then run `ldconfig` so that it picks up the libraries correctly.

## The Source Files

From the make file, we see that there's four dependencies for `devpkg` which are:

`bstrlib.o`

Comes from `bstrlib.c` and header file `bstrlib.h` which you already have.

`db.o`

From `db.c` and header file `db.h`, and it will contain code for our little "database" routines.

`shell.o`

From `shell.c` and header `shell.h`, with a couple functions that make running other commands like `curl` easier.

`commands.o`

From `command.c` and header `command.h`, and contains all the commands that `devpkg` needs to be useful.

`devpkg`

It's not explicitly mentioned, but instead is the target (on the left) in this part of the `Makefile`. It comes from `devpkg.c` which contains the `main` function for the whole program.

Your job is to now create each of these files and type in their code and get them correct.

## NOTE

You may read this description and think, "Man! How is it that Zed is so smart he just sat down and typed these files out like this!? I could never do that." I didn't magically craft `devpkg` in this form with my awesome code powers. Instead, what I did is this:

- I wrote a quick little README to get an idea of how I wanted it to work.
- I created a simple bash script (like the one you did) to figure out all the pieces that you need.
- I made one .c file and hacked on it for a few days working through the idea and figuring it out.
- I got it mostly working and debugged, *then* I started breaking up the one big file into these four files.
- After getting these files laid down, I renamed and refined the functions and data structures so they'd be more logical and "pretty".
- Finally, after I had it working the *exact same* but with the new structure, I added a few features like the `-F` and `-B` options.

You're reading this in the order I want to teach it to you, but don't think this is how I always build software. Sometimes I already know the subject and I use more planning. Sometimes I just hack up an idea and see how well it'd work. Sometimes I write one, then throw it away and plan out a better one. It all depends on what my experience tells me is best, or where my inspiration takes me.

If you run into an "expert" who tries to tell you there's only one way to solve a programming problem, then they're lying to you. Either they actually use multiple tactics, or they're not very good.

## The DB Functions

There must be a way to record URLs that have been installed, list these URLs, and check if something has already been installed so we can skip it. I'll use a simple flat file database and the `bstrlib.h` library to do it.

First, create the `db.h` header file so you know what you'll be implementing.

```
#ifndef _db_h
#define _db_h

#define DB_FILE "/usr/local/.devpkg/db"
#define DB_DIR "/usr/local/.devpkg"

int DB_init();
int DB_list();
int DB_update(const char *url);
int DB_find(const char *url);
```

```
#endif
```

Then implement those functions in `db.c`, as you build this, use `make` like you've been to get it to compile cleanly.

```
#include <unistd.h>
#include <apr_errno.h>
#include <apr_file_io.h>

#include "db.h"
#include "bstrlib.h"
#include "dbg.h"

static FILE *DB_open(const char *path, const char *mode)
{
    return fopen(path, mode);
}

static void DB_close(FILE *db)
{
    fclose(db);
}

static bstring DB_load()
{
    FILE *db = NULL;
    bstring data = NULL;

    db = DB_open(DB_FILE, "r");
    check(db, "Failed to open database: %s", DB_FILE);

    data = bread((bNread)fread, db);
    check(data, "Failed to read from db file: %s", DB_FILE);

    DB_close(db);
    return data;
}

error:
    if(db) DB_close(db);
    if(data) bdestroy(data);
    return NULL;
}

int DB_update(const char *url)
{
    if(DB_find(url)) {
        log_info("Already recorded as installed: %s", url);
    }

    FILE *db = DB_open(DB_FILE, "a+");
    check(db, "Failed to open DB file: %s", DB_FILE);

    bstring line = bfromcstr(url);
    bconchar(line, '\n');
    int rc = fwrite(line->data, blength(line), 1, db);
    check(rc == 1, "Failed to append to the db.");

    return 0;
}

error:
    if(db) DB_close(db);
    return -1;
}

int DB_find(const char *url)
{

```

```

bstring data = NULL;
bstring line = bfromcstr(url);
int res = -1;

data = DB_load();
check(data, "Failed to load: %s", DB_FILE);

if(binstr(data, 0, line) == BSTR_ERR) {
    res = 0;
} else {
    res = 1;
}

error: // fallthrough
if(data) bdestroy(data);
if(line) bdestroy(line);

return res;
}

int DB_init()
{
    apr_pool_t *p = NULL;
    apr_pool_initialize();
    apr_pool_create(&p, NULL);

    if(access(DB_DIR, W_OK | X_OK) == -1) {
        apr_status_t rc = apr_dir_make_recursive(DB_DIR,
            APR_UREAD | APR_UWRITE | APR_UEXECUTE |
            APR_GREAD | APR_GWRITE | APR_GEXECUTE, p);
        check(rc == APR_SUCCESS, "Failed to make database dir: %s", DB_DIR);
    }

    if(access(DB_FILE, W_OK) == -1) {
        FILE *db = DB_open(DB_FILE, "w");
        check(db, "Cannot open database: %s", DB_FILE);
        DB_close(db);
    }

    apr_pool_destroy(p);
    return 0;

error:
    apr_pool_destroy(p);
    return -1;
}

int DB_list()
{
    bstring data = DB_load();
    check(data, "Failed to read load: %s", DB_FILE);

    printf("%s", bdata(data));
    bdestroy(data);
    return 0;

error:
    return -1;
}

```

## Challenge 1: Code Review

Before continuing, read every line of these files carefully and confirm that you have them entered in *exactly*. Read them line-by-line backwards to practice that. Also trace each function call and make sure you are using `check` to validate the return codes. Finally, look up *every* function

that you don't recognize either on the APR web site documentation, or in the `bstrlib.h` and `bstrlib.c` source.

## The Shell Functions

A key design decision for `devpkg` is to do most of the work using external tools like `curl`, `tar`, and `git`. We could find libraries to do all of this internally, but it's pointless if we just need the base features of these programs. There is no shame in running another command in Unix.

To do this I'm going to use the `apr_thread_proc.h` functions to run programs, but I also want to make a simple kind of "template" system. I'll use a `struct Shell` that holds all the information needed to run a program, but has "holes" in the arguments list where I can replace them with values.

Look at the `shell.h` file to see the structure and the commands I'll use. You can see I'm using `extern` to indicate that other `.c` files can access variables I'm defining in `shell.c`.

```
#ifndef _shell_h
#define _shell_h

#define MAX_COMMAND_ARGS 100

#include <apr_thread_proc.h>

typedef struct Shell {
    const char *dir;
    const char *exe;

    apr_procattr_t *attr;
    apr_proc_t proc;
    apr_exit_why_e exit_why;
    int exit_code;

    const char *args[MAX_COMMAND_ARGS];
} Shell;

int Shell_run(apr_pool_t *p, Shell *cmd);
int Shell_exec(Shell cmd, ...);

extern Shell CLEANUP_SH;
extern Shell GIT_SH;
extern Shell TAR_SH;
extern Shell CURL_SH;
extern Shell CONFIGURE_SH;
extern Shell MAKE_SH;
extern Shell INSTALL_SH;

#endif
```

Make sure you've created `shell.h` exactly, and that you've got the same names and number of `extern Shell` variables. Those are used by the `Shell_run` and `Shell_exec` functions to run commands. I define these two functions, and create the real variables in `shell.c`.

```
#include "shell.h"
#include "dbg.h"
#include <stdarg.h>

int Shell_exec(Shell template, ...)
{
    apr_pool_t *p = NULL;
```

```

int rc = -1;
apr_status_t rv = APR_SUCCESS;
va_list argp;
const char *key = NULL;
const char *arg = NULL;
int i = 0;

rv = apr_pool_create(&p, NULL);
check(rv == APR_SUCCESS, "Failed to create pool.");

va_start(argp, template);

for(key = va_arg(argp, const char *);
    key != NULL;
    key = va_arg(argp, const char *))
{
    arg = va_arg(argp, const char *);

    for(i = 0; template.args[i] != NULL; i++) {
        if(strcmp(template.args[i], key) == 0) {
            template.args[i] = arg;
            break; // found it
        }
    }
}

rc = Shell_run(p, &template);
apr_pool_destroy(p);
va_end(argp);
return rc;
}

error:
if(p) {
    apr_pool_destroy(p);
}
return rc;
}

int Shell_run(apr_pool_t *p, Shell *cmd)
{
    apr_procattr_t *attr;
    apr_status_t rv;
    apr_proc_t newproc;

    rv = apr_procattr_create(&attr, p);
    check(rv == APR_SUCCESS, "Failed to create proc attr.");

    rv = apr_procattr_io_set(attr, APR_NO_PIPE, APR_NO_PIPE,
        APR_NO_PIPE);
    check(rv == APR_SUCCESS, "Failed to set IO of command.");

    rv = apr_procattr_dir_set(attr, cmd->dir);
    check(rv == APR_SUCCESS, "Failed to set root to %s", cmd->dir);

    rv = apr_procattr_cmdtype_set(attr, APR_PROGRAM_PATH);
    check(rv == APR_SUCCESS, "Failed to set cmd type.");

    rv = apr_proc_create(&newproc, cmd->exe, cmd->args, NULL, attr, p);
    check(rv == APR_SUCCESS, "Failed to run command.");

    rv = apr_proc_wait(&newproc, &cmd->exit_code, &cmd->exit_why, APR_WAIT);
    check(rv == APR_CHILD_DONE, "Failed to wait.");

    check(cmd->exit_code == 0, "%s exited badly.", cmd->exe);
    check(cmd->exit_why == APR_PROC_EXIT, "%s was killed or crashed", cmd->exe);
};

return 0;
}

error:

```

```

        return -1;
    }

    Shell CLEANUP_SH = {
        .exe = "rm",
        .dir = "/tmp",
        .args = {"rm", "-rf", "/tmp/pkg-build", "/tmp/pkg-src.tar.gz",
                "/tmp/pkg-src.tar.bz2", "/tmp/DEPENDS", NULL}
    };

    Shell GIT_SH = {
        .dir = "/tmp",
        .exe = "git",
        .args = {"git", "clone", "URL", "pkg-build", NULL}
    };

    Shell TAR_SH = {
        .dir = "/tmp/pkg-build",
        .exe = "tar",
        .args = {"tar", "-xzf", "FILE", "--strip-components", "1", NULL}
    };

    Shell CURL_SH = {
        .dir = "/tmp",
        .exe = "curl",
        .args = {"curl", "-L", "-o", "TARGET", "URL", NULL}
    };

    Shell CONFIGURE_SH = {
        .exe = "./configure",
        .dir = "/tmp/pkg-build",
        .args = {"configure", "OPTS", NULL},
    };

    Shell MAKE_SH = {
        .exe = "make",
        .dir = "/tmp/pkg-build",
        .args = {"make", "OPTS", NULL}
    };

    Shell INSTALL_SH = {
        .exe = "sudo",
        .dir = "/tmp/pkg-build",
        .args = {"sudo", "make", "TARGET", NULL}
    };
};

```

Read the `shell.c` from the bottom to the top (which is a common C source layout) and you see I've created the actual `Shell` variables that you indicated were `extern` in `shell.h`. They live here, but are available to the rest of the program. This is how you make global variables that live in one `.o` file but are used everywhere. You shouldn't make many of these, but they are handy for things like this.

Continuing up the file we get to the `Shell_run` function, which is a "base" function that just runs a command based on what's in a `Shell` struct. It uses many of the functions defined in `apr_thread_proc.h` so go look up each one to see how it works. This seems like a lot of work compared to just using the `system` function call, but this also gives you more control over the other program's execution. For example, in our `Shell` struct we have a `.dir` attribute which forces the program to be in a specific directory before running.

Finally, I have the `Shell_exec` function, which is a "variable arguments" function. You've seen this before, but make sure you grasp the `stdarg.h` functions and how to write one of these. In the challenge for this section you are going to analyze this function.



## Challenge 2: Analyze Shell\_exec

Challenge for these files (in addition to a full code review just like you did in Challenge 1) is to fully analyze `Shell_exec` and break down exactly how it works. You should be able to understand each line, how the two `for-loops` work, and how arguments are being replaced.

Once you have it analyzed, add a field to `struct Shell` that gives the number of variable `args` that must be replaced. Update all the commands to have the right count of args, and then have an error check that confirms these args have been replaced and error exit.

## The Command Functions

Now you get to make the actual commands that do the work. These commands will use functions from APR, `db.h` and `shell.h` to do the real work of downloading and building software you want it to build. This is the most complex set of files, so do them carefully. As before, you start by making the `commands.h` file, then implementing its functions in the `commands.c` file.

```
#ifndef _commands_h
#define _commands_h

#include <apr_pools.h>

#define DEPENDS_PATH "/tmp/DEPENDS"
#define TAR_GZ_SRC "/tmp/pkg-src.tar.gz"
#define TAR_BZ2_SRC "/tmp/pkg-src.tar.bz2"
#define BUILD_DIR "/tmp/pkg-build"
#define GIT_PAT "*.git"
#define DEPEND_PAT "*DEPENDS"
#define TAR_GZ_PAT "*.tar.gz"
#define TAR_BZ2_PAT "*.tar.bz2"
#define CONFIG_SCRIPT "/tmp/pkg-build/configure"

enum CommandType {
    COMMAND_NONE, COMMAND_INSTALL, COMMAND_LIST, COMMAND_FETCH,
    COMMAND_INIT, COMMAND_BUILD
};

int Command_fetch(apr_pool_t *p, const char *url, int fetch_only);

int Command_install(apr_pool_t *p, const char *url, const char *configure_opts,
    const char *make_opts, const char *install_opts);

int Command_depends(apr_pool_t *p, const char *path);

int Command_build(apr_pool_t *p, const char *url, const char *configure_opts,
    const char *make_opts, const char *install_opts);

#endif
```

There's not much in `commands.h` that you haven't seen already. You should see that there's some defines for strings that are used everywhere. The real interesting code is in `commands.c`.

```
#include <apr_uri.h>
#include <apr_fnmatch.h>
#include <unistd.h>
```

```

#include "commands.h"
#include "dbg.h"
#include "bstrlib.h"
#include "db.h"
#include "shell.h"

int Command_depends(apr_pool_t *p, const char *path)
{
    FILE *in = NULL;
    bstring line = NULL;

    in = fopen(path, "r");
    check(in != NULL, "Failed to open downloaded depends: %s", path);

    for(line = bgets((bNgetc)fgetc, in, '\n'); line != NULL;
        line = bgets((bNgetc)fgetc, in, '\n'))
    {
        btrimws(line);
        log_info("Processing depends: %s", bdata(line));
        int rc = Command_install(p, bdata(line), NULL, NULL, NULL);
        check(rc == 0, "Failed to install: %s", bdata(line));
        bdestroy(line);
    }

    fclose(in);
    return 0;
error:
    if(line) bdestroy(line);
    if(in) fclose(in);
    return -1;
}

int Command_fetch(apr_pool_t *p, const char *url, int fetch_only)
{
    apr_uri_t info = {.port = 0};
    int rc = 0;
    const char *depends_file = NULL;
    apr_status_t rv = apr_uri_parse(p, url, &info);

    check(rv == APR_SUCCESS, "Failed to parse URL: %s", url);

    if(apr_fnmatch(GIT_PAT, info.path, 0) == APR_SUCCESS) {
        rc = Shell_exec(GIT_SH, "URL", url, NULL);
        check(rc == 0, "git failed.");
    } else if(apr_fnmatch(DEPEND_PAT, info.path, 0) == APR_SUCCESS) {
        check(!fetch_only, "No point in fetching a DEPENDS file.");

        if(info.scheme) {
            depends_file = DEPENDS_PATH;
            rc = Shell_exec(CURL_SH, "URL", url, "TARGET", depends_file, NULL)
;
            check(rc == 0, "Curl failed.");
        } else {
            depends_file = info.path;
        }

        // recursively process the devpkg list
        log_info("Building according to DEPENDS: %s", url);
        rv = Command_depends(p, depends_file);
        check(rv == 0, "Failed to process the DEPENDS: %s", url);

        // this indicates that nothing needs to be done
        return 0;
    } else if(apr_fnmatch(TAR_GZ_PAT, info.path, 0) == APR_SUCCESS) {
        if(info.scheme) {
            rc = Shell_exec(CURL_SH,
                "URL", url,

```

```

        "TARGET", TAR_GZ_SRC, NULL);
        check(rc == 0, "Failed to curl source: %s", url);
    }

    rv = apr_dir_make_recursive(BUILD_DIR,
        APR_UREAD | APR_UWRITE | APR_UEXECUTE, p);
    check(rv == APR_SUCCESS, "Failed to make directory %s", BUILD_DIR);

    rc = Shell_exec(TAR_SH, "FILE", TAR_GZ_SRC, NULL);
    check(rc == 0, "Failed to untar %s", TAR_GZ_SRC);
} else if(apr_fnmatch(TAR_BZ2_PAT, info.path, 0) == APR_SUCCESS) {
    if(info.scheme) {
        rc = Shell_exec(CURL_SH, "URL", url, "TARGET", TAR_BZ2_SRC, NULL);
        check(rc == 0, "Curl failed.");
    }

    apr_status_t rc = apr_dir_make_recursive(BUILD_DIR,
        APR_UREAD | APR_UWRITE | APR_UEXECUTE, p);

    check(rc == 0, "Failed to make directory %s", BUILD_DIR);
    rc = Shell_exec(TAR_SH, "FILE", TAR_BZ2_SRC, NULL);
    check(rc == 0, "Failed to untar %s", TAR_BZ2_SRC);
} else {
    sentinel("Don't now how to handle %s", url);
}

// indicates that an install needs to actually run
return 1;
error:
return -1;
}

int Command_build(apr_pool_t *p, const char *url, const char *configure_opts,
    const char *make_opts, const char *install_opts)
{
    int rc = 0;

    check(access(BUILD_DIR, X_OK | R_OK | W_OK) == 0,
        "Build directory doesn't exist: %s", BUILD_DIR);

    // actually do an install
    if(access(CONFIG_SCRIPT, X_OK) == 0) {
        log_info("Has a configure script, running it.");
        rc = Shell_exec(CONFIGURE_SH, "OPTS", configure_opts, NULL);
        check(rc == 0, "Failed to configure.");
    }

    rc = Shell_exec(MAKE_SH, "OPTS", make_opts, NULL);
    check(rc == 0, "Failed to build.");

    rc = Shell_exec(INSTALL_SH,
        "TARGET", install_opts ? install_opts : "install",
        NULL);
    check(rc == 0, "Failed to install.");

    rc = Shell_exec(CLEANUP_SH, NULL);
    check(rc == 0, "Failed to cleanup after build.");

    rc = DB_update(url);
    check(rc == 0, "Failed to add this package to the database.");

    return 0;
error:
return -1;
}

int Command_install(apr_pool_t *p, const char *url, const char *configure_opts,
    const char *make_opts, const char *install_opts)

```

```

{
    int rc = 0;
    check(Shell_exec(CLEANUP_SH, NULL) == 0, "Failed to cleanup before building.");

    rc = DB_find(url);
    check(rc != -1, "Error checking the install database.");

    if(rc == 1) {
        log_info("Package %s already installed.", url);
        return 0;
    }

    rc = Command_fetch(p, url, 0);

    if(rc == 1) {
        rc = Command_build(p, url, configure_opts, make_opts, install_opts);
        check(rc == 0, "Failed to build: %s", url);
    } else if(rc == 0) {
        // no install needed
        log_info("Depends successfully installed: %s", url);
    } else {
        // had an error
        sentinel("Install failed: %s", url);
    }

    Shell_exec(CLEANUP_SH, NULL);
    return 0;

error:
    Shell_exec(CLEANUP_SH, NULL);
    return -1;
}

```

After you have this entered in and compiling, you can analyze it. If you've done the challenges until now, you should see how the `shell.c` functions are being used to run shells and how the arguments are being replaced. If not then go back and make sure you *truly* understand how `Shell_exec` actually works.

### Challenge 3: Critique My Design

As before, do a complete review of this code and make sure it's exactly the same. Then go through each function and make sure you know how it works and what it's doing. You also should trace how each function calls the other functions you've written in this file and other files. Finally, confirm that you understand all the functions you're calling from APR here.

Once you have the file correct and analyzed, go back through and assume I'm an idiot. Then, criticize the design I have to see how you can improve it if you can. Don't *actually* change the code, just create a little `notes.txt` file and write down your thoughts and what you might change.

## The `devpkg` Main Function

The last and most important file, but probably the simplest, is `devpkg.c` where the `main` function lives. There's no `.h` file for this, since this one includes all the others. Instead this just creates the executable `devpkg` when combined with the other `.o` files from our `Makefile`. Enter in the code for this file, and make sure it's correct.

```

#include <stdio.h>
#include <apr_general.h>
#include <apr_getopt.h>
#include <apr_strings.h>
#include <apr_lib.h>

#include "dbg.h"
#include "db.h"
#include "commands.h"

int main(int argc, const char const *argv[])
{
    apr_pool_t *p = NULL;
    apr_pool_initialize();
    apr_pool_create(&p, NULL);

    apr_getopt_t *opt;
    apr_status_t rv;

    char ch = '\\0';
    const char *optarg = NULL;
    const char *config_opts = NULL;
    const char *install_opts = NULL;
    const char *make_opts = NULL;
    const char *url = NULL;
    enum CommandType request = COMMAND_NONE;

    rv = apr_getopt_init(&opt, p, argc, argv);

    while(apr_getopt(opt, "I:Lc:m:i:d:SF:B:", &ch, &optarg) == APR_SUCCESS) {
        switch (ch) {
            case 'I':
                request = COMMAND_INSTALL;
                url = optarg;
                break;

            case 'L':
                request = COMMAND_LIST;
                break;

            case 'c':
                config_opts = optarg;
                break;

            case 'm':
                make_opts = optarg;
                break;

            case 'i':
                install_opts = optarg;
                break;

            case 'S':
                request = COMMAND_INIT;
                break;

            case 'F':
                request = COMMAND_FETCH;
                url = optarg;
                break;

            case 'B':
                request = COMMAND_BUILD;
                url = optarg;
                break;

        }
    }

    switch(request) {

```

```

    case COMMAND_INSTALL:
        check(url, "You must at least give a URL.");
        Command_install(p, url, config_opts, make_opts, install_opts);
        break;

    case COMMAND_LIST:
        DB_list();
        break;

    case COMMAND_FETCH:
        check(url != NULL, "You must give a URL.");
        Command_fetch(p, url, 1);
        log_info("Downloaded to %s and in /tmp/", BUILD_DIR);
        break;

    case COMMAND_BUILD:
        check(url, "You must at least give a URL.");
        Command_build(p, url, config_opts, make_opts, install_opts);
        break;

    case COMMAND_INIT:
        rv = DB_init();
        check(rv == 0, "Failed to make the database.");
        break;

    default:
        sentinel("Invalid command given.");
}

return 0;

error:
return 1;
}

```

## Challenge 4: The README And Test Files

The challenge for this file is to understand how the arguments are being processed, what the arguments are, and then create the `README` file with instructions on how to use it. As you write the README, also write a simple `test.sh` that runs `./devpkg` to check that each command is actually working against real live code. Use the `set -e` at the top of your script so that it aborts on the first error.

Finally, run the program under valgrind and make sure it's all working before moving on to the mid-term exam.

## The Mid-Term Exam

Your final challenge is the mid-term exam and it involves three things:

- Compare your code to my code available online and starting with 100%, remove 1% for each line you got wrong.
- Take your notes.txt on how you would improve the code and functionality of `devpkg` and implement your improvements.
- Write an alternative version of `devpkg` using your other favorite language or the one you think can do this the best. Compare the two, then improve your C version of `devpkg` based on what you've learned.

To compare your code with mine, do the following:

```
cd .. # get one directory above your current one
git clone git://gitorious.org/devpkg/devpkg.git
devpkgzed
diff -r devpkg devpkgzed
```

This will clone my version of `devpkg` into a directory `devpkgzed` and then use the tool `diff` to compare what you've done to what I did. The files you're working with in this book come directly from this project, so if you get different lines then that's an error.

Keep in mind that there's no real pass or fail on this exercise, just a way for you to challenge yourself to be as exact and meticulous as possible.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 27: Creative And Defensive Programming

You have now learned most of the basics of C programming and are ready to start becoming a serious programmer. This is where you go from beginner to expert, both with C and hopefully with core computer science concepts. I will be teaching you a few of the core data structures and algorithms that every programmer should know, and then a few very interesting ones I've used in real software for years.

Before I can do that I have to teach you some basic skills and ideas that will help you make better software. Exercises 27 through 31 will teach you advanced concepts and feature more talking than code, but after those you'll apply what you learn to making a core library of useful data structures.

The first step in getting better at writing C code (and really any language) is to learn a new mindset called "defensive programming". Defensive programming assumes that you are going to make many mistakes and then attempts to prevent them at every possible step. In this exercise I'm going to teach you how to think about programming defensively.

### The Creative Programmer Mindset

It's not possible to tell you how to be creative in a short exercise like this, but I will tell you that creativity involves taking risks and being open minded. Fear will quickly kill creativity, so the mindset I adopt, and many programmers adopt on, accident is designed to make me unafraid of taking chances and looking like an idiot:

- I can't make a mistake.
- It doesn't matter what people think.
- Whatever my brain comes up with is going to be a great idea.

I only adopt this mindset temporarily, and even have little tricks to turn it on. By doing this I can come up with ideas, find creative solutions, open my thoughts to odd connections, and just generally invent weirdness without fear. In this mindset I will typically write a horrible first version of something just to get the idea out.

However, when I've finished my creative prototype I will throw it out and get serious about making it solid. Where other people make a mistake is carrying the creative mindset into their implementation phase. This then leads to a very different destructive mindset that is the dark side of the creative mindset:

- It is possible to write perfect software.
- My brain tells me the truth, and it can't find any errors, therefore I have written perfect software.



- My code is who I am and people who criticize its perfection are criticizing me.

These are lies. You will frequently run into programmers who feel intense pride about what they've created, which is natural, but this pride gets in the way of their ability to objectively improve their craft. Because of pride and attachment to what they've written, they can continue to believe that what they write is perfect. As long as they ignore other people's criticism of their code they can protect their fragile ego and never improve.

The trick to being creative *and* making solid software is to also be able to adopt a defensive programming mindset.

## The Defensive Programmer Mindset

After you have a working creative prototype and you're feeling good about the idea, it's time to switch to being a defensive programmer. The defensive programmer basically hates your code and believes these things:

- Software has errors.
- You are not your software, yet you are responsible for the errors.
- You can never remove the errors, only reduce their probability.

This mindset lets you be honest about your work and critically analyze it for improvements. Notice that it doesn't say *you* are full of errors? It says your *code* is full of errors. This is a significant thing to understand because it gives you the power of objectivity for the next implementation.

Just like the creative mindset, the defensive programming mindset has a dark side as well. The defensive programmer is a paranoid who is afraid of everything, and this fear prevents them from possibly being wrong or making mistakes. That's great when you are trying to be ruthlessly consistent and correct, but it is murder on creative energy and concentration.

## The Eight Defensive Programmer Strategies

Once you've adopted this mindset, you can then rewrite your prototype and follow a set of eight strategies I use to make my code as solid as I can. While I work on the "real" version I ruthlessly follow these strategies and try to remove as many errors as I can, thinking like someone who wants to break the software.

### Never Trust Input

Never trust the data you are given and always validate it.

### Prevent Errors

If an error is possible, no matter how probable, try to prevent it.

### Fail Early And Openly

Fail early, cleanly, and openly, stating what happened, where and how to fix it.

#### Document Assumptions

Clearly state the pre-conditions, post-conditions, and invariants.

#### Prevention Over Documentation

Do not do with documentation, that which can be done with code or avoided completely.

#### Automate Everything

Automate everything, especially testing.

#### Simplify And Clarify

Always simplify the code to the smallest, cleanest form that works without sacrificing safety.

#### Question Authority

Do not blindly follow or reject rules.

These aren't the only ones, but they're the core things I feel programmers have to focus on when trying to make good solid code. Notice that I don't really say exactly how to do these. I'll go into each of these in more detail, and some of the exercises actually cover them extensively.

## Applying The Eight Strategies

These ideas are all great pop-psychology platitudes, but how do you actually apply them to working code? I'm now going to give you a set of things to always do in this book's code that demonstrate each one with a concrete example. The ideas aren't limited to these examples, and you should use these as a guide to making your own code tougher.

## Never Trust Input

Let's look at an example of bad design and "better" design. I won't say good design because this could be done even better. Take a look at two functions that both copy a string and a simple `main` to test out the better one.

```
#undef NDEBUG
#include "dbg.h"
#include <stdio.h>
#include <assert.h>

/*
 * Naive copy that assumes all inputs are always valid
 * taken from K&R C and cleaned up a bit.
 */
void copy(char to[], char from[])
{
    int i = 0;

    // while loop will not end if from isn't '\0' terminated
    while((to[i] = from[i]) != '\0') {
        ++i;
    }
}

/*
 * A safer version that checks for many common errors using the
 * length of each string to control the loops and termination.
 */
int safercopy(int from_len, char *from, int to_len, char *to)
{
    assert(from != NULL && to != NULL && "from and to can't be NULL");
    int i = 0;
```

```

int max = from_len > to_len - 1 ? to_len - 1 : from_len;

// to_len must have at least 1 byte
if(from_len < 0 || to_len <= 0) return -1;

for(i = 0; i < max; i++) {
    to[i] = from[i];
}

to[to_len - 1] = '\0';

return i;
}

int main(int argc, char *argv[])
{
    // careful to understand why we can get these sizes
    char from[] = "0123456789";
    int from_len = sizeof(from);

    // notice that it's 7 chars + \0
    char to[] = "0123456";
    int to_len = sizeof(to);

    debug("Copying '%s':%d to '%s':%d", from, from_len, to, to_len);

    int rc = safercopy(from_len, from, to_len, to);
    check(rc > 0, "Failed to safercopy.");
    check(to[to_len - 1] == '\0', "String not terminated.");

    debug("Result is: '%s':%d", to, to_len);

    // now try to break it
    rc = safercopy(from_len * -1, from, to_len, to);
    check(rc == -1, "safercopy should fail #1");
    check(to[to_len - 1] == '\0', "String not terminated.");

    rc = safercopy(from_len, from, 0, to);
    check(rc == -1, "safercopy should fail #2");
    check(to[to_len - 1] == '\0', "String not terminated.");

    return 0;
error:
    return 1;
}

```

The `copy` function is typical C code and it's the source of a huge number of buffer overflows. It is flawed because it assumes that it will always receive a validly terminated C string (with `'\0'`) and just uses a while-loop to process it. Problem is, ensuring that is incredibly difficult, and if not handled right it causes the while-loop to loop infinitely. *A cornerstone of writing solid code is never writing loops that can possibly loop forever.*

The `safercopy` function tries to solve this by requiring the caller to give the lengths of the two strings it must deal with. By doing this it can make certain checks about these strings that the `copy` function can't. It can check the lengths are right, that the `to` string has enough space, and it will *always* terminate. It's impossible for this function to run on forever like the `copy` function.

This is the idea behind never trusting the inputs you receive. If you assume that your function is going to get a string that's not terminated (which is common) then you design your function to not rely on that to function properly. If you need the arguments to never be `NULL` then you

should check for that too. If the sizes should be within sane levels, then check that. You simply assume that whoever is calling you got it wrong and try to make it difficult for them to give you bad state.

This then extends out to software you write that gets input from the external universe. The famous last words of the programmer are, "Nobody's going to do that." I've seen them say that and then the *next* day someone does exactly that, crashing or hacking their application. If you say nobody is going to do that, just throw in the code to make sure they simply can't hack your application. You'll be glad you did.

There is a diminishing returns on this, but here's a list of things I try to do with all of my functions I write in C:

- For each parameter identify what its preconditions are, and whether the precondition should cause a failure or return an error. If you are writing a library, favor errors over failures.
- Add `assert` calls at the beginning that checks for each failure precondition using `assert(test && "message");` This little hack does the test, and when it fails the OS will typically print the assert line for you, which then includes that message. Very helpful when you're trying to figure out why that `assert` is there.
- For the other preconditions, return the error code or use my `check` macro to do that and give an error message. I didn't use `check` in this example since it would confuse the comparison.
- Document *why* these preconditions exist so that when a programmer hits the error they can figure out if they are really necessary or not.
- If you are modifying the inputs, make sure that they are correctly formed when the function exits, or abort if they aren't.
- Always check the error codes of functions you use. For example, people frequently forget to check the return codes from `fopen` or `fread` which causes them to use the resources they give despite the error. This causes your program to crash or gives an avenue for an attack.
- You also need to be returning consistent error codes so that you can do this for all of your functions too. Once you get in this habit you will then understand why my `check` macros work the way they do.

Just doing these simple things will improve your resource handling and prevent quite a few errors.

## Prevent Errors

In the previous example you may hear people say, "Well it's not very likely someone will use `copy` wrong." Despite the mountain of attacks made against this very kind of function they still believe that the probability of this error is very low. Probability is a funny thing because people are incredibly bad at guessing the probability of any event. People are however much better at determining if something is *possible*. They may say the error in `copy` is not *probably*, but they can't deny that it's *possible*.

The key reason is that for something to be probable, it first has to be

possible. Determining the possibility is easy, since we can all imagine something happening. What's not so easy is determining its possibility after that. Is the chance that someone might use `copy` wrong 20%, 10%, or 1%? Who knows, and to determine that you'd need to gather evidence, look at rates of failure in many software packages, and probably survey real programmers and how they use the function.

This means, if you're going to prevent errors then you need to try to prevent what is possible, but focus your energies on what's most probable first. It may not be feasible to handle all the possible ways your software can be broken, but you have to attempt it. But, at the same time, if you don't constrain your efforts to the most probably events with the least effort then you'll be wasting time on irrelevant attacks.

Here's a process for determining what to prevent in your software:

- List all the possible errors that can happen, no matter how probable. Within reason of course. No point listing aliens sucking your memories out to steal your passwords.
- Give each one a probability that's a percentage of operations that can be vulnerable. If you are handling requests from the internet, then it's the percentage of requests that can cause the error. If it's function calls, then it's what percentage of function calls can cause it.
- Give each one an effort in number of hours or amount of code to prevent it. You could also just give an easy or hard metric. Any metric that prevents you from working on the impossible when there's easier things to fix still on the list.
- Rank them by effort (lowest to highest), and probability (highest to lowest). This is now your task list.
- Prevent all the errors you can in this list, aiming for removing the possibility, then reducing the probability if you can't make it impossible.
- If there are errors you can't fix, then document them so someone else can fix it.

This little process will give you a nice list of things to do, but more importantly keep you from working on useless things when there's other more important things to work on. You can also be more or less formal with this process. If you're doing a full security audit this will be better done with a whole team and a nice spreadsheet. If you're just writing a function then simply reviewing the code and scratching out these into some comments is good enough. What's important is you stop assuming that errors don't happen, and you work on removing them when you can without wasting effort.

## Fail Early And Openly

If you encounter an error in C you have two choices:

- Return an error code.
- Abort the process.

This is just how it is, so what you need to do is make sure the failures happen quickly, are clearly documented, give an error message, and are easy for the programmer to avoid. This is why the `check` macros I've

given you work the way they do. For every error you find it prints a message, the file and line number where it happened, and force a return code. If you just use my macros you'll end up doing the right thing anyway.

I tend to prefer returning error code to aborting the program. If it's catastrophic then I will, but very few errors are truly catastrophic. A good example of when I'll abort a program is if I'm given an invalid pointer, as I did in `safercopy`. Instead of having the programmer experience a segmentation fault explosion "somewhere", I catch it right away and abort. However, if it's common to pass in a NULL then I'll probably change that to a `check` instead so that the caller can adapt and keep running.

In libraries however, I try my hardest to *never* abort. The software using my library can decide if it should abort, and typically I'll only abort if the library is very badly used.

Finally, a big part of being "open" about errors is not using the same message or error code for more than one possible error. You typically see this with errors on external resources. A library will receive an error on a socket, and then simply report "bad socket". What they should do is return exactly what the error was on the socket so it can be debugged properly and fixed. When designing your error reporting, make sure you give a different error message for the different possible errors.

## Document Assumptions

If you're following along and doing this advice then what you'll be doing is building a "contract" of how your functions expect the world to be. You've created preconditions for each argument, you've handled possible errors, and you're failing elegantly. The next step is to complete the contract and add "invariants" and "postconditions".

An invariant is some condition that must be held true in some state while the function runs. This isn't very common in simple functions, but when you're dealing with complex structures it becomes more necessary. A good example of an invariant is that a structure is always initialized properly while it's being used. Another would be that a sorted data structure is always sorted during processing.

A postcondition is a guarantee on the exit value or result of a function running. This can blend together with invariants, but this is something as simple as "function always returns 0 or -1 on error". Usually these are documented, but if your function returns an allocated resource, you can add a postcondition that checks to make sure it's returning something and not NULL. Or, you can use NULL to indicate an error, so in that case your postcondition is now checking the resource is deallocated on any errors.

In C programming invariants and postconditions are usually more documentation than actual code and assertions. The best way to handle them is add `assert` calls for the ones you can, then document the rest. If you do that then when people hit an error they can see what assumptions you made when writing the function.

## Prevention Over Documentation

A common problem when programmers write code is they will document a common bug rather than simply fix it. My favorite is when the Ruby on Rails system simply assumed that all months had 30 days. Calendars are hard, so rather than fix it they threw a tiny little comment somewhere that said this was on purpose, and then they refused to fix it for years. Every time someone would complain they would then bluster and yell, "But it's documented!"

Documentation doesn't matter if you can actually fix the problem, and if the function has a fatal flaw then simply don't include it until you can fix it. In the case of Ruby on Rails, not having date functions would have been better than including purposefully broken ones that nobody could use.

As you go through your defensive programming cleanups, try to fix everything you can. If you find yourself documenting more and more problems you can't fix, then consider redesigning the feature or simply removing it. If you *really* have to keep this horribly broken feature, then I suggest you write it, document it and find a new job before you are blamed for it.

## Automate Everything

You are a programmer, and that means your job is putting other people out of jobs with automation. The pinnacle of this is putting yourself out of a job with your own automation. Obviously you won't completely remove what you do, but if are spending your whole day rerunning manual tests in your terminal, then your job is not programming. You are doing QA, and you should automate yourself out of this QA job you probably don't really want anyway.

The easiest way to do this is to write automated tests, or unit tests. In this book I'm going to get into how to do this easily, and I'll avoid most of the dogma of when you should write tests. I'll focus on how to write them, what to test, and how to be efficient at the testing.

Common things programmers fail to automate but they should:

- Testing and validation.
- Build processes.
- Deployment of software.
- System administration.
- Error reporting.

Try to devote some of your time to automating this and you'll have more time to work on the fun stuff. Or, if this is fun to you, then maybe you should work on software that makes automating these things easier.

## Simplify And Clarify

The concept of "simplicity" is a slippery one to many people, especially smart people. They generally confuse "comprehension" with "simplicity". If they understand it well, clearly it's simple. The actual test of simplicity



is by comparison with something else that could be simpler. But, you'll see people who write code go running to the most complex obtuse structures possible because they think the simpler version of the same thing is "dirty". A love affair with complexity is a programming sickness.

You can fight this disease by first telling yourself, "Simple and clear is not dirty, no matter what everyone else is doing." If everyone else is writing insane visitor patterns involving 19 classes over 12 interfaces and you can do it with two string operations, then you win. They are wrong, no matter how "elegant" they think their complex monstrosity is.

The simplest test of which function to use is:

- Make sure both functions have no errors. It doesn't matter how fast or simple a function is if it has errors.
- If you can't fix one, then pick the other.
- Do they produce the same result? If not then pick the one that has the result you need.
- If they produce the same result, then pick the one that either has fewer features, fewer branches, or you just think is simpler.
- Make sure you're not just picking the one that is most impressive. Simple and dirty beats complex and clean any day.

You'll notice that I mostly give up at the end and tell you to use your judgment. Simplicity is ironically a very complex thing, so using your tastes as a guide is the best way to go. Just make sure you adjust your view of what's "good" as you grow and gain more experience.

## Question Authority

The final strategy is the most important because it breaks you out of the defensive programming mindset and lets you transition into the creative mindset. Defensive programming is authoritarian and it can be cruel. The job of this mindset is to make you follow rules because without them you'll miss something or get distracted.

This authoritarian attitude has the disadvantage of disabling independent creative thought. Rules are necessary for getting things done, but being a slave to them will kill your creativity.

This final strategy means you should question the rules you follow periodically and assume that they could be wrong, just like the software you are reviewing. What I will typically do is, after a session of defensive programming, I'll go take a non-programming break and let the rules go. Then I'll be ready to do some creative work or do more defensive coding if need to.

## Order Is Not Important

The final thing I'll say on this philosophy is that I'm not telling you to do this in a strict order of "CREATE! DEFEND! CREATE! DEFEND!" At first you may want to do that, but I will actually do either in varying amounts depend on what I want to do, and I may even meld them together with no defined boundary.

I also don't think one mindset is better than another, or that there are



strict separation between them. You need both creativity and strictness to do programming well, so work on both if you want to improve.

## Extra Credit

- The code in the book up to this point (and for the rest of it) potentially violates these rules. Go back through and apply what you've learned to one exercise to see if you can improve it or find bugs.
- Find an open source project and give some of the files a similar code review. Submit a patch that fixes a bug if you find it.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 28: Intermediate Makefiles

In the next three Exercises you'll create a skeleton project directory to use in building your C programs later. This skeleton directory will be used in the rest of the book, and in this exercise I'll cover just the `Makefile` so you can understand it.

The purpose of this structure is to make it easy to build medium sized programs without having to resort to configure tools. If done right you can get very far with just GNU make and some small shell scripts.

### The Basic Project Structure

The first thing to do is make a `c-skeleton` directory and then put a set of basic files and directories in it that many projects have. Here's my starter:

```
$ mkdir c-skeleton
$ cd c-skeleton/
$ touch LICENSE README.md Makefile
$ mkdir bin src tests
$ cp dbg.h src/ # this is from Ex20
$ ls -l
total 8
-rw-r--r-- 1 zedshaw staff 0 Mar 31 16:38 LICENSE
-rw-r--r-- 1 zedshaw staff 1168 Apr 1 17:00 Makefile
-rw-r--r-- 1 zedshaw staff 0 Mar 31 16:38 README.md
drwxr-xr-x 2 zedshaw staff 68 Mar 31 16:38 bin
drwxr-xr-x 2 zedshaw staff 68 Apr 1 10:07 build
drwxr-xr-x 3 zedshaw staff 102 Apr 3 16:28 src
drwxr-xr-x 2 zedshaw staff 68 Mar 31 16:38 tests
$ ls -l src
total 8
-rw-r--r-- 1 zedshaw staff 982 Apr 3 16:28 dbg.h
$
```

At the end you see me do an `ls -l` so you can see the final results.

Here's what each of these does:

#### `LICENSE`

If you release the source of your projects you'll want to include a license. If you don't though, the code is copyright by you and nobody has rights to it by default.

#### `README.md`

Basic instructions for using your project go here. It ends in `.md` so that it will be interpreted as markdown.

#### `Makefile`

The main build file for the project.

#### `bin/`

Where programs that users can run go. This is usually empty and the Makefile will create it if it's not there.

#### `build/`

Where libraries and other build artifacts go. Also empty and the Makefile will create it if it's not there.

`src/`

Where the source code goes, usually `.c` and `.h` files.

`tests/`

Where automated tests go.

`src/dbg.h`

I copied the `dbg.h` from Exercise 20 into `src/` for later.

I'll now break down each of the components of this skeleton project so you can understand how it works.

## Makefile

The first thing I'll cover is the Makefile because from that you can understand how everything else works. The Makefile in this exercise is much more detailed than ones you've used so far, so I'm going to break it down after you type it in:

```
CFLAGS=-g -O2 -Wall -Wextra -Isrc -rdynamic -DDEBUG $(OPTFLAGS)
LIBS=-ldl $(OPTLIBS)
PREFIX?=/usr/local

SOURCES=$(wildcard src/**/*.c src/*.c)
OBJECTS=$(patsubst %.c,%.o,$(SOURCES))

TEST_SRC=$(wildcard tests/*_tests.c)
TESTS=$(patsubst %.c,%, $(TEST_SRC))

TARGET=build/libYOUR_LIBRARY.a
SO_TARGET=$(patsubst %.a,%.so,$(TARGET))

# The Target Build
all: $(TARGET) $(SO_TARGET) tests

dev: CFLAGS=-g -Wall -Isrc -Wall -Wextra $(OPTFLAGS)
dev: all

$(TARGET): CFLAGS += -fPIC
$(TARGET): build $(OBJECTS)
    ar rcs $@ $(OBJECTS)
    ranlib $@

$(SO_TARGET): $(TARGET) $(OBJECTS)
    $(CC) -shared -o $@ $(OBJECTS)

build:
    @mkdir -p build
    @mkdir -p bin

# The Unit Tests
.PHONY: tests
tests: CFLAGS += $(TARGET)
tests: $(TESTS)
    sh ./tests/runtests.sh

valgrind:
    VALGRIND="valgrind --log-file=/tmp/valgrind-%p.log" $(MAKE)

# The Cleaner
clean:
    rm -rf build $(OBJECTS) $(TESTS)
    rm -f tests/tests.log
    find . -name "*.gc*" -exec rm {} \;
    rm -rf `find . -name "*.dSYM" -print`
```

```
# The Install
install: all
    install -d $(DESTDIR)/$(PREFIX)/lib/
    install $(TARGET) $(DESTDIR)/$(PREFIX)/lib/

# The Checker
BADFUNCS='[^_>a-zA-Z0-9](str(n?cpy|n?cat|xfrm|n?dup|str|pbrk|tok|_)|stpn?cpy|a?sn?printf|byte_)'
check:
    @echo Files with potentially dangerous functions.
    @egrep $(BADFUNCS) $(SOURCES) || true
```

Remember that you need to indent the Makefile consistently with tab characters. Your editor should know that and do the right thing, but if it doesn't then get a different text editor. No programmer should use an editor that fails at something so simple.

## The Header

This makefile is designed to build a library we'll be working on later and to do so reliably on almost any platform by using special features of `GNU make`. I'll break down each part in sections, starting with the header.

### Makefile:1

These are the usual `CFLAGS` that you set in all of your projects, but with a few others that may be needed to build libraries. *You may need to adjust these for different platforms.* Notice the `OPTFLAGS` variable at the end which lets people augment the build options as needed.

### Makefile:2

Options used when linking a library, and allows someone else to augment the linking options using the `OPTLIBS` variable.

### Makefile:3

Setting an *optional* variable called `PREFIX` that will only have this value if the person running the Makefile didn't already give a `PREFIX` setting. That's what the `?=` does.

### Makefile:5

This fancy line of awesome *dynamically* creates the `SOURCES` variable by doing a `wildcard` search for all `*.c` files in the `src/` directory. You have to give both `src/**/*.c` and `src/*.c` so that GNU make will include the files in `src` and also the ones below it.

### Makefile:6

Once you have the list of source files, you can then use the `patsubst` to take the `SOURCES` list of `*.c` files and make a *new* list of all the object files. You do this by telling `patsubst` to change all `%.c` extensions to `%.o` and then those are assigned to `OBJECTS`.

### Makefile:8

Using the `wildcard` again to find all the test source files for the unit tests. These are separate from the library's source files.

### Makefile:9

Then using the same `patsubst` trick to dynamically get all the `TEST` targets. In this case I'm stripping away the `.c` extension so that a full program will be made with the same name. Previously I had replaced the `.c` with `{.o}` so an object file is

created.

Makefile:11

Finally, we say the ultimate target is `build/libYOUR_LIBRARY.a`, which you will change to be whatever library you are actually trying to build.

This completes the top of the Makefile, but I should explain what I mean by "lets people augment the build". When you run make you can do this:

```
# WARNING! Just a demonstration, won't really work
right now.
# this installs the library into /tmp
$ make PREFIX=/tmp install
# this tells it to add pthreads
$ make OPTFLAGS=-pthread
```

If you pass in options that match the same kind of variables you have in your `Makefile`, then those will show up in your build. You can then use this to change how the `Makefile` runs. The first one alters the `PREFIX` so that it installs into `/tmp` instead. The second one sets `OPTFLAGS` so that the `-pthread` option is present.

## The Target Build

Continuing with the breakdown of the `Makefile` I have actually building the object files and targets:

Makefile:14

Remember that the first target is what `make` will run by default when no target is given. In this case it's called `all:` and it gives `$(TARGET) tests` as the targets to build. Look up at the `TARGET` variable and you see that's the library, so `all:` will first build the library. The `tests` target is then further down in the `Makefile` and builds the unit tests.

Makefile:16

Another target for making "developer builds" that introduces a technique for changing options for just one target. If I do a "dev build" I want the `CFLAGS` to include options like `-Wextra` that are useful for finding bugs. If you place them on the target line as options like this, then give another line that says the original target (in this case `all`) then it will change the options you set. I use this for setting different flags on different platforms that need it.

Makefile:19

Builds the `TARGET` library, whatever that is, and also uses the same trick from line 15 of giving a target with just options changes to alter them for this run. In this case I'm adding `-fPIC` just for the library build using the `+=` syntax to add it on.

Makefile:20

Now the real target where I say first make the `build` directory, then compile all the `OBJECTS`.

Makefile:21

Runs the `ar` command which actually makes the `TARGET`. The syntax `$(@) $(OBJECTS)` is a way of saying, "put the target for this Makefile source here and all the OBJECTS after that". In this case the `$(@)` maps back to the `$(TARGET)` on line 19, which

maps to `build/libYOUR_LIBRARY.a`. It seems like a lot to keep track of this indirection, and it can be, but once you get it working this means you just change `TARGET` at the top and build a whole new library.

Makefile:22

Finally, to make the library you run `ranlib` on the `TARGET` and it's built.

Makefile:24-24

This just makes the `build/` or `bin/` directories if they don't exist. This is then referenced from line 19 when it gives the `build` target to make sure the `build/` directory is made.

You now have all the stuff you need to build the software, so we'll create a way to build and run unit tests to do automated testing.

## The Unit Tests

C is different from other languages because it's easier to create one tiny little program for each thing you're testing. Some testing frameworks try to emulate the module concept other languages have and do dynamic loading, but this doesn't work well in C. It's also unnecessary because you can just make a single program that's run for each test instead.

I'll cover this part of the Makefile, and then later you'll see the contents of the `tests/` directory that make it actually work.

Makefile:29

If you have a target that's not "real", but there is a directory or file with that name, then you need to tag the target with `.PHONY:` so `make` will ignore the file and always run.

Makefile:30

I use the same trick for modifying the `CFLAGS` variable to add the `TARGET` to the build so that each of the test programs will be linked with the `TARGET` library. In this case it will add `build/libYOUR_LIBRARY.a` to the linking.

Makefile:31

Then I have the actual `tests:` target which depends on all the programs listed in the `TESTS` variable we created in the header. This one line actually says, "Make, use what you know about building programs and the current CFLAGS settings to build each program in `TESTS`."

Makefile:32

Finally, when all of the `TESTS` are built, there's a simple shell script I'll create later that knows how to run them all and report their output. This line actually runs it so you can see the test results.

Makefile:34-35

In order to be able to dynamically rerun the tests with Valgrind there's a `valgrind:` target that sets the right variable and runs itself again. This puts the valgrind logs into `/tmp/valgrind-*.log` so you can go look and see what might be going on. The `tests/runtests.sh` then knows to run the test programs under Valgrind when it sees this `VALGRIND` variable.

For the unit testing to work you'll need to create a little shell script that

knows how to run the programs. Go ahead and create this

`tests/runtests.sh` script:

```
echo "Running unit tests:"

for i in tests/*_tests
do
    if test -f $i
    then
        if $VALGRIND ./$i 2>> tests/tests.log
        then
            echo $i PASS
        else
            echo "ERROR in test $i: here's tests/tests.log"
            echo "-----"
            tail tests/tests.log
            exit 1
        fi
    fi
done

echo ""
```

I'll be using this later when I cover how unit tests work.

## The Cleaner

I now have fully working unit tests, so next up is making things clean when I need to reset everything.

Makefile:38

The `clean:` target starts things off whenever we need to clean up the project.

Makefile:39-42

This cleans out most of the junk that various compilers and tools leave behind. It also gets rid of the `build/` directory and uses a trick at the end to cleanly erase the weird `*.dSYM` directories Apple's XCode leaves behind for debugging purposes.

If you run into junk that you need to clean out, simply augment the list of things being deleted in this target.

## The Install

After that I'll need a way to install the project, and for a `Makefile` that's building a library I just need to put something in the common `PREFIX` directory, which is usually `/usr/local/lib`.

Makefile:45

This makes `install:` depend on the `all:` target so that when you run `make install` it will be sure to build everything.

Makefile:46

I then use the program `install` to create the target `lib` directory if it doesn't exist. In this case I'm trying to make the install as flexible as possible by using two variables that are conventions for installers. `DESTDIR` is handed to make by installers that do their builds in secure or odd locations to build packages. `PREFIX` is used when people want the project to be

installed in someplace other than `/usr/local`.

Makefile:47

After that I'm just using `install` to actually install the library where it needs to go.

The purpose of the `install` program is to make sure things have the right permissions set. When you run `make install` you usually have to do it as the root user, so the typical build process is `make && sudo make install`.

## The Checker

The very last part of this `Makefile` is a bonus that I include in my C projects to help me dig out any attempts to use the "bad" functions in C. Namely the string functions and other "unprotected buffer" functions.

Makefile:50

Sets a variable which is a big regex looking for bad functions like `strcpy`.

Makefile:51

The `check:` target so you can run a check whenever you need.

Makefile:52

Just a way to print a message, but doing `@echo` tells `make` to not print the command, just its output.

Makefile:53

Run the `egrep` command on the source files looking for any bad patterns. The `|| true` at the end is a way to prevent `make` from thinking that `egrep` not finding things is a failure.

When you run this it will have the odd effect that you'll get an error when there is nothing bad going on.

## What You Should See

I have two more exercises to go before I'm done building the project skeleton directory, but here's me testing out the features of the `Makefile`.

```
$ make clean
rm -rf build
rm -f tests/tests.log
find . -name "*.gc*" -exec rm {} \;
rm -rf `find . -name "*.dSYM" -print`
$ make check
Files with potentially dangerous functions.
^Cmake: *** [check] Interrupt: 2

$ make
ar rcs build/libYOUR_LIBRARY.a
ar: no archive members specified
usage: ar -d [-TLsv] archive file ...
       ar -m [-TLsv] archive file ...
       ar -m [-abiTLsv] position archive file ...
       ar -p [-TLsv] archive [file ...]
       ar -q [-cTLsv] archive file ...
       ar -r [-cuTLsv] archive file ...
       ar -r [-abciuTLsv] position archive file ...
       ar -t [-TLsv] archive [file ...]
       ar -x [-ouTLsv] archive [file ...]
```



```

make: *** [build/libYOUR_LIBRARY.a] Error 1
$ make valgrind
VALGRIND="valgrind --log-file=/tmp/valgrind-%p.log" make
ar rcs build/libYOUR_LIBRARY.a
ar: no archive members specified
usage: ar -d [-TLsv] archive file ...
       ar -m [-TLsv] archive file ...
       ar -m [-abiTLsv] position archive file ...
       ar -p [-TLsv] archive [file ...]
       ar -q [-cTLsv] archive file ...
       ar -r [-cuTLsv] archive file ...
       ar -r [-abciuTLsv] position archive file ...
       ar -t [-TLsv] archive [file ...]
       ar -x [-ouTLsv] archive [file ...]
make[1]: *** [build/libYOUR_LIBRARY.a] Error 1
make: *** [valgrind] Error 2
$

```

When I run the `clean:` target that works, but because I don't have any source files in the `src/` directory none of the other commands really work. I'll finish that up in the next exercises.

## Extra Credit

- Try to get the `Makefile` to actually work by putting a source and header file in `src/` and making the library. You shouldn't need a `main` function in the source file.
- Research what functions the `check:` target is looking for in the `BADFUNCS` regular expression it's using.
- If you don't do automated unit testing, then go read about it so you're prepared later.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 29: Libraries And Linking

A central part of any C program is the ability to link it to libraries that your operating system provides. Linking is how you get additional features for your program that someone else created and packaged on the system. You've been using some standard libraries that are automatically included, but I'm going to explain the different types of libraries and what they do.

First off, libraries are poorly designed in every programming language. I have no idea why, but it seems language designers think of linking as something they just slap on later. They are usually confusing, hard to deal with, can't do versioning right, and end up being linked differently everywhere.

C is no different, but the way linking and libraries are done in C is an artifact of how the Unix operating system and executable formats were designed years ago. Learning how C links things helps you understand how your OS works and how it runs your programs.

To start off there are two basic types of libraries:

### static

You've made one of these when you used `ar` and `ranlib` to create the `libYOUR_LIBRARY.a` in the last exercise. This kind of library is nothing more than a container for a set of `.o` object files and their functions, and you can treat it like one big `.o` file when building your programs.

### dynamic

These typically end in `.so`, `.dll` or about 1 million other endings on OSX depending on the version and who happened to be working that day. Seriously though, OSX adds `.dylib`, `.bundle`, and `.framework` with not much distinction between the three. These files are built and then placed in a common location. When you run your program the OS dynamically loads these files and links them to your program on the fly.

I tend to like static libraries for small to medium sized projects because they are easier to deal with and work on more operating systems. I also like to put all of the code I can into a static library so that I can then link it to unit tests and to the file programs as needed.

Dynamic libraries are good for larger systems, when space is tight, or if you have a large number of programs that use common functionality. In this case you don't want to statically link all of the code for the common features to every program, so you put it in a dynamic library so that it is loaded only once for all of them.

In the previous exercise I laid out how to make a static library (a `.a` file), and that's what I'll use in the rest of the book. In this exercise I'm going to show you how to make a simple `.so` library, and how to dynamically load

it with the Unix `dlopen` system. I'll have you do this manually so that you understand everything that's actually happening, then the Extra Credit will be to use the `c-skeleton` skeleton to create it.

## Dynamically Loading A Shared Library

To do this I will create two source files. One will be used to make a `libex29.so` library, the other will be a program called `ex29` that can load this library and run functions from it.

```
#include <stdio.h>
#include <ctype.h>
#include "dbg.h"

int print_a_message(const char *msg)
{
    printf("A STRING: %s\n", msg);

    return 0;
}

int uppercase(const char *msg)
{
    int i = 0;

    // BUG: \0 termination problems
    for(i = 0; msg[i] != '\0'; i++) {
        printf("%c", toupper(msg[i]));
    }

    printf("\n");

    return 0;
}

int lowercase(const char *msg)
{
    int i = 0;

    // BUG: \0 termination problems
    for(i = 0; msg[i] != '\0'; i++) {
        printf("%c", tolower(msg[i]));
    }

    printf("\n");

    return 0;
}

int fail_on_purpose(const char *msg)
{
    return 1;
}
```

There's nothing fancy in there, although there's some bugs I'm leaving in on purpose to see if you've been paying attention. You'll fix those later.

What we want to do is use the functions `dlopen`, `dlsym` and `dlclose` to work with the above functions.

```
#include <stdio.h>
```

```

#include "dbg.h"
#include <dlfcn.h>

typedef int (*lib_function)(const char *data);

int main(int argc, char *argv[])
{
    int rc = 0;
    check(argc == 4, "USAGE: ex29 libex29.so function data");

    char *lib_file = argv[1];
    char *func_to_run = argv[2];
    char *data = argv[3];

    void *lib = dlopen(lib_file, RTLD_NOW);
    check(lib != NULL, "Failed to open the library %s: %s", lib_file, dlerror());

    lib_function func = dlsym(lib, func_to_run);
    check(func != NULL, "Did not find %s function in the library %s: %s", func_to_run, lib_file, dlerror());

    rc = func(data);
    check(rc == 0, "Function %s return %d for data: %s", func_to_run, rc, data);

    rc = dlclose(lib);
    check(rc == 0, "Failed to close %s", lib_file);

    return 0;
error:
    return 1;
}

```

I'll now break this down so you can see what's going on in this small bit of useful code:

ex29.c:5

I'll use this function pointer definition later to call functions in the library. This is nothing new, but make sure you understand what it's doing.

ex29.c:17

After the usual setup for a small program, I use the `dlopen` function to load up the library indicated by `lib_file`. This function returns a handle that we use later and works a lot like opening a file.

ex29.c:18

If there's an error, I do the usual check and exit, but notice at then end that I'm using `dlerror` to find out what the library related error was.

ex29.c:20

I use `dlsym` to get a function out of the `lib` by it's *string* name in `func_to_run`. This is the powerful part, since I'm dynamically getting a pointer to a function based on a string I got from the command line `argv`.

ex29.c:23

I then call the `func` function that was returned, and check its return value.

ex29.c:26

Finally, I close the library up just like I would a file. Usually you keep these open the whole time the program is running, so

closing at the end isn't as useful, but I'm demonstrating it here.

## What You Should See

Now that you know what this file does, here's a shell session of me building the `libex29.so`, `ex29` and then working with it. Follow along so you learn how these things are built manually.

```
# compile the lib file and make the .so
# you may need -fPIC here on some platforms. add that if you get an error
$ cc -c libex29.c -o libex29.o
$ cc -shared -o libex29.so libex29.o

# make the loader program
$ cc -Wall -g -DDEBUG ex29.c -ldl -o ex29

# try it out with some things that work
$ ex29 ./libex29.so print_a_message "hello there"
-bash: ex29: command not found
$ ./ex29 ./libex29.so print_a_message "hello there"
A STRING: hello there
$ ./ex29 ./libex29.so uppercase "hello there"
HELLO THERE
$ ./ex29 ./libex29.so lowercase "HELLO tHeRe"
hello there
$ ./ex29 ./libex29.so fail_on_purpose "i fail"
[ERROR] (ex29.c:23: errno: None) Function fail_on_purpose return 1 for data: i
fail

# try to give it bad args
$ ./ex29 ./libex29.so fail_on_purpose
[ERROR] (ex29.c:11: errno: None) USAGE: ex29 libex29.so function data

# try calling a function that is not there
$ ./ex29 ./libex29.so adfasfasdf asdfadff
[ERROR] (ex29.c:20: errno: None) Did not find adfasfasdf
function in the library libex29.so: dlsym(0x1076009b0, adfasfasdf): symbol n
ot found

# try loading a .so that is not there
$ ./ex29 ./libex.so adfasfasdf asdfadfas
[ERROR] (ex29.c:17: errno: No such file or directory) Failed to open
the library libex.so: dlopen(libex.so, 2): image not found
$
```

One thing that you may run into is that every OS, every version of every OS, and every compiler on every version of every OS, seems to want to change the way you build a shared library every other month that some new programmer thinks it's wrong. If the line I use to make the `libex29.so` file is wrong, then let me know and I'll add some comments for other platforms.

### NOTE

Sometimes you'll do what you think is normal and run this command `cc -Wall -g -DDEBUG -ldl ex29.c -o ex29` thinking everything will work, but nope. You see, on some platforms the order of where libraries goes makes them work or not, and for no real reason. On Debian or Ubuntu you have to do `cc -Wall`

`-g -DNDEBUG ex29.c -ldl -o ex29` for no reason at all. It's just the way it is, so since this works on OSX I'm doing it here, but in the future, if you link against a dynamic library and it can't find a function, try shuffling things around.

The irritation here is there is an actual platform difference on nothing more than order of command line arguments. On no rational planet should putting an `-ldl` at one position be different from another. It's an option, and having to know these things is incredibly annoying.

## How To Break It

Open `libex29.so` and edit it with an editor that can handle binary files. Change a couple bytes, then close it. Try to see if you can get the `dlopen` function to load it even though you've corrupted it.

## Extra Credit

- Were you paying attention to the bad code I have in the `libex29.c` functions? See how, even though I use a for-loop they still check for `'\0'` endings? Fix this so the functions always take a length for the string to work with inside the function.
- Take the `c-skeleton` skeleton, and create a new project for this exercise. Put the `libex29.c` file in the `src/` directory. Change the Makefile so that it builds this as `build/libex29.so`.
- Take the `ex29.c` file and put it in `tests/ex29_tests.c` so that it runs as a unit test. Make this all work, which means you have to change it so that it loads the `build/libex29.so` file and runs tests similar to what I did manually above.
- Read the `man dlopen` documentation and read about all the related functions. Try some of the other options to `dlopen` beside `RTLD_NOW`.

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 30: Automated Testing

Automated testing is used frequently in other languages like Python and Ruby, but rarely used in C. Part of the reason comes from the difficulty of automatically loading and testing pieces of C code. In this chapter we'll create a very small little testing "framework" and get your skeleton directory building an example test case.

The frameworks I'm going to use, and which you'll include in your `c-skeleton` skeleton is called "minunit" which started with code from a tiny snippet of code by [Jera Design](#). I then evolved it further, to be this:

```
#undef NDEBUG
#ifdef _minunit_h
#define _minunit_h

#include <stdio.h>
#include <dbg.h>
#include <stdlib.h>

#define mu_suite_start() char *message = NULL

#define mu_assert(test, message) if (!(test)) { log_err(message); return message; }
#define mu_run_test(test) debug("\n-----%s", " " #test); \
    message = test(); tests_run++; if (message) return message;

#define RUN_TESTS(name) int main(int argc, char *argv[]) {\
    argc = 1; \
    debug("----- RUNNING: %s", argv[0]);\
    printf("-----\nRUNNING: %s\n", argv[0]);\
    char *result = name();\
    if (result != 0) {\
        printf("FAILED: %s\n", result);\
    }\
    else {\
        printf("ALL TESTS PASSED\n");\
    }\
    printf("Tests run: %d\n", tests_run);\
    exit(result != 0);\
}

int tests_run;

#endif
```

There's mostly nothing left of the original, as now I'm using the `dbg.h` macros and I've created a large macro at the end for the boilerplate test runner. Even with this tiny amount of code we'll create a fully functioning unit test system you can use in your C code once it's combined with a shell script to run the tests.

## Wiring Up The Test Framework

To continue this exercise, you should have your `src/libex29.c` working

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

⚙  
HELP

Follow  
@lzssth

and that you completed the Exercise 29 extra credit where you got the `ex29.c` loader program to properly run. In Exercise 29 I had an extra credit to make it work like a unit test, but I'm going to start over and show you how to do that with `minunit.h`.

The first thing to do is create a simple empty unit test name `tests/libex29_tests.c` with this in it:

```
#include "minunit.h"

char *test_dlopen()
{
    return NULL;
}

char *test_functions()
{
    return NULL;
}

char *test_failures()
{
    return NULL;
}

char *test_dlclose()
{
    return NULL;
}

char *all_tests() {
    mu_suite_start();

    mu_run_test(test_dlopen);
    mu_run_test(test_functions);
    mu_run_test(test_failures);
    mu_run_test(test_dlclose);

    return NULL;
}

RUN_TESTS(all_tests);
```

This code is demonstrating the `RUN_TESTS` macro in `tests/minunit.h` and how to use the other test runner macros. I have the actual test functions stubbed out so that you can see how to structure a unit test. I'll break this file down first:

libex29\_tests.c:1

Include the `minunit.h` framework.

libex29\_tests.c:3-7

A first test. Tests are structured so they take no arguments and return a `char *` which is `NULL` on *success*. This is important because the other macros will be used to return an error message to the test runner.

libex29\_tests.c:9-25

More tests that are the same as the first one.

libex29\_tests.c:27

The runner function that will control all the other tests. It has the same form as any other test case, but it gets configured



with some additional gear.

libex29\_tests.c:28

Sets up some common stuff for a test with `mu_suite_start`.

libex29\_tests.c:30

This is how you say what test to run, using the `mu_run_test` macro.

libex29\_tests.c:35

After you say what tests to run, you then return `NULL` just like a normal test function.

libex29\_tests.c:38

Finally, you just use the big `RUN_TESTS` macro to wire up the `main` method with all the goodies and tell it to run the `all_tests` starter.

That's all there is to running a test, now you should try getting just this to run within the project skeleton. Here's what it looks like when I do it:

*not printable*

I first did a `make clean` and then I ran the build, which remade the template `libYOUR_LIBRARY.a` and `libYOUR_LIBRARY.so` files. Remember that you had to do this in the extra credit for Exercise 29, but just in case you didn't figure it out, here's the diff for the `Makefile` I'm using now:

```
diff --git a/code/c-skeleton/Makefile b/code/c-skeleton/Makefile
index 135d538..21b92bf 100644
--- a/code/c-skeleton/Makefile
+++ b/code/c-skeleton/Makefile
@@ -9,9 +9,10 @@ TEST_SRC=$(wildcard tests/*_tests.c)
 TESTS=$(patsubst %.c,%, $(TEST_SRC))

 TARGET=build/libYOUR_LIBRARY.a
+SO_TARGET=$(patsubst %.a,%.so, $(TARGET))

# The Target Build
-all: $(TARGET) tests
+all: $(TARGET) $(SO_TARGET) tests

dev: CFLAGS=-g -Wall -Isrc -Wall -Wextra $(OPTFLAGS)
dev: all
@@ -21,6 +22,9 @@ $(TARGET): build $(OBJECTS)
    ar rcs $@ $(OBJECTS)
    ranlib $@

+$(SO_TARGET): $(TARGET) $(OBJECTS)
+    $(CC) -shared -o $@ $(OBJECTS)
+
build:
    @mkdir -p build
    @mkdir -p bin
```

With those changes you should be now building everything and you can finally fill in the remaining unit test functions:

```
#include "minunit.h"
#include <dlfcn.h>

typedef int (*lib_function)(const char *data);
char *lib_file = "build/libYOUR_LIBRARY.so";
void *lib = NULL;

int check_function(const char *func_to_run, const char *data, int expected)
{
    lib_function func = dlsym(lib, func_to_run);
    check(func != NULL, "Did not find %s function in the library %s: %s", func
```

```

_to_run, lib_file, dlerror());

    int rc = func(data);
    check(rc == expected, "Function %s return %d for data: %s", func_to_run, r
c, data);

    return 1;
error:
    return 0;
}

char *test_dlopen()
{
    lib = dlopen(lib_file, RTLD_NOW);
    mu_assert(lib != NULL, "Failed to open the library to test.");

    return NULL;
}

char *test_functions()
{
    mu_assert(check_function("print_a_message", "Hello", 0), "print_a_message
failed.");
    mu_assert(check_function("uppercase", "Hello", 0), "uppercase failed.");
    mu_assert(check_function("lowercase", "Hello", 0), "lowercase failed.");

    return NULL;
}

char *test_failures()
{
    mu_assert(check_function("fail_on_purpose", "Hello", 1), "fail_on_purpose
should fail.");

    return NULL;
}

char *test_dlclose()
{
    int rc = dlcclose(lib);
    mu_assert(rc == 0, "Failed to close lib.");

    return NULL;
}

char *all_tests() {
    mu_suite_start();

    mu_run_test(test_dlopen);
    mu_run_test(test_functions);
    mu_run_test(test_failures);
    mu_run_test(test_dlclose);

    return NULL;
}

RUN_TESTS(all_tests);

```

Hopefully by now you can figure out what's going on, since there's nothing new in this except for the `check_function` function. This is a common pattern where I see that I'll be doing a chunk of code repeatedly, and then simply automate it either by creating a function or a macro for it. In this case I'm going to run functions in the `.so` I load so I just made a little function to do it.

## Extra Credit

- This works but it's probably a bit messy. Clean the `c-skeleton` directory up so that it has all these files, but remove any of the code related to Exercise 29. You should be able to copy this directory over and kickstart new projects without much editing.
- Study the `runtests.sh` and go read about `bash` syntax so you know what it does. Think you could write a C version of this script?

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 31: Debugging Code

I've already taught you about my awesome debug macros and you've been using them. When I debug code I use the `debug()` macro almost exclusively to analyze what's going on and track down the problem. In this exercise I'm going to teach you the basics of using gdb to inspect a simple program that runs and doesn't exit. You'll learn how to use gdb to attach to a running process, stop it, and see what's happening. After that I'll give you some little tips and tricks that you can use with gdb.

## Debug Printing Vs. GDB Vs. Valgrind

I approach debugging primarily with a "scientific method" style, where I come up with possible causes and then rule them out or prove they cause the defect. The problem many programmers have though is their panic and rush to solve a bug makes them feel like this approach will "slow them down". In their rush to solve they fail to notice that they're really just flailing around and gathering no useful information. I find that logging (debug printing) forces me to solve a bug scientifically and it's also just easier to gather information in more situations.

In addition to that, I also have these reasons for using debug printing as my primary debugging tool:

- You see an entire tracing of a program's execution with debug printing of variables which lets you track how things are going wrong. With gdb you have to place watch and debug statements all over for every thing you want and it's difficult to get a solid trace of the execution.
- The debug prints can stay in the code, and when you need them you can recompile and they come back. With gdb you have to configure the same information uniquely for every defect you have to hunt down.
- It's easier to turn on debug logging on a server that's not working right and then inspect the logs while it runs to see what's going on. System administrators know how to handle logging, they don't know how to use gdb.
- Printing things is just easier. Debuggers are always obtuse and weird with their own quirky interface and inconsistencies. There's nothing complicated about `debug("Yo, dis right? %d", my_stuff);`.
- Writing debug prints to find a defect forces you to actually analyze the code and use the scientific method. You can think of a debug usage as, "I hypothesize that the code is broken here." Then when you run it you get your hypothesis tested and if it's not broken then you can move to another part where it could be. This may seem like it takes longer, but it's actually faster because you go through a process of "differential diagnosis" and rule out possible causes until you find the real

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

⚙️  
HELP

Follow  
@lzssthw

one.

- Debug printing works better with unit testing. You can actually just compile the debugs in all the time while you work, and when a unit test explodes just go look at the logs any time. With gdb you'd have to rerun the unit test under gdb and then trace through it to see what's going on.
- With valgrind you get the equivalent of debug prints for many memory related errors, so you don't need to use something like gdb to find those defects anymore.

Despite all these reasons that I rely on `debug` over `gdb`, I still use `gdb` in a few situations and I think you should have any tool that helps you get your work done. Sometimes, you just have to connect to a broken program and poke around. Or, maybe you've got a server that's crashing and you can only get at core files to see why. In these and a few other cases, gdb is the way to go, and it's always good to have as many tools as possible to help solve problems.

I then break down when I use gdb vs. valgrind vs. debug printing like this:

- Valgrind is used to catch all memory errors. I use gdb if valgrind is having problems or if using valgrind would slow the program down too much.
- Print with debug to diagnose and fix defects related to logic or usage. This amounts to about 90% of the defects after you start using Valgrind.
- Use gdb for the remaining "mystery weird stuff" or emergency situations to gather information. If Valgrind isn't turning anything up and I can't even print out the information I need, then I bust out gdb and start poking around. My use of gdb in this case is entirely to gather information. Once I have an idea of what's going on I go back to writing a unit test to cause the defect, and then do print statements to find out why.

## A Debugging Strategy

This process will actually work with any debugging technique you're going to use, whether that's Valgrind, debug printing, or using a debugger. I'm going to describe it in terms of using `gdb` since it seems people skip this process the most when using debuggers, but use this for every bug until you only need it on the very difficult ones.

- Start a little text file called `notes.txt` and use it as a kind of "lab notes" for ideas, bugs, problems, etc.
- Before you use `gdb`, write out the bug you're going to fix and what could be causing it.
- For each cause, write out the files and functions where you think the cause is coming from, or just write that you don't know.
- Now start `gdb` and pick the first possible cause with good `file:function` possibilities and set breakpoints there.
- Use `gdb` to then run the program and confirm if that is the cause. The best way is to see if you can use the `set` command to either fix the program easily or cause the error immediately.
- If this isn't the cause, then mark in the `notes.txt` that it wasn't

and why. Move on to the next possible cause that's easiest to debug, and keep adding information you gather.

In case you haven't noticed, this is basically the scientific method. You write down a set of hypotheses, then you use debugging to prove or disprove them. This gives you insight into more possible causes and then eventually you find it. This process helps you avoid going over the same possible causes repeatedly even though you've found they aren't possible.

You can also do this with debug printing, the only difference is you actually write out your hypotheses in the source code where you think the problem is instead of the `notes.txt`. In a way, debug printing forces you to tackle bugs scientifically since you have to write out hypotheses as print statements.

## Using GDB

The program I'll debug in this exercise is just a while-loop that doesn't terminate correctly. I'm putting a small `usleep` call in it so that there's something interesting to troll through as well.

```
#include <unistd.h>

int main(int argc, char *argv[])
{
    int i = 0;

    while(i < 100) {
        usleep(3000);
    }

    return 0;
}
```

Compile this like normal and then start it under `gdb` like this: `gdb ./ex31`

Once it's running I want you to play around with these `gdb` commands to see what they do and how to use them.

`help COMMAND`

Get a short help with COMMAND.

`break file.c:(line|function)`

Sets a break point where you want to pause execution. You can give lines or function names to break at after the file.

`run ARGS`

Runs the program, using the ARGS as arguments to the program.

`cont`

Continues execution until a new breakpoint or error.

`step`

Step through the code, but move *into functions*. Use this to trace into a function and see what it's doing.

`next`

Just like `step`, but go *over functions* by just running them.

backtrace (or bt)

Does a "backtrace", which dumps the trace of function calls leading to the current point in the program. Very useful for figuring out how you got there, since it also prints the parameters that were passed to each function. It's also similar to what Valgrind reports when you have a memory error.

set var X = Y

Set variable X equal to Y.

print X

Prints out the value of X, and you can usually use C syntax to access the values of pointers and contents of structs.

ENTER

The ENTER key just repeats the last command.

quit

Exits `gdb`

Those are the majority of commands I use with `gdb`. Your job is to now play with these and `ex31` so you can get familiar with the output.

Once you're familiar with `gdb` you'll want to play with it some more. Try using it on more complicated programs like `devpkg` to see if you can alter the program's execution or analyze what it's doing.

## Process Attaching

The most useful thing about `gdb` is the ability to attach to a running program and debug it right there. When you have a crashing server or a GUI program, you can't usually start it under `gdb` like you just did. Instead, you have to start it, hope it doesn't crash right away, then attach to it and set a breakpoint. In this part of the exercise I'll show you how to do that.

After you exit `gdb` I want you to restart `ex31` if you stopped it, and then start another Terminal window so you can process attach to it. Process attaching is where you tell `gdb` to connect to a program that's already running so you can inspect it live. It stops the program and then you can walk through it, and when you're done it'll continue just like normal.

Here's a session of me doing it to `ex31`, stepping through it, then fixing the while-loop to make it exit.

```
$ ps ax | grep ex31
10026 s000  S+      0:00.11 ./ex31
10036 s001  R+      0:00.00 grep ex31

$ gdb ./ex31 10026
GNU gdb 6.3.50-20050815 (Apple version gdb-1705) (Fri Jul  1 10:50:06 UTC 2011)
Copyright 2004 Free Software Foundation, Inc.
GDB is free software, covered by the GNU General Public License, and you are
welcome to change it and/or distribute copies of it under certain conditions.
Type "show copying" to see the conditions.
There is absolutely no warranty for GDB.  Type "show warranty" for details.
This GDB was configured as "x86_64-apple-darwin"...Reading symbols for shared
libraries .. done

/Users/zedshaw/projects/books/learn-c-the-hard-way/code/10026: No such file or
directory
Attaching to program: `/Users/zedshaw/projects/books/learn-c-the-hard-way/code
/ex31', process 10026.
```

```

Reading symbols for shared libraries + done
Reading symbols for shared libraries ++..... done
Reading symbols for shared libraries + done
0x00007ffff862c9e42 in __semwait_signal ()

(gdb) break 8
Breakpoint 1 at 0x107babf14: file ex31.c, line 8.

(gdb) break ex31.c:11
Breakpoint 2 at 0x107babf1c: file ex31.c, line 12.

(gdb) cont
Continuing.

Breakpoint 1, main (argc=1, argv=0x7fff677aabd8) at ex31.c:8
8      while(i < 100) {

(gdb) p i
$1 = 0

(gdb) cont
Continuing.

Breakpoint 1, main (argc=1, argv=0x7fff677aabd8) at ex31.c:8
8      while(i < 100) {

(gdb) p i
$2 = 0

(gdb) list
3
4  int main(int argc, char *argv[])
5  {
6      int i = 0;
7
8      while(i < 100) {
9          usleep(3000);
10     }
11
12     return 0;

(gdb) set var i = 200

(gdb) p i
$3 = 200

(gdb) next

Breakpoint 2, main (argc=1, argv=0x7fff677aabd8) at ex31.c:12
12     return 0;

(gdb) cont
Continuing.

Program exited normally.
(gdb) quit
$

```

## NOTE

On OSX you may see a GUI prompt for the root password, and even after you give it you still get an error from `gdb` saying "Unable to access task for process-id XXX: (os/kern) failure." In that case stop both `gdb` and the `ex31` program, then start over and it should work as long as you



successfully entered the root password.

I'll walk through this session and explain what I did:

gdb:1

I use `ps` to find out what the process id is of the `ex31` I want to attach.

gdb:5

I'm attaching using `gdb ./ex31 PID` replacing PID with the process id I have.

gdb:6-19

`gdb` prints out a bunch of information about it's license and then all the things it's reading.

gdb:21

The program is attached and stopped at this point, so now I set a breakpoint at line 8 in the file with `break`. I'm assuming that I'm already in the file I want to break when I do this.

gdb:24

A better way to do a `break`, is give `file.c:line` format so you can be sure you did the right location. I do that in this `break`.

gdb:27

I use `cont` to continue processing until I hit a breakpoint.

gdb:30-31

The breakpoint is reached so `gdb` prints out variables I need to know about (`argc` and `argv`) and where it's stopped, then the line of code for the breakpoint.

gdb:33-34

I use the abbreviation for `print` "p" to print out the value of the `i` variable. It's 0.

gdb:36

Continue again to see if `i` changes.

gdb:42

Print out `i` again, and nope it's not changing.

gdb:45-55

Use `list` to see what the code is, and then I realize it's not exiting because I'm not incrementing `i`.

gdb:57

Confirm my hypothesis that `i` needs to change by using the `set` command to change it to be `i = 200`. This is one of the best features of `gdb` as it lets you "fix" a program really quick to see if you're right.

gdb:59

Print out `i` just to make sure it changed.

gdb:62

Use `next` to move to the next piece of code, and I see that the breakpoint at `ex31.c:12` is hit, so that means the while-loop exited. My hypothesis is correct, I need to make `i` change.

gdb:67

Use `cont` to continue and the program exits like normal.

gdb:71

I finally use `quit` to get out of `gdb`.

# GDB Tricks

Here's a list of simple tricks you can do with GDB:

`gdb --args`

Normally `gdb` takes arguments you give it and assumes they are for itself. Using `--args` passes them to the program.

`thread apply all bt`

Dumps a backtrace for *all* threads. Very useful.

`gdb --batch --ex r --ex bt --ex q --args`

Runs the program so that, if it bombs you get a backtrace.

?

Got one? Leave it in the comments.

## Extra Credit

- Find a graphical debugger and compare using it to raw `gdb`. These are useful when the program you're looking at is local, but they are pointless if you have to debug a program on a server.
- You can enable "core dumps" on your OS, and when a program crashes you'll get a core file. This core file is like a post-mortem of the program so you can load up what happened right at the crash and see what caused it. Change `ex31.c` so that it crashes after a few iterations, then try to get a core dump and analyze it.

### Interested In Python?

Python is also a great language.

**Learn Python The Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 32: Double Linked Lists

The purpose of this book is to teach you how your computer really works, and included in that is how various data structures and algorithms function. Computers by themselves don't do a lot of useful processing. To make them do useful things you need to structure the data and then organize processing on these structures. Other programming languages either include libraries that implement all of these structures, or they have direct syntax for them. C makes you implement all the data structures you need yourself, which makes it the perfect language to learn how they actually work.

My goal in teaching you these data structures and these algorithms is to help you do three things:

- Understand what is really going on in Python, Ruby, or JavaScript code like: `data = {"name": "Zed"}`
- Get even better at C code by applying what you know to a set of solved problems using the data structures.
- Learn a core set of data structures and algorithms so that you are better informed about what ones work best in certain situations.

## What Are Data Structures

The name "data structure" is self-explanatory. It is an organization of data that fits a certain model. Maybe the model is designed to allow processing the data in a new way. Maybe it's just organized to store it on disk efficiently. In this book I'll follow a simple pattern for making data structures that works reliably:

- Define a struct for the main "outer structure".
- Define a struct for the contents, usually nodes with links between them.
- Create functions that operate on these two.

There's other styles of data structures in C, but this pattern works well and is consistent for most data structures you'll make.

## Making The Library

For the rest of this book you'll be creating a library that you can use when you're done with this book. This library will have the following elements:

- Header (.h) files for each data structure.
- Implementation (.c) files for the algorithms.
- Unit tests that test all of them to make sure they keep working.
- Documentation we'll autogenerate from the header files.

You already have the `c-skeleton` so use it to create a `liblcthw` project:

```
$ cp -r c-skeleton liblcthw
$ cd liblcthw/
$ ls
LICENSE          Makefile          README.md         bin               build
    src              tests
$ vim Makefile
$ ls src/
dbg.h              libex29.c          libex29.o
$ mkdir src/lcthw
$ mv src/dbg.h src/lcthw
$ vim tests/minunit.h
$ rm src/libex29.* tests/libex29*
$ make clean
rm -rf build tests/libex29_tests
rm -f tests/tests.log
find . -name "*.gc*" -exec rm {} \;
rm -rf `find . -name "*.dSYM" -print`
$ ls tests/
minunit.h  runtests.sh
$
```

In this session I'm doing the following:

- Copy the `c-skeleton` over.
- Edit the Makefile to change `libYOUR_LIBRARY.a` to `liblcthw.a` as the new `TARGET`.
- Make the `src/lcthw` directory where we'll put our code.
- Move the `src/dbg.h` into this new directory.
- Edit `tests/minunit.h` so that it uses `#include <lcthw/dbg.h>` as the include.
- Get rid of the source and test files we don't need for `libex29.*`.
- Clean up everything that's left over.

With that you're ready to start building the library, and the first data structure I'll build is the Double Linked List.

## Double Linked Lists

The first data structure we'll add to `liblcthw` is a double linked list. This is the simplest data structure you can make, and it has useful properties for certain operations. A linked list works by nodes having pointers to their next or previous element. A "double linked list" contains pointers to both, while a "single linked list" only points at the next element.

Because each node has pointers to the next and previous, and because you keep track of the first and last element of the list, you can do some operations very quickly. Anything that involves inserting or deleting an element will be very fast. They are also easy to implement by most people.

The main disadvantage of a linked list is that traversing it involves processing every single pointer along the way. This means that searching, most sorting, or iterating over the elements will be slow. It also means that you can't really jump to random parts of the list. If you had an array of elements you could just index right into the middle of the list, but a linked list uses a stream of pointers. That means if you want the 10th element, you have to go through elements 1-9.

## Definition

As I said in the introduction to this exercise, the process to follow is to first write a header file with the right C struct statements in it.

```
#ifndef lcthw_List_h
#define lcthw_List_h

#include <stdlib.h>

struct ListNode;

typedef struct ListNode {
    struct ListNode *next;
    struct ListNode *prev;
    void *value;
} ListNode;

typedef struct List {
    int count;
    ListNode *first;
    ListNode *last;
} List;

List *List_create();
void List_destroy(List *list);
void List_clear(List *list);
void List_clear_destroy(List *list);

#define List_count(A) ((A)->count)
#define List_first(A) ((A)->first != NULL ? (A)->first->value : NULL)
#define List_last(A) ((A)->last != NULL ? (A)->last->value : NULL)

void List_push(List *list, void *value);
void *List_pop(List *list);

void List_unshift(List *list, void *value);
void *List_shift(List *list);

void *List_remove(List *list, ListNode *node);

#define LIST_FOREACH(L, S, M, V) ListNode *_node = NULL;\
    ListNode *V = NULL;\
    for(V = _node = L->S; _node != NULL; V = _node = _node->M)

#endif
```

The first thing I do is create two structs for the `ListNode` and the `List` that will contain those nodes. This creates the data structure I'll use in the functions and macros I define after that. If you read through these functions they seem rather simple. I'll be explaining them when I cover the implementation, but hopefully you can guess what they do.

How the data structure works is each `ListNode` has three components:

- A value, which is a pointer to anything and stores the thing we want to put in the list.
- A `ListNode *next` pointer which points at another `ListNode` that holds the next element in the list.
- A `ListNode *prev` that holds the previous element. Complex right? Calling the previous thing "previous". I could have used "anterior" and "posterior" but only a jerk would do that.

The `List` struct is then nothing more than a container for these `ListNode` structs that have been linked together in a chain. It keeps

track of the `count`, `first` and `last` element of the list.

Finally, take a look at `src/lcthw/list.h:37` where I define the `LIST_FOREACH` macro. This is a common idiom where you make a macro that generates iteration code so people can't mess it up. Getting this kind of processing right can be difficult with data structures, so writing macros helps people out. You'll see how I use this when I talk about the implementation.

## Implementation

Once you understand that, you mostly understand how a double linked list works. It is nothing more than nodes with two pointers to the next and previous element of the list. You can then write the `src/lcthw/list.c` code to see how each operation is implemented.

```
#include <lcthw/list.h>
#include <lcthw/dbg.h>

List *List_create()
{
    return calloc(1, sizeof(List));
}

void List_destroy(List *list)
{
    LIST_FOREACH(list, first, next, cur) {
        if(cur->prev) {
            free(cur->prev);
        }
    }

    free(list->last);
    free(list);
}

void List_clear(List *list)
{
    LIST_FOREACH(list, first, next, cur) {
        free(cur->value);
    }
}

void List_clear_destroy(List *list)
{
    List_clear(list);
    List_destroy(list);
}

void List_push(List *list, void *value)
{
    ListNode *node = calloc(1, sizeof(ListNode));
    check_mem(node);

    node->value = value;

    if(list->last == NULL) {
        list->first = node;
        list->last = node;
    } else {
        list->last->next = node;
        node->prev = list->last;
        list->last = node;
    }
}
```

```

    }

    list->count++;

error:
    return;
}

void *List_pop(List *list)
{
    ListNode *node = list->last;
    return node != NULL ? List_remove(list, node) : NULL;
}

void List_unshift(List *list, void *value)
{
    ListNode *node = calloc(1, sizeof(ListNode));
    check_mem(node);

    node->value = value;

    if(list->first == NULL) {
        list->first = node;
        list->last = node;
    } else {
        node->next = list->first;
        list->first->prev = node;
        list->first = node;
    }

    list->count++;

error:
    return;
}

void *List_shift(List *list)
{
    ListNode *node = list->first;
    return node != NULL ? List_remove(list, node) : NULL;
}

void *List_remove(List *list, ListNode *node)
{
    void *result = NULL;

    check(list->first && list->last, "List is empty.");
    check(node, "node can't be NULL");

    if(node == list->first && node == list->last) {
        list->first = NULL;
        list->last = NULL;
    } else if(node == list->first) {
        list->first = node->next;
        check(list->first != NULL, "Invalid list, somehow got a first that is NULL.");
        list->first->prev = NULL;
    } else if (node == list->last) {
        list->last = node->prev;
        check(list->last != NULL, "Invalid list, somehow got a next that is NULL.");
        list->last->next = NULL;
    } else {
        ListNode *after = node->next;
        ListNode *before = node->prev;
        after->prev = before;
        before->next = after;
    }

    list->count--;
}

```

```

        result = node->value;
        free(node);

error:
    return result;
}

```

I then implement all of the operations on a double linked list that can't be done with simple macros. Rather than cover every tiny little line of this file, I'm going to give high-level overview of every operation in both the `list.h` and `list.c` file, then leave you to read the code.

`list.h:List_count`

Returns the number of elements in the list, which is maintained as elements are added and removed.

`list.h:List_first`

Returns the first element of the list, but does not remove it.

`list.h:List_last`

Returns the last element of the list, but does not remove it.

`list.h:LIST_FOREACH`

Iterates over the elements in the list.

`list.c:List_create`

Simply creates the main `List` struct.

`list.c:List_destroy`

Destroys a `List` and any elements it might have.

`list.c:List_clear`

Convenience function for freeing the *values* in each node, not the nodes.

`list.c:List_clear_destroy`

Clears and destroys a list. It's not very efficient since it loops through them twice.

`list.c:List_push`

The first operation that demonstrates the advantage of a linked list. It adds a new element to the end of the list, and because that's just a couple of pointer assignments, does it very fast.

`list.c:List_pop`

The inverse of `List_push`, this takes the last element off and returns it.

`list.c:List_unshift`

The other thing you can easily do to a linked list is add elements to the *front* of the list very fast. In this case I call that `List_unshift` for lack of a better term.

`list.c:List_shift`

Just like `List_pop`, this removes the first element and returns it.

`list.c:List_remove`

This is actually doing all of the removal when you do `List_pop` or `List_shift`. Something that seems to always be difficult in data structures is removing things, and this function is no different. It has to handle quite a few conditions depending on if the element being removed is at the front; the end; both front and end; or middle.

Most of these functions are nothing special, and you should be able to easily digest this and understand it from just the code. You should definitely focus on how the `LIST_FOREACH` macro is used in `List_destroy` so you can understand how much it simplifies this



common operation.

## Tests

After you have those compiling it's time to create the test that makes sure they operate correctly.

```
#include "minunit.h"
#include <lcthw/list.h>
#include <assert.h>

static List *list = NULL;
char *test1 = "test1 data";
char *test2 = "test2 data";
char *test3 = "test3 data";

char *test_create()
{
    list = List_create();
    mu_assert(list != NULL, "Failed to create list.");

    return NULL;
}

char *test_destroy()
{
    List_clear_destroy(list);

    return NULL;
}

char *test_push_pop()
{
    List_push(list, test1);
    mu_assert(List_last(list) == test1, "Wrong last value.");

    List_push(list, test2);
    mu_assert(List_last(list) == test2, "Wrong last value.");

    List_push(list, test3);
    mu_assert(List_last(list) == test3, "Wrong last value.");
    mu_assert(List_count(list) == 3, "Wrong count on push.");

    char *val = List_pop(list);
    mu_assert(val == test3, "Wrong value on pop.");

    val = List_pop(list);
    mu_assert(val == test2, "Wrong value on pop.");

    val = List_pop(list);
    mu_assert(val == test1, "Wrong value on pop.");
    mu_assert(List_count(list) == 0, "Wrong count after pop.");

    return NULL;
}

char *test_unshift()
{
    List_unshift(list, test1);
    mu_assert(List_first(list) == test1, "Wrong first value.");

    List_unshift(list, test2);
    mu_assert(List_first(list) == test2, "Wrong first value.");
}
```

```

List_unshift(list, test3);
mu_assert(List_first(list) == test3, "Wrong last value.");
mu_assert(List_count(list) == 3, "Wrong count on unshift.");

return NULL;
}

char *test_remove()
{
    // we only need to test the middle remove case since push/shift
    // already tests the other cases

    char *val = List_remove(list, list->first->next);
    mu_assert(val == test2, "Wrong removed element.");
    mu_assert(List_count(list) == 2, "Wrong count after remove.");
    mu_assert(List_first(list) == test3, "Wrong first after remove.");
    mu_assert(List_last(list) == test1, "Wrong last after remove.");

    return NULL;
}

char *test_shift()
{
    mu_assert(List_count(list) != 0, "Wrong count before shift.");

    char *val = List_shift(list);
    mu_assert(val == test3, "Wrong value on shift.");

    val = List_shift(list);
    mu_assert(val == test1, "Wrong value on shift.");
    mu_assert(List_count(list) == 0, "Wrong count after shift.");

    return NULL;
}

char *all_tests() {
    mu_suite_start();

    mu_run_test(test_create);
    mu_run_test(test_push_pop);
    mu_run_test(test_unshift);
    mu_run_test(test_remove);
    mu_run_test(test_shift);
    mu_run_test(test_destroy);

    return NULL;
}

RUN_TESTS(all_tests);

```

This test simply goes through every operation and makes sure it works. I use a simplification in the test where I create just one `List *list` for the whole program, then have the tests work on it. This saves the trouble of building a `List` for every test, but it could mean that some tests only pass because of how the previous test ran. In this case I try to make each test keep the list clear or actually use the previous test's results.

## What You Should See

If you did everything right, then when you do a build and run the unit tests it should look like this:

```
$ make
```

```
cc -g -O2 -Wall -Wextra -Isrc -rdynamic -DNDEBUG -fPIC -c -o src/lcthw/list.o src/lcthw/list.c
ar rcs build/liblcthw.a src/lcthw/list.o
ranlib build/liblcthw.a
cc -shared -o build/liblcthw.so src/lcthw/list.o
cc -g -O2 -Wall -Wextra -Isrc -rdynamic -DNDEBUG build/liblcthw.a tests/list_tests.c -o tests/list_tests
sh ./tests/runtests.sh
Running unit tests:
----
RUNNING: ./tests/list_tests
ALL TESTS PASSED
Tests run: 6
tests/list_tests PASS
$
```

Make sure 6 tests ran, that it builds without warnings or errors, and that it's making the `build/liblcthw.a` and `build/liblcthw.so` files.

## How To Improve It

Instead of breaking this, I'm going to tell you how to improve the code:

- You can make `List_clear_destroy` more efficient by using `LIST_FOREACH` and doing both `free` calls inside one loop.
- You can add asserts for preconditions that it isn't given a `NULL` value for the `List *list` parameters.
- You can add invariants that check the list's contents are always correct, such as `count` is never `< 0`, and if `count > 0` then `first` isn't `NULL`.
- You can add documentation to the header file in the form of comments before each struct, function, and macro that describes what it does.

These amount to going through the defensive programming practices I talked about and "hardening" this code against flaws or improving usability. Go ahead and do these things, then find as many other ways to improve the code.

## Extra Credit

- Research double vs. single linked lists and when one is preferred over the other.
- Research the limitations of a double linked list. For example, while they are efficient for inserting and deleting elements, they are very slow for iterating over them all.
- What operations are missing that you can imagine needing? Some examples are copying, joining, splitting. Implement these operations and write the unit tests for them.

**Interested  
In  
Python?**

Python is also a great language.

**Learn  
Python The  
Hard Way**

**Interested In Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 33: Linked List Algorithms

I'm going to cover two algorithms you can do on a linked list that involve sorting. I'm going to warn you first that if you need to sort the data, then don't use a linked list. These are horrible for sorting things, and there's much better data structures you can use if that's a requirement. I'm covering these two algorithms because they are slightly difficult to pull off with a linked list and get you thinking about manipulating them efficiently.

In the interest of writing this book, I'm going to put the algorithms in two different files `list_algos.h` and `list_algos.c` then write a test in `list_algos_test.c`. For now just follow my structure, as it does keep things clean, but if you ever work on other libraries remember this isn't a common structure.

In this exercise I'm going to also give you an extra challenge and I want you to try not to cheat. I'm going to give you the *unit test* first, and I want you to type it in. Then I want you to try and implement the two algorithms based on their descriptions in Wikipedia before seeing if your code is like my code.

## Bubble And Merge Sort

You know what's awesome about the Internet? I can just link you to the [Bubble Sort page](#) [Merge Sort page](#) and tell you to read that. Man, that saves me a boat load of typing. Now I can tell you how to actually implement each of these using the pseudo-code they have there. Here's how you can tackle an algorithm like this:

- Read the description and look at any visualizations it has.
- Either draw the algorithm on paper using boxes and lines, or actually take a deck of numbered cards (like Poker Cards) and try to do the algorithm manually. This gives you a concrete demonstration of how the algorithm works.
- Create the skeleton functions in your `list_algos.c` file and make a working `list_algos.h` file, then setup your test harness.
- Write your first failing test and get everything to compile.
- Go back to the Wikipedia page and copy-paste the pseudo-code (not the C code!) into the guts of the first function you're making.
- Translate this pseudo-code into good C code like I've taught you, using your unit test to make sure it's working.
- Fill out some more tests for edge cases like, empty lists, already sorted lists, etc.
- Repeat for the next algorithm and test.

I just gave you the secret to figuring out most of the algorithms out there,

that is until you get to some of the more insane ones. In this case you're just doing the Bubble and Merge Sorts from Wikipedia, but those will be good starters.

## The Unit Test

Here is the unit test you should try to get passing:

```
#include "minunit.h"
#include <lcthw/list_algos.h>
#include <assert.h>
#include <string.h>

char *values[] = {"XXXX", "1234", "abcd", "xjvef", "NDSS"};
#define NUM_VALUES 5

List *create_words()
{
    int i = 0;
    List *words = List_create();

    for(i = 0; i < NUM_VALUES; i++) {
        List_push(words, values[i]);
    }

    return words;
}

int is_sorted(List *words)
{
    LIST_FOREACH(words, first, next, cur) {
        if(cur->next && strcmp(cur->value, cur->next->value) > 0) {
            debug("%s %s", (char *)cur->value, (char *)cur->next->value);
            return 0;
        }
    }

    return 1;
}

char *test_bubble_sort()
{
    List *words = create_words();

    // should work on a list that needs sorting
    int rc = List_bubble_sort(words, (List_compare)strcmp);
    mu_assert(rc == 0, "Bubble sort failed.");
    mu_assert(is_sorted(words), "Words are not sorted after bubble sort.");

    // should work on an already sorted list
    rc = List_bubble_sort(words, (List_compare)strcmp);
    mu_assert(rc == 0, "Bubble sort of already sorted failed.");
    mu_assert(is_sorted(words), "Words should be sort if already bubble sorted.");

    List_destroy(words);

    // should work on an empty list
    words = List_create(words);
    rc = List_bubble_sort(words, (List_compare)strcmp);
    mu_assert(rc == 0, "Bubble sort failed on empty list.");
    mu_assert(is_sorted(words), "Words should be sorted if empty.");

    List_destroy(words);

    return NULL;
}
```

```

char *test_merge_sort()
{
    List *words = create_words();

    // should work on a list that needs sorting
    List *res = List_merge_sort(words, (List_compare)strcmp);
    mu_assert(is_sorted(res), "Words are not sorted after merge sort.");

    List *res2 = List_merge_sort(res, (List_compare)strcmp);
    mu_assert(is_sorted(res), "Should still be sorted after merge sort.");
    List_destroy(res2);
    List_destroy(res);

    List_destroy(words);
    return NULL;
}

char *all_tests()
{
    mu_suite_start();

    mu_run_test(test_bubble_sort);
    mu_run_test(test_merge_sort);

    return NULL;
}

RUN_TESTS(all_tests);

```

I suggest that you start with the bubble sort and get that working, then move on to the merge sort. What I would do is lay out the function prototypes and skeletons that get all three files compiling, but not passing the test. Then just fill in the implementation until it starts working.

## The Implementation

Are you cheating? In future exercises I will do exercises where I just give you a unit test and tell you to implement it, so it'll be good practice for you to not look at this code until you get your own working. Here's the code for the `list_algos.c` and `list_algos.h`:

```

#ifndef lcthw_List_algos_h
#define lcthw_List_algos_h

#include <lcthw/list.h>

typedef int (*List_compare)(const void *a, const void *b);

int List_bubble_sort(List *list, List_compare cmp);

List *List_merge_sort(List *list, List_compare cmp);

#endif

#include <lcthw/list_algos.h>
#include <lcthw/dbg.h>

inline void ListNode_swap(ListNode *a, ListNode *b)
{
    void *temp = a->value;
    a->value = b->value;
    b->value = temp;
}

```

```

int List_bubble_sort(List *list, List_compare cmp)
{
    int sorted = 1;

    if(List_count(list) <= 1) {
        return 0; // already sorted
    }

    do {
        sorted = 1;
        LIST_FOREACH(list, first, next, cur) {
            if(cur->next) {
                if(cmp(cur->value, cur->next->value) > 0) {
                    ListNode_swap(cur, cur->next);
                    sorted = 0;
                }
            }
        }
    } while(!sorted);

    return 0;
}

inline List *List_merge(List *left, List *right, List_compare cmp)
{
    List *result = List_create();
    void *val = NULL;

    while(List_count(left) > 0 || List_count(right) > 0) {
        if(List_count(left) > 0 && List_count(right) > 0) {
            if(cmp(List_first(left), List_first(right)) <= 0) {
                val = List_shift(left);
            } else {
                val = List_shift(right);
            }

            List_push(result, val);
        } else if(List_count(left) > 0) {
            val = List_shift(left);
            List_push(result, val);
        } else if(List_count(right) > 0) {
            val = List_shift(right);
            List_push(result, val);
        }
    }

    return result;
}

List *List_merge_sort(List *list, List_compare cmp)
{
    if(List_count(list) <= 1) {
        return list;
    }

    List *left = List_create();
    List *right = List_create();
    int middle = List_count(list) / 2;

    LIST_FOREACH(list, first, next, cur) {
        if(middle > 0) {
            List_push(left, cur->value);
        } else {
            List_push(right, cur->value);
        }

        middle--;
    }
}

```

```

List *sort_left = List_merge_sort(left, cmp);
List *sort_right = List_merge_sort(right, cmp);

if(sort_left != left) List_destroy(left);
if(sort_right != right) List_destroy(right);

return List_merge(sort_left, sort_right, cmp);
}

```

The bubble sort isn't too bad to figure out, although it is really slow. The merge sort is much more complicated, and honestly I could probably spend a bit more time optimizing this code if I wanted to sacrifice clarity.

There is another way to implement merge sort using a "bottom up" method, but it's a little harder to understand so I didn't do it. As I've already said, sorting algorithms on linked lists are entirely pointless. You could spend all day trying to make this faster and it will still be slower than almost any other sortable data structure. The nature of linked lists is such that you simply don't use them if you need to sort things.

## What You Should See

If everything works then you should get something like this:

```

$ make clean all
rm -rf build src/lcthw/list.o src/lcthw/list_algos.o tests/list_algos_tests te
sts/list_tests
rm -f tests/tests.log
find . -name "*.gc*" -exec rm {} \;
rm -rf `find . -name "*.dSYM" -print`
cc -g -O2 -Wall -Wextra -Isrc -rdynamic -DNDEBUG -fPIC -c -o src/lcthw/list
.o src/lcthw/list.c
cc -g -O2 -Wall -Wextra -Isrc -rdynamic -DNDEBUG -fPIC -c -o src/lcthw/list
_algos.o src/lcthw/list_algos.c
ar rcs build/liblcthw.a src/lcthw/list.o src/lcthw/list_algos.o
ranlib build/liblcthw.a
cc -shared -o build/liblcthw.so src/lcthw/list.o src/lcthw/list_algos.o
cc -g -O2 -Wall -Wextra -Isrc -rdynamic -DNDEBUG build/liblcthw.a tests/li
st_algos_tests.c -o tests/list_algos_tests
cc -g -O2 -Wall -Wextra -Isrc -rdynamic -DNDEBUG build/liblcthw.a tests/li
st_tests.c -o tests/list_tests
sh ./tests/runtests.sh
Running unit tests:
----
RUNNING: ./tests/list_algos_tests
ALL TESTS PASSED
Tests run: 2
tests/list_algos_tests PASS
----
RUNNING: ./tests/list_tests
ALL TESTS PASSED
Tests run: 6
tests/list_tests PASS
$

```

After this exercise I'm not going to show you this output unless it's necessary to show you how it works. From now on you should know that I ran the tests and they all passed and everything compiled.

## How To Improve It

Going back to the description of the algorithms, there's several ways to improve these implementations, and there's a few obvious ones:



- The merge sort does a crazy amount of copying and creating lists, find ways to reduce this.
- The bubble sort Wikipedia description mentions a few optimizations, implement them.
- Can you use the `List_split` and `List_join` (if you implemented them) to improve merge sort?
- Go through all the defensive programming checks and improve the robustness of this implementation, protecting against bad `NULL` pointers, and create an optional debug level invariant that does what `is_sorted` does after a sort.

## Extra Credit

- Create a unit test that compares the performance of the two algorithms. You'll want to look at `man 3 time` for a basic timer function, and you'll want to run enough iterations to at least have a few seconds of samples.
- Play with the amount of data in the lists that need to be sorted and see if that changes your timing.
- Find a way to simulate filling different sized random lists and measuring how long they take, then graph it and see how it compares to the description of the algorithm.
- Try to explain why sorting linked lists is a really bad idea.
- Implement a `List_insert_sorted` that will take a given value, and using the `List_compare`, insert the element at the right position so that the list is always sorted. How does using this method compare to sorting a list after you've built it?
- Try implementing the "bottom up" merge sort on the wikipedia page. The code there is already C so it should be easy to recreate, but try to understand how it's working compared to the slower one I have here.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 34: Dynamic Array

This is an array that grows on its own and has most of the same features as a linked list. It will usually take up less space, run faster, and has other beneficial properties. This exercise will cover a few of the disadvantages like very slow removal from the front, with a solution (just do it at the end).

A dynamic array is simply an array of `void **` pointers that is pre-allocated in one shot and that point at the data. In the linked list you had a full struct that stored the `void *value` pointer, but in a dynamic array there's just a single array with all of them. This means you don't need any other pointers for next and previous records since you can just index into it directly.

To start, I'll give you the header file you should type up for the implementation:

```
#ifndef _DArray_h
#define _DArray_h
#include <stdlib.h>
#include <assert.h>
#include <lcthw/dbg.h>

typedef struct DArray {
    int end;
    int max;
    size_t element_size;
    size_t expand_rate;
    void **contents;
} DArray;

DArray *DArray_create(size_t element_size, size_t initial_max);

void DArray_destroy(DArray *array);

void DArray_clear(DArray *array);

int DArray_expand(DArray *array);

int DArray_contract(DArray *array);

int DArray_push(DArray *array, void *el);

void *DArray_pop(DArray *array);

void DArray_clear_destroy(DArray *array);

#define DArray_last(A) ((A)->contents[(A)->end - 1])
#define DArray_first(A) ((A)->contents[0])
#define DArray_end(A) ((A)->end)
#define DArray_count(A) DArray_end(A)
#define DArray_max(A) ((A)->max)

#define DEFAULT_EXPAND_RATE 300

static inline void DArray_set(DArray *array, int i, void *el)
```

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@l2sthw

```

{
    check(i < array->max, "darray attempt to set past max");
    if(i > array->end) array->end = i;
    array->contents[i] = el;
error:
    return;
}

static inline void *DArray_get(DArray *array, int i)
{
    check(i < array->max, "darray attempt to get past max");
    return array->contents[i];
error:
    return NULL;
}

static inline void *DArray_remove(DArray *array, int i)
{
    void *el = array->contents[i];

    array->contents[i] = NULL;

    return el;
}

static inline void *DArray_new(DArray *array)
{
    check(array->element_size > 0, "Can't use DArray_new on 0 size darrays.");

    return calloc(1, array->element_size);
error:
    return NULL;
}

#define DArray_free(E) free((E))

#endif

```

This header file is showing you a new technique where I put `static inline` functions right in the header. These function definitions will work similar to the `#define` macros you've been making, but they're cleaner and easier to write. If you need to create a block of code for a macro and you don't need code generation, then use a `static inline` function.

Compare this technique to the `LIST_FOREACH` that *generates* a proper for-loop for a list. This would be impossible to do with a `static inline` function because it actually has to generate the inner block of code for the loop. The only way to do that is with a callback function, but that's not as fast and is harder to use.

I'll then change things up and have you create the unit test for `DArray`:

```

#include "minunit.h"
#include <lcthw/darray.h>

static DArray *array = NULL;
static int *val1 = NULL;
static int *val2 = NULL;

char *test_create()
{
    array = DArray_create(sizeof(int), 100);
    mu_assert(array != NULL, "DArray_create failed.");
    mu_assert(array->contents != NULL, "contents are wrong in darray");
    mu_assert(array->end == 0, "end isn't at the right spot");
    mu_assert(array->element_size == sizeof(int), "element size is wrong.");
    mu_assert(array->max == 100, "wrong max length on initial size");
}

```

```

    return NULL;
}

char *test_destroy()
{
    DArray_destroy(array);

    return NULL;
}

char *test_new()
{
    val1 = DArray_new(array);
    mu_assert(val1 != NULL, "failed to make a new element");

    val2 = DArray_new(array);
    mu_assert(val2 != NULL, "failed to make a new element");

    return NULL;
}

char *test_set()
{
    DArray_set(array, 0, val1);
    DArray_set(array, 1, val2);

    return NULL;
}

char *test_get()
{
    mu_assert(DArray_get(array, 0) == val1, "Wrong first value.");
    mu_assert(DArray_get(array, 1) == val2, "Wrong second value.");

    return NULL;
}

char *test_remove()
{
    int *val_check = DArray_remove(array, 0);
    mu_assert(val_check != NULL, "Should not get NULL.");
    mu_assert(*val_check == *val1, "Should get the first value.");
    mu_assert(DArray_get(array, 0) == NULL, "Should be gone.");
    DArray_free(val_check);

    val_check = DArray_remove(array, 1);
    mu_assert(val_check != NULL, "Should not get NULL.");
    mu_assert(*val_check == *val2, "Should get the first value.");
    mu_assert(DArray_get(array, 1) == NULL, "Should be gone.");
    DArray_free(val_check);

    return NULL;
}

char *test_expand_contract()
{
    int old_max = array->max;
    DArray_expand(array);
    mu_assert((unsigned int)array->max == old_max + array->expand_rate, "Wrong
size after expand.");

    DArray_contract(array);
    mu_assert((unsigned int)array->max == array->expand_rate + 1, "Should stay
at the expand_rate at least.");

    DArray_contract(array);
    mu_assert((unsigned int)array->max == array->expand_rate + 1, "Should stay
at the expand_rate at least.");
}

```

```

        return NULL;
    }

    char *test_push_pop()
    {
        int i = 0;
        for(i = 0; i < 1000; i++) {
            int *val = DArray_new(array);
            *val = i * 333;
            DArray_push(array, val);
        }

        mu_assert(array->max == 1201, "Wrong max size.");

        for(i = 999; i >= 0; i--) {
            int *val = DArray_pop(array);
            mu_assert(val != NULL, "Shouldn't get a NULL.");
            mu_assert(*val == i * 333, "Wrong value.");
            DArray_free(val);
        }

        return NULL;
    }

    char * all_tests() {
        mu_suite_start();

        mu_run_test(test_create);
        mu_run_test(test_new);
        mu_run_test(test_set);
        mu_run_test(test_get);
        mu_run_test(test_remove);
        mu_run_test(test_expand_contract);
        mu_run_test(test_push_pop);
        mu_run_test(test_destroy);

        return NULL;
    }

    RUN_TESTS(all_tests);

```

This shows you how all of the operations are used, which then makes implementing the `DArray` much easier:

```

#include <lcthw/darray.h>
#include <assert.h>

DArray *DArray_create(size_t element_size, size_t initial_max)
{
    DArray *array = malloc(sizeof(DArray));
    check_mem(array);
    array->max = initial_max;
    check(array->max > 0, "You must set an initial_max > 0.");

    array->contents = calloc(initial_max, sizeof(void *));
    check_mem(array->contents);

    array->end = 0;
    array->element_size = element_size;
    array->expand_rate = DEFAULT_EXPAND_RATE;

    return array;
error:
    if(array) free(array);
    return NULL;
}

```

```

void DArray_clear(DArray *array)
{
    int i = 0;
    if(array->element_size > 0) {
        for(i = 0; i < array->max; i++) {
            if(array->contents[i] != NULL) {
                free(array->contents[i]);
            }
        }
    }
}

static inline int DArray_resize(DArray *array, size_t newsize)
{
    array->max = newsize;
    check(array->max > 0, "The newsize must be > 0.");

    void *contents = realloc(array->contents, array->max * sizeof(void *));
    // check contents and assume realloc doesn't harm the original on error

    check_mem(contents);

    array->contents = contents;

    return 0;
error:
    return -1;
}

int DArray_expand(DArray *array)
{
    size_t old_max = array->max;
    check(DArray_resize(array, array->max + array->expand_rate) == 0,
        "Failed to expand array to new size: %d",
        array->max + (int)array->expand_rate);

    memset(array->contents + old_max, 0, array->expand_rate + 1);
    return 0;
error:
    return -1;
}

int DArray_contract(DArray *array)
{
    int new_size = array->end < (int)array->expand_rate ? (int)array->expand_r
ate : array->end;

    return DArray_resize(array, new_size + 1);
}

void DArray_destroy(DArray *array)
{
    if(array) {
        if(array->contents) free(array->contents);
        free(array);
    }
}

void DArray_clear_destroy(DArray *array)
{
    DArray_clear(array);
    DArray_destroy(array);
}

int DArray_push(DArray *array, void *el)
{
    array->contents[array->end] = el;
    array->end++;
}

```

```

        if(DArray_end(array) >= DArray_max(array)) {
            return DArray_expand(array);
        } else {
            return 0;
        }
    }

void *DArray_pop(DArray *array)
{
    check(array->end - 1 >= 0, "Attempt to pop from empty array.");

    void *e1 = DArray_remove(array, array->end - 1);
    array->end--;

    if(DArray_end(array) > (int)array->expand_rate && DArray_end(array) % arra
y->expand_rate) {
        DArray_contract(array);
    }

    return e1;
error:
    return NULL;
}

```

This shows you another way to tackle complex code. Instead of diving right into the `.c` implementation, look at the header file, then read the unit test. This gives you an "abstract to concrete" understanding how the pieces work together and making it easier to remember.

## Advantages And Disadvantages

A `DArray` is better when you need to optimize these operations:

- Iteration. You can just use a basic for-loop and `DArray_count` with `DArray_get` and you're done. No special macros needed, and it's faster because you aren't walking pointers.
- Indexing. You can use `DArray_get` and `DArray_set` to access any element at random, but with a `List` you have to go through N elements to get to N+1.
- Destroying. You just free the struct and the `contents` in two operations. A `List` requires a series of `free` calls and also walking every element.
- Cloning. You can also clone it in just two operations (plus whatever it's storing) by copying the struct and `contents`. A list again requires walking the whole thing and copying every `ListNode` plus its value.
- Sorting. As you saw, `List` is horrible if you need to keep the data sorted. A `DArray` opens up a whole class of great sorting algorithms because now you can access elements randomly.
- Large Data. If you need to keep around a lot of data, then a `DArray` wins since its base `contents` takes up less memory than the same number of `ListNode` structs.

The `List` however wins on these operations:

- Insert and remove on the front (what I called shift). A `DArray` needs special treatment to be able to do this efficiently, and usually has to do some copying.
- Splitting or joining. A `List` can just copy some pointers and it's

done, but with a `DArray` you have to do copying of the arrays involved.

- Small Data. If you only need to store a few elements, then typically the storage will be less in a `List` than a generic `DArray` because the `DArray` needs to expand the backing store to accommodate future inserts, but a `List` only makes what it needs.

With this, I prefer to use a `DArray` for most of the things you see other people use a `List`. I reserve using `List` for any data structure that requires small number of nodes that are inserted and removed from either end. I'll show you two similar data structures called a `Stack` and `Queue` where this is important.

## How To Improve It

As usual, go through each function and operation and add the defensive programming checks, pre-conditions, invariants, and anything else you can find to make the implementation more bulletproof.

## Extra Credit

- Improve the unit tests to cover more of the operations and test that using a for-loop to iterate works.
- Research what it would take to implement bubble sort and merge sort for `DArray`, but don't do it yet. I'll be implementing `DArray` algorithms next and you'll do this then.
- Write some performance tests for common operations and compare them to the same operations in `List`. You did some of this, but this time, write a unit test that repeatedly does the operation in question, then in the main runner do the timing.
- Look at how the `DArray_expand` is implemented using a constant increase ( $\text{size} + 300$ ). Typically dynamic arrays are implemented with a multiplicative increase ( $\text{size} * 2$ ), but I've found this to cost needless memory for no real performance gain. Test my assertion and see when you'd want a multiplied increase instead of a constant increase.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 35: Sorting And Searching

In this exercise I'm going to cover four sorting algorithms and one search algorithm. The sorting algorithms are going to be quick sort, heap sort, merge sort, and radix sort. I'm then going to show you how to binary search after you've done a radix sort.

However, I'm a lazy guy, and in most standard C libraries you have existing implementations of the heapsort, quicksort, and mergesort algorithms. Here's how you use them:

```
#include <lcthw/darray_algos.h>
#include <stdlib.h>

int DArray_qsort(DArray *array, DArray_compare cmp)
{
    qsort(array->contents, DArray_count(array), sizeof(void *), cmp);
    return 0;
}

int DArray_heapsort(DArray *array, DArray_compare cmp)
{
    return heapsort(array->contents, DArray_count(array), sizeof(void *), cmp);
}

int DArray_mergesort(DArray *array, DArray_compare cmp)
{
    return mergesort(array->contents, DArray_count(array), sizeof(void *), cmp);
}
```

That's the whole implementation of the `darray_algos.c` file, and it should work on most modern Unix systems. What each of these does is sort the `contents` store of void pointers using the `DArray_compare` you give it. I'll show you the header file for this too:

```
#ifndef darray_algos_h
#define darray_algos_h

#include <lcthw/darray.h>

typedef int (*DArray_compare)(const void *a, const void *b);

int DArray_qsort(DArray *array, DArray_compare cmp);

int DArray_heapsort(DArray *array, DArray_compare cmp);

int DArray_mergesort(DArray *array, DArray_compare cmp);

#endif
```

About the same size and should be what you expect. Next you can see how these functions are used in the unit test for these three:

```
#include "minunit.h"
#include <lcthw/darray_algos.h>
```

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzsthw

```

int testcmp(char **a, char **b)
{
    return strcmp(*a, *b);
}

DArray *create_words()
{
    DArray *result = DArray_create(0, 5);
    char *words[] = {"asdfasfd", "werwar", "13234", "asdfasfd", "oioj"};
    int i = 0;

    for(i = 0; i < 5; i++) {
        DArray_push(result, words[i]);
    }

    return result;
}

int is_sorted(DArray *array)
{
    int i = 0;

    for(i = 0; i < DArray_count(array) - 1; i++) {
        if(strcmp(DArray_get(array, i), DArray_get(array, i+1)) > 0) {
            return 0;
        }
    }

    return 1;
}

char *run_sort_test(int (*func)(DArray *, DArray_compare), const char *name)
{
    DArray *words = create_words();
    mu_assert(!is_sorted(words), "Words should start not sorted.");

    debug("--- Testing %s sorting algorithm", name);
    int rc = func(words, (DArray_compare)testcmp);
    mu_assert(rc == 0, "sort failed");
    mu_assert(is_sorted(words), "didn't sort it");

    DArray_destroy(words);

    return NULL;
}

char *test_qsort()
{
    return run_sort_test(DArray_qsort, "qsort");
}

char *test_heapsort()
{
    return run_sort_test(DArray_heapsort, "heapsort");
}

char *test_mergesort()
{
    return run_sort_test(DArray_mergesort, "mergesort");
}

char * all_tests()
{
    mu_suite_start();

    mu_run_test(test_qsort);
    mu_run_test(test_heapsort);
    mu_run_test(test_mergesort);
}

```

```

    return NULL;
}

RUN_TESTS(all_tests);

```

The thing to notice, and actually what tripped me up for a whole day, is the definition of `testcmp` on line 4. You have to use a `char **` and *not* a `char *` because `qsort` is going to give you a pointer to *the pointers* in the `contents` array. The reason is `qsort` and friends are scanning the array, and handing *pointers* to each element in the array to your comparison function. Since what I have in the `contents` array is pointers, that means you get a pointer to a pointer.

With that out of the way you have to just implemented three difficult sorting algorithms in about 20 lines of code. You could stop there, but part of this book is learning how these algorithms work so the extra credit is going to involve implementing each of these.

## Radix Sort And Binary Search

Since you're going to implement quicksort, heapsort, and mergesort on your own, I'm going to show you a funky algorithm called Radix Sort. It has a slightly narrow usefulness in sorting arrays of integers, and seems to work like magic. In this case I'm going to create a special data structure called a `RadixMap` that is used to map one integer to another.

Here's the header file for the new algorithm that is both algorithm and data structure in one:

```

#ifndef _radixmap_h
#include <stdint.h>

typedef union RMElement {
    uint64_t raw;
    struct {
        uint32_t key;
        uint32_t value;
    } data;
} RMElement;

typedef struct RadixMap {
    size_t max;
    size_t end;
    uint32_t counter;
    RMElement *contents;
    RMElement *temp;
} RadixMap;

RadixMap *RadixMap_create(size_t max);

void RadixMap_destroy(RadixMap *map);

void RadixMap_sort(RadixMap *map);

RMElement *RadixMap_find(RadixMap *map, uint32_t key);

int RadixMap_add(RadixMap *map, uint32_t key, uint32_t value);

int RadixMap_delete(RadixMap *map, RMElement *el);

#endif

```

You see I have a lot of the same operations as in a `Dynamic Array` or a `List` data structure, the difference is I'm working only with fixed size 32 bit `uint32_t` integers. I'm also introducing you to a new C concept called the `union` here.

## C Unions

A union is a way to refer to the same piece of memory in a number of different ways. How they work is you define them like a `struct` except every element is sharing the same space with all of the others. You can think of a union as a picture of the memory, and the elements in the union as different colored lenses to view the picture.

What they are used for is to either save memory, or to convert chunks of memory between formats. The first usage is typically done with "variant types", where you create a struct that has "tag" for the type, and then a union inside it for each type. When used for converting between formats of memory, you simply define the two structures, and then access the right one.

First let me show you how to make a variant type with C unions:

```
#include <stdio.h>

typedef enum {
    TYPE_INT,
    TYPE_FLOAT,
    TYPE_STRING,
} VariantType;

struct Variant {
    VariantType type;
    union {
        int as_integer;
        float as_float;
        char *as_string;
    } data;
};

typedef struct Variant Variant;

void Variant_print(Variant *var)
{
    switch(var->type) {
        case TYPE_INT:
            printf("INT: %d\n", var->data.as_integer);
            break;
        case TYPE_FLOAT:
            printf("FLOAT: %f\n", var->data.as_float);
            break;
        case TYPE_STRING:
            printf("STRING: %s\n", var->data.as_string);
            break;
        default:
            printf("UNKNOWN TYPE: %d", var->type);
    }
}

int main(int argc, char *argv[])
{
    Variant a_int = {.type = TYPE_INT, .data.as_integer = 100};
    Variant a_float = {.type = TYPE_FLOAT, .data.as_float = 100.34};
    Variant a_string = {.type = TYPE_STRING, .data.as_string = "YO DUDE!"};

    Variant_print(&a_int);
}
```

```

Variant_print(&a_float);
Variant_print(&a_string);

// here's how you access them
a_int.data.as_integer = 200;
a_float.data.as_float = 2.345;
a_string.data.as_string = "Hi there.";

Variant_print(&a_int);
Variant_print(&a_float);
Variant_print(&a_string);

return 0;
}

```

You find this in many implementations of dynamic languages. The language will define some base variant type with tags for all the base types of the language, and then usually there's a generic "object" tag for the types you create. The advantage of doing this is that the `Variant` only takes up as much space as the `VariantType` tag and the largest member of the union. This is because C is "layering" each element of the `Variant.data` union together so they overlap, and to do that it sizes it big enough to hold the largest element.

In the `radixmap.h` file I have the `RMElement` union which demonstrates using a union to convert blocks of memory between types. In this case, I want to store a `uint64_t` sized integer for sorting purposes, but I want a two `uint32_t` integers for the data to represent a `key` and `value` pair. By using a union I'm able to access the same block of memory in the two different ways I need cleanly.

## The Implementation

I next have the actual `RadixMap` implementation for each of these operations:

```

/*
 * Based on code by Andre Reinald then heavily modified by Zed A. Shaw.
 */

#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <lcthw/radixmap.h>
#include <lcthw/dbg.h>

RadixMap *RadixMap_create(size_t max)
{
    RadixMap *map = calloc(sizeof(RadixMap), 1);
    check_mem(map);

    map->contents = calloc(sizeof(RMElement), max + 1);
    check_mem(map->contents);

    map->temp = calloc(sizeof(RMElement), max + 1);
    check_mem(map->temp);

    map->max = max;
    map->end = 0;

    return map;
error:
    return NULL;
}

```

```

void RadixMap_destroy(RadixMap *map)
{
    if(map) {
        free(map->contents);
        free(map->temp);
        free(map);
    }
}

#define ByteOf(x,y) (((uint8_t *)x)[(y)])

static inline void radix_sort(short offset, uint64_t max, uint64_t *source, uint64_t *dest)
{
    uint64_t count[256] = {0};
    uint64_t *cp = NULL;
    uint64_t *sp = NULL;
    uint64_t *end = NULL;
    uint64_t s = 0;
    uint64_t c = 0;

    // count occurrences of every byte value
    for (sp = source, end = source + max; sp < end; sp++) {
        count[ByteOf(sp, offset)]++;
    }

    // transform count into index by summing elements and storing into same array
    for (s = 0, cp = count, end = count + 256; cp < end; cp++) {
        c = *cp;
        *cp = s;
        s += c;
    }

    // fill dest with the right values in the right place
    for (sp = source, end = source + max; sp < end; sp++) {
        cp = count + ByteOf(sp, offset);
        dest[*cp] = *sp;
        ++(*cp);
    }
}

void RadixMap_sort(RadixMap *map)
{
    uint64_t *source = &map->contents[0].raw;
    uint64_t *temp = &map->temp[0].raw;

    radix_sort(0, map->end, source, temp);
    radix_sort(1, map->end, temp, source);
    radix_sort(2, map->end, source, temp);
    radix_sort(3, map->end, temp, source);
}

RMElement *RadixMap_find(RadixMap *map, uint32_t to_find)
{
    int low = 0;
    int high = map->end - 1;
    RMElement *data = map->contents;

    while (low <= high) {
        int middle = low + (high - low)/2;
        uint32_t key = data[middle].data.key;

        if (to_find < key) {
            high = middle - 1;
        } else if (to_find > key) {
            low = middle + 1;
        } else {
            return &data[middle];
        }
    }
}

```

```

    }
}

return NULL;
}

int RadixMap_add(RadixMap *map, uint32_t key, uint32_t value)
{
    check(key < UINT32_MAX, "Key can't be equal to UINT32_MAX.");

    RMElement element = {.data = {.key = key, .value = value}};
    check(map->end + 1 < map->max, "RadixMap is full.");

    map->contents[map->end++] = element;

    RadixMap_sort(map);

    return 0;
error:
    return -1;
}

int RadixMap_delete(RadixMap *map, RMElement *el)
{
    check(map->end > 0, "There is nothing to delete.");
    check(el != NULL, "Can't delete a NULL element.");

    el->data.key = UINT32_MAX;

    if(map->end > 1) {
        // don't bother resorting a map of 1 length
        RadixMap_sort(map);
    }

    map->end--;

    return 0;
error:
    return -1;
}

```

As usual enter this in and get it working along with the unit test then I'll explain what's happening. Take *special* care with the `radix_sort` function as it's very particular in how it's implemented.

```

#include "minunit.h"
#include <lcsth/radixmap.h>
#include <time.h>

static int make_random(RadixMap *map)
{
    size_t i = 0;

    for (i = 0; i < map->max - 1; i++) {
        uint32_t key = (uint32_t)(rand() | (rand() << 16));
        check(RadixMap_add(map, key, i) == 0, "Failed to add key %u.", key);
    }

    return i;
error:
    return 0;
}

static int check_order(RadixMap *map)
{
    RMElement d1, d2;
    unsigned int i = 0;

```

```

// only signal errors if any (should not be)
for (i = 0; map->end > 0 && i < map->end-1; i++) {
    d1 = map->contents[i];
    d2 = map->contents[i+1];

    if(d1.data.key > d2.data.key) {
        debug("FAIL:i=%u, key: %u, value: %u, equals max? %d\n", i, d1.data
a.key, d1.data.value,
        d2.data.key == UINT32_MAX);
        return 0;
    }
}

return 1;
}

static int test_search(RadixMap *map)
{
    unsigned i = 0;
    RMElement *d = NULL;
    RMElement *found = NULL;

    for(i = map->end / 2; i < map->end; i++) {
        d = &map->contents[i];
        found = RadixMap_find(map, d->data.key);
        check(found != NULL, "Didn't find %u at %u.", d->data.key, i);
        check(found->data.key == d->data.key, "Got the wrong result: %p:%u loo
king for %u at %u",
            found, found->data.key, d->data.key, i);
    }

    return 1;
error:
    return 0;
}

// test for big number of elements
static char *test_operations()
{
    size_t N = 200;

    RadixMap *map = RadixMap_create(N);
    mu_assert(map != NULL, "Failed to make the map.");
    mu_assert(make_random(map), "Didn't make a random fake radix map.");

    RadixMap_sort(map);
    mu_assert(check_order(map), "Failed to properly sort the RadixMap.");

    mu_assert(test_search(map), "Failed the search test.");
    mu_assert(check_order(map), "RadixMap didn't stay sorted after search.");

    while(map->end > 0) {
        RMElement *el = RadixMap_find(map, map->contents[map->end / 2].data.ke
y);
        mu_assert(el != NULL, "Should get a result.");

        size_t old_end = map->end;

        mu_assert(RadixMap_delete(map, el) == 0, "Didn't delete it.");
        mu_assert(old_end - 1 == map->end, "Wrong size after delete.");

        // test that the end is now the old value, but uint32 max so it trails
off
        mu_assert(check_order(map), "RadixMap didn't stay sorted after delete.
");
    }

    RadixMap_destroy(map);

    return NULL;
}

```



```

}

char *all_tests()
{
    mu_suite_start();
    srand(time(NULL));

    mu_run_test(test_operations);

    return NULL;
}

RUN_TESTS(all_tests);

```

I shouldn't have to explain too much about the test. It's simply simulating placing random integers into the `RadixMap` and then making sure it can get them out reliably. Not too interesting.

In the `radixmap.c` file most of the operations are easy to understand if you read the code. Here's a description of what the basic functions are doing and how they work:

#### RadixMap\_create

As usual I'm allocating all the memory needed for the structures defined in `radixmap.h`. I'll be using the `temp` and `contents` later when I talk about `radix_sort`.

#### RadixMap\_destroy

Again, just destroying what was created.

#### radix\_sort

The meat of the data structure, but I'll explain what it's doing in the next section.

#### RadixMap\_sort

This uses the `radix_sort` function to actually sort the `contents`. It does this by sorting between the `contents` and `temp` until finally `contents` is sorted. You'll see how this works when I describe `radix_sort` later.

#### RadixMap\_find

This is using a binary search algorithm to find a key you give it. I'll explain how this works shortly.

#### RadixMap\_add

Using the `RadixMap_sort` function, this will add the key and value you request at the end, then simply sort it again so that everything is in the right place. Once everything is sorted, the `RadixMap_find` will work properly because it's a binary search.

#### RadixMap\_delete

Works the same as `RadixMap_add` except "deletes" elements of the structure by setting their values to the max for a unsigned 32 bit integer, `UINT32_MAX`. This means you can't use that value as an key value, but it makes deleting elements easy. Simply set it to that and then sort and it'll get moved to the end. Now it's deleted.

Study the code for the ones I described, and then that just leaves `RadixMap_sort`, `radix_sort`, and `RadixMap_find` to understand.

## RadixMap\_find And Binary Search

I'll start with how the binary search is implemented. Binary search is

simple algorithm that most people can understand intuitively. In fact, you could take a deck of playing cards (or cards with numbers) and do this manually. Here's how this function works, and how a binary search works:

- Set a high and low mark based on the size of the array.
- Get the middle element between the low and high marks.
- If the key is less-than, then the key must be below the middle. Set high to one less than middle.
- If the key is greater-than, then the key must be above the middle. Set the low mark one greater than the middle.
- If it's equal then you found it, stop.
- Keep looping until low and high pass each other. You don't find it if you exit the loop.

What you are effectively doing is guessing where the key might be by picking the middle and comparing it. Since the data is sorted, you know that the the key has to be above or below this. If it's below, then you just divided the search space in half. You keep going until you either find it or you overlap the boundaries and exhaust the search space.

## RadixMap\_sort And radix\_sort

A radix sort is easy to understand if you try to do it manually first. What this algorithm does is exploit the fact that numbers are stored with a sequence of digits that go from "least significant" to "most significant". It then takes the numbers and buckets them by the digit, and when it has processed all the digits the numbers come out sorted. At first it seems like magic, and honestly looking at the code sure seems like it is, but try doing it manually once.

To do this algorithm write out a bunch of three digit numbers, in a random order, let's say we do 223, 912, 275, 100, 633, 120, and 380.

- Place the number in buckets by their 1's digit: [380, 100, 120], [912], [633, 223], [275].
- I now have to go through each of these buckets in order, and then sort it into 10's buckets: [100], [912], [120, 223], [633], [275], [380].
- Now each bucket contains numbers that are sorted by the 1's then 10's digit. I need to then go through these in order and fill the final 100's buckets: [100, 120], [223, 275], [380], [633], [912].
- At this point each bucket is sorted by 100's, 10's, then 1's and if I take each bucket in order I get the final sorted list: 100, 120, 223, 275, 380, 633, 912.

Make sure you do this a few times so you understand how it works. It really is a slick little algorithm and most importantly it will work on numbers of arbitrary size, so you can sort really huge numbers because you are just doing them one byte at a time.

In my situation the "digits" are individual 8 bit bytes, so I need 256 buckets to store the distribution of the numbers by their digits. I also need a way to store them such that I don't use too much space. If you look at `radix_sort` first thing I do is build a `count` histogram so I know

how many occurrences of each digit there are for the given `offset`.

Once I know the counts for each digit (all 256 of them) I can then use that as distribution points into a target array. For example, if I have 10 bytes that are 0x00, then I know I can place them in the first 10 slots of the target array. This gives me an index for where they go in the target array, which is the second for-loop in `radix_sort`.

Finally, once I know where they can go in the target array, I simply go through all the digits in the `source` array, for this `offset` and place the numbers in their slots in order. Using the `ByteOf` macro helps keep the code clean since there's a bit of pointer hackery to make it work, but the end result is all of the integers will be placed in the bucket for their digit when the final for-loop is done.

What becomes interesting is then how I use this in `RadixMap_sort` to sort these 64 bit integers by just the first 32 bits. Remember how I have the key and value in a union for the `RMElement` type? That means to sort this array by the key I only need to sort the first 4 bytes (32 bits / 8 bits per byte) of every integer.

If you look at the `RadixMap_sort` you see that I grab a quick pointer to the `contents` and `temp` to for source and target arrays, and then I call `radix_sort` four times. Each time I call it, I alternate source and target and do the next byte. When I'm done, the `radix_sort` has done its job and the final copy has been done into the `contents`.

## How To Improve It

There is a big disadvantage to this implementation because it has to process the entire array four times on every insertion. It does do it fast, but it'd be better if you could limit the amount of sorting by the size of what needs to be sorted.

There's two ways you can improve this implementation:

- Use a binary search to find the minimum position for the new element, then only sort from there to the end. You find the minimum, put the new element on the end, then just sort from the minimum on. This will cut your sort space down considerably most of the time.
- Keep track of the biggest key currently being used, and then only sort enough digits to handle that key. You can also keep track of the smallest number, and then only sort the digits necessary for the range. To do this you'll have to start caring about CPU integer ordering (endianess).

Try these optimizations, but after you augment the unit test with some timing information so you can see if you're actually improving the speed of the implementation.

## Extra Credit

- Implement quicksort, heapsort, and mergesort and provide a `#define` that lets you pick between the two, or create a second set of functions you can call. Use the technique I taught you to

read the Wikipedia page for the algorithm and then implement it with the psuedo-code.

- Compare the performance of your implementations to the original ones.
- Use these sorting functions to create a `DArray_sort_add` that adds elements to the `DArray` but sorts the array after.
- Write a `DArray_find` that uses the binary search algorithm from `RadixMap_find` and the `DArray_compare` to find elements in a sorted `DArray`.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 36: Safer Strings

I've already introduced you to the `Better String` library in Exercise 26 when we made `devpkg`. This exercise is designed to get you into using `bstring` from now on, why C's strings are an incredibly bad idea, and then have you change the `liblcthw` code to use `bstring`.

## Why C Strings Were A Horrible Idea

When people talk about problems with C, it's concept of a "string" is one of the top flaws. You've been using these extensively, and I've talked about the kinds of flaws they have, but there's not much that explains exactly why C strings are flawed and always will be. I'll try to explain that right now, but part of my explanation will just be that after decades of using C's strings there's enough evidence that they are just a bad idea.

It is impossible to confirm that any given C string is valid:

- A C string is invalid if it does not end in `'\0'`.
- Any loop that processes an invalid C string will loop infinitely (or, just buffer overflow).
- C strings do not have a known length, so the only way to check if it's terminated correctly is to loop through it.
- Therefore, it is not possible to validate a C string without possibly looping infinitely.

This is simple logic. You can't write a loop that checks if a C string is valid because invalid C strings cause loops to never terminate. That's it, and the only solution is to *include the size*. Once you know the size you can avoid the infinite loop problem. If you look at the two functions I showed you from Exercise 27 you can see this:

```
void copy(char to[], char from[])
{
    int i = 0;

    // while loop will not end if from isn't '\0' terminated
    while((to[i] = from[i]) != '\0') {
        ++i;
    }
}

int safercopy(int from_len, char *from, int to_len, char *to)
{
    int i = 0;
    int max = from_len > to_len - 1 ? to_len - 1 : from_len;

    // to_len must have at least 1 byte
    if(from_len < 0 || to_len <= 0) return -1;

    for(i = 0; i < max; i++) {
        to[i] = from[i];
    }
}
```

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssthw

```

    to[to_len - 1] = '\0';

    return i;
}

```

Imagine you want to add a check to the `copy` function to confirm that the `from` string is valid. How would you do that? Why you'd write a loop that checked that the string ended in `'\0'`. Oh wait, if the string doesn't end in `'\0'` then how does the checking loop end? It doesn't. Checkmate.

No matter what you do, you can't check that a C string is valid without knowing the length of the underlying storage, and in this case the `safercopy` includes those lengths. This function doesn't have the same problem as it's loops will always terminate, even if you lie to it about the size, you still have to give it a finite size.

What the Better String library does is create a struct that always includes the length of the string's storage. Because the length is always available to a `bstring` then all of its operations can be safer. The loops will terminate, the contents can be validated, and it will not have this major flaw. The `bstring` library also comes with a ton of operations you need with strings, like splitting, formatting, searching, and they are most likely done right and safer.

There could be flaws in `bstring`, but it's been around a long time so those are probably minimal. They still find flaws in `glibc` so what's a programmer to do right?

## Using bstrlib

There's quite a few improved string libraries, but I like `bstrlib` because it fits in one file for the basics and has most of the stuff you need to deal with strings. You've already used it a bit, so in this exercise you'll go get the two files `bstrlib.c` and `bstrlib.h` from the [Better String](#)

Here's me doing this in the `liblcthw` project directory:

```

$ mkdir bstrlib
$ cd bstrlib/
$ unzip ~/Downloads/bstrlib-05122010.zip
Archive:  /Users/zedshaw/Downloads/bstrlib-05122010.zip
...
$ ls
bsafe.c          bstraux.c        bstrlib.h        bstrwrap.h       license.tx
t    test.cpp
bsafe.h          bstraux.h        bstrlib.txt      cpptest.cpp      porting.tx
t    testaux.c
bctest.c  bstrlib.c        bstrwrap.cpp    gpl.txt          security.txt
$ mv bstrlib.h bstrlib.c ../src/lcthw/
$ cd ../
$ rm -rf bstrlib
# make the edits
$ vim src/lcthw/bstrlib.c
$ make clean all
...
$

```

On line 14 you seem me edit the `bstrlib.c` file to move it to a new location and to fix a bug on OSX. Here's the diff:

```

25c25
< #include "bstrlib.h"
---
> #include <lcthw/bstrlib.h>
2759c2759
< #ifdef __GNUC__
---
> #if defined(__GNUC__) && !defined(__APPLE__)

```

That is, change the include to be `<lcthw/bstrlib.h>`, and then fix one of the `ifdef` at line 2759.

## Learning The Library

This exercise is short and simply getting you ready for the remaining exercises that use the library. In the next two exercises I'll use `bstrlib.c` to create a `HashMap` data structure.

You should now get familiar with this library by reading the header file, the implementations, and then write a `tests/bstr_tests.c` that tests out the following functions:

`bfromcstr`

Create a bstring from a C style constant.

`blk2bstr`

Same but give the length of the buffer.

`bstrcpy`

Copy a bstring.

`bassign`

Set one bstring to another.

`bassigncstr`

Set a bstring to a C string's contents.

`bassignblk`

Set a bstring to a C string but give the length.

`bdestroy`

Destroy a bstring.

`bconcat`

Concatenate one bstring onto another.

`bstrcmp`

Compare two bstrings returning the same result as strcmp.

`biseq`

Tests if two bstrings are equal.

`binstr`

Tells if one bstring is in another.

`bfindreplace`

Find one bstring in another then replace it with a third.

`bsplit`

How to split a bstring into a bstrList.

`bformat`

Doing a format string, super handy.

`blength`

Getting the length of a bstring.

`bdata`

Getting the data from a bstring.

`bchar`

Getting a char from a bstring.

Your test should try out all of these operations, and a few more that you find interesting from the header file. Make sure to run the test under `valgrind` to make sure you use the memory correctly.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 37: Hashmaps

Hash Maps (Hashmaps, Hashes, or sometimes Dictionaries) are used frequently in many dynamic programming for storing key/value data. A Hashmap works by performing a "hashing" calculation on the keys to produce an integer, then uses that integer to find a bucket to get or set the value. It is a very fast practical data structure since it works on nearly any data and they are easy to implement.

Here's an example of using a Hashmap (aka dict) in Python:

```
fruit_weights = {'Apples': 10, 'Oranges': 100, 'Grapes': 1.0}

for key, value in fruit_weights.items():
    print key, "=", value
```

Almost every modern language has something like this, so many people end up writing code and never understand how this actually works. By creating the `Hashmap` data structure in C I'll show you how this works. I'll start with the header file so I can talk about the data structure.

```
#ifndef _lcthw_Hashmap_h
#define _lcthw_Hashmap_h

#include <stdint.h>
#include <lcthw/darray.h>

#define DEFAULT_NUMBER_OF_BUCKETS 100

typedef int (*Hashmap_compare)(void *a, void *b);
typedef uint32_t (*Hashmap_hash)(void *key);

typedef struct Hashmap {
    DArray *buckets;
    Hashmap_compare compare;
    Hashmap_hash hash;
} Hashmap;

typedef struct HashmapNode {
    void *key;
    void *data;
    uint32_t hash;
} HashmapNode;

typedef int (*Hashmap_traverse_cb)(HashmapNode *node);

Hashmap *Hashmap_create(Hashmap_compare compare, Hashmap_hash);
void Hashmap_destroy(Hashmap *map);

int Hashmap_set(Hashmap *map, void *key, void *data);
void *Hashmap_get(Hashmap *map, void *key);

int Hashmap_traverse(Hashmap *map, Hashmap_traverse_cb traverse_cb);

void *Hashmap_delete(Hashmap *map, void *key);

#endif
```

The structure consists of a `Hashmap` that contains any number of

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

🛠  
HELP

Follow  
@lzsthw

`HashMapNode` structs. Looking at `HashMap` you can see that it is structured like this:

`DArray *buckets`

A dynamic array that will be set to a fixed size of 100 buckets. Each bucket will in turn contain a `DArray` that will actually hold `HashMapNode` pairs.

`HashMap_compare compare`

This is a comparison function that the `HashMap` uses to actually find elements by their key. It should work like all of the other compare functions, and defaults to using `bstrcmp` so that keys are just bstrings.

`HashMap_hash hash`

This is the hashing function and it's responsible for taking a key, processing its contents, and producing a single `uint32_t` index number. You'll see the default one soon.

This almost tells you how the data is stored, but the `buckets` `DArray` isn't created yet. Just remember that it's kind of a two level mapping:

- There are 100 buckets that make up the first level, and things are in these buckets based on their hash.
- Each bucket is a `DArray` that then contains `HashMapNode` structs simply appended to the end as they're added.

The `HashMapNode` is then composed of these three elements:

`void *key`

The key for this key=value pair.

`void *value`

The value.

`uint32_t hash`

The calculated hash, which makes finding this

node quicker since we can just check the hash and skip any that don't match, only checking they key if it's equal.

The rest of the header file is nothing new, so now I can show you the implementation `hashmap.c` file:

```
#undef NDEBUG
#include <stdint.h>
#include <lcthw/hashmap.h>
#include <lcthw/dbg.h>
#include <lcthw/bstrlib.h>

static int default_compare(void *a, void *b)
{
    return bstrcmp((bstring)a, (bstring)b);
}

/**
 * Simple Bob Jenkins's hash algorithm taken from the
 * wikipedia description.
 */
static uint32_t default_hash(void *a)
{
    size_t len = blength((bstring)a);
    char *key = bdata((bstring)a);
    uint32_t hash = 0;
    uint32_t i = 0;
```

```

    for(hash = i = 0; i < len; ++i)
    {
        hash += key[i];
        hash += (hash << 10);
        hash ^= (hash >> 6);
    }

    hash += (hash << 3);
    hash ^= (hash >> 11);
    hash += (hash << 15);

    return hash;
}

Hashmap *Hashmap_create(Hashmap_compare compare, Hashmap_hash hash)
{
    Hashmap *map = calloc(1, sizeof(Hashmap));
    check_mem(map);

    map->compare = compare == NULL ? default_compare : compare;
    map->hash = hash == NULL ? default_hash : hash;
    map->buckets = DArray_create(sizeof(DArray *), DEFAULT_NUMBER_OF_BUCKETS);
    map->buckets->end = map->buckets->max; // fake out expanding it
    check_mem(map->buckets);

    return map;
}

error:
    if(map) {
        Hashmap_destroy(map);
    }

    return NULL;
}

void Hashmap_destroy(Hashmap *map)
{
    int i = 0;
    int j = 0;

    if(map) {
        if(map->buckets) {
            for(i = 0; i < DArray_count(map->buckets); i++) {
                DArray *bucket = DArray_get(map->buckets, i);
                if(bucket) {
                    for(j = 0; j < DArray_count(bucket); j++) {
                        free(DArray_get(bucket, j));
                    }
                    DArray_destroy(bucket);
                }
            }
            DArray_destroy(map->buckets);
        }

        free(map);
    }
}

static inline HashmapNode *Hashmap_node_create(int hash, void *key, void *data)
{
    HashmapNode *node = calloc(1, sizeof(HashmapNode));
    check_mem(node);

    node->key = key;
    node->data = data;
    node->hash = hash;
}

```

```

        return node;

error:
    return NULL;
}

static inline DArray *Hashmap_find_bucket(Hashmap *map, void *key,
                                          int create, uint32_t *hash_out)
{
    uint32_t hash = map->hash(key);
    int bucket_n = hash % DEFAULT_NUMBER_OF_BUCKETS;
    check(bucket_n >= 0, "Invalid bucket found: %d", bucket_n);
    *hash_out = hash; // store it for the return so the caller can use it

    DArray *bucket = DArray_get(map->buckets, bucket_n);

    if(!bucket && create) {
        // new bucket, set it up
        bucket = DArray_create(sizeof(void *), DEFAULT_NUMBER_OF_BUCKETS);
        check_mem(bucket);
        DArray_set(map->buckets, bucket_n, bucket);
    }

    return bucket;

error:
    return NULL;
}

int Hashmap_set(Hashmap *map, void *key, void *data)
{
    uint32_t hash = 0;
    DArray *bucket = Hashmap_find_bucket(map, key, 1, &hash);
    check(bucket, "Error can't create bucket.");

    HashmapNode *node = Hashmap_node_create(hash, key, data);
    check_mem(node);

    DArray_push(bucket, node);

    return 0;

error:
    return -1;
}

static inline int Hashmap_get_node(Hashmap *map, uint32_t hash, DArray *bucket,
                                   void *key)
{
    int i = 0;

    for(i = 0; i < DArray_end(bucket); i++) {
        debug("TRY: %d", i);
        HashmapNode *node = DArray_get(bucket, i);
        if(node->hash == hash && map->compare(node->key, key) == 0) {
            return i;
        }
    }

    return -1;
}

void *Hashmap_get(Hashmap *map, void *key)
{
    uint32_t hash = 0;
    DArray *bucket = Hashmap_find_bucket(map, key, 0, &hash);

```

```

    if(!bucket) return NULL;

    int i = Hashmap_get_node(map, hash, bucket, key);
    if(i == -1) return NULL;

    HashmapNode *node = DArray_get(bucket, i);
    check(node != NULL, "Failed to get node from bucket when it should exist."
);

    return node->data;

error: // fallthrough
    return NULL;
}

int Hashmap_traverse(Hashmap *map, Hashmap_traverse_cb traverse_cb)
{
    int i = 0;
    int j = 0;
    int rc = 0;

    for(i = 0; i < DArray_count(map->buckets); i++) {
        DArray *bucket = DArray_get(map->buckets, i);
        if(bucket) {
            for(j = 0; j < DArray_count(bucket); j++) {
                HashmapNode *node = DArray_get(bucket, j);
                rc = traverse_cb(node);
                if(rc != 0) return rc;
            }
        }
    }

    return 0;
}

void *Hashmap_delete(Hashmap *map, void *key)
{
    uint32_t hash = 0;
    DArray *bucket = Hashmap_find_bucket(map, key, 0, &hash);
    if(!bucket) return NULL;

    int i = Hashmap_get_node(map, hash, bucket, key);
    if(i == -1) return NULL;

    HashmapNode *node = DArray_get(bucket, i);
    void *data = node->data;
    free(node);

    HashmapNode *ending = DArray_pop(bucket);

    if(ending != node) {
        // alright looks like it's not the last one, swap it
        DArray_set(bucket, i, ending);
    }

    return data;
}

```

There's nothing very complicated in the implementation, but the `default_hash` and `Hashmap_find_bucket` functions will need some explanation. When you use `Hashmap_create` you can pass in any compare and hash functions you want, but if you don't it uses the `default_compare` and `default_hash` functions.

The first thing to look at is how `default_hash` does its thing. This is a simple hash function called a "Jenkins hash" after Bob Jenkins. I got it from the [Wikipedia page](#) for the algorithm. It simply goes through each

byte of the key to hash (a bstring) and works the bits so that the end result is a single `uint32_t`. It does this with some adding and xor operations.

There are many different hash functions, all with different properties, but once you have one you need a way to use it to find the right buckets. The `Hashmap_find_bucket` does it like this:

- First it calls `map->hash(key)` to get the hash for the key.
- It then finds the bucket using `hash % DEFAULT_NUMBER_OF_BUCKETS`, that way every hash will always find some bucket no matter how big it is.
- It then gets the bucket, which is also a `DArray`, and if it's not there it will create it. That depends on if the `create` variable says too.
- Once it has found the `DArray` bucket for the right hash, it returns it, and also the `hash_out` variable is used to give the caller the hash that was found.

All of the other functions then use `Hashmap_find_bucket` to do their work:

- Setting a key/value involves finding the bucket, then making a `HashmapNode`, and then adding it to the bucket.
- Getting a key involves finding the bucket, then finding the `HashmapNode` that matches the `hash` and `key` you want.
- Deleting an item again finds the bucket, finds where the requested node is, and then removes it by swapping the last node into its place.

The only other function that you should study is the `Hashmap_traverse`. This simply walks every bucket, and for any bucket that has possible values, it calls the `traverse_cb` on each value. This is how you scan a whole `Hashmap` for its values.

## The Unit Test

Finally you have the unit test that is testing all of these operations:

```
#include "minunit.h"
#include <lcthw/hashmap.h>
#include <assert.h>
#include <lcthw/bstrlib.h>

Hashmap *map = NULL;
static int traverse_called = 0;
struct tagbstring test1 = bsStatic("test data 1");
struct tagbstring test2 = bsStatic("test data 2");
struct tagbstring test3 = bsStatic("xest data 3");
struct tagbstring expect1 = bsStatic("THE VALUE 1");
struct tagbstring expect2 = bsStatic("THE VALUE 2");
struct tagbstring expect3 = bsStatic("THE VALUE 3");

static int traverse_good_cb(HashmapNode *node)
{
    debug("KEY: %s", bdata((bstring)node->key));
    traverse_called++;
    return 0;
}

static int traverse_fail_cb(HashmapNode *node)
```

```

{
    debug("KEY: %s", bdata((bstring)node->key));
    traverse_called++;

    if(traverse_called == 2) {
        return 1;
    } else {
        return 0;
    }
}

char *test_create()
{
    map = Hashmap_create(NULL, NULL);
    mu_assert(map != NULL, "Failed to create map.");

    return NULL;
}

char *test_destroy()
{
    Hashmap_destroy(map);

    return NULL;
}

char *test_get_set()
{
    int rc = Hashmap_set(map, &test1, &expect1);
    mu_assert(rc == 0, "Failed to set &test1");
    bstring result = Hashmap_get(map, &test1);
    mu_assert(result == &expect1, "Wrong value for test1.");

    rc = Hashmap_set(map, &test2, &expect2);
    mu_assert(rc == 0, "Failed to set test2");
    result = Hashmap_get(map, &test2);
    mu_assert(result == &expect2, "Wrong value for test2.");

    rc = Hashmap_set(map, &test3, &expect3);
    mu_assert(rc == 0, "Failed to set test3");
    result = Hashmap_get(map, &test3);
    mu_assert(result == &expect3, "Wrong value for test3.");

    return NULL;
}

char *test_traverse()
{
    int rc = Hashmap_traverse(map, traverse_good_cb);
    mu_assert(rc == 0, "Failed to traverse.");
    mu_assert(traverse_called == 3, "Wrong count traverse.");

    traverse_called = 0;
    rc = Hashmap_traverse(map, traverse_fail_cb);
    mu_assert(rc == 1, "Failed to traverse.");
    mu_assert(traverse_called == 2, "Wrong count traverse for fail.");

    return NULL;
}

char *test_delete()
{
    bstring deleted = (bstring)Hashmap_delete(map, &test1);
    mu_assert(deleted != NULL, "Got NULL on delete.");
    mu_assert(deleted == &expect1, "Should get test1");
    bstring result = Hashmap_get(map, &test1);
    mu_assert(result == NULL, "Should delete.");
}

```

```

deleted = (bstring)Hashmap_delete(map, &test2);
mu_assert(deleted != NULL, "Got NULL on delete.");
mu_assert(deleted == &expect2, "Should get test2");
result = Hashmap_get(map, &test2);
mu_assert(result == NULL, "Should delete.");

deleted = (bstring)Hashmap_delete(map, &test3);
mu_assert(deleted != NULL, "Got NULL on delete.");
mu_assert(deleted == &expect3, "Should get test3");
result = Hashmap_get(map, &test3);
mu_assert(result == NULL, "Should delete.");

return NULL;
}

char *all_tests()
{
    mu_suite_start();

    mu_run_test(test_create);
    mu_run_test(test_get_set);
    mu_run_test(test_traverse);
    mu_run_test(test_delete);
    mu_run_test(test_destroy);

    return NULL;
}

RUN_TESTS(all_tests);

```

The only thing to learn about this unit test is that at the top I use a feature of `bstring` to create static strings to work with in the tests. I use the `tagbstring` and `bsStatic` to create them on lines 7-13.

## How To Improve It

This is a very simple implementation of `Hashmap` as are most of the other data structures in this book. My goal isn't to give you insanely great hyper speed well tuned data structures. Usually those are much too complicated to discuss and only distract you from the real basic data structure at work. My goal is to give you an understandable starting point to then improve it or understand how they are implemented.

In this case, there's some things you can do with this implementation:

- You can use a sort on each bucket so that they are always sorted. This increases your insert time, but decreases your find time because you can then use a binary search to find each node. Right now it's looping through all of the nodes in a bucket just to find one.
- You can dynamically size the number of buckets, or let the caller specify the number for each `Hashmap` created.
- You can use a better `default_hash`. There are tons of them.
- This (and nearly every `Hashmap`) is vulnerable to someone picking keys that will fill only one bucket, and then tricking your program into processing them. This then makes your program run slower because it changes from processing a `Hashmap` to effectively processing a single `DArray`. If you sort the nodes in the bucket this helps, but you can also use better hashing functions, and for the really paranoid add a random salt so that keys can't be predicted.



- You could have it delete buckets that are empty of nodes to save space, or put empty buckets into a cache so you save on creating and destroying them.
- Right now it just adds elements even if they already exist. Write an alternative set method that only adds it if it isn't set already.

As usual you should go through each function and make it bullet proof. The `HashMap` could also use a debug setting for doing an invariant check.

## Extra Credit

- Research the `HashMap` implementation of your favorite programming language to see what features they have.
- Find out what the major disadvantages of a `HashMap` are and how to avoid them. For example, they do not preserve order without special changes and they don't work when you need to find things based on parts of keys.
- Write a unit test that demonstrates the defect of filling a `HashMap` with keys that land in the same bucket, then test how this impact performance. A good way to do this is to just reduce the number of buckets to something stupid like 5.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 38: Hashmap Algorithms

There are three hash functions that you'll implement in this exercise:

### FNV-1a

Named after the creators Glenn Fowler, Phong Vo, and Landon Curt Noll. This hash produces good numbers and is reasonably fast.

### Adler-32

Named after Mark Adler, is a horrible hash algorithm, but it's been around a long time and it's good for studying.

### DJB Hash

This hash algorithm is attributed to Dan J. Bernstein (DJB) but it's difficult to find his discussion of the algorithm. It's shown to be fast, but possibly not great numbers.

You've already seen the Jenkins hash as the default hash for the Hashmap data structure, so this exercise will be looking at these three new ones. The code for them is usually small, and it's not optimized at all. As usual I'm going for understanding and not blinding latest speed.

The header file is very simple, so I'll start with that:

```
#ifndef hashmap_algos_h
#define hashmap_algos_h

#include <stdint.h>

uint32_t Hashmap_fnv1a_hash(void *data);
uint32_t Hashmap_adler32_hash(void *data);
uint32_t Hashmap_djb_hash(void *data);

#endif
```

I'm just declaring the three functions I'll implement in the `hashmap_algos.c` file:

```
#include <lcthw/hashmap_algos.h>
#include <lcthw/bstrlib.h>

// settings taken from
// http://www.isthe.com/chongo/tech/comp/fnv/index.html#FNV-param
const uint32_t FNV_PRIME = 16777619;
const uint32_t FNV_OFFSET_BASIS = 2166136261;

uint32_t Hashmap_fnv1a_hash(void *data)
{
    bstring s = (bstring)data;
    uint32_t hash = FNV_OFFSET_BASIS;
    int i = 0;

    for(i = 0; i < blength(s); i++) {
        hash ^= bchare(s, i, 0);
        hash *= FNV_PRIME;
    }
}
```

```

        return hash;
    }

    const int MOD_ADLER = 65521;

    uint32_t Hashmap_adler32_hash(void *data)
    {
        bstring s = (bstring)data;
        uint32_t a = 1, b = 0;
        int i = 0;

        for (i = 0; i < blength(s); i++)
        {
            a = (a + bchare(s, i, 0)) % MOD_ADLER;
            b = (b + a) % MOD_ADLER;
        }

        return (b << 16) | a;
    }

    uint32_t Hashmap_djb_hash(void *data)
    {
        bstring s = (bstring)data;
        uint32_t hash = 5381;
        int i = 0;

        for(i = 0; i < blength(s); i++) {
            hash = ((hash << 5) + hash) + bchare(s, i, 0); /* hash * 33 + c */
        }

        return hash;
    }
}

```

This file then has the three hash algorithms. You should notice that I'm defaulting to just using a `bstring` for the key, and I'm using the `bchare` function to get a character from the bstring, but return 0 if that character is outside the string's length.

Each of these algorithms are found online so go search for them and read about them. Again I used Wikipedia primarily and then followed it to other sources.

I then have a unit test that tests out each algorithm, but also tests that it will distribute well across a number of buckets:

```

#include <lcthw/bstrlib.h>
#include <lcthw/hashmap.h>
#include <lcthw/hashmap_algos.h>
#include <lcthw/darray.h>
#include "minunit.h"

struct tagbstring test1 = bsStatic("test data 1");
struct tagbstring test2 = bsStatic("test data 2");
struct tagbstring test3 = bsStatic("xest data 3");

char *test_fnv1a()
{
    uint32_t hash = Hashmap_fnv1a_hash(&test1);
    mu_assert(hash != 0, "Bad hash.");

    hash = Hashmap_fnv1a_hash(&test2);
    mu_assert(hash != 0, "Bad hash.");

    hash = Hashmap_fnv1a_hash(&test3);
    mu_assert(hash != 0, "Bad hash.");

    return NULL;
}

```

```

}

char *test_adler32()
{
    uint32_t hash = Hashmap_adler32_hash(&test1);
    mu_assert(hash != 0, "Bad hash.");

    hash = Hashmap_adler32_hash(&test2);
    mu_assert(hash != 0, "Bad hash.");

    hash = Hashmap_adler32_hash(&test3);
    mu_assert(hash != 0, "Bad hash.");

    return NULL;
}

char *test_djb()
{
    uint32_t hash = Hashmap_djb_hash(&test1);
    mu_assert(hash != 0, "Bad hash.");

    hash = Hashmap_djb_hash(&test2);
    mu_assert(hash != 0, "Bad hash.");

    hash = Hashmap_djb_hash(&test3);
    mu_assert(hash != 0, "Bad hash.");

    return NULL;
}

#define BUCKETS 100
#define BUFFER_LEN 20
#define NUM_KEYS BUCKETS * 1000
enum { ALGO_FNV1A, ALGO_ADLER32, ALGO_DJB};

int gen_keys(DArray *keys, int num_keys)
{
    int i = 0;
    FILE *urand = fopen("/dev/urandom", "r");
    check(urand != NULL, "Failed to open /dev/urandom");

    struct bStream *stream = bsopen((bNread)fread, urand);
    check(stream != NULL, "Failed to open /dev/urandom");

    bstring key = bfromcstr("");
    int rc = 0;

    // FNV1a histogram
    for(i = 0; i < num_keys; i++) {
        rc = bsread(key, stream, BUFFER_LEN);
        check(rc >= 0, "Failed to read from /dev/urandom.");

        DArray_push(keys, bstrcpy(key));
    }

    bsclose(stream);
    fclose(urand);
    return 0;

error:
    return -1;
}

void destroy_keys(DArray *keys)
{
    int i = 0;
    for(i = 0; i < NUM_KEYS; i++) {
        bdestroy(DArray_get(keys, i));
    }
}

```

```

    DArray_destroy(keys);
}

void fill_distribution(int *stats, DArray *keys, Hashmap_hash hash_func)
{
    int i = 0;
    uint32_t hash = 0;

    for(i = 0; i < DArray_count(keys); i++) {
        hash = hash_func(DArray_get(keys, i));
        stats[hash % BUCKETS] += 1;
    }
}

char *test_distribution()
{
    int i = 0;
    int stats[3][BUCKETS] = {{0}};
    DArray *keys = DArray_create(0, NUM_KEYS);

    mu_assert(gen_keys(keys, NUM_KEYS) == 0, "Failed to generate random keys.");

    fill_distribution(stats[ALGO_FNV1A], keys, Hashmap_fnv1a_hash);
    fill_distribution(stats[ALGO_ADLER32], keys, Hashmap_adler32_hash);
    fill_distribution(stats[ALGO_DJB], keys, Hashmap_djb_hash);

    fprintf(stderr, "FNV\tA32\tDJB\n");

    for(i = 0; i < BUCKETS; i++) {
        fprintf(stderr, "%d\t%d\t%d\n",
            stats[ALGO_FNV1A][i],
            stats[ALGO_ADLER32][i],
            stats[ALGO_DJB][i]);
    }

    destroy_keys(keys);

    return NULL;
}

char *all_tests()
{
    mu_suite_start();

    mu_run_test(test_fnv1a);
    mu_run_test(test_adler32);
    mu_run_test(test_djb);
    mu_run_test(test_distribution);

    return NULL;
}

RUN_TESTS(all_tests);

```

I have the number of `BUCKETS` in this code set fairly high, since I have a fast enough computer, but if it runs slow just lower them, and also lower `NUM_KEYS`. What this test lets me do is run the test and then look at the distribution of keys for each hash function using a bit of analysis with a language called R.

How I do this is I craft a big list of keys using the `gen_keys` function. These keys are taken out of the `/dev/urandom` device they are random byte keys. I then use these keys to have the `fill_distribution` function fill up the `stats` array with where those keys would hash in a theoretical set of buckets. All this function does is go through all the keys, do the

hash, then do what the `HashMap` would do to find its bucket.

Finally I'm simply printing out a three column table of the final count for each bucket, showing how many keys managed to get into each bucket randomly. I can then look at these numbers to see if the hash functions are distributing keys mostly evenly.

## What You Should See

Teaching you R is outside the scope of this book, but if you want to get it and try this then it can be found at [r-project.org](http://r-project.org).

Here is an abbreviated shell session showing me run the `tests/hashmap_algos_test` to get the table produced by `test_distribution` (not shown here), and then use R to see what the summary statistics are:

```
$ tests/hashmap_algos_tests
# copy-paste the table it prints out
$ vim hash.txt
$ R
> hash <- read.table("hash.txt", header=T)
> summary(hash)
```

FNV		A32		DJB	
Min.	: 945	Min.	: 908.0	Min.	: 927
1st Qu.:	980	1st Qu.:	980.8	1st Qu.:	979
Median	: 998	Median	:1000.0	Median	: 998
Mean	:1000	Mean	:1000.0	Mean	:1000
3rd Qu.:	1016	3rd Qu.:	1019.2	3rd Qu.:	1021
Max.	:1072	Max.	:1075.0	Max.	:1082

```
>
```

First I just run the test, which on your screen will print the table. Then I just copy-paste it out of my terminal and use `vim hash.txt` to save the data. If you look at the data it has the header `FNV A32 DJB` for each of the three algorithms.

Second I run R and load the data using the `read.table` command. This is a smart function that works with this kind of tab-delimited data and I only have to tell it `header=T` so it knows the data has a header.

Finally, I have the data loaded and I can use `summary` to print out its summary statistics for each column. Here you can see that each function actually does alright with this random data. I'll explain what each of these rows means:

Min.

This is the minimum value found for the data in that column. FNV seems to win this on this run since it has the largest number, meaning it has a tighter range at the low end.

1st Qu.

The point where the first quarter of the data ends.

Median

This is the number that is in the middle if you sorted them. Median is most useful when compared to mean.

Mean

Mean is the "average" most people think of, and is the sum/count of the data. If you look, all of them are 1000, which is great. If you compare this to the median you see that all

three have really close medians to the mean. What this means is the data isn't "skewed" in one direction, so you can trust the mean.

3rd Qu.

The point where the last quarter of the data starts and represents the tail end of the numbers.

Max.

This is the maximum number of the data, and presents the upper bound on all of them.

Looking at this data, you see that all of these hashes seem to do good on random keys, and that the means match the `NUM_KEYS` setting I made. What I'm looking for is that if I make 1000 keys per buckets (`BUCKETS * 1000`), then on average each bucket should have 1000 keys in it. If the hash function isn't working then you'll see these summary statistics show a Mean that's not 1000, and really high ranges at the 1st quarter and 3rd quarter. A good hash function should have a dead on 1000 mean, and as tight as possible range.

You should also know that you will get mostly different numbers from mine, and even between different runs of this unit test.

## How To Break It

I'm finally going to have you do some breaking in this exercise. I want you to write the worst hash function you can, and then use the data to prove that it's really bad. You can use R to do the stats, just like I did, but maybe you have another tool you can use to give you the same summary statistics.

The goal is to make a hash function that seems normal to an untrained eye, but when actually run has a bad mean and is all over the place. That means you can't just have it return 1, but have to give a stream of numbers that seem alright, but really are all over the place and loading up some buckets too much.

Extra points if you can make a minimal change to one of the four hash algorithms I gave you to do this.

The purpose of this exercise is to imagine that some "friendly" coder comes to you and offers to improve your hash function, but actually just makes a nice little backdoor that screws up your `Hashmap`.

As the Royal Society says, "Nullius in verba."

## Extra Credit

- Take the `default_hash` out of the `hashmap.c`, make it one of the algorithms in `hashmap_algos.c` and then make all the tests work again.
- Add the `default_hash` to the `hashmap_algos_tests.c` test and compare its statistics to the other hash functions.
- Find a few more hash functions and add them too. You can never have too many hash functions!

**Interested  
In  
Python?**

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

**Interested In Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 39: String Algorithms

In this exercise I'm going to show you one of the supposedly faster string search algorithms, and compare it to the one that exists in `bstrlib.c` call `binstr`. The documentation for `binstr` says that it uses a simple "brute force" string search to find the first instance. The one I'll implement will use the Boyer-Moore-Horspool (BMH) algorithm, which is supposed to be faster if you analyze the theoretical time. You'll see that, assuming my implementation isn't flawed, that the practical time for BMH is much worse than the simple brute force of `binstr`.

The point of this exercise isn't really to explain the algorithm because it's simple enough for you to go to [the Boyer-Moore-Horspool Wikipedia page](#) and read it. The gist of this algorithm is that it calculates a "skip characters list" as a first operation, then it uses this list to quickly scan through the string. It is supposed to be faster than brute force, so let's get the code into the right files and see.

First, I have the header:

```
#ifndef string_algos_h
#define string_algos_h

#include <lcsth/bstrlib.h>
#include <lcsth/darray.h>

typedef struct StringScanner {
    bstring in;
    const unsigned char *haystack;
    ssize_t hlen;
    const unsigned char *needle;
    ssize_t nlen;
    size_t skip_chars[ UCHAR_MAX + 1 ];
} StringScanner;

int String_find(bstring in, bstring what);

StringScanner *StringScanner_create(bstring in);

int StringScanner_scan(StringScanner *scan, bstring tofind);

void StringScanner_destroy(StringScanner *scan);

#endif
```

In order to see the effects of this "skip characters list" I'm going to make two versions of the BMH algorithm:

`String_find`

Simply find the first instance of one string in another, doing the entire algorithm in one shot.

`StringScanner_scan`

Uses a `StringScanner` state structure to separate the skip list build from the actual find. This will let me see what impact that has on performance. This model also has the advantage that I

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@lzssth

can incrementally scan for one string in another and find all instances quickly.

Once you have that, here's the implementation:

```
#include <lcthw/string_algos.h>
#include <limits.h>

static inline void String_setup_skip_chars(
    size_t *skip_chars,
    const unsigned char *needle, ssize_t nlen)
{
    size_t i = 0;
    size_t last = nlen - 1;

    for(i = 0; i < UCHAR_MAX + 1; i++) {
        skip_chars[i] = nlen;
    }

    for (i = 0; i < last; i++) {
        skip_chars[needle[i]] = last - i;
    }
}

static inline const unsigned char *String_base_search(
    const unsigned char *haystack, ssize_t hlen,
    const unsigned char *needle, ssize_t nlen,
    size_t *skip_chars)
{
    size_t i = 0;
    size_t last = nlen - 1;

    assert(haystack != NULL && "Given bad haystack to search.");
    assert(needle != NULL && "Given bad needle to search for.");

    check(nlen > 0, "nlen can't be <= 0");
    check(hlen > 0, "hlen can't be <= 0");

    while (hlen >= nlen)
    {
        for (i = last; haystack[i] == needle[i]; i--) {
            if (i == 0) {
                return haystack;
            }
        }

        hlen -= skip_chars[haystack[last]];
        haystack += skip_chars[haystack[last]];
    }

    error: // fallback
    return NULL;
}

int String_find(bstring in, bstring what)
{
    const unsigned char *found = NULL;

    const unsigned char *haystack = (const unsigned char *)bdata(in);
    ssize_t hlen = blength(in);
    const unsigned char *needle = (const unsigned char *)bdata(what);
    ssize_t nlen = blength(what);
    size_t skip_chars[UCHAR_MAX + 1] = {0};

    String_setup_skip_chars(skip_chars, needle, nlen);

    found = String_base_search(haystack, hlen, needle, nlen, skip_chars);

    return found != NULL ? found - haystack : -1;
}
```

```

}

StringScanner *StringScanner_create(bstring in)
{
    StringScanner *scan = calloc(1, sizeof(StringScanner));
    check_mem(scan);

    scan->in = in;
    scan->haystack = (const unsigned char *)bdata(in);
    scan->hlen = blength(in);

    assert(scan != NULL && "fuck");
    return scan;

error:
    free(scan);
    return NULL;
}

static inline void StringScanner_set_needle(StringScanner *scan, bstring tofind)
{
    scan->needle = (const unsigned char *)bdata(tofind);
    scan->nlen = blength(tofind);

    String_setup_skip_chars(scan->skip_chars, scan->needle, scan->nlen);
}

static inline void StringScanner_reset(StringScanner *scan)
{
    scan->haystack = (const unsigned char *)bdata(scan->in);
    scan->hlen = blength(scan->in);
}

int StringScanner_scan(StringScanner *scan, bstring tofind)
{
    const unsigned char *found = NULL;
    ssize_t found_at = 0;

    if(scan->hlen <= 0) {
        StringScanner_reset(scan);
        return -1;
    }

    if((const unsigned char *)bdata(tofind) != scan->needle) {
        StringScanner_set_needle(scan, tofind);
    }

    found = String_base_search(
        scan->haystack, scan->hlen,
        scan->needle, scan->nlen,
        scan->skip_chars);

    if(found) {
        found_at = found - (const unsigned char *)bdata(scan->in);
        scan->haystack = found + scan->nlen;
        scan->hlen -= found_at - scan->nlen;
    } else {
        // done, reset the setup
        StringScanner_reset(scan);
        found_at = -1;
    }

    return found_at;
}

void StringScanner_destroy(StringScanner *scan)
{
    if(scan) {

```

```

        free(scan);
    }
}

```

The entire algorithm is in two `static inline` functions called `String_setup_skip_chars` and `String_base_search`. These are then used in the other functions to actually implement the searching styles I want. Study these first two functions and compare them to the Wikipedia description so you know what's going on.

The `String_find` then just uses these two functions to do a find and return the position found. It's very simple and I'll use it to see how this "build skip chars" phase impacts real practical performance. Keep in mind that you could maybe make this faster, but I'm teaching you how to confirm theoretical speed after you implement an algorithm.

The `StringScanner_scan` function is then following the common pattern I use of "create, scan, destroy" and is used to incrementally scan a string for another string. You'll see how this is used when I show you the unit test that will test this out.

Finally, I have the unit test that first confirms this is all working, then runs simple performance tests for all three finding algorithms in a *commented out section*.

```

#include "minunit.h"
#include <lcthw/string_algos.h>
#include <lcthw/bstrlib.h>
#include <time.h>

struct tagbstring IN_STR = bsStatic("I have ALPHA beta ALPHA and oranges ALPHA");
struct tagbstring ALPHA = bsStatic("ALPHA");
const int TEST_TIME = 1;

char *test_find_and_scan()
{
    StringScanner *scan = StringScanner_create(&IN_STR);
    mu_assert(scan != NULL, "Failed to make the scanner.");

    int find_i = String_find(&IN_STR, &ALPHA);
    mu_assert(find_i > 0, "Failed to find 'ALPHA' in test string.");

    int scan_i = StringScanner_scan(scan, &ALPHA);
    mu_assert(scan_i > 0, "Failed to find 'ALPHA' with scan.");
    mu_assert(scan_i == find_i, "find and scan don't match");

    scan_i = StringScanner_scan(scan, &ALPHA);
    mu_assert(scan_i > find_i, "should find another ALPHA after the first");

    scan_i = StringScanner_scan(scan, &ALPHA);
    mu_assert(scan_i > find_i, "should find another ALPHA after the first");

    mu_assert(StringScanner_scan(scan, &ALPHA) == -1, "shouldn't find it");

    StringScanner_destroy(scan);

    return NULL;
}

char *test_binstr_performance()
{
    int i = 0;
    int found_at = 0;
    unsigned long find_count = 0;
    time_t elapsed = 0;

```

```

time_t start = time(NULL);

do {
    for(i = 0; i < 1000; i++) {
        found_at = bistr(&IN_STR, 0, &ALPHA);
        mu_assert(found_at != BSTR_ERR, "Failed to find!");
        find_count++;
    }

    elapsed = time(NULL) - start;
} while(elapsed <= TEST_TIME);

debug("BINSTR COUNT: %lu, END TIME: %d, OPS: %f",
      find_count, (int)elapsed, (double)find_count / elapsed);
return NULL;
}

char *test_find_performance()
{
    int i = 0;
    int found_at = 0;
    unsigned long find_count = 0;
    time_t elapsed = 0;
    time_t start = time(NULL);

    do {
        for(i = 0; i < 1000; i++) {
            found_at = String_find(&IN_STR, &ALPHA);
            find_count++;
        }

        elapsed = time(NULL) - start;
    } while(elapsed <= TEST_TIME);

    debug("FIND COUNT: %lu, END TIME: %d, OPS: %f",
          find_count, (int)elapsed, (double)find_count / elapsed);

    return NULL;
}

char *test_scan_performance()
{
    int i = 0;
    int found_at = 0;
    unsigned long find_count = 0;
    time_t elapsed = 0;
    StringScanner *scan = StringScanner_create(&IN_STR);

    time_t start = time(NULL);

    do {
        for(i = 0; i < 1000; i++) {
            found_at = 0;

            do {
                found_at = StringScanner_scan(scan, &ALPHA);
                find_count++;
            } while(found_at != -1);
        }

        elapsed = time(NULL) - start;
    } while(elapsed <= TEST_TIME);

    debug("SCAN COUNT: %lu, END TIME: %d, OPS: %f",
          find_count, (int)elapsed, (double)find_count / elapsed);

    StringScanner_destroy(scan);

    return NULL;
}

```

```

char *all_tests()
{
    mu_suite_start();

    mu_run_test(test_find_and_scan);

    // this is an idiom for commenting out sections of code
    #if 0
        mu_run_test(test_scan_performance);
        mu_run_test(test_find_performance);
        mu_run_test(test_binstr_performance);
    #endif

    return NULL;
}

RUN_TESTS(all_tests);

```

I have it written here with `#if 0` which is a way to use the CPP to comment out a section of code. Type it in like this, and then remove that and the `#endif` so you can see these performance tests run. When you continue with the book, simply comment these out so that the test doesn't waste development time.

There's nothing amazing in this unit test, it just runs each of the different functions in loops that last long enough to get a few seconds of sampling. The first test (`test_find_and_scan`) just confirms that what I've written works, because there's no point in testing the speed of something that doesn't work. Then the next three functions run a large number of searches using each of the three functions.

The trick to notice is that I grab the starting time in `start`, and then I loop until at least `TEST_TIME` seconds have passed. This makes sure that I get enough samples to work with in comparing the three. I'll then run this test with different `TEST_TIME` settings and analyze the results.

## What You Should See

When I run this test on my laptop, I get number that look like this:

```

$ ./tests/string_algos_tests
DEBUG tests/string_algos_tests.c:124: ----- RUNNING: ./tests/string_algos_test
s
-----
RUNNING: ./tests/string_algos_tests
DEBUG tests/string_algos_tests.c:116:
----- test_find_and_scan
DEBUG tests/string_algos_tests.c:117:
----- test_scan_performance
DEBUG tests/string_algos_tests.c:105: SCAN COUNT: 110272000, END TIME: 2, OPS:
55136000.000000
DEBUG tests/string_algos_tests.c:118:
----- test_find_performance
DEBUG tests/string_algos_tests.c:76: FIND COUNT: 12710000, END TIME: 2, OPS: 6
355000.000000
DEBUG tests/string_algos_tests.c:119:
----- test_binstr_performance
DEBUG tests/string_algos_tests.c:54: BINSTR COUNT: 72736000, END TIME: 2, OPS:
36368000.000000
ALL TESTS PASSED
Tests run: 4
$

```

I look at this and I sort of want to do more than 2 seconds of each run, and I want to run this many times then use R to check it out like I did before. Here's what I get for 10 samples of about 10 seconds each:

```
scan find binstr
71195200 6353700 37110200
75098000 6358400 37420800
74910000 6351300 37263600
74859600 6586100 37133200
73345600 6365200 37549700
74754400 6358000 37162400
75343600 6630400 37075000
73804800 6439900 36858700
74995200 6384300 36811700
74781200 6449500 37383000
```

The way I got this is with a little bit of shell help and then editing the output:

```
$ for i in 1 2 3 4 5 6 7 8 9 10; do echo "RUN --- $i" >> times.log; ./tests/st
ring_algos_tests 2>&1 | grep COUNT >> times.log ; done
$ less times.log
$ vim times.log
```

Right away you can see that the scanning system beats the pants off both of the others, but I'll open this in R and confirm the results:

```
> times <- read.table("times.log", header=T)
> summary(times)
      scan      find      binstr
Min.   :71195200  Min.   :6351300  Min.   :36811700
1st Qu.:74042200  1st Qu.:6358100  1st Qu.:37083800
Median :74820400  Median :6374750  Median :37147800
Mean   :74308760  Mean   :6427680  Mean   :37176830
3rd Qu.:74973900  3rd Qu.:6447100  3rd Qu.:37353150
Max.   :75343600  Max.   :6630400  Max.   :37549700
>
```

To understand why I'm getting the summary statistics I have to explain some statistics for you. What I'm looking for in these numbers can be said simply to be, "Are these three functions (scan, find, bsinter) actually different?" I know that each time I run my tester function I get slightly different numbers, and that those numbers can cover a certain range. You see here that the 1st and 3rd quarters do that for each sample.

What I look at first is the mean and I want to see if each sample's mean is different from the others. I can see that, and clearly the `scan` beats `binstr` which also beats `find`. However, I have a problem, if I use just the mean, there's a *chance* that the `ranges` of each sample might overlap.

What if I have means that are different, but the 1st and 3rd quarters overlap? In that case I could say that there's a chance that if I ran the samples again the means might not be different. The more overlap I have in the ranges the higher probability that my two samples (and my two functions) are *not* actually different. Any difference I'm seeing in the two (in this case three) is just random chance.

Statistics has many tools to solve this problem, but in our case I can just look at the 1st and 3rd quarters as well as the mean for all three samples. If the means are different and the quarters are way off never possibly overlapping, then it's alright to say they are different.

In my three samples I can say that `scan`, `find` and `binstr` are different, don't overlap in range, and that I can trust the sample (for the most part).

## Analyzing The Results

Looking at the results I can see that `String_find` is much slower than the other two. In fact, so slow I'd think there's something wrong with how I implemented it. However when I compare it with `StringScanner_scan` I can see that it's the part that builds the skip list that is most likely costing the time. Not only is `find` slower, it's also doing *less* than `scan` because it's just finding the first string while `scan` finds all of them.

I can also see that `scan` beats `binstr` as well by quite a large margin. Again I can say that not only does `scan` do more than both of these, but it's also much faster.

There's a few caveats with this analysis:

- I may have messed up this implementation or the test. At this point I would go research all the possible ways to do a BMH algorithm and try to improve it. I would also confirm that I'm doing the test right.
- If you alter the time the test runs, you get different results. There is a "warm up" period I'm not investigating.
- The `test_scan_performance` unit test isn't quite the same as the others, but it is doing more than the other tests so it's probably alright.
- I'm only doing the test by searching for one string in another. I could randomize the strings to find to remove their position and length as a confounding factor.
- Maybe `binstr` is implemented better than "simple" brute force.
- I could be running these in an unfortunate order and maybe randomizing which test runs first will give better results.

One thing to gather from this is you need to confirm real performance even if you implement an algorithm "correctly". In this case the claim is that the BMH algorithm should have beaten the `binstr` algorithm, but a simple test proved it didn't. Had I not done this I would have been using an inferior algorithm implementation without knowing it. With these metrics I can start to tune my implementation, or simply scrap it and find another one.

## Extra Credit

- See if you can make the `Scan_find` faster. Why is my implementation here slow?
- Try some different scan times and see if you get different numbers. What impact does the length of time that you run the test have on the `scan` times? What can you say about that result?
- Alter the unit test so that it runs each function for a short burst in the beginning to clear out any "warm up" period, then start the timing portion. Does that change the dependence on the length of time the test runs and how many operations / second are



possible?

- Make the unit test randomize the strings to find and then measure the performance you get. One way to do this is use the `bssplit` function from `bstrlib.h` to split the `IN_STR` on spaces. Then use the `bstrList` struct you get to access each string it returns. This will also teach you how to use `bstrList` operations for string processing.
- Try some runs with the tests in different orders and see if you get different results.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 40: Binary Search Trees

The binary tree is the simplest tree based data structure and while it has been replaced by Hash Maps in many languages is still useful for many applications. Variants on the binary tree exist for very useful things like database indexes, search algorithm structures, and even graphics processing.

I'm calling my binary tree a `BSTree` for "binary search tree" and the best way to describe it is that it's another way to do a `HashMap` style key/value store. The difference is that instead of hashing the key to find a location, the `BSTree` compares the key to nodes in a tree, and then walks the tree to find the best place to store it based on how it compares to other nodes.

Before I really explain how this works, let me show you the `bstree.h` header file so you can see the data structures, then I can use that to explain how it's built.

```
#ifndef _lcthw_BSTree_h
#define _lcthw_BSTree_h

typedef int (*BSTree_compare)(void *a, void *b);

typedef struct BSTreeNode {
    void *key;
    void *data;

    struct BSTreeNode *left;
    struct BSTreeNode *right;
    struct BSTreeNode *parent;
} BSTreeNode;

typedef struct BSTree {
    int count;
    BSTree_compare compare;
    BSTreeNode *root;
} BSTree;

typedef int (*BSTree_traverse_cb)(BSTreeNode *node);

BSTree *BSTree_create(BSTree_compare compare);
void BSTree_destroy(BSTree *map);

int BSTree_set(BSTree *map, void *key, void *data);
void *BSTree_get(BSTree *map, void *key);

int BSTree_traverse(BSTree *map, BSTree_traverse_cb traverse_cb);

void *BSTree_delete(BSTree *map, void *key);

#endif
```

This follows the same pattern I've been using this whole time where I have a base "container" named `BSTree` and that then has nodes names `BSTreeNode` that make up the actual contents. Bored yet? Good, there's

no reason to be clever with this kind of structure.

The important part is how the `BSTreeNode` is configured and how it gets used to do each operation: set, get, and delete. I'll cover get first since it's the easiest operation and I'll pretend I'm doing it manually against the data structure:

- I take the key you're looking for and I start at the root. First thing I do is compare your key with that node's key.
- If your key is less-than the `node.key`, then I traverse down the tree using the `left` pointer.
- If your key is greater-than the `node.key`, then I go down with `right`.
- I repeat step 2 and 3 until either I find a matching `node.key`, or I get to a node that has no left and right. In the first case I return the `node.data`, in the second I return `NULL`.

That's all there is to `get`, so now to do `set` it's nearly the same thing, except you're looking for where to put a new node:

- If there is no `BSTree.root` then I just make that and we're done. That's the first node.
- After that I compare your key to `node.key`, starting at the root.
- If your key is less-than or equal to the `node.key` then I want to go left. If your key is greater-than (not equal) then I want to go right.
- I keep repeating 3 until I reach a node where the left or right doesn't exist, but that's the direction I need to go.
- Once there I set that direction (left or right) to a new node for the key and data I want, and set this new node's parent to the previous node I came from. I'll use the parent node when I do delete.

This also makes sense given how `get` works. If finding a node involves going left or right depending on how they key compares, well then setting a node involves the same thing until I can set the left or right for a new node.

Take some time to draw out a few trees on paper and go through some setting and getting nodes so you understand how it work. After that you are ready to look at the implementation so that I can explain delete. Deleting in trees is a *major* pain, and so it's best explained by doing a line-by-line code breakdown.

```
#include <lcthw/dbg.h>
#include <lcthw/bstree.h>
#include <stdlib.h>
#include <lcthw/bstrlib.h>

static int default_compare(void *a, void *b)
{
    return bstrcmp((bstring)a, (bstring)b);
}

BSTree *BSTree_create(BSTree_compare compare)
{
    BSTree *map = calloc(1, sizeof(BSTree));
    check_mem(map);

    map->compare = compare == NULL ? default_compare : compare;
```

```

        return map;

error:
    if(map) {
        BSTree_destroy(map);
    }
    return NULL;
}

static int BSTree_destroy_cb(BSTreeNode *node)
{
    free(node);
    return 0;
}

void BSTree_destroy(BSTree *map)
{
    if(map) {
        BSTree_traverse(map, BSTree_destroy_cb);
        free(map);
    }
}

static inline BSTreeNode *BSTreeNode_create(BSTreeNode *parent, void *key, void
id *data)
{
    BSTreeNode *node = calloc(1, sizeof(BSTreeNode));
    check_mem(node);

    node->key = key;
    node->data = data;
    node->parent = parent;
    return node;

error:
    return NULL;
}

static inline void BSTree_setnode(BSTree *map, BSTreeNode *node, void *key, vo
id *data)
{
    int cmp = map->compare(node->key, key);

    if(cmp <= 0) {
        if(node->left) {
            BSTree_setnode(map, node->left, key, data);
        } else {
            node->left = BSTreeNode_create(node, key, data);
        }
    } else {
        if(node->right) {
            BSTree_setnode(map, node->right, key, data);
        } else {
            node->right = BSTreeNode_create(node, key, data);
        }
    }
}

int BSTree_set(BSTree *map, void *key, void *data)
{
    if(map->root == NULL) {
        // first so just make it and get out
        map->root = BSTreeNode_create(NULL, key, data);
        check_mem(map->root);
    } else {
        BSTree_setnode(map, map->root, key, data);
    }
}

```

```

    }

    return 0;
error:
    return -1;
}

static inline BSTreeNode *BSTree_getnode(BSTree *map, BSTreeNode *node, void *
key)
{
    int cmp = map->compare(node->key, key);

    if(cmp == 0) {
        return node;
    } else if(cmp < 0) {
        if(node->left) {
            return BSTree_getnode(map, node->left, key);
        } else {
            return NULL;
        }
    } else {
        if(node->right) {
            return BSTree_getnode(map, node->right, key);
        } else {
            return NULL;
        }
    }
}

void *BSTree_get(BSTree *map, void *key)
{
    if(map->root == NULL) {
        return NULL;
    } else {
        BSTreeNode *node = BSTree_getnode(map, map->root, key);
        return node == NULL ? NULL : node->data;
    }
}

static inline int BSTree_traverse_nodes(BSTreeNode *node, BSTree_traverse_cb t
raverse_cb)
{
    int rc = 0;

    if(node->left) {
        rc = BSTree_traverse_nodes(node->left, traverse_cb);
        if(rc != 0) return rc;
    }

    if(node->right) {
        rc = BSTree_traverse_nodes(node->right, traverse_cb);
        if(rc != 0) return rc;
    }

    return traverse_cb(node);
}

int BSTree_traverse(BSTree *map, BSTree_traverse_cb traverse_cb)
{
    if(map->root) {
        return BSTree_traverse_nodes(map->root, traverse_cb);
    }

    return 0;
}

static inline BSTreeNode *BSTree_find_min(BSTreeNode *node)
{
    while(node->left) {

```

```

        node = node->left;
    }

    return node;
}

static inline void BSTree_replace_node_in_parent(BSTree *map, BSTreeNode *node
, BSTreeNode *new_value)
{
    if(node->parent) {
        if(node == node->parent->left) {
            node->parent->left = new_value;
        } else {
            node->parent->right = new_value;
        }
    } else {
        // this is the root so gotta change it
        map->root = new_value;
    }

    if(new_value) {
        new_value->parent = node->parent;
    }
}

static inline void BSTree_swap(BSTreeNode *a, BSTreeNode *b)
{
    void *temp = NULL;
    temp = b->key; b->key = a->key; a->key = temp;
    temp = b->data; b->data = a->data; a->data = temp;
}

static inline BSTreeNode *BSTree_node_delete(BSTree *map, BSTreeNode *node, vo
id *key)
{
    int cmp = map->compare(node->key, key);

    if(cmp < 0) {
        if(node->left) {
            return BSTree_node_delete(map, node->left, key);
        } else {
            // not found
            return NULL;
        }
    } else if(cmp > 0) {
        if(node->right) {
            return BSTree_node_delete(map, node->right, key);
        } else {
            // not found
            return NULL;
        }
    } else {
        if(node->left && node->right) {
            // swap this node for the smallest node that is bigger than us
            BSTreeNode *successor = BSTree_find_min(node->right);
            BSTree_swap(successor, node);

            // this leaves the old successor with possibly a right child
            // so replace it with that right child
            BSTree_replace_node_in_parent(map, successor, successor->right);

            // finally it's swapped, so return successor instead of node
            return successor;
        } else if(node->left) {
            BSTree_replace_node_in_parent(map, node, node->left);
        } else if(node->right) {
            BSTree_replace_node_in_parent(map, node, node->right);
        } else {
            BSTree_replace_node_in_parent(map, node, NULL);
        }
    }
}

```

```

        return node;
    }
}

void *BSTree_delete(BSTree *map, void *key)
{
    void *data = NULL;

    if(map->root) {
        BSTreeNode *node = BSTree_node_delete(map, map->root, key);

        if(node) {
            data = node->data;
            free(node);
        }
    }

    return data;
}

```

Before getting into how `BSTree_delete` works I want to explain a pattern I'm using for doing recursive function calls in a sane way. You'll find that many tree based data structures are easy to write if you use recursion, but that formulating a single recursive function is difficult. Part of the problem is that you need to setup some initial data for the first operation, *then* recurse into the data structure, which is hard to do with one function.

The solution is to use two functions. One function "sets up" the data structure and initial recursion conditions so that a second function can do the real work. Take a look at `BSTree_get` first to see what I mean:

- I have an initial condition to handle that if `map->root` is `NULL` then return `NULL` and don't recurse.
- I then setup a call to the real recursion, which is in `BSTree_getnode`. I create the initial condition of the root node to start with, the key, and the `map`.
- In the `BSTree_getnode` then I do the actual recursive logic. I compare the keys with `map->compare(node->key, key)` and go left, right, or equal depending on that.
- Since this function is "self-similar" and doesn't have to handle any initial conditions (because `BSTree_get` did) then I can structure it very simply. When it's done it returns to the caller, and that return then comes back to `BSTree_get` the result.
- At the end, the `BSTree_get` then handles getting the `node.data` element but only if the result isn't `NULL`.

This way of structuring a recursive algorithm matches the way I structure my recursive data structures. I have an initial "base function" that handles initial conditions and some edge cases, then it calls a clean recursive function that does the work. Compare that with how I have a "base struct" in `BSTree` combined with recursive `BSTreeNode` structures that all reference each other in a tree. Using this pattern makes it easy to deal with recursion and keep it straight.

Next, go look at `BSTree_set` and `BSTree_setnode` to see the exact same pattern going on. I use `BSTree_set` to configure the initial conditions and edge cases. A common edge case is that there's no root node, so I have to make one to get things started.

This pattern will work with nearly any recursive algorithm you have to figure out. The way I do this is I follow this pattern:

- Figure out the initial variables, how they change, and what the stopping conditions are for each recursive step.
- Write a recursive function that calls itself, with arguments for each stopping condition and initial variable.
- Write a setup function to set initial starting conditions for the algorithm and handle edge cases, then it calls the recursive function.
- Finally, the setup function returns the final result and possibly alters it if the recursive function can't handle final edge cases.

Which leads me finally to `BSTree_delete` and `BSTree_node_delete`. First you can just look at `BSTree_delete` and see that it's the setup function, and what it is doing is grabbing the resulting node data and freeing the node that's found. In `BSTree_node_delete` is where things get complex because to delete a node at any point in the tree, I have to *rotate* that node's children up to the parent. I'll break this function down and the ones it uses:

bstree.c:190

I run the compare function to figure out which direction I'm going.

bstree.c:192-198

This is the usual "less-than" branch where I want to go left. I'm handling the case that left doesn't exist here and returning `NULL` to say "not found". This handles deleting something that isn't in the `BSTree`.

bstree.c:199-205

The same thing but for the right branch of the tree. Just keep recursing down into the tree just like in the other functions, and return `NULL` if it doesn't exist.

bstree.c:206

This is where I have found the node since the key is equal (`compare` return 0).

bstree.c:207

This node has both a `left` and `right` branch, so it's deeply embedded in the tree.

bstree.c:209

To remove this node I need to first find the smallest node that's greater than this node, which means I call `BSTree_find_min` on the right child.

bstree.c:210

Once I have this node, I will do a swap on its `key` and `data` with the current node's. This will effectively take this node that was down at the bottom of the tree, and put it's contents here so that I don't have to try and shuffle this node out by its pointers.

bstree.c:214

The `successor` is now this dead branch that has the current node's values. It could just be removed, but there's a chance that it has a right node value, which means I need to do a single rotate so that the successor's right node gets moved up to completely detach it.

bstree.c:217



At this point, the successor is removed from the tree, its values replaced the current node's values, and any children it had are moved up into the parent. I can return the `successor` as if it were the `node`.

bstree.c:218

At this branch I know that the node has a left but no right, so I want to replace this node with its left child.

bstree.c:219

I again use `BSTree_replace_node_in_parent` to do the replace, rotating the left child up.

bstree.c:220

This branch of the if-statement means I have a right child but no left child, so I want to rotate the right child up.

bstree.c:221

Again, use the function to do the rotate, but this time of the right node.

bstree.c:222

Finally, the only thing that's left is the condition that I've found the node, and it has no children (no left or right). In this case, I simply replace this node with NULL using the same function I did with all the others.

bstree.c:210

After all that, I have the current node rotated out of the tree and replaced with some child element that will fit in the tree. I just return this to the caller so it can be freed and managed.

This operation is very complex, and to be honest, in some tree data structures I just don't bother doing deletes and treat them like constant data in my software. If I need to do heavy insert and delete, I use a `Hashmap` instead.

Finally, you can look at the unit test to see how I'm testing it:

```
#include "minunit.h"
#include <lcthw/bstree.h>
#include <assert.h>
#include <lcthw/bstrlib.h>
#include <stdlib.h>
#include <time.h>

BSTree *map = NULL;
static int traverse_called = 0;
struct tagbstring test1 = bsStatic("test data 1");
struct tagbstring test2 = bsStatic("test data 2");
struct tagbstring test3 = bsStatic("test data 3");
struct tagbstring expect1 = bsStatic("THE VALUE 1");
struct tagbstring expect2 = bsStatic("THE VALUE 2");
struct tagbstring expect3 = bsStatic("THE VALUE 3");

static int traverse_good_cb(BSTreeNode *node)
{
    debug("KEY: %s", bdata((bstring)node->key));
    traverse_called++;
    return 0;
}

static int traverse_fail_cb(BSTreeNode *node)
{
    debug("KEY: %s", bdata((bstring)node->key));
    traverse_called++;
}
```

```

        if(traverse_called == 2) {
            return 1;
        } else {
            return 0;
        }
    }

char *test_create()
{
    map = BSTree_create(NULL);
    mu_assert(map != NULL, "Failed to create map.");

    return NULL;
}

char *test_destroy()
{
    BSTree_destroy(map);

    return NULL;
}

char *test_get_set()
{
    int rc = BSTree_set(map, &test1, &expect1);
    mu_assert(rc == 0, "Failed to set &test1");
    bstring result = BSTree_get(map, &test1);
    mu_assert(result == &expect1, "Wrong value for test1.");

    rc = BSTree_set(map, &test2, &expect2);
    mu_assert(rc == 0, "Failed to set test2");
    result = BSTree_get(map, &test2);
    mu_assert(result == &expect2, "Wrong value for test2.");

    rc = BSTree_set(map, &test3, &expect3);
    mu_assert(rc == 0, "Failed to set test3");
    result = BSTree_get(map, &test3);
    mu_assert(result == &expect3, "Wrong value for test3.");

    return NULL;
}

char *test_traverse()
{
    int rc = BSTree_traverse(map, traverse_good_cb);
    mu_assert(rc == 0, "Failed to traverse.");
    mu_assert(traverse_called == 3, "Wrong count traverse.");

    traverse_called = 0;
    rc = BSTree_traverse(map, traverse_fail_cb);
    mu_assert(rc == 1, "Failed to traverse.");
    mu_assert(traverse_called == 2, "Wrong count traverse for fail.");

    return NULL;
}

char *test_delete()
{
    bstring deleted = (bstring)BSTree_delete(map, &test1);
    mu_assert(deleted != NULL, "Got NULL on delete.");
    mu_assert(deleted == &expect1, "Should get test1");
    bstring result = BSTree_get(map, &test1);
    mu_assert(result == NULL, "Should delete.");

    deleted = (bstring)BSTree_delete(map, &test1);
    mu_assert(deleted == NULL, "Should get NULL on delete");

    deleted = (bstring)BSTree_delete(map, &test2);

```

```

mu_assert(deleted != NULL, "Got NULL on delete.");
mu_assert(deleted == &expect2, "Should get test2");
result = BSTree_get(map, &test2);
mu_assert(result == NULL, "Should delete.");

deleted = (bstring)BSTree_delete(map, &test3);
mu_assert(deleted != NULL, "Got NULL on delete.");
mu_assert(deleted == &expect3, "Should get test3");
result = BSTree_get(map, &test3);
mu_assert(result == NULL, "Should delete.");

// test deleting non-existent stuff
deleted = (bstring)BSTree_delete(map, &test3);
mu_assert(deleted == NULL, "Should get NULL");

return NULL;
}

char *test_fuzzing()
{
    BSTree *store = BSTree_create(NULL);
    int i = 0;
    int j = 0;
    bstring numbers[100] = {NULL};
    bstring data[100] = {NULL};
    srand((unsigned int)time(NULL));

    for(i = 0; i < 100; i++) {
        int num = rand();
        numbers[i] = bformat("%d", num);
        data[i] = bformat("data %d", num);
        BSTree_set(store, numbers[i], data[i]);
    }

    for(i = 0; i < 100; i++) {
        bstring value = BSTree_delete(store, numbers[i]);
        mu_assert(value == data[i], "Failed to delete the right number.");

        mu_assert(BSTree_delete(store, numbers[i]) == NULL, "Should get nothing.");

        for(j = i+1; j < 99 - i; j++) {
            bstring value = BSTree_get(store, numbers[j]);
            mu_assert(value == data[j], "Failed to get the right number.");
        }

        bdestroy(value);
        bdestroy(numbers[i]);
    }

    BSTree_destroy(store);

    return NULL;
}

char *all_tests()
{
    mu_suite_start();

    mu_run_test(test_create);
    mu_run_test(test_get_set);
    mu_run_test(test_traverse);
    mu_run_test(test_delete);
    mu_run_test(test_destroy);
    mu_run_test(test_fuzzing);

    return NULL;
}

RUN_TESTS(all_tests);

```

I'll point you at the `test_fuzzing` function, which is an interesting technique for testing complex data structures. It is difficult to create a set of keys that cover all the branches in `BSTree_node_delete`, and chances are I would miss some edge case. A better way is to create a "fuzz" function that does all the operations, but does them in as horrible and random a way as possible. In this case I'm inserting a set of random string keys, then I'm deleting them and trying to get the rest after each delete.

Doing this prevents the situation where you test only what you know to work, which means you'll miss things you don't know. By throwing random junk at your data structures you'll hit things you didn't expect and work out any bugs you have.

## How To Improve It

Do *not* do any of these yet since in the next exercise I'll be using this unit test to teach you some more performance tuning tricks. You'll come back and do these after you do Exercise 41.

- As usual, you should go through all the defensive programming checks and add `assert`s for conditions that shouldn't happen. For example, you should not be getting `NULL` values for the recursion functions, so assert that.`
- The traverse function traverses the tree in order by traversing left, then right, then the current node. You can create traverse functions for reverse order as well.
- It does a full string compare on every node, but I could use the `Hashmap` hashing functions to speed this up. I could hash the keys, then keep the hash in the `BSTreeNode`. Then in each of the set up functions I can hash the key ahead of time, and pass it down to the recursive function. Using this hash I can then compare each node much quicker similar to I do in `Hashmap`.

## Extra Credit

Again, do *not* do these yet, wait until Exercise 41 when you can use performance tuning features of Valgrind to do them.

- There's an alternative way to do this data structure without using recursion. The Wikipedia page shows alternatives that don't use recursion but do the same thing. Why would this be better or worse?
- Read up on all the different similar trees you can find. There's AVL trees, Red-Black trees, and some non-tree structures like Skip Lists.

**Interested  
In  
Python?**

Python is also a  
great language.

**Learn  
Python The**

**Interested In Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 41: Using Cachegrind And Callgrind For Performance Tuning

In this exercise I'm going to give you a quick course in using two tools for `Valgrind` called `callgrind` and `cachegrind`. These two tools will analyze your program's execution and tell you what parts are running slow. The results are accurate because of the way `Valgrind` works and help you spot problems such as lines of code that execute too much, hot spots, memory access problems, and even CPU cache misses.

To do this exercise I'm going to use the `bstree_tests` unit tests you just did to look for places to improve the algorithms used. Make sure your versions of these programs are running without any `valgrind` errors and that it is exactly like my code. I'll be using dumps of my code to talk about how `cachegrind` and `callgrind` work.

### Running Callgrind

To run callgrind you pass the `--tool=callgrind` option to `valgrind` and it will produce a `callgrind.out.PID` file (where PID is replace with the process ID of the program that ran). Once you run it you can analyze this `callgrind.out` file with a tool called `callgrind_annotate` which will tell you which functions used the most instructions to run. Here's an example of me running `callgrind` on `bstree_tests` and then getting its information:

```
$ valgrind --dsymutil=yes --tool=callgrind tests/bstree_tests
...
$ callgrind_annotate callgrind.out.1232
-----
--
Profile data file 'callgrind.out.1232' (creator: callgrind-3.7.0.SVN)
-----
--
I1 cache:
D1 cache:
LL cache:
Timerange: Basic block 0 - 1098689
Trigger: Program termination
Profiled target: tests/bstree_tests (PID 1232, part 1)
Events recorded: Ir
Events shown: Ir
Event sort order: Ir
Thresholds: 99
Include dirs:
User annotated:
Auto-annotation: off

-----
--
Ir
-----
--
4,605,808 PROGRAM TOTALS
```

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶ NEXT

⛶  
HELP

Follow  
@l2sthw

```

--
Ir file:function
--
670,486 src/lcthw/bstrlib.c:bstrcmp [tests/bstree_tests]
194,377 src/lcthw/bstree.c:BSTree_get [tests/bstree_tests]
65,580 src/lcthw/bstree.c:default_compare [tests/bstree_tests]
16,338 src/lcthw/bstree.c:BSTree_delete [tests/bstree_tests]
13,000 src/lcthw/bstrlib.c:bformat [tests/bstree_tests]
11,000 src/lcthw/bstrlib.c:bfromcstralloc [tests/bstree_tests]
7,774 src/lcthw/bstree.c:BSTree_set [tests/bstree_tests]
5,800 src/lcthw/bstrlib.c:bdestroy [tests/bstree_tests]
2,323 src/lcthw/bstree.c:BSTreeNode_create [tests/bstree_tests]
1,183 /private/tmp/pkg-build/coregrind/vg_preloaded.c:vg_cleanup_env [/u
sr/local/lib/valgrind/vgpreload_core-amd64-darwin.so]

$

```

I've removed the unit test run and the `valgrind` output since it's not very useful for this exercise. What you should look at is the `callgrind_annotate` output. What this shows you is the number of instructions run (which `valgrind` calls `Ir`) for each function, and the functions sorted highest to lowest. You can usually ignore the header data and just jump to the list of functions.

## NOTE

In if you get a ton of "???Image" lines and things that are not in your program then you're picking up junk from the OS. Just add `| grep -v "???"` at the end to filter those out, like this.

I can now do a quick breakdown of this output to figure out where to look next:

- Each line lists the number of `Ir` and the `file:function` that executed them. The `Ir` is just the instruction count, and if you make that lower then you have made it faster. There's some complexity to that, but at first just focus on getting the `Ir` down.
- The way to attack this is to look at your top functions, and then see which one you think you can improve first. In this case, I'd look at improving `bstrcmp` or `BSTree_get`. It's probably easier to start with `BSTree_get`.
- Some of these functions are just called from the unit test, so I would just ignore those. Functions like `bformat`, `bfromcstralloc`, and `bdestroy` fit this description.
- I would also look for functions I can simply avoid calling. For example, maybe I can just say `BSTree` only works with `bstring` keys, and then I can just not use the callback system and remove `default_compare` entirely.

At this point though, I only know that I want to look at `BSTree_get` to improve it, and not the reason `BSTree_get` is slow. That is phase two of the analysis.

# Callgrind Annotating Source

I will next tell `callgrind_annotate` to dump out the `bstree.c` file and annotate each line with the number of `Ir` it took. You get the annotated source file by running:

```
$ callgrind_annotate callgrind.out.1232 src/lcthw/bstree.c
...
```

Your output will have a big dump of the file's source, but I've cut out the parts for `BSTree_get` and `BSTree_getnode`:

```
-----
-- User-annotated source: src/lcthw/bstree.c
-----

    Ir

    2,453 static inline BSTreeNode
*BSTree_getnode(BSTree *map, BSTreeNode *node, void
*key)
    . {
61,853     int cmp = map->compare(node->key, key);
663,908 => src/lcthw/bstree.c:default_compare
(14850x)
    .
14,850     if(cmp == 0) {
    .         return node;
24,794     } else if(cmp < 0) {
30,623         if(node->left) {
    .             return BSTree_getnode(map, node-
>left, key);
    .         } else {
    .             return NULL;
    .         }
    .     } else {
13,146         if(node->right) {
    .             return BSTree_getnode(map, node-
>right, key);
    .         } else {
    .             return NULL;
    .         }
    .     }
    . }
    .
    . void *BSTree_get(BSTree *map, void *key)
    4,912 {
24,557     if(map->root == NULL) {
14,736         return NULL;
    .     } else {
    .         BSTreeNode *node = BSTree_getnode(map,
map->root, key);
    2,453         return node == NULL ? NULL : node-
>data;
    .     }
    . }
```

Each line is shown with either the number of `Ir` (instructions) it ran, or a period (.) to show that it's not important. What I'm looking for is hotspots,



or lines that have huge numbers of `Ir` that I can possibly bring down. In this case, line 10 of the output above shows that what makes `BSTree_getnode` so expensive is that it calls `default_comapre` which calls `bstrcmp`. I already know that `bstrcmp` is the worst running function, so if I want to improve the speed of `BSTree_getnode` I should work on that first.

I'll then look at `bstrcmp` the same way:

```
98,370  int bstrcmp (const_bstring b0, const_bstring
b1) {
    .   int i, v, n;
    .
196,740      if (b0 == NULL || b1 == NULL || b0->data
== NULL || b1->data == NULL ||
32,790          b0->slen < 0 || b1->slen < 0)
return SHRT_MIN;
65,580      n = b0->slen; if (n > b1->slen) n = b1-
>slen;
89,449      if (b0->slen == b1->slen && (b0->data ==
b1->data || b0->slen == 0))
    .          return BSTR_OK;
    .
23,915      for (i = 0; i < n; i++) {
163,642          v = ((char) b0->data[i]) - ((char)
b1->data[i]);
    .          if (v != 0) return v;
    .          if (b0->data[i] == (unsigned char)
'\0') return BSTR_OK;
    .      }
    .
    .      if (b0->slen > n) return 1;
    .      if (b1->slen > n) return -1;
    .      return BSTR_OK;
    .  }
```

The `Ir` for this function shows two lines that take up most of the execution. First, `bstrcmp` seems to go through a lot of trouble to make sure that it is not given a `NULL` value. That's a good thing so I want to leave that alone, but I'd consider writing a different compare function that was more "risky" and assumed it was never given a `NULL`. The next one is the loop that does the actual comparison. It seems that there's some optimization that could be done in comparing the characters of the two data buffers.

## Analyzing Memory Access With Cachegrind

What I want to do next is see how many times this `bstrcmp` function access memory to either read it or write it. The tool for doing that (and other things) is `cachegrind` and you use it like this:

```
$ valgrind --tool=cachegrind tests/bstree_tests
...
$ cg_annotate --show=Dr,Dw cachegrind.out.1316 | grep -v "???"
-----
--
I1 cache:          32768 B, 64 B, 8-way associative
D1 cache:          32768 B, 64 B, 8-way associative
```

```

LL cache:          4194304 B, 64 B, 16-way associative
Command:           tests/bstree_tests
Data file:         cachegrind.out.1316
Events recorded:   Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLMw
Events shown:      Dr Dw
Event sort order:  Ir I1mr ILmr Dr D1mr DLmr Dw D1mw DLMw
Thresholds:        0.1 100 100 100 100 100 100 100 100
Include dirs:
User annotated:
Auto-annotation:  off

```

```

-----
--
      Dr      Dw
-----
--
997,124 349,058  PROGRAM TOTALS
-----

```

```

--
      Dr      Dw  file:function
-----
--
169,754 19,430  src/lcthw/bstrlib.c:bstrcmp
 67,548 27,428  src/lcthw/bstree.c:BSTree_get
19,430 19,430  src/lcthw/bstree.c:default_compare
 5,420  2,383  src/lcthw/bstree.c:BSTree_delete
 2,000  4,200  src/lcthw/bstrlib.c:bformat
 1,600  2,800  src/lcthw/bstrlib.c:bfromcstralloc
 2,770  1,410  src/lcthw/bstree.c:BSTree_set
 1,200  1,200  src/lcthw/bstrlib.c:bdestroy

```

\$

I tell `valgrind` to use the `cachegrind` tool, which runs `bstree_tests` and then produces a `cachegrind.out.PID` file just like `callgrind` did. I then use the program `cg_annotate` to get a similar output, but notice that I'm telling it to do `--show=Dr,Dw`. This option says that I only want the memory read `Dr` and write `Dw` counts for each function.

After that you get your usual header and then the counts for `Dr` and `Dw` for each `file:function` combination. I've edited this down so it shows the files and also removed any OS junk with `| grep -v "???"` so your output may be a little different. What you see in my output is that `bstrcmp` is the worst function for memory usage too, which is to be expected since that's mostly the only thing it does. I'm going to now dump it's annotated source to see where.

```

-----
-- User-annotated source: src/lcthw/bstrlib.c
-----
--
      Dr      Dw
-----
--
      0 19,430  int bstrcmp (const_bstring b0,
const_bstring b1) {
      .      .  int i, v, n;
      .      .
77,720      0      if (b0 == NULL || b1 == NULL ||
b0->data == NULL || b1->data == NULL ||
38,860      0      b0->slen < 0 || b1->slen <
0) return SHRT_MIN;
      0      0      n = b0->slen; if (n > b1->slen) n

```

```

= b1->slen;
    0      0      if (b0->slen == b1->slen && (b0->
>data == b1->data || b0->slen == 0))
    .      .      return BSTR_OK;
    .      .
    0      0      for (i = 0; i < n; i++) {
53,174    0      v = ((char) b0->data[i]) -
((char) b1->data[i]);
    .      .      if (v != 0) return v;
    .      .      if (b0->data[i] ==
(unsigned char) '\0') return BSTR_OK;
    .      .      }
    .      .
    .      .      if (b0->slen > n) return 1;
    .      .      if (b1->slen > n) return -1;
    .      .      return BSTR_OK;
    .      .      }

```

The surprising thing about this output is that the worst line of `bstrcmp` isn't the character comparison like I thought. For memory access it's that protective if-statement at the top that checks every possible bad variable it could receive. That one if statement does more than twice as many memory accesses compared to the line that's comparing the characters on line 17 of this output. If I were to make `bstrcmp` then I would definitely just ditch that or do it once somewhere else.

Another option is to turn this check into an `assert` that only exists when running in development, and then compile it out in production. I now have enough evidence to say that this line is bad for this data structure, so I can justify removing it.

What I don't want to do however is justify making this function less defensive to just gain a few more cycles. In a real performance improvement situation I would simply put this on a list and then dig for other gains I can make in the program.

## Judo Tuning

"We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil."

—Donald Knuth

In my opinion, this quote seems to miss a major point about performance tuning. In this Knuth is saying that when you performance tune matters, in that if you do it in the beginning, then you'll cause all sorts of problems. According to him optimization should happen "sometime later", or at least that's my guess. Who knows these days really.

I'm going to declare this quote not necessarily wrong, but missing the point, and instead I'm going to officially give my quote. You can quote me on this:

"Use evidence to find the biggest optimizations that take the least effort."

—Zed A. Shaw

It doesn't matter when you try to optimize something, but instead it's

how you figure out if your optimization actually improved the software, and how much effort you put into doing them. With evidence you can find the places in the code where just a little effort gets you big improvements. Usually these places are just dumb decisions, as in `bstrcmp` trying to check everything possible for a `NULL` value.

At a certain point you have tuned the code to where the only thing that remains is tiny little micro-optimizations such as reorganizing if-statements and special loops like Duff's Device. At this point, just stop because there's a good chance that you'd gain more by redesigning the software to just *not do things*.

This is something that programmers who are optimizing simply fail to see. Many times the best way to do something fast is to find out ways to not do them. In the above analysis, I wouldn't try to make `bstrcmp` faster, I'd try to find a way to not use `bstrcmp` so much. Maybe there's a hashing scheme I can use that let's me do a sortable hash instead of constantly doing `bstrcmp`. Maybe I can optimize it by trying the first char first, and if it's comparable just don't call `bstrcmp`.

If after all that you can't do a redesign then start looking for little micro-optimizations, but as you do them *constantly confirm they improve speed*. Remember that the goal is to cause the biggest impact with the least effort possible.

## Using KCachegrind

The final section of this exercise is going to point you at a tool called `KCachegrind` is a *fantastic* GUI for analyzing `callgrind` and `cachegrind` output. I use it almost exclusively when I'm working on a Linux or BSD computer, and I've actually switched to just coding on Linux for projects because of `KCachegrind`.

Teaching you how to use it is outside the scope of this exercise, but you should be able to understand how to use it after this exercise. The output is nearly the same except `KCachegrind` lets you do the following:

- Graphically browse the source and execution times doing various sorts to find things to improve.
- Analyze different graphs to visually see what's taking up the most time and also what it is calling.
- Look at the actual machine code assembler output so you can see possible instructions that are happening, giving you more clues.
- Visualize the jump patterns for loops and branches in the source code, helping you find wayso to optimize the code easier.

You should spend some time getting `KCachegrind` installed and play with it.

## Extra Credit

- Read the `callgrind manual` and try some advanced options.
- Read the `cachegrind manual` and also try some advanced options.

- Use `callgrind` and `cachegrind` on all the unit tests and see if you can find optimizations to make. Did you find some things that surprised you? If not you probably aren't looking hard enough.
- Use KCachegrind and see how it compares to doing the terminal output like I'm doing here.
- Now use these tools to do the Exercise 40 extra credits and improvements.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 42: Stacks and Queues

At this point in the book you know most of the data structures that are used to build all the other data structures. If you have some kind of `List`, `DArray`, `Hashmap`, and `Tree` then you can build most anything else that's out there. Everything else you run into either uses these or is some variant on these. If it's not then it's most likely an exotic data structure that you probably will not need.

`Stacks` and `Queues` are very simple data structures that are really variants of the `List` data structure. All they are is using a `List` with a "discipline" or convention that says you'll always place elements on one end of the `List`. For a `Stack`, you always push and pop. For a `Queue`, you always shift to the front, but pop from the end.

I can implement both data structures using nothing but the CPP and two header files. My header files are 21 lines long and do all the `Stack` and `Queue` operations without any fancy defines.

To see if you've been paying attention, I'm going to show you the unit tests, and then *you* have to implement the header files needed to make them work. To pass this exercise you can't create any `stack.c` or `queue.c` implementation files. Use only the `stack.h` and `queue.h` files to make the tests runs.

```
#include "minunit.h"
#include <lcsthw/stack.h>
#include <assert.h>

static Stack *stack = NULL;
char *tests[] = {"test1 data", "test2 data", "test3 data"};
#define NUM_TESTS 3

char *test_create()
{
    stack = Stack_create();
    mu_assert(stack != NULL, "Failed to create stack.");

    return NULL;
}

char *test_destroy()
{
    mu_assert(stack != NULL, "Failed to make stack #2");
    Stack_destroy(stack);

    return NULL;
}

char *test_push_pop()
{
    int i = 0;
    for(i = 0; i < NUM_TESTS; i++) {
        Stack_push(stack, tests[i]);
        mu_assert(Stack_peek(stack) == tests[i], "Wrong next value.");
    }
}
```

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

🛠  
HELP

Follow  
@lzssthw

```

mu_assert(Stack_count(stack) == NUM_TESTS, "Wrong count on push.");

STACK_FOREACH(stack, cur) {
    debug("VAL: %s", (char *)cur->value);
}

for(i = NUM_TESTS - 1; i >= 0; i--) {
    char *val = Stack_pop(stack);
    mu_assert(val == tests[i], "Wrong value on pop.");
}

mu_assert(Stack_count(stack) == 0, "Wrong count after pop.");

return NULL;
}

char *all_tests() {
    mu_suite_start();

    mu_run_test(test_create);
    mu_run_test(test_push_pop);
    mu_run_test(test_destroy);

    return NULL;
}

RUN_TESTS(all_tests);

```

Then the `queue_tests.c` is almost the same just using `Queue`:

```

#include "minunit.h"
#include <lcthw/queue.h>
#include <assert.h>

static Queue *queue = NULL;
char *tests[] = {"test1 data", "test2 data", "test3 data"};
#define NUM_TESTS 3

char *test_create()
{
    queue = Queue_create();
    mu_assert(queue != NULL, "Failed to create queue.");

    return NULL;
}

char *test_destroy()
{
    mu_assert(queue != NULL, "Failed to make queue #2");
    Queue_destroy(queue);

    return NULL;
}

char *test_send_recv()
{
    int i = 0;
    for(i = 0; i < NUM_TESTS; i++) {
        Queue_send(queue, tests[i]);
        mu_assert(Queue_peek(queue) == tests[0], "Wrong next value.");
    }

    mu_assert(Queue_count(queue) == NUM_TESTS, "Wrong count on send.");

    QUEUE_FOREACH(queue, cur) {
        debug("VAL: %s", (char *)cur->value);
    }
}

```

```

for(i = 0; i < NUM_TESTS; i++) {
    char *val = Queue_recv(queue);
    mu_assert(val == tests[i], "Wrong value on recv.");
}

mu_assert(Queue_count(queue) == 0, "Wrong count after recv.");

return NULL;
}

char *all_tests() {
    mu_suite_start();

    mu_run_test(test_create);
    mu_run_test(test_send_recv);
    mu_run_test(test_destroy);

    return NULL;
}

RUN_TESTS(all_tests);

```

## What You Should See

Your unit test should run without you changing the tests, and it should pass `valgrind` with no memory errors. Here's what it looks like if I run `stack_tests` directly:

```

$ ./tests/stack_tests
DEBUG tests/stack_tests.c:60: ----- RUNNING: ./tests/stack_tests
-----
RUNNING: ./tests/stack_tests
DEBUG tests/stack_tests.c:53:
----- test_create
DEBUG tests/stack_tests.c:54:
----- test_push_pop
DEBUG tests/stack_tests.c:37: VAL: test3 data
DEBUG tests/stack_tests.c:37: VAL: test2 data
DEBUG tests/stack_tests.c:37: VAL: test1 data
DEBUG tests/stack_tests.c:55:
----- test_destroy
ALL TESTS PASSED
Tests run: 3
$

```

The `queue_test` is basically the same kind of output so I shouldn't have to show it to you at this stage.

## How To Improve It

The only real improvement you could make to this is to switch from using a `List` to using a `DArray`. The `Queue` data structure is more difficult to do with a `DArray` because it works at both ends of the list of nodes.

A disadvantage of doing this entirely in a header file is that you can't easily performance tune it. Mostly what you're doing with this technique is establishing a "protocol" for how to use a `List` in a certain style. When performance tuning, if you make `List` fast then these two should improve as well.



## Extra Credit

- Implement `Stack` using `DArray` instead of `List` without changing the unit test. That means you'll have to create your own `STACK_FOREACH`.

### Interested In Python?

Python is also a great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 43: A Simple Statistics Engine

This is a simple algorithm I use for collecting summary statistics "online", or without storing all of the samples. I use this in any software that needs to keep some statistics such as mean, standard deviation, and sum, but where I can't store all the samples needed. Instead I can just store the rolling results of the calculations which is only 5 numbers.

## Rolling Standard Deviation And Mean

The first thing you need is a sequence of samples. This can be anything from time to complete a task, numbers of times someone accesses something, or even star ratings on a website. Doesn't really matter what, just so long as you have a stream of numbers and you want to know the following summary statistics about them:

sum

This is the total of all the numbers added together.

sum squared (sumsq)

This is the sum of the square of each number.

count (n)

This is the number samples you've taken.

min

This is the smallest sample you've seen.

max

This is the largest sample you've seen.

mean

This is the most likely middle number. It's not actually the middle, since that's the median, but it's an accepted approximation for it.

stddev

Calculated using  $\sqrt{\text{sumsq} - (\text{sum} * \text{mean}) / (n - 1)}$   
where `sqrt` is the square root function in the `math.h` header.

I will confirm this calculation works using R since I know R gets these right:

```
> s <- runif(n=10, max=10)
> s
[1] 6.1061334 9.6783204 1.2747090 8.2395131 0.3333483 6.9755066 1.0626275
[8] 7.6587523 4.9382973 9.5788115
> summary(s)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.3333  2.1910  6.5410  5.5850  8.0940  9.6780
> sd(s)
[1] 3.547868
> sum(s)
```

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

⛑  
HELP

Follow  
@lzssth

```

[1] 55.84602
> sum(s * s)
[1] 425.1641
> sum(s) * mean(s)
[1] 311.8778
> sum(s * s) - sum(s) * mean(s)
[1] 113.2863
> (sum(s * s) - sum(s) * mean(s)) / (length(s) - 1)
[1] 12.58737
> sqrt((sum(s * s) - sum(s) * mean(s)) / (length(s) - 1))
[1] 3.547868
>

```

You don't need to know R, just follow along while I explain how I'm breaking this down to check my math:

lines 1-4

I use the function `runif` to get a "random uniform" distribution of numbers, then print them out. I'll use these in the unit test later.

lines 5-7

Here's the summary, so you can see the values that R calculates for these.

lines 8-9

This is the `stddev` using the `sd` function.

lines 10-11

Now I begin to build this calculation manually, first by getting the

`sum`.

lines 12-13

Next piece of the `stdev` formula is the `sumsq`, which I can get with `sum(s * s)` which tells R to multiple the whole `s` list by itself and then `sum` those. The power of R is being able to do math on entire data structures like this.

lines 14-15

Looking at the formula, I then need the `sum` multiplied by `mean`, so I do `sum(s) * mean(s)`.

lines 16-17

I then combine the `sumsq` with this to get `sum(s * s) - sum(s) * mean(s)`.

lines 18-19

That needs to be divided by  $n-1$ , so I do `(sum(s * s) - sum(s) * mean(s)) / (length(s) - 1)`.

lines 20-21

Finally, I `sqrt` that and I get 3.547868 which matches the number R gave me for `sd` above.

## Implementation

That's how you calculate the `stddev`, so now I can make some simple code to implement this calculation.

```

#ifndef lcthw_stats_h
#define lcthw_stats_h

typedef struct Stats {

```

```

    double sum;
    double sumsq;
    unsigned long n;
    double min;
    double max;
} Stats;

Stats *Stats_recreate(double sum, double sumsq, unsigned long n, double min, double max);

Stats *Stats_create();

double Stats_mean(Stats *st);

double Stats_stddev(Stats *st);

void Stats_sample(Stats *st, double s);

void Stats_dump(Stats *st);

#endif

```

Here you can see I've put the calculations I need to store in a `struct` and then I have functions for sampling and getting the numbers. Implementing this is then just an exercise in converting the math:

```

#include <math.h>
#include <lcsth/stats.h>
#include <stdlib.h>
#include <lcsth/dbg.h>

Stats *Stats_recreate(double sum, double sumsq, unsigned long n, double min, double max)
{
    Stats *st = malloc(sizeof(Stats));
    check_mem(st);

    st->sum = sum;
    st->sumsq = sumsq;
    st->n = n;
    st->min = min;
    st->max = max;

    return st;
error:
    return NULL;
}

Stats *Stats_create()
{
    return Stats_recreate(0.0, 0.0, 0L, 0.0, 0.0);
}

double Stats_mean(Stats *st)
{
    return st->sum / st->n;
}

double Stats_stddev(Stats *st)
{
    return sqrt( (st->sumsq - (st->sum * st->sum / st->n)) / (st->n - 1) );
}

void Stats_sample(Stats *st, double s)
{
    st->sum += s;
    st->sumsq += s * s;

    if(st->n == 0) {

```

```

        st->min = s;
        st->max = s;
    } else {
        if(st->min > s) st->min = s;
        if(st->max < s) st->max = s;
    }

    st->n += 1;
}

void Stats_dump(Stats *st)
{
    fprintf(stderr, "sum: %f, sumsq: %f, n: %ld, min: %f, max: %f, mean: %f, s
    tdddev: %f",
            st->sum, st->sumsq, st->n, st->min, st->max,
            Stats_mean(st), Stats_stddev(st));
}

```

Here's what each function in `stats.c` does:

**Stats\_recreate**

I'll want to load these numbers from some kind of database, and this function let's me recreate a `Stats` struct.

**Stats\_create**

Simply called `Stats_recreate` with all 0 values.

**Stats\_mean**

Using the `sum` and `n` it gives the mean.

**Stats\_stddev**

Implements the formula I worked out, with the only difference being that I calculate the mean with `st->sum / st->n` in this formula instead of calling `Stats_mean`.

**Stats\_sample**

This does the work of maintaining the numbers in the `Stats` struct. When you give it the first value it sees that `n` is 0 and sets `min` and `max` accordingly. Every call after that keeps increasing `sum`, `sumsq`, and `n`. It then figures out if this new sample is a new `min` or `max`.

**Stats\_dump**

Simple debug function that dumps the stats so you can view them.

The last thing I need to do is confirm that this math is correct. I'm going to use my numbers and calculations from my R session to create a unit test that confirms I'm getting the right results.

```

#include "minunit.h"
#include <lcthw/stats.h>
#include <math.h>

const int NUM_SAMPLES = 10;
double samples[] = {
    6.1061334, 9.6783204, 1.2747090, 8.2395131, 0.3333483,
    6.9755066, 1.0626275, 7.6587523, 4.9382973, 9.5788115
};

Stats expect = {
    .sumsq = 425.1641,
    .sum = 55.84602,
    .min = 0.333,
    .max = 9.678,
    .n = 10,
};

double expect_mean = 5.584602;
double expect_stddev = 3.547868;

```

```

#define EQ(X,Y,N) (round((X) * pow(10, N)) == round((Y) * pow(10, N)))

char *test_operations()
{
    int i = 0;
    Stats *st = Stats_create();
    mu_assert(st != NULL, "Failed to create stats.");

    for(i = 0; i < NUM_SAMPLES; i++) {
        Stats_sample(st, samples[i]);
    }

    Stats_dump(st);

    mu_assert(EQ(st->sumsq, expect.sumsq, 3), "sumsq not valid");
    mu_assert(EQ(st->sum, expect.sum, 3), "sum not valid");
    mu_assert(EQ(st->min, expect.min, 3), "min not valid");
    mu_assert(EQ(st->max, expect.max, 3), "max not valid");
    mu_assert(EQ(st->n, expect.n, 3), "max not valid");
    mu_assert(EQ(expect_mean, Stats_mean(st), 3), "mean not valid");
    mu_assert(EQ(expect_stddev, Stats_stddev(st), 3), "stddev not valid");

    return NULL;
}

char *test_recreate()
{
    Stats *st = Stats_recreate(expect.sum, expect.sumsq, expect.n, expect.min,
    expect.max);

    mu_assert(st->sum == expect.sum, "sum not equal");
    mu_assert(st->sumsq == expect.sumsq, "sumsq not equal");
    mu_assert(st->n == expect.n, "n not equal");
    mu_assert(st->min == expect.min, "min not equal");
    mu_assert(st->max == expect.max, "max not equal");
    mu_assert(EQ(expect_mean, Stats_mean(st), 3), "mean not valid");
    mu_assert(EQ(expect_stddev, Stats_stddev(st), 3), "stddev not valid");

    return NULL;
}

char *all_tests()
{
    mu_suite_start();

    mu_run_test(test_operations);
    mu_run_test(test_recreate);

    return NULL;
}

RUN_TESTS(all_tests);

```

There's nothing new in this unit test, except maybe the `EQ` macro. I felt lazy and didn't want to look up the standard way to tell if two `double` values are close, so I made this macro. The problem with `double` is that equality assumes totally equal, but I'm using two different systems with slightly different rounding errors. The solution is to say I want the numbers to be "equal to X decimal places".

I do this with `EQ` by raising the number to a power of 10, then using the `round` function to get an integer. This is a simple way to round to N decimal places and compare the results as an integer. I'm sure there's a billion other ways to do the same thing, but this works for now.

The expected results are then in a `Stats` `struct` and then I simply make

sure that the number I get is close to the number R gave me.

## How To Use It

You can use the standard deviation and mean to determine if a new sample is "interesting", or you can use this to collect statistics on statistics. The first one is easy for people to understand so I'll explain that quickly using an example for login times.

Imagine you're tracking how long users spend on a server and you're using stats to analyze it. Every time someone logs in, you keep track of how long they are there, then you call `Stats_sample`. I'm looking for people are a on "too long" and also people who seem to be on "too quickly".

Instead of setting specific levels, what I'd do is compare how long someone is on with the `mean (plus or minus) 2 * stddev` range. I get the `mean` and `2 * stddev`, and consider login times to be "interesting" if they are outside these two ranges. Since I'm keeping these statistics using a rolling algorithm this is a very fast calculation and I can then have the software flag the users who are outside of this range.

This doesn't necessarily point out people who are behaving badly, but instead it flags potential problems that you can review to see what's going on. It's also doing it based on the behavior of all the users, which avoids the problem where you pick some arbitrary number that's not based on what's really happening.

The general rule you can get from this is that the `mean (plus or minus) 2 * stddev` is an estimate of where 90% of the values are expected to fall, and that anything outside those ranges is interesting.

The second way to use these statistics is to go meta and calculate them for other `Stats` calculations. You basically do your `Stats_sample` like normal, but then you run `Stats_sample` on the `min`, `max`, `n`, `mean`, and `stddev` on that sample. This gives a two-level measurement, and let's you compare samples of samples.

Confusing right? I'll continue my example above and add that you have 100 servers that each hold a different application. You are already tracking user's login times for each application server, but you want to compare all 100 applications and flag any users that are logging in "too much" on all of them. Easiest way to do that is each time someone logs in, calculate the new login stats, then add *that* `Stats structs` elements to a second `Stat`.

What you end up with is a series of stats that can be named like this:

mean of means

This is a full `Stats struct` that gives you `mean` and `stddev` of the means of all the servers. Any server or user who is outside of this is work looking at on a global level.

mean of stddevs

Another `Stats struct` that produces the statistics of how *all* of the servers range. You can then analyze each server and see if any of them have unusually wide ranging numbers by

comparing their `stddev` to this `mean of stddevs` statistic.

You could do them all, but these are the most useful. If you wanted to then monitor servers for erratic login times you'd do this:

- User John logs into and out of server A. Grab server A's stats, update them.
- Grab the `mean of means` stats, and take A's mean and add it as a sample. I'll call this `m_of_m`.
- Grab the `mean of stddevs` stats, and add A's stddev to it as a sample. I'll call this `m_of_s`.
- If A's `mean` is outside of `m_of_m.mean + 2 * m_of_m.stddev` then flag it as possibly having a problem.
- If A's `stddev` is outside of `m_of_s.mean + 2 * m_of_s.stddev` then flag it as possible behaving too erratically.
- Finally, if John's login time is outside of A's range, or A's `m_of_m` range, then flag it as interesting.

Using this "mean of means" and "mean of stddevs" calculation you can do efficient tracking of many metrics with a minimal amount of processing and storage.

## Extra Credit

- Convert the `Stats_stddev` and `Stats_mean` to `static inline` functions in the `stats.h` file instead of in the `stats.c` file.
- Use this code to write a performance test of the `string_algos_test.c`. Make it optional and have it run the base test as a series of samples then report the results.
- Write a version of this in another programming language you know. Confirm that this version is correct based on what I have here.
- Write a little program that can take a file full of numbers and spit these statistics out for them.
- Make the program accept a table of data that has headers on one line, then all the other numbers on lines after it separated by any number of spaces. Your program should then print out these stats for each column by the header name.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**





## Exercise 44: Ring Buffer

Ring buffers are incredibly useful when processing asynchronous IO. They allow one side to receive data in random intervals of random sizes, but feed cohesive chunks to another side in set sizes or intervals. They are a variant on the `Queue` data structure but it focuses on blocks of bytes instead of a list of pointers. In this exercise I'm going to show you the `RingBuffer` code, and then you have to make a full unit test for it.

```
#ifndef _lcthw_RingBuffer_h
#define _lcthw_RingBuffer_h

#include <lcthw/bstrlib.h>

typedef struct {
    char *buffer;
    int length;
    int start;
    int end;
} RingBuffer;

RingBuffer *RingBuffer_create(int length);

void RingBuffer_destroy(RingBuffer *buffer);

int RingBuffer_read(RingBuffer *buffer, char *target, int amount);

int RingBuffer_write(RingBuffer *buffer, char *data, int length);

int RingBuffer_empty(RingBuffer *buffer);

int RingBuffer_full(RingBuffer *buffer);

int RingBuffer_available_data(RingBuffer *buffer);

int RingBuffer_available_space(RingBuffer *buffer);

bstring RingBuffer_gets(RingBuffer *buffer, int amount);

#define RingBuffer_available_data(B) (((B)->end + 1) % (B)->length - (B)->start - 1)

#define RingBuffer_available_space(B) ((B)->length - (B)->end - 1)

#define RingBuffer_full(B) (RingBuffer_available_data((B)) - (B)->length == 0)

#define RingBuffer_empty(B) (RingBuffer_available_data((B)) == 0)

#define RingBuffer_puts(B, D) RingBuffer_write((B), bdata((D)), blength((D)))

#define RingBuffer_get_all(B) RingBuffer_gets((B), RingBuffer_available_data((B)))

#define RingBuffer_starts_at(B) ((B)->buffer + (B)->start)

#define RingBuffer_ends_at(B) ((B)->buffer + (B)->end)

#define RingBuffer_commit_read(B, A) ((B)->start = ((B)->start + (A)) % (B)->length)
```

1  
2  
3  
4  
MAIN

◀  
PREVIOUS

▶  
NEXT

🛠  
HELP

Follow  
@lzsthw

```
#define RingBuffer_commit_write(B, A) ((B)->end = ((B)->end + (A)) % (B)->length)
```

```
#endif
```

Looking at the data structure you see I have a `buffer`, `start` and `end`. A `RingBuffer` does nothing more than move the `start` and `end` around the buffer so that it "loops" whenever it reaches the buffer's end. Doing this gives the illusion of an infinite read device in a small space. I then have a bunch of macros that do various calculations based on this.

Here's the implementation which is a much better explanation of how this works:

```
#undef NDEBUG
#include <assert.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <lcthw/dbg.h>
#include <lcthw/ringbuffer.h>
```

```
RingBuffer *RingBuffer_create(int length)
{
    RingBuffer *buffer = calloc(1, sizeof(RingBuffer));
    buffer->length = length + 1;
    buffer->start = 0;
    buffer->end = 0;
    buffer->buffer = calloc(buffer->length, 1);

    return buffer;
}

void RingBuffer_destroy(RingBuffer *buffer)
{
    if(buffer) {
        free(buffer->buffer);
        free(buffer);
    }
}

int RingBuffer_write(RingBuffer *buffer, char *data, int length)
{
    if(RingBuffer_available_data(buffer) == 0) {
        buffer->start = buffer->end = 0;
    }

    check(length <= RingBuffer_available_space(buffer),
           "Not enough space: %d request, %d available",
           RingBuffer_available_data(buffer), length);

    void *result = memcpy(RingBuffer_ends_at(buffer), data, length);
    check(result != NULL, "Failed to write data into buffer.");

    RingBuffer_commit_write(buffer, length);

    return length;
error:
    return -1;
}

int RingBuffer_read(RingBuffer *buffer, char *target, int amount)
{
    check_debug(amount <= RingBuffer_available_data(buffer),
                "Not enough in the buffer: has %d, needs %d",
                RingBuffer_available_data(buffer), amount);

    void *result = memcpy(target, RingBuffer_starts_at(buffer), amount);
    check(result != NULL, "Failed to write buffer into data.");
}
```

```

RingBuffer_commit_read(buffer, amount);

if(buffer->end == buffer->start) {
    buffer->start = buffer->end = 0;
}

return amount;
error:
return -1;
}

bstring RingBuffer_gets(RingBuffer *buffer, int amount)
{
    check(amount > 0, "Need more than 0 for gets, you gave: %d ", amount);
    check_debug(amount <= RingBuffer_available_data(buffer),
        "Not enough in the buffer.");

    bstring result = blk2bstr(RingBuffer_starts_at(buffer), amount);
    check(result != NULL, "Failed to create gets result.");
    check(blength(result) == amount, "Wrong result length.");

    RingBuffer_commit_read(buffer, amount);
    assert(RingBuffer_available_data(buffer) >= 0 && "Error in read commit.");

    return result;
error:
return NULL;
}

```

This is all there is to a basic `RingBuffer` implementation. You can read and write blocks of data to it. You can ask how much is in it and how much space it has. There are some fancier ring buffers that use tricks in the OS to create an imaginary infinite store, but those aren't portable.

Since my `RingBuffer` deals with reading and writing blocks of memory, I'm making sure that any time `end == start` then I reset them to 0 (zero) so that they go to the beginning of the buffer. In the Wikipedia version it wasn't writing blocks of data, so it only had to move `end` and `start` around in a circle. To better handle blocks you have to drop to the beginning of the internal buffer whenever the data is empty.

## The Unit Test

For your unit test, you'll want to test as many possible conditions as you can. Easiest way to do that is to preconstruct different `RingBuffer` structs and then manually check that the functions and math work right. For example, you could make one where `end` is right at the end of the buffer and `start` is right before it, then see how it fails.

## What You Should See

Here's my `ringbuffer_tests` run:

```

$ ./tests/ringbuffer_tests
DEBUG tests/ringbuffer_tests.c:60: ----- RUNNING: ./tests/ringbuffer_tests
----
RUNNING: ./tests/ringbuffer_tests
DEBUG tests/ringbuffer_tests.c:53:
----- test_create
DEBUG tests/ringbuffer_tests.c:54:
----- test_read_write

```

```
DEBUG tests/ringbuffer_tests.c:55:
----- test_destroy
ALL TESTS PASSED
Tests run: 3
$
```

You should have at least three tests that confirm all the basic operations, and then see how much more you can test beyond what I've done.

## How To Improve It

As usual you should go back and add the defensive programming checks to this exercise. Hopefully you've been doing this because the base code in most of `liblcthw` doesn't check for common defensive programming that I'm teaching you. I leave this to you so that you get used to improving code with these extra checks.

For example, in this ring buffer there's not a lot of checking that an access will actually be inside the buffer.

If you read the [Ring Buffer Wikipedia page](#) you'll see the "Optimized POSIX implementation" that uses POSIX specific calls to create an infinite space. Study that as I'll have you try it in the extra credit.

## Extra Credit

- Create an alternative implementation of `RingBuffer` that uses the POSIX trick and a unit test for it.
- Add a performance comparison test to this unit test that compares the two versions by fuzzing them with random data and random read/write operations. Make sure that you setup this fuzzing so that the same operations are done to each so you can compare them between runs.
- Use `callgrind` and `cachegrind` to compare the performance of these two.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 45: A Simple TCP/IP Client

I'm going to use the `RingBuffer` to create a very simplistic little network testing tool called `netclient`. To do this I have to add some stuff to the `Makefile` to handle little programs in the `bin/` directory.

### Augment The Makefile

First, add a variable for the programs just like the unit tests `TESTS` and `TEST_SRC` variables:

```
PROGRAMS_SRC=$(wildcard bin/*.c)
PROGRAMS=$(patsubst %.c,%, $(PROGRAMS_SRC))
```

Then you want to add the `PROGRAMS` to the all target:

```
all: $(TARGET) $(SO_TARGET) tests $(PROGRAMS)
```

Then add `PROGRAMS` to the `rm` line in the `clean` target:

```
rm -rf build $(OBJECTS) $(TESTS) $(PROGRAMS)
```

Finally you just need a target at the end to build them all:

```
$(PROGRAMS): CFLAGS += $(TARGET)
```

With these changes you can drop simple `.c` files into `bin` and `make` will build them and link them to the library just like the tests are done.

### The netclient Code

The code for the little netclient looks like this:

```
#undef NDEBUG
#include <stdlib.h>
#include <sys/select.h>
#include <stdio.h>
#include <lcthw/ringbuffer.h>
#include <lcthw/dbg.h>
#include <sys/socket.h>
#include <sys/types.h>
#include <sys/uio.h>
#include <arpa/inet.h>
#include <netdb.h>
#include <unistd.h>
#include <fcntl.h>

struct tagbstring NL = bsStatic("\n");
struct tagbstring CRLF = bsStatic("\r\n");

int nonblock(int fd) {
    int flags = fcntl(fd, F_GETFL, 0);
    check(flags >= 0, "Invalid flags on nonblock.");
```

```

    int rc = fcntl(fd, F_SETFL, flags | O_NONBLOCK);
    check(rc == 0, "Can't set nonblocking.");

    return 0;
error:
    return -1;
}

int client_connect(char *host, char *port)
{
    int rc = 0;
    struct addrinfo *addr = NULL;

    rc = getaddrinfo(host, port, NULL, &addr);
    check(rc == 0, "Failed to lookup %s:%s", host, port);

    int sock = socket(AF_INET, SOCK_STREAM, 0);
    check(sock >= 0, "Cannot create a socket.");

    rc = connect(sock, addr->ai_addr, addr->ai_addrlen);
    check(rc == 0, "Connect failed.");

    rc = nonblock(sock);
    check(rc == 0, "Can't set nonblocking.");

    freeaddrinfo(addr);
    return sock;
error:
    freeaddrinfo(addr);
    return -1;
}

int read_some(RingBuffer *buffer, int fd, int is_socket)
{
    int rc = 0;

    if(RingBuffer_available_data(buffer) == 0) {
        buffer->start = buffer->end = 0;
    }

    if(is_socket) {
        rc = recv(fd, RingBuffer_starts_at(buffer), RingBuffer_available_space(buffer), 0);
    } else {
        rc = read(fd, RingBuffer_starts_at(buffer), RingBuffer_available_space(buffer));
    }

    check(rc >= 0, "Failed to read from fd: %d", fd);

    RingBuffer_commit_write(buffer, rc);

    return rc;
error:
    return -1;
}

int write_some(RingBuffer *buffer, int fd, int is_socket)
{
    int rc = 0;
    bstring data = RingBuffer_get_all(buffer);

    check(data != NULL, "Failed to get from the buffer.");
    check(bfindreplace(data, &NL, &CRLF, 0) == BSTR_OK, "Failed to replace NL.");
    );

    if(is_socket) {

```

```

        rc = send(fd, bdata(data), blength(data), 0);
    } else {
        rc = write(fd, bdata(data), blength(data));
    }

    check(rc == blength(data), "Failed to write everything to fd: %d.", fd);
    bdestroy(data);

    return rc;
error:
    return -1;
}

int main(int argc, char *argv[])
{
    fd_set allreads;
    fd_set readmask;

    int socket = 0;
    int rc = 0;
    RingBuffer *in_rb = RingBuffer_create(1024 * 10);
    RingBuffer *sock_rb = RingBuffer_create(1024 * 10);

    check(argc == 3, "USAGE: netclient host port");

    socket = client_connect(argv[1], argv[2]);
    check(socket >= 0, "connect to %s:%s failed.", argv[1], argv[2]);

    FD_ZERO(&allreads);
    FD_SET(socket, &allreads);
    FD_SET(0, &allreads);

    while(1) {
        readmask = allreads;
        rc = select(socket + 1, &readmask, NULL, NULL, NULL);
        check(rc >= 0, "select failed.");

        if(FD_ISSET(0, &readmask)) {
            rc = read_some(in_rb, 0, 0);
            check_debug(rc != -1, "Failed to read from stdin.");
        }

        if(FD_ISSET(socket, &readmask)) {
            rc = read_some(sock_rb, socket, 0);
            check_debug(rc != -1, "Failed to read from socket.");
        }

        while(!RingBuffer_empty(sock_rb)) {
            rc = write_some(sock_rb, 1, 0);
            check_debug(rc != -1, "Failed to write to stdout.");
        }

        while(!RingBuffer_empty(in_rb)) {
            rc = write_some(in_rb, socket, 1);
            check_debug(rc != -1, "Failed to write to socket.");
        }
    }

    return 0;
error:
    return -1;
}

```

This code uses `select` to handle events from both `stdin` (file descriptor 0) and from the `socket` it uses to talk to a server. It uses `RingBuffers` to store the data and copy it around, and you can consider

the functions `read_some` and `write_some` early prototypes for similar functions in the `RingBuffer` library.

In this little bit of code are quite a few networking functions you may not know. As you hit a function you don't know, look it up in the man pages and make sure you understand it. This one little file could actually get you to research all the APIs required to write a little server in C.

## What You Should See

If you have everything building then the quickest way to test it is see if you can get a special file off [learncodethehardway.org](http://learncodethehardway.org):

```
$
$ ./bin/netclient learncodethehardway.org 80
GET /ex45.txt HTTP/1.1
Host: learncodethehardway.org

HTTP/1.1 200 OK
Date: Fri, 27 Apr 2012 00:41:25 GMT
Content-Type: text/plain
Content-Length: 41
Last-Modified: Fri, 27 Apr 2012 00:42:11 GMT
ETag: 4f99eb63-29
Server: Mongrel2/1.7.5

Learn C The Hard Way, Exercise 45 works.
^C
$
```

What I did there is I type in the syntax needed to make the HTTP request for the file `/ex45.txt`, then the `Host:` header line, then hit ENTER to get an empty line. I then get the response, with headers and the content. After that I just hit CTRL-c to exit.

## How To Break It

This code definitely could have bugs, and currently in the draft of the book I'm going to have to keep working on this. In the meantime, try analyzing the code I have here and thrashing it against other servers. There's a tool called `netcat` that is great for setting up these kinds of servers. Another thing to do is use a language like `Python` or `Ruby` to create a simple "junk server" that spews out junk and bad data, closes connections randomly, and other nasty things.

If you find bugs, report them in the comments and I'll fix them up.

## Extra Credit

- As I mentioned, there's quite a few functions you may not know, so look them up. In fact, look them all up even if you think you know them.
- Run this under `valgrind` and look for errors.
- Go back through and add various defensive programming checks to the functions to improve them.
- Use the `getopt` function to allow the user to give this the option to *not* translate `\n` to `\r\n`. This is only needed on protocols that require it for line endings, like HTTP. Sometimes



you don't want the translation, so give the user an option.

**Interested  
In  
Python?**

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

**Interested In Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 46: Ternary Search Tree

The final data structure I'll show you is called the *TSTree* and it's similar to the *BSTree* except it has three branches *low*, *equal*, and *high*. It's primarily used to just like *BSTree* and *HashMap* to store key/value data, but it is keyed off of the individual characters in the keys. This gives the *TSTree* some abilities that neither *BSTree* or *HashMap* have.

How a *TSTree* works is every key is a string, and it's inserted by walking and building a tree based on the equality of the characters in the string. Start at the root, look at the character for that node, and if lower, equal to, or higher than that then go in that direction. You can see this in the header file:

```
#ifndef _lcthw_TSTree_h
#define _lcthw_TSTree_h

#include <stdlib.h>
#include <lcthw/darray.h>

typedef struct TSTree {
    char splitchar;
    struct TSTree *low;
    struct TSTree *equal;
    struct TSTree *high;
    void *value;
} TSTree;

void *TSTree_search(TSTree *root, const char *key, size_t len);

void *TSTree_search_prefix(TSTree *root, const char *key, size_t len);

typedef void (*TSTree_traverse_cb)(void *value, void *data);

TSTree *TSTree_insert(TSTree *node, const char *key, size_t len, void *value);

void TSTree_traverse(TSTree *node, TSTree_traverse_cb cb, void *data);

void TSTree_destroy(TSTree *root);

#endif
```

The *TSTree* has the following elements:

*splitchar*

The character at this point in the tree.

*low*

The branch that is lower than *splitchar*.

*equal*

The branch that is equal to *splitchar*.

*high*

The branch that is higher than *splitchar*.

*value*

The value set for a string at that point with that *splitchar*.

You can see this implementation has the following operations:

1  
2  
3  
4  
MAIN

PREVIOUS

NEXT

HELP

Follow  
@l2sthw

search

Typical "find a value for this `key`" operation.

search\_prefix

Finds the first value that has this as a prefix of its key. This is the an operation that you can't easily do in a `BSTree` or `Hashmap`.

insert

Breaks the `key` down by each character and inserts it into the tree.

traverse

Walks the tree allowing you to collect or analyze all the keys and values it contains.

The only thing missing is a `TSTree_delete`, and that's because it is a horribly expensive operation, even more than `BSTree_delete` was. When I use `TSTree` structures I treat them as constant data that I plan on traversing many times and not removing anything from them. They are very fast for this, but are not good if you need to insert and delete quickly. For that I use `Hashmap` since it beats both `BSTree` and `TSTree`.

The implementation for the `TSTree` is actually simple, but it might be hard to follow at first. I'll break it down after you enter it in:

```
#include <stdlib.h>
#include <stdio.h>
#include <assert.h>
#include <lcthw/dbg.h>
#include <lcthw/tstree.h>

static inline TSTree *TSTree_insert_base(TSTree *root, TSTree *node,
    const char *key, size_t len, void *value)
{
    if(node == NULL) {
        node = (TSTree *) calloc(1, sizeof(TSTree));

        if(root == NULL) {
            root = node;
        }

        node->splitchar = *key;
    }

    if(*key < node->splitchar) {
        node->low = TSTree_insert_base(root, node->low, key, len, value);
    } else if(*key == node->splitchar) {
        if(len > 1) {
            node->equal = TSTree_insert_base(root, node->equal, key+1, len - 1
, value);
        } else {
            assert(node->value == NULL && "Duplicate insert into tst.");
            node->value = value;
        }
    } else {
        node->high = TSTree_insert_base(root, node->high, key, len, value);
    }

    return node;
}

TSTree *TSTree_insert(TSTree *node, const char *key, size_t len, void *value)
{
    return TSTree_insert_base(node, node, key, len, value);
}

void *TSTree_search(TSTree *root, const char *key, size_t len)
```

```

{
    TSTree *node = root;
    size_t i = 0;

    while(i < len && node) {
        if(key[i] < node->splitchar) {
            node = node->low;
        } else if(key[i] == node->splitchar) {
            i++;
            if(i < len) node = node->equal;
        } else {
            node = node->high;
        }
    }

    if(node) {
        return node->value;
    } else {
        return NULL;
    }
}

void *TSTree_search_prefix(TSTree *root, const char *key, size_t len)
{
    if(len == 0) return NULL;

    TSTree *node = root;
    TSTree *last = NULL;
    size_t i = 0;

    while(i < len && node) {
        if(key[i] < node->splitchar) {
            node = node->low;
        } else if(key[i] == node->splitchar) {
            i++;
            if(i < len) {
                if(node->value) last = node;
                node = node->equal;
            }
        } else {
            node = node->high;
        }
    }

    node = node ? node : last;

    // traverse until we find the first value in the equal chain
    // this is then the first node with this prefix
    while(node && !node->value) {
        node = node->equal;
    }

    return node ? node->value : NULL;
}

void TSTree_traverse(TSTree *node, TSTree_traverse_cb cb, void *data)
{
    if(!node) return;

    if(node->low) TSTree_traverse(node->low, cb, data);

    if(node->equal) {
        TSTree_traverse(node->equal, cb, data);
    }

    if(node->high) TSTree_traverse(node->high, cb, data);

    if(node->value) cb(node->value, data);
}

```

```

void TSTree_destroy(TSTree *node)
{
    if(node == NULL) return;

    if(node->low) TSTree_destroy(node->low);

    if(node->equal) {
        TSTree_destroy(node->equal);
    }

    if(node->high) TSTree_destroy(node->high);

    free(node);
}

```

For `TSTree_insert` I'm using the same pattern for recursive structures where I have a small function that calls the real recursive function. I'm not doing any additional check here but you should add the usual defensive programming to it. One thing to keep in mind is it is using a slightly different design where you don't have a separate `TSTree_create` function, and instead if you pass it a `NULL` for the `node` then it will create it, and returns the final value.

That means I need to break down `TSTree_insert_base` for you to understand the insert operation:

tstree.c:10-18

As I mentioned, if I'm given a `NULL`, then I need to make this node and assign the `*key` (current char) to it. This is used to build the tree as we insert keys.

tstree.c:20-21

If the `*key` < this then recurse, but go to the `low` branch.

tstree.c:22

This `splitchar` is equal, so I want to go to deal with equality. This will happen if we just created this node, so we'll be building the tree at this point.

tstree.c:23-24

There's still characters to handle, so recurse down the `equal` branch, but go to the next `*key` char.

tstree.c:26-27

This is the last char, so I set the value and that's it. I have an `assert` here in case of a duplicate.

tstree.c:29-30

The last condition is that this `*key` is greater than `splitchar` so I need to recurse down the `high` branch.

The key to some of the properties of this data structure is the fact that I'm only incrementing the character I analyze when a `splitchar` is equal. The other two conditions I just walk the tree until I hit an equal character to recurse into next. What this does is it makes it very fast to *not* find a key. I can get a bad key, and simply walk a few `high` and `low` nodes until I hit a dead end to know that this key doesn't exist. I don't need to process every character of the key, or every node of the tree.

Once you understand that then move onto analyzing how `TSTree_search` works:

tstree.c:46

I don't need to process the tree recursively in the `TSTree`, I can just use a while loop and a `node` for where I am currently.

tstree.c:47-48

If the current char is less than the node `splitchar`, then go low.

tstree.c:49-51

If it's equal, then increment `i` and go equal as long as it's not the last character. That's why the `if(i < len)` is there, so that I don't go too far past the final `value`.

tstree.c:52-53

Otherwise I go `high` since the char is greater.

tstree.c:57-61

If after the loop I have a node, then return its `value`, otherwise return `NULL`.

This isn't too difficult to understand, and you can then see that it's almost exactly the same algorithm for the `TSTree_search_prefix` function. The only difference is I'm trying to not find an exact match, but the longest prefix I can. To do that I keep track of the `last` node that was equal, and then after the search loop, walk that node until I find a `value`.

Looking at `TSTree_search_prefix` you can start to see the second advantage a `TSTree` has over the `BSTree` and `Hashmap` for finding strings. Given any key of X length, you can find any key in X moves. You can also find the first prefix in X moves, plus N more depending on how big the matching key is. If the biggest key in the tree is 10 characters long, then you can find any prefix in that key in 10 moves. More importantly, you can do all of this by only comparing each character of the key *once*.

In comparison, to do the same with a `BSTree` you would have to check the prefixes of each character in every possibly matching node in the `BSTree` against the characters in the prefix. It's the same for finding keys, or seeing if a key doesn't exist. You have to compare each character against most of the characters in the `BSTree` to find or not find a match.

A `Hashamp` is even worse for finding prefixes since you can't hash just the prefix. You basically can't do this efficiently in a `Hashmap` unless the data is something you can parse like a URL. Even then that usually requires whole trees of `Hashmaps`.

The last two functions should be easy for you to analyze as they are the typical traversing and destroying operations you've seen already for other data structures.

Finally, I have a simple unit test for the whole thing to make sure it works right:

```
#include "minunit.h"
#include <lcthw/tstree.h>
#include <string.h>
#include <assert.h>
#include <lcthw/bstrlib.h>
```

```
TSTree *node = NULL;
char *valueA = "VALUEA";
char *valueB = "VALUEB";
char *value2 = "VALUE2";
char *value4 = "VALUE4";
```

```

char *reverse = "VALUER";
int traverse_count = 0;

struct tagbstring test1 = bsStatic("TEST");
struct tagbstring test2 = bsStatic("TEST2");
struct tagbstring test3 = bsStatic("TSET");
struct tagbstring test4 = bsStatic("T");

char *test_insert()
{
    node = TSTree_insert(node, bdata(&test1), blength(&test1), valueA);
    mu_assert(node != NULL, "Failed to insert into tst.");

    node = TSTree_insert(node, bdata(&test2), blength(&test2), value2);
    mu_assert(node != NULL, "Failed to insert into tst with second name.");

    node = TSTree_insert(node, bdata(&test3), blength(&test3), reverse);
    mu_assert(node != NULL, "Failed to insert into tst with reverse name.");

    node = TSTree_insert(node, bdata(&test4), blength(&test4), value4);
    mu_assert(node != NULL, "Failed to insert into tst with second name.");

    return NULL;
}

char *test_search_exact()
{
    // tst returns the last one inserted
    void *res = TSTree_search(node, bdata(&test1), blength(&test1));
    mu_assert(res == valueA, "Got the wrong value back, should get A not B.");

    // tst does not find if not exact
    res = TSTree_search(node, "TESTNO", strlen("TESTNO"));
    mu_assert(res == NULL, "Should not find anything.");

    return NULL;
}

char *test_search_prefix()
{
    void *res = TSTree_search_prefix(node, bdata(&test1), blength(&test1));
    debug("result: %p, expected: %p", res, valueA);
    mu_assert(res == valueA, "Got wrong valueA by prefix.");

    res = TSTree_search_prefix(node, bdata(&test1), 1);
    debug("result: %p, expected: %p", res, valueA);
    mu_assert(res == value4, "Got wrong value4 for prefix of 1.");

    res = TSTree_search_prefix(node, "TE", strlen("TE"));
    mu_assert(res != NULL, "Should find for short prefix.");

    res = TSTree_search_prefix(node, "TE--", strlen("TE--"));
    mu_assert(res != NULL, "Should find for partial prefix.");

    return NULL;
}

void TSTree_traverse_test_cb(void *value, void *data)
{
    assert(value != NULL && "Should not get NULL value.");
    assert(data == valueA && "Expecting valueA as the data.");
    traverse_count++;
}

char *test_traverse()
{
    traverse_count = 0;
    TSTree_traverse(node, TSTree_traverse_test_cb, valueA);
    debug("traverse count is: %d", traverse_count);
}

```

```

mu_assert(traverse_count == 4, "Didn't find 4 keys.");

return NULL;
}

char *test_destroy()
{
    TSTree_destroy(node);

    return NULL;
}

char * all_tests() {
    mu_suite_start();

    mu_run_test(test_insert);
    mu_run_test(test_search_exact);
    mu_run_test(test_search_prefix);
    mu_run_test(test_traverse);
    mu_run_test(test_destroy);

    return NULL;
}

RUN_TESTS(all_tests);

```

## Advantages And Disadvantages

There's other interesting practical things you can do with a `TSTree`:

- In addition to finding prefixes, you can reverse all the keys you insert, and then find by *suffix*. I use this to lookup host names, since I want to find `*.learncodethehardway.com` so if I go backwards I can match them quickly.
- You can do "approximate" matching, where you gather all the nodes that have most of the same characters as the key, or using other algorithms for what's a close match.
- You can find all the keys that have a part in the middle.

I've already talked about some of the things `TSTrees` can do, but they aren't the best data structure all the time. The disadvantages of the `TSTree` are:

- As I mentioned, deleting from them is murder. They are better for data that needs to be looked up fast and you rarely remove from. If you need to delete then simply disable the `value` and then periodically rebuild the tree when it gets too big.
- It uses a ton of memory compared to `BSTree` and `Hashmaps` for the same key space. Think about it, it's using a full node for each character in every key. It might do better for smaller keys, but if you put a lot in a `TSTree` it will get huge.
- They also do not work well with large keys, but "large" is subjective so as usual test first. If you're trying to store 10k character sized keys then use a `HashMap`.

## How To Improve It

As usual, go through and improve this by adding the defensive preconditions, asserts, and checks to each function. There's some other possible improvements, but you don't necessarily have to implement all



of these:

- You could allow duplicates by using a `DArray` instead of the `value`.
- As I mentioned deleting is hard, but you could simulate it by setting the values to `NULL` so they are effectively gone.
- There are no ways to collect all the possible matching values. I'll have you implement that in an extra credit.
- There are other algorithms that are more complex but have slightly better properties. Take a look at Suffix Array, Suffix Tree, and Radix Tree structures.

## Extra Credit

- Implement a `TSTree_collect` that returns a `DArray` containing all the keys that match the given prefix.
- Implement `TSTree_search_suffix` and a `TSTree_insert_suffix` so you can do suffix searches and inserts.
- Use `valgrind` to see how much memory this structure uses to store data compared to the `BSTree` and `Hashmap`.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



## Exercise 47: A Fast URL Router

I'm going to now show you how I use the `TSTree` to do fast URL routing in web servers I've written. This works for simple URL routing you might use at the edge of an application, not really for the more complex (and sometimes unnecessary) routing found in many web application frameworks.

To play with routing I'm going to make a little command line tool I'm calling `urlor` that reads a simple file of routes, and then prompts the user to enter in URLs to look up.

```
#include <lcthw/tstree.h>
#include <lcthw/bstrlib.h>

TSTree *add_route_data(TSTree *routes, bstring line)
{
    struct bstrList *data = bsplit(line, ' ');
    check(data->qty == 2, "Line '%s' does not have 2 columns",
          bdata(line));

    routes = TSTree_insert(routes,
                          bdata(data->entry[0]), blength(data->entry[0]),
                          bstrcpy(data->entry[1]));

    bstrListDestroy(data);

    return routes;
error:
    return NULL;
}

TSTree *load_routes(const char *file)
{
    TSTree *routes = NULL;
    bstring line = NULL;
    FILE *routes_map = NULL;

    routes_map = fopen(file, "r");
    check(routes_map != NULL, "Failed to open routes: %s", file);

    while((line = bgets((bNgetc)fgetc, routes_map, '\n')) != NULL) {
        check(btrimws(line) == BSTR_OK, "Failed to trim line.");
        routes = add_route_data(routes, line);
        check(routes != NULL, "Failed to add route.");
        bdestroy(line);
    }

    fclose(routes_map);
    return routes;
error:
    if(routes_map) fclose(routes_map);
    if(line) bdestroy(line);

    return NULL;
}

bstring match_url(TSTree *routes, bstring url)
```

```

{
    bstring route = TSTree_search(routes, bdata(url), blength(url));

    if(route == NULL) {
        printf("No exact match found, trying prefix.\n");
        route = TSTree_search_prefix(routes, bdata(url), blength(url));
    }

    return route;
}

bstring read_line(const char *prompt)
{
    printf("%s", prompt);

    bstring result = bgets((bNgetc)fgetc, stdin, '\n');
    check_debug(result != NULL, "stdin closed.");

    check(btrimws(result) == BSTR_OK, "Failed to trim.");

    return result;

error:
    return NULL;
}

void bdestroy_cb(void *value, void *ignored)
{
    (void)ignored;
    bdestroy((bstring)value);
}

void destroy_routes(TSTree *routes)
{
    TSTree_traverse(routes, bdestroy_cb, NULL);
    TSTree_destroy(routes);
}

int main(int argc, char *argv[])
{
    bstring url = NULL;
    bstring route = NULL;
    check(argc == 2, "USAGE: urlor <urlfile>");

    TSTree *routes = load_routes(argv[1]);
    check(routes != NULL, "Your route file has an error.");

    while(1) {
        url = read_line("URL> ");
        check_debug(url != NULL, "goodbye.");

        route = match_url(routes, url);

        if(route) {
            printf("MATCH: %s == %s\n", bdata(url), bdata(route));
        } else {
            printf("FAIL: %s\n", bdata(url));
        }

        bdestroy(url);
    }

    destroy_routes(routes);
    return 0;

error:
    destroy_routes(routes);
    return 1;
}

```

I'll then make a simple file with some fake routes to play with:

```
/ MainApp /hello Hello /hello/ Hello /signup Signup /logout Logout
/album/ Album
```

## What You Should See

Once you have `urlor` working and a routes file, you can try it out:

```
$ ./bin/urlor urls.txt
URL> /
MATCH: / == MainApp
URL> /hello
MATCH: /hello == Hello
URL> /hello/zed
No exact match found, trying prefix.
MATCH: /hello/zed == Hello
URL> /album
No exact match found, trying prefix.
MATCH: /album == Album
URL> /album/12345
No exact match found, trying prefix.
MATCH: /album/12345 == Album
URL> asdfasfdasfd
No exact match found, trying prefix.
FAIL: asdfasfdasfd
URL> /asdfasfdasf
No exact match found, trying prefix.
MATCH: /asdfasfdasf == MainApp
URL>
$
```

You can see that the routing system first tries an exact match, and then if it can't find one it will give a prefix match. This is mostly to try out the difference between the two. Depending on the semantics of your URLs you may want to always match exactly, always to prefixes, or do both and pick the "best" one.

## How To Improve It

URLs are weird because people want them to magically handle all of the insane things their web applications do, even if that's not very logical. In this simple demonstration of how to use the `TSTree` to do routing, it has some flaws that people wouldn't be able to articulate. For example, it will match `/al` to `Album`, which generally isn't what they want. They want `/album/*` to match `Album` and `/al` to be a 404 error.

This isn't difficult to implement though, since you could change the prefix algorithm to match any way you want. If you change the matching algorithm to find *all* matching prefixes, and then pick the "best" one, you'll be able to do it easily. In this case, `/al` could match `MainApp` or `Album`. Take those results then do a little logic on which is "best".

Another thing you can do in a real routing system is use the `TSTree` to find all possible matches, but that these matches are a small set of patterns to check. In many web applications there's a list of regex that have to be matched against URLs on each request. Running all the regex can be time consuming, so you can use a `TSTree` to find all the possible ones by their prefixes. Then you narrow the patterns to try down to a few

very quickly.

Using this method, your URLs will match exactly since you are actually running real regex patterns, and they'll match much faster since you're finding them by possible prefixes.

This kind of algorithm also works for anything else that needs to have flexible user-visible routing mechanisms. Domain names, IP address, registries and directories, files, or URLs.

## Extra Credit

- Instead of just storing the string for the handler, create an actual engine that uses an `Handler` struct to store the application. The struct would store the URL it is attached to, the name, and anything else you'd need to make an actual routing system.
- Instead of mapping URLs to arbitrary names, map them to .so files and use the `dlopen` system to load handlers on the fly and call callbacks they contain. Put these callbacks in your `Handler` struct and then you have yourself a fully dynamic callback handler system in C.

### Interested In Python?

Python is also a  
great language.

**Learn  
Python The  
Hard Way**

### Interested In Ruby?

Ruby is also a great language.

**Learn Ruby The Hard Way**



# LEARN C

## THE HARD WAY

BETA  
EDITION

## Exercise 48: A Tiny Virtual Machine Part 1

The rest of the book will be implementing a version of the DCPU16 virtual machine using the algorithms created so far. This will be done in 5 parts so it's broken down and understandable. It will apply nearly everything taught so far.

## What You Should See

## How To Break It

## Extra Credit

1 —  
2 —  
3 —  
4 — MAIN

◀  
PREVIOUS

▶ NEXT

🚑 HELP

Follow  
@l2sthw

Interested  
In  
Python?

Python is  
also a great  
language.

Learn  
Python  
The Hard  
Way

Interested In  
Ruby?

Ruby is also a great language.

Learn Ruby The Hard  
Way



# LEARN C

## THE HARD WAY

BETA  
EDITION

## Exercise 48: A Tiny Virtual Machine Part 2

### What You Should See

### How To Break It

### Extra Credit

1 \_\_\_\_\_  
2 \_\_\_\_\_  
3 \_\_\_\_\_  
4 \_\_\_\_\_ MAIN

◀  
PREVIOUS

▶  
NEXT

🚚  
HELP

Follow  
@l2sthw

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Learn  
Python  
The Hard  
Way**

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**



# LEARN C

## THE HARD WAY

BETA  
EDITION

### Exercise 50: A Tiny Virtual Machine Part 3

### What You Should See

### How To Break It

### Extra Credit

1 —  
2 —  
3 — MAIN  
4 —

◀  
PREVIOUS

▶ NEXT

🚚 HELP

Follow  
@l2sthw

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Learn  
Python  
The Hard  
Way**

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**





# LEARN C

## THE HARD WAY

BETA  
EDITION

## Exercise 51: A Tiny Virtual Machine Part 4

### What You Should See

### How To Break It

### Extra Credit

1 —  
2 —  
3 —  
4 — MAIN

◀  
PREVIOUS

▶  
NEXT

🚚  
HELP

Follow  
@l2sthw

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Learn  
Python  
The Hard  
Way**

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**



# LEARN C

## THE HARD WAY

BETA  
EDITION

## Exercise 52: A Tiny Virtual Machine Part 5

### What You Should See

### How To Break It

### Extra Credit

1 —  
2 —  
3 —  
4 — MAIN

◀  
PREVIOUS

▶  
NEXT

🚚  
HELP

Follow  
@l2sthw

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Learn  
Python  
The Hard  
Way**

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**



# LEARN C

## THE HARD WAY

BETA  
EDITION

## Deconstructing **K&RC** Is Dead

I have lost. I am giving up after years of trying to figure out how I can get the message out that the way C has been written since its invention is flawed. Originally I had a section of my book called Deconstructing K&R C. The purpose of the section is to teach people to never assume that their code is correct, or that the code of anyone, no matter how famous, is free of defects. This doesn't seem to be a revolutionary idea, and to me is just part of how you analyze code for defects and get better at making your own work solid.

Over the years, this one piece of writing has tanked the book and received more criticism and more insults than any other thing I've written. Not only that, but the criticisms of this part of the book end up being along the lines of, "You're right, but you're wrong that their code is bad." I cannot fathom how a group of people who are supposedly so intelligent and geared toward rational thought can hold in their head the idea that I can be wrong, and also right at the same time. I've had to battle pedants on ##c IRC channels, email chains, comments, and in every case they come up with minor tiny weird little pedantic jabs that require ever more logical modifications to my prose to convince them.

The interesting data point is that before I wrote that part of the book I received positive comments about the book. It was a work in progress so I felt it'd need to be improved for sure. I even setup a bounty at one point to get people to help with that. Sadly, once they were blinded by their own hero worship the tone changed dramatically. I became actually hated. For doing nothing more than trying to teach people how to use an error prone shitty language like C safely. Something I'm pretty good at.

It didn't matter that most of these detractors admitted to me that they don't code C anymore, that they don't teach it, and that they just memorized the standard so they could "help" people. It didn't matter that I was entirely open to trying to fix things and even offered to pay people bounties to help fix it. It

1 —  
2 —  
3 —  
4 — MAIN

◀  
PREVIOUS

▶ NEXT

🚑 HELP

Follow  
@l2sthw

didn't matter that this could get more people to love C and help others get into programming. All that mattered was I "insulted" their heroes and that means everything I said is permanently broken, never to be deemed worthy ever again.

Frankly, this is the deep dark ugly evil side of programming culture. They talk all day long of how, "We're all in this together" but if you don't bow to the great altar of their vast knowledge and beg them for permission to question what they believe you are suddenly the enemy. Programmers consistently go out of their way to set themselves up in positions of power that require others to pay homage to their brilliant ability to memorize standards or know obscure trivia, and will do their very best to destroy anyone who dares question that.

It's disgusting, and there's nothing I can do about it. I *cannot* help old programmers. They are all doomed. Destined to have all the knowledge they accumulated through standards memorization evaporate at the next turn of the worm. They have no interest in questioning the way things are and potentially improving things, or helping teach their craft to others unless that education involves a metric ton of ass kissing to make them feel good. Old programmers are just screwed.

I can't do anything about their current power over younger new programmers. I can't prevent the slander by incompetent people who haven't worked as professional C coders ever. And I'd rather make the book useful for people who can learn C and how to make it solid than fight a bunch of closed minded conservatives who's only joy in life is feeling like they know more about a pedantic pathetically small topic like C undefined behavior.

With that in mind, I'm removing the K&R C part of the book and I finally have my new theme. I've wanted to rewrite the book but couldn't figure out how to do it. I was held in limbo because I was personally too attached to something I felt was important, but that I couldn't advance forward. I now realize this was wrong because it prevented me from teaching many new programmers important skills unrelated to C. Skills in rigor, code analysis, defects, security flaws, and how to learn any programming language.

Now I know that I will make the book a course in writing the best secure C code possible and breaking C code as a way to

learn both C and also rigorous programming. I will fill it with pandering to the pedants saying that my humble book is merely a gateway to C and that all should go read K&R C and worship at the feet of the great golden codes held within. I will make it clear that my version of C is limited and odd on purpose because it makes my code safe. I will be sure to mention all of the pedantic things that every pedant demands about NULL pointers on a PDP-11 computer from the 1960s.

And then I will also tell people to never write another C program again. It won't be obvious. It won't be outright, but my goal will be to move people right off C onto other languages that are doing it better. Go, Rust, and Swift, come to mind as recent entrants that can handle the majority of tasks that C does now, so I will push people there. I will tell them that their skills at finding defects, and rigorous analysis of C code will pay massive dividends in every language and make learning any other language possible.

But C? C's dead. It's the language for old programmers who want to debate section A.6.2 paragraph 4 of the undefined behavior of pointers. Good riddance. I'm going to go learn Go (or Rust, or Swift, or anything else).

**Interested  
In  
Python?**

Python is  
also a great  
language.

**Learn  
Python  
The Hard  
Way**

**Interested In  
Ruby?**

Ruby is also a great language.

**Learn Ruby The Hard  
Way**